

Kube-OVN Document

v1.14.4

Kube-OVN Team

2025 Kube-OVN Team

Table of contents

1.	Kube-OVN	5
1.1	What is Kube-OVN?	5
1.2	Why Kube-OVN?	5
1.3	CNI	6
1.4	OVN/ovn-kubernetes/Kube-OVN	6
1.5		7
2.		8
2.1		8
2.2		10
2.3	Underlay	12
2.4	Talos	17
2.5		18
2.6		20
3.		21
3.1		21
3.2		28
3.3	IP	31
3.4		32
3.5	EIP SNAT	33
3.6	QoS	36
3.7		38
3.8	Webhook	39
3.9		41
3.10	NetworkPolicy	47
3.11	LoadBalancer Service	49
3.12		54
4.	KubeVirt	58
4.1	VM IP	58
4.2		61
4.3		63
4.4		65
4.5	DHCP	70
5.	VPC	71
5.1	VPC	71
5.2	VPC Egress Gateway	80

5.3	VPC QoS	90
5.4	VPC	93
5.5	VPC DNS	97
5.6	SecurityGroup	100
5.7	OVN EIP FIP SNAT DNAT	103
5.8	OVN SNAT ECMP BFD L3 HA	112
5.9	VPC	116
6.		118
6.1	kubectl	118
6.2		128
6.3	ovn-central	130
6.4	OVN	134
6.5	CIDR	136
6.6	Join CIDR	137
6.7		138
6.8		139
7.		143
7.1		143
7.2		149
7.3	FastPath	154
7.4	eBPF TCP	155
7.5	OVN-IC	159
7.6	Submariner	165
7.7	Overlay	167
7.8	BGP	169
7.9	MetalLB Kube-OVN Underlay	173
7.10	Cilium	177
7.11	Cilium NetworkPolicy	180
7.12	Cilium	184
7.13		189
7.14	VIP IP	190
7.15	Mellanox Offload	194
7.16	Offload	203
7.17	Offload	207
7.18	Offload	210
7.19	DPDK	213
7.20	OpenStack	217
7.21	IPsec	221

7.22 OVN	222
7.23 DNS Kube-OVN	225
7.24 VPC NAT	227
8.	228
8.1	228
8.2 What's Next	231
8.3	232
8.4	236
8.5	238
8.6 Underlay	242
8.7 Iptables	249
8.8	254
8.9 OVS/OVN	256
8.10	258
8.11 Kube-OVN	260
8.12 Kube-OVN	266
8.13 Annotation	287
8.14	291
9.	294
9.1	294

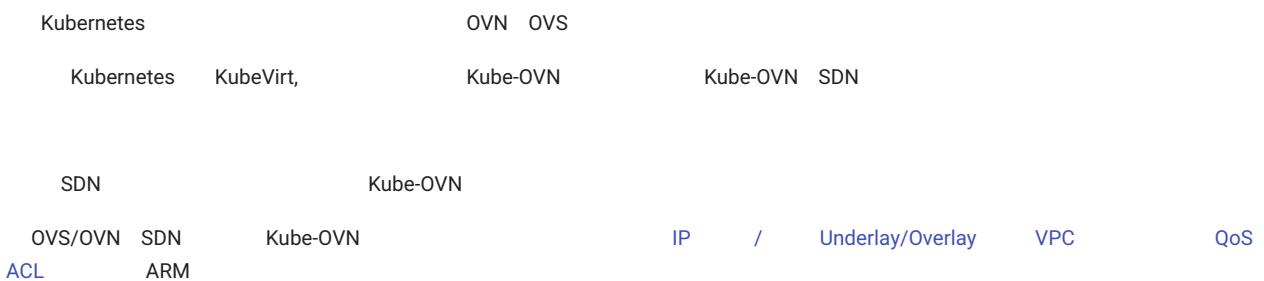
1. Kube-OVN



1.1 What is Kube-OVN?



1.2 Why Kube-OVN?



1.3 CNI



1.3.1 eBPF

[Cilium](#) [Calico eBPF](#)
[Kube-OVN](#) [Open vSwitch](#)

1.3.2 CNI, Ingress, Service Mesh Observability All in One

[Cilium](#)
[Kube-OVN](#) CNI

1.3.3 OpenShift

[ovn-kubernetes](#)
[OpenShift](#) CNI [Cluster Network Operator](#) Kube-OVN RedHat Kubernetes

1.3.4 Kubernetes EKS/AKS/GKE

Kubernetes CNI

1.3.5 AI

Hostnetwork [host-device](#)
AI

1.4 OVN/ovn-kubernetes/Kube-OVN

1.4.1 OVN

[OVN](#) Open vSwitch
Kubernetes OpenStack, Kubernetes CMS ovn-kubernetes Kube-OVN OVN

1.4.2 ovn-kubernetes

[ovn-kubernetes](#) OVN OVN Kubernetes CNI OpenShift OpenShift [UDN, Multihoming, Hardware Acceleration](#)

1.4.3 Kube-OVN

Kube-OVN IP Namespace
Underlay VPC KubeVirt OVN ovn-kubernetes annotation Pod join



⌚2025 9 10

⌚2022 5 20



1.5

2.

2.1

Kube-OVN CNI Kubernetes

2.1.1

- Kubernetes >= 1.29
- Docker >= 1.12.6, Containerd >= 1.3.4
- : CentOS 7/8, Ubuntu 16.04/18.04/20.04
- Linux geneve, openvswitch, ip_tables, iptable_nat Kube-OVN

1.	3.10.0-862	netfilter	bug	Kube-OVN	CentOS	bug	Floating IPs broken after kernel upgrade to Centos/RHEL 7.5 - DNAT not working
2.	Rocky Linux 8.6	4.18.0-372.9.1.el8.x86_64	TCP	TCP connection failed in Rocky Linux 8.6		4.18.0-372.13.1.el8_6.x86_64	
3.	4.4	openvswitch		openvswitch			
4.	Geneve	IPv6	cat /proc/cmdline		bug	Geneve tunnels don't work when ipv6 is disabled	

2.1.2

- IPv6 ipv6.disable=1 0
- kube-proxy Kube-OVN Service ClusterIP kube-apiserver
- kubelet CNI , kubelet --network-plugin=cni --cni-bin-dir=/opt/cni/bin --cni-conf-dir=/etc/cni/net.d /etc/cni/net.d/

2.1.3

ovn-central	6641/tcp	ovn nb db server
ovn-central	6642/tcp	ovn sb db server
ovn-central	6643/tcp	ovn northd server
ovn-central	6644/tcp	ovn raft server
ovn-ic	6645/tcp	ovn ic nb db server
ovn-ic	6646/tcp	ovn ic sb db server
ovs-ovn	Geneve 6081/udp, STT 7471/tcp, Vxlan 4789/udp	
kube-ovn-controller	10660/tcp	
kube-ovn-daemon	10665/tcp	
kube-ovn-monitor	10661/tcp	

firewalld Packet Forwarding Masquerade

```
# Packet Forwarding
firewall-cmd --add-forward --permanent
```

```
#   IPv4 Masquerade
firewall-cmd --add-masquerade --permanent
#   Kube-OVN IPv6/           Masquerade
firewall-cmd --permanent --add-rich-rule 'rule family="ipv6" source address="fd00:10:16::/112" masquerade'

firewall-cmd --reload
```

[!\[\]\(7e49c700e4adaed94ad5398cf2e7059e_img.jpg\) PDF](#)[!\[\]\(5ebcf382a6ee952d6c5b8b948415801e_img.jpg\) Slack](#)[!\[\]\(71ceb62b681518c82e95d615e7265d66_img.jpg\) Support](#)

⌚2025 9 10

⌚2022 5 20



2.1.4

2.2



2.2.1

release

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/refs/tags/v1.14.9/dist/images/install.sh
```

master

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/images/install.sh
```

```

REGISTRY="kubeovn"
VERSION="v1.14.9"
#           # /Tag
POD_CIDR="10.16.0.0/16"      # CIDR      SVC/NODE/JOIN CIDR
SVC_CIDR="10.96.0.0/12"       # apiserver  service-cluster-ip-range
JOIN_CIDR="100.64.0.0/16"     # Pod        CIDR      SVC/NODE/POD CIDR
LABEL="node-role.kubernetes.io/master" # OVN DB
IFACE=""                      #                 Kubernetes   Node IP
TUNNEL_TYPE="geneve"          #             geneve, vxlan   stt stt   ovs

```

IFACE=enp6s0f0,eth.*

root

```
bash install.sh
```

Kube-OVN

1. [Step 4/6] Pod
2. Kube-OVN

2.2.2 Helm Chart

Kube-OVN Helm Kube-OVN

IP

```
# kubectl get node -o wide
NAME           STATUS    ROLES      AGE   VERSION   INTERNAL-IP   EXTERNAL-IP   OS-IMAGE      KERNEL-VERSION   CONTAINER-RUNTIME
kube-ovn-control-plane  NotReady  control-plane  20h  v1.26.0  172.18.0.3  <none>        Ubuntu 22.04.1 LTS  5.10.104-linuxkit  containerd://1.6.9
kube-ovn-worker    NotReady  <none>      20h  v1.26.0  172.18.0.2  <none>        Ubuntu 22.04.1 LTS  5.10.104-linuxkit  containerd://1.6.9
```

label

```
# kubectl label node -lbeta.kubernetes.io/os=linux kubernetes.io/os=linux --overwrite
node/kube-ovn-control-plane not labeled
node/kube-ovn-worker not labeled

# kubectl label node -lnode-role.kubernetes.io/control-plane kube-ovn/role=master --overwrite
node/kube-ovn-control-plane labeled

#   label    dpdk      dpdk
# kubectl label node -lvn.kubernetes.io/ovs_dp_type!=userspace ovn.kubernetes.io/ovs_dp_type=kernel --overwrite
node/kube-ovn-control-plane labeled
node/kube-ovn-worker labeled
```

Helm Repo

```
# helm repo add kubeovn https://kubeovn.github.io/kube-ovn/
" kubeovn " has been added to your repositories

# helm repo list
NAME          URL
kubeovn      https://kubeovn.github.io/kube-ovn/

# helm repo update kubeovn
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the " kubeovn " chart repository
Update Complete. *Happy Helm-ing!*

# helm search repo kubeovn
NAME          CHART VERSION   APP VERSION   DESCRIPTION
kubeovn/kube-ovn   v1.14.9       v1.14.9       Helm chart for Kube-OVN
```

helm install Kube-OVN

Chart values.yaml

```
# helm install kube-ovn kubeovn/kube-ovn --wait -n kube-system --version v1.14.9
NAME: kube-ovn
LAST DEPLOYED: Thu Apr 24 08:30:13 2025
NAMESPACE: kube-system
STATUS: deployed
REVISION: 1
TEST SUITE: None
```

Helm values.yaml

```
helm upgrade -f values.yaml kube-ovn kubeovn/kube-ovn --wait -n kube-system --version v1.14.9
```

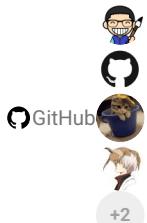
 PDF

 Slack

 Support

 2025 9 17

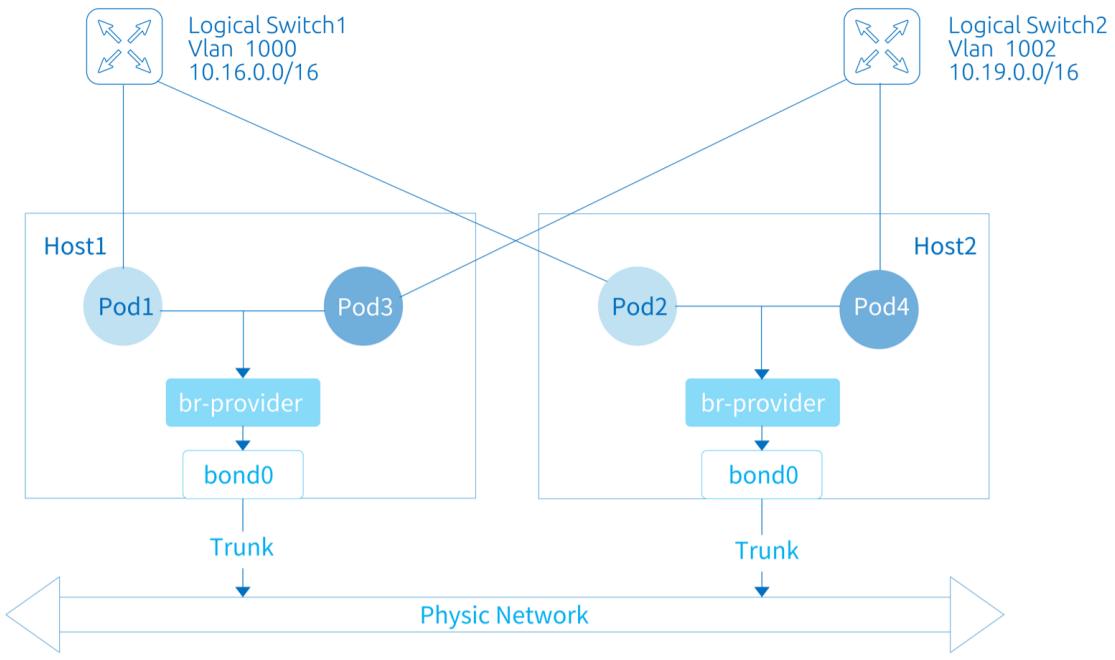
 2022 5 20



2.2.3

2.3 Underlay

Kube-OVN Geneve Overlay
 Kube-OVN Underlay



2.3.1

Overlay SNAT/EIP / L3 VPC Underlay

2.3.2 Macvlan

Kube-OVN	Underlay	Macvlan				
1. Macvlan	OVS	Macvlan				
2. Kube-OVN	arp-proxy	arp				
3. Macvlan	netfilter Service	NetworkPolicy	Kube-OVN	OVS	Service	NetworkPolicy
4. Kube-OVN	Underlay	Macvlan	IP	QoS		

2.3.3

Underlay	OVS	OVS	L2/L3	Vlan
1. OpenStack VM		PortSecurity		
2. VMware vSwitch		MAC Address Changes, Forged Transmits	Promiscuous Mode Operation	allow
3. Hyper-V		MAC Address Spoofing		
4. AWS GCE	Mac	Underlay	Underlay	VPC-CNI
5. Linux Bridge				

	Kube-OVN	Mac	IP	MTU	OVS Bridge	Underlay	OVS Bridge	br-PROVIDER_NAME
PROVIDER_NAME	Provider	provider						

2.3.4

Underlay Pod Underlay

```
 wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install.sh
```

```
ENABLE_ARP_DETECT_IP_CONFLICT #           vlan     arp
NETWORK_TYPE          #           vlan
VLAN_INTERFACE_NAME   #           eth1
VLAN_ID               #           VLAN Tag    0     VLAN
POD_CIDR              #           CIDR      192.168.1.0/24
POD_GATEWAY           #           192.168.1.1
EXCLUDE_IPS           #           IP        192.168.1.1..192.168.1.100
ENABLE_LB              #           Underlay  Service   true
EXCHANGE_LINK_NAME    #           provider-network OVS       false
LS_DNAT_MOD_DL_DST   #           DNAT      MAC       Service   true
```

```
bash install.sh
```

2.3.5 CRD Underlay

Underlay Pod ProviderNetwork Vlan Subnet

ProviderNetwork

ProviderNetwork Underlay

ProviderNetwork :

```
apiVersion: kubeovn.io/v1
kind: ProviderNetwork
metadata:
  name: net1
spec:
  defaultInterface: eth1
  customInterfaces:
    - interface: eth2
      nodes:
        - node1
  excludeNodes:
    - node2
```

ProviderNetwork 12

- ```
• defaultInterface: ProviderNetwork excludeNodes: br-net1 br-NAME OVS
• customInterfaces:
• excludeNodes: net1.provider-network.own.kubernetes.io/exclude=true
```

| Key                                               | Value |
|---------------------------------------------------|-------|
| net1.provider-network.ovn.kubernetes.io/ready     | true  |
| net1.provider-network.ovn.kubernetes.io/interface | eth1  |
| net1.provider-network.ovn.kubernetes.io/mtu       | 1500  |
|                                                   | MTU   |

IP IP OVS

**VLAN**

Vlan Vlan Tag ProviderNetwork

**VLAN**

```
apiVersion: kubeovn.io/v1
kind: Vlan
metadata:
 name: vlan1
spec:
 id: 0
 provider: net1
```

- id: VLAN ID/Tag Kube-OVN Vlan Vlan 0 vlan localnet
- provider: ProviderNetwork VLAN ProviderNetwork

**Subnet**

Vlan

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: subnet1
spec:
 protocol: IPv4
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 vlan: vlan1
 disableGatewayCheck: false
```

- vlan VLAN Subnet VLAN
- disableGatewayCheck Underlay true

**2.3.6**

IP

IP Pod IP Mac

**2.3.7**

Kube-OVN Underlay spec.logicalGateway true

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: subnet1
spec:
 protocol: IPv4
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 vlan: vlan1
 logicalGateway: true
```

|     |          |                |
|-----|----------|----------------|
| Pod | Kube-OVN | Logical Router |
|-----|----------|----------------|

### 2.3.8 Underlay Overlay

|          |         |          |                    |         |          |        |             |          |             |
|----------|---------|----------|--------------------|---------|----------|--------|-------------|----------|-------------|
| Underlay | Overlay | Overlay  | Pod                | NAT     | Underlay | Pod IP | Underlay    | Pod      | Overlay     |
|          | Overlay | Underlay | Pod                | Pod IP  | Underlay | Pod    |             |          |             |
| Underlay | Overlay |          | u2oInterconnection | true    | Kube-OVN |        | Underlay IP | Underlay | ovn-cluster |
|          |         | Kube-OVN | Underlay           | Overlay |          |        |             |          |             |

#### IP

|        |    |                 |                      |
|--------|----|-----------------|----------------------|
| subnet | IP | Underlay Subnet | u2oInterconnectionIP |
|--------|----|-----------------|----------------------|

|                 |     |
|-----------------|-----|
| Underlay Subnet | VPC |
|-----------------|-----|

|                 |     |                |     |                    |      |                 |     |
|-----------------|-----|----------------|-----|--------------------|------|-----------------|-----|
| Underlay Subnet | VPC | Overlay Subnet | VPC | u2oInterconnection | true | subnet.spec.vpc | VPC |
|-----------------|-----|----------------|-----|--------------------|------|-----------------|-----|

### 2.3.9

|    |         |        |         |          |                |    |      |
|----|---------|--------|---------|----------|----------------|----|------|
| IP | Netplan | Ubuntu | Netplan | renderer | NetworkManager | IP | DHCP |
|----|---------|--------|---------|----------|----------------|----|------|

```
network:
 renderer: NetworkManager
 ethernets:
 eth0:
 dhcp4: no
 addresses:
 - 172.16.143.129/24
 version: 2
```

|    |         |
|----|---------|
| IP | netplan |
|----|---------|

```
netplan generate

nmcli connection reload netplan-eth0
nmcli device set eth0 managed yes
```

|          |    |     |
|----------|----|-----|
| Kube-OVN | IP | OVS |
|----------|----|-----|

|                |        |
|----------------|--------|
| NetworkManager | CentOS |
|----------------|--------|

```
nmcli connection reload eth0
nmcli device set eth0 managed yes
nmcli -t -f GENERAL.STATE device show eth0 | grep -qw unmanaged || nmcli device reapply eth0
```

|    |     |
|----|-----|
| IP | MAC |
|----|-----|

### 2.3.10

#### hairpin Pod

|         |     |          |     |     |
|---------|-----|----------|-----|-----|
| hairpin | Pod | Pod      | OVS | MAC |
|         |     | Kube-OVN |     |     |

#### Pod Pod

|     |     |     |     |          |
|-----|-----|-----|-----|----------|
| Pod | 300 | ARP | OVS | resubmit |
|-----|-----|-----|-----|----------|

```
2022-11-13T08:43:46.782Z|00222|ofproto_dpif_upcall(handler5)|WARN|Flow: arp,in_port=331,vlan_tci=0x0000,dl_src=00:00:00:25:eb:39,dl_dst=ff:ff:ff:ff:ff:ff,arp_spa=10.213.131.240,arp_tpa=10.213.159.254,arp_op=1,arp_sha=00:00:00:25:eb:39,arp_tha=ff:ff:ff:ff:ff:ff
bridge("br-int")

0. No match.
```

```
>>> received packet on unknown port 331 <<<
drop

Final flow: unchanged
Megaflow: recirc_id=0,eth,arp,in_port=331,dl_src=00:00:00:25:eb:39
Datapath actions: drop
2022-11-13T08:44:34.077Z|00224|ofproto_dpif_xlate(handler5)|WARN|over 4096 resubmit actions on bridge br-int while processing
arp,in_port=13483,vlan_tci=0x0000,dl_src=00:00:00:59:ef:
13,dl_dst=ff:ff:ff:ff:ff:ff,arp_spa=10.213.152.3,arp_tpa=10.213.159.254,arp_op=1,arp_sha=00:00:00:59:ef:13,arp_tha=ff:ff:ff:ff:ff:ff
```

OVN NB    bcast\_arp\_req\_flood    false

```
kubectl ko nbctl set NB_Global . options:bcast_arp_req_flood=false
```

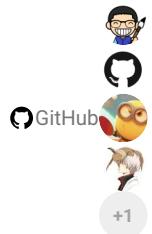
[!\[\]\(397cc4c04b5e7ea225dbaa029a5dee1f\_img.jpg\) PDF](#)

[!\[\]\(115eff7009a76771e6b7adb966005e4c\_img.jpg\) Slack](#)

[!\[\]\(a6eac08c103efb51b40f958fe35f07bb\_img.jpg\) Support](#)

 2025 9 10

 2022 5 20



 GitHub

+1

2.3.11

---

## 2.4 Talos

Talos Linux    Kubernetes    Linux

### 2.4.1 Helm Chart    Kube-OVN

Talos Linux    Kube-OVN

```
helm install kubeovn kubeovn/kube-ovn --wait \
-n kube-system \
--version v1.14.9 \
--set OVN_DIR=/var/lib/ovn \
--set OPENVSWITCH_DIR=/var/lib/openvswitch \
--set DISABLE_MODULES_MANAGEMENT=true \
--set cni_conf.MOUNT_LOCAL_BIN_DIR=false
```

Underlay

Helm

Chart

```
helm install kubeovn kubeovn/kube-ovn --wait \
-n kube-system \
--version v1.14.9 \
--set OVN_DIR=/var/lib/ovn \
--set OPENVSWITCH_DIR=/var/lib/openvswitch \
--set DISABLE_MODULES_MANAGEMENT=true \
--set cni_conf.MOUNT_LOCAL_BIN_DIR=false \
--set networking.NETWORK_TYPE=vlan \
--set networking.vlan.VLAN_INTERFACE_NAME=enp0s5f1 \
--set networking.vlan.VLAN_ID=0 \
--set networking.NET_STACK=ipv4 \
--set-json networking.EXCLUDE_IPS='["172.99.99.11..172.99.99.99"]' \
--set-json ipv4.POD_CIDR='["172.99.99.8/24"]' \
--set-json ipv4.POD_GATEWAY='["172.99.99.1"]'
```



#### Note

VLAN Bond Bridge    Underlay    Underlay    Talos    ignore=true

```
machine:
 network:
 interfaces:
 - interface: enp0s5f1
 ignore: true
```

PDF

Slack

Support

⌚2025 7 1

⌚2025 4 16



GitHub



### 2.4.2

## 2.5

---

Kube-OVN

pre-v0.0.1

Kube-OVN

### 2.5.1

---

- [Linkerd](#)
- [Elasticsearch](#)
- [EMQX](#)
- [KubeSphere](#)

### 2.5.2 OpenVswitch/OVN

---

Kube-OVN

OpenVswitch OVN

- [OVN](#)
- [OpenVswitch](#)

ovn-architecture

### 2.5.3 Kube-OVN

---

Kube-OVN

Kube-OVN      kubectl      bash

Kube-OVN

### 2.5.4

---

Kube-OVN

Kube-OVN

Kube-OVN E2E

### 2.5.5

---

Kube-OVN

[OpenTelemetry](#)[DeepFlow](#)

### 2.5.6

---

Kube-OVN

7\*24

Github Issue

Github Issue

AI

AI

[!\[\]\(d544e576856cad8f6b886171f7e2e480\_img.jpg\) PDF](#)
[!\[\]\(700de1146ce8d8d31c70e1dbe6152e98\_img.jpg\) Slack](#)
[!\[\]\(20bc8a2224e0b0c62d32464117ca7264\_img.jpg\) Support](#)

⌚2025 9 22

⌚2025 8 29



2.5.7

---

## 2.6

Kube-OVN

Kube-OVN

OVS

issue

Kube-OVN

### 2.6.1 Kubernetes

[Script Uninstall](#)    [Helm Uninstall](#)

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/cleanup.sh
bash cleanup.sh

helm uninstall kube-ovn -n kube-system
```

### 2.6.2

ovsdb openvswitch

```
rm -rf /var/run/openvswitch
rm -rf /var/run/ovn
rm -rf /etc/origin/openvswitch/
rm -rf /etc/origin/ovn/
rm -rf /etc/cni/net.d/00-kube-ovn.conflist
rm -rf /etc/cni/net.d/01-kube-ovn.conflist
rm -rf /var/log/openvswitch
rm -rf /var/log/ovn
rm -fr /var/log/kube-ovn
```

### 2.6.3

iptable/ipset

reboot

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 5 24



### 2.6.4

## 3.

## 3.1



Overlay Underlay

## 3.1.1

| Kube-OVN | Namespace | IP      | CIDR |
|----------|-----------|---------|------|
| Overlay  | NAT       | Flannel |      |
| Underlay | arping    |         |      |

spec default true ovn-default

```
kubectl get subnet ovn-default -o yaml
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 creationTimestamp: "2019-08-06T09:33:43Z"
 generation: 1
 name: ovn-default
 resourceVersion: "1571334"
```

```

selfLink: /apis/kubeovn.io/v1/subnets/ovn-default
uid: 7e2451f8-fb44-4f7f-b3e0-cfd27f6fd5d6
spec:
 cidrBlock: 10.16.0.0/16
 default: true
 excludeIps:
 - 10.16.0.1
 gateway: 10.16.0.1
 gatewayType: distributed
 natOutgoing: true
 private: false
 protocol: IPv4

```

### 3.1.2 Join

|                |       |      |         |          |      |      |      |      |
|----------------|-------|------|---------|----------|------|------|------|------|
| Kubernetes Pod | Node  | Pod  | Overlay | Kube-OVN | join | Node | ovn0 | join |
| Pods           | Nodes | ovn0 | Node    | Pod      | ovn0 |      | ovn0 |      |
|                |       | join | CIDR    |          | Join |      |      |      |

#### Join

|     |          |                              |          |         |
|-----|----------|------------------------------|----------|---------|
| Pod | hostport | externalTrafficPolicy: Local | NodePort | Service |
|-----|----------|------------------------------|----------|---------|

|      |      |
|------|------|
| join | CIDR |
|------|------|

```

kubectl get subnet join -o yaml
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 creationTimestamp: "2019-08-06T09:33:43Z"
 generation: 1
 name: join
 resourceVersion: "1571333"
 selfLink: /apis/kubeovn.io/v1/subnets/join
 uid: 9c744810-c678-4d50-8a7d-b8ec12ef91b8
spec:
 cidrBlock: 100.64.0.0/16
 default: false
 excludeIps:
 - 100.64.0.1
 gateway: 100.64.0.1
 gatewayNode: ""
 gatewayType: ""
 natOutgoing: false
 private: false
 protocol: IPv4

```

Node ovn0

```

ifconfig ovn0
ovn0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1420
 inet 100.64.0.4 netmask 255.255.0.0 broadcast 100.64.255.255
 inet6 fe80::800:ff:fe40:5 prefixlen 64 scopeid 0x20<link>
 ether 0a:00:00:40:00:05 txqueuelen 1000 (Ethernet)
 RX packets 18 bytes 1428 (1.3 KiB)
 RX errors 0 dropped 0 overruns 0 frame 0
 TX packets 19 bytes 1810 (1.7 KiB)
 TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

```

### 3.1.3

#### Namespace

```

cat <<EOF | kubectl create -f -
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: subnet1
spec:
 protocol: IPv4
 cidrBlock: 10.66.0.0/16
 excludeIps:

```

```

- 10.66.0.1..10.66.0.10
- 10.66.0.101..10.66.0.151
gateway: 10.66.0.1
gatewayType: distributed
natOutgoing: true
routeTable: ""
namespaces:
- ns1
- ns2
EOF

```

- cidrBlock : CIDR      VPC      Subnet CIDR
- excludeIps :              IP      Underlay
- gateway :      Overlay      Kube-OVN      Underlay
- namespaces :      Namespace      Namespace      Pod
- routeTable :

```

kubectl create ns ns1
namespace/ns1 created

kubectl run nginx --image=docker.io/library/nginx:alpine -n ns1
deployment.apps/nginx created

kubectl get pod -n ns1 -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
nginx-74d5899f46-n8wtg 1/1 Running 0 10s 10.66.0.11 node1 <none> <none>

```

## Workload

| Pod | Namespace | IP | Namespace | Workload | Pod | Annotation ovn.kubernetes.io/logical_switch |
|-----|-----------|----|-----------|----------|-----|---------------------------------------------|
|-----|-----------|----|-----------|----------|-----|---------------------------------------------|

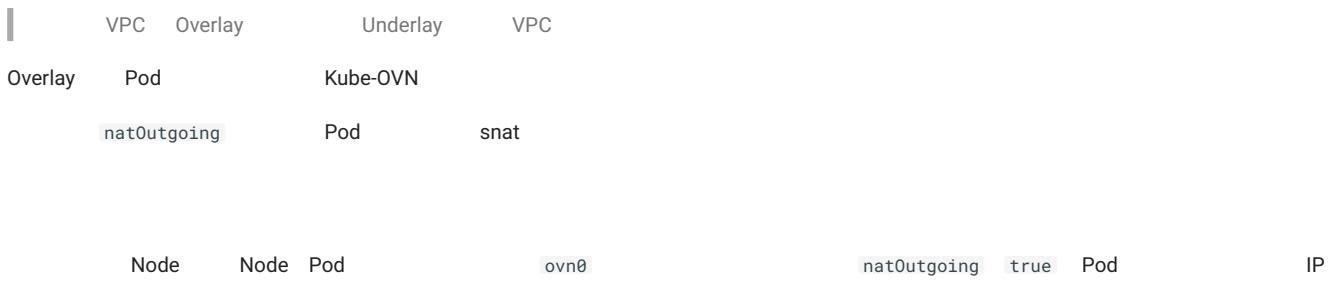
```

apiVersion: v1
kind: Pod
metadata:
 name: another-subnet
 annotations:
 ovn.kubernetes.io/logical_switch: subnet1
spec:
 containers:
 - name: another-subnet
 image: docker.io/library/nginx:alpine

```

| Workload                           | Deployment | StatefulSet | ovn.kubernetes.io/logical_switch Annotation |
|------------------------------------|------------|-------------|---------------------------------------------|
| spec.template.metadata.annotations |            |             |                                             |

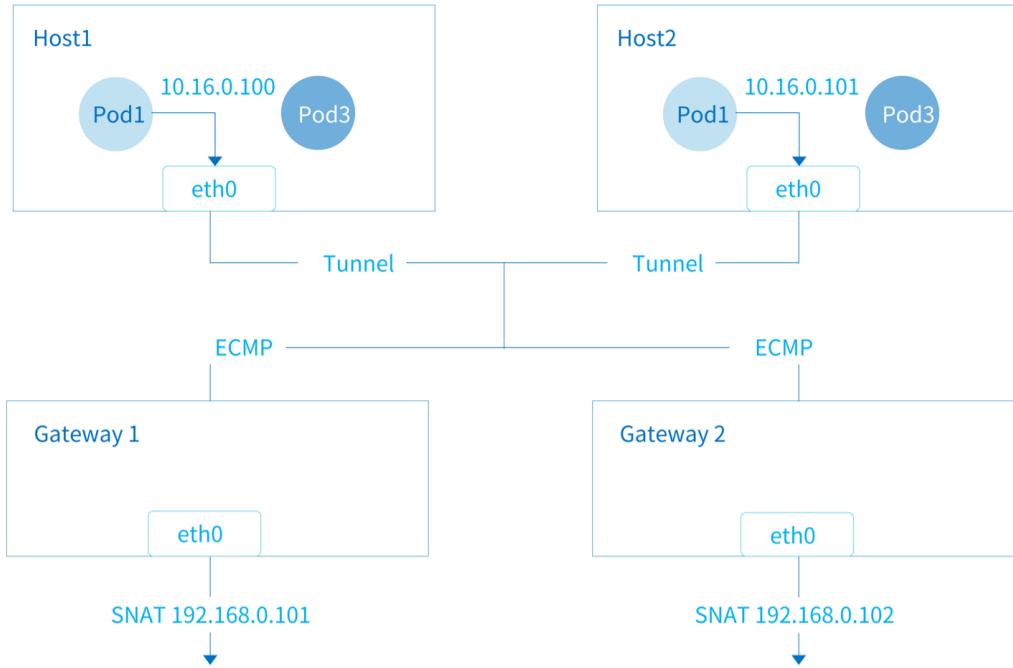
## 3.1.4 Overlay





```
gatewayType distributed
```

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: distributed
spec:
 protocol: IPv4
 cidrBlock: 10.166.0.0/16
 default: false
 excludeIps:
 - 10.166.0.1
 gateway: 10.166.0.1
 gatewayType: distributed
 natOutgoing: true
```



| Pod                                                                                                                                                                                                                                                                               | IP          | Pod                                                           | ovn0       | natOutgoing | true      |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|---------------------------------------------------------------|------------|-------------|-----------|
|                                                                                                                                                                                                                                                                                   | IP          | Kubernetes                                                    | NodeName   | gatewayNode |           |
| <pre>gatewayType      centralized   gatewayNode</pre>                                                                                                                                                                                                                             |             |                                                               |            |             |           |
|                                                                                                                                                                                                                                                                                   |             |                                                               |            |             |           |
| <pre>apiVersion: kubeovn.io/v1 kind: Subnet metadata:   name: centralized spec:   protocol: IPv4   cidrBlock: 10.166.0.0/16   default: false   excludeIps:   - 10.166.0.1   gateway: 10.166.0.1   gatewayType: centralized   gatewayNode: "node1,node2"   natOutgoing: true</pre> |             |                                                               |            |             |           |
| •                                                                                                                                                                                                                                                                                 | gatewayNode | kube-ovn-worker:172.18.0.2, kube-ovn-control-plane:172.18.0.3 |            |             |           |
| •                                                                                                                                                                                                                                                                                 | ECMP        | ECMP                                                          |            |             |           |
| • Kube-OVN v1.12.0                                                                                                                                                                                                                                                                | subnet crd  | spec                                                          | enableEcmp | ECMP        | ECMP      |
| controller Deployment                                                                                                                                                                                                                                                             | enable-ecmp |                                                               | v1.12.0    |             | kube-ovn- |

### 3.1.5 ACL

#### ⚠ Warning

Kube-OVN NetworkPolicy Network Policy API Subnet ACL Security Group OVN ACL NetworkPolicy NetworkPolicy API

|        |          |        |     |                                                  |
|--------|----------|--------|-----|--------------------------------------------------|
| ACL    | Kube-OVN | Subnet | ACL |                                                  |
| Subnet | ACL      | OVN    | ACL | ovn-nb ACL Table match ovn-sb Logical Flow Table |

IP 10.10.0.2 Pod

ACL

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: acl
spec:
 allowEWTraffic: false
 acls:
 - action: drop
 direction: to-lport
 match: ip4.dst == 10.10.0.2 && ip
 priority: 1002
 - action: allow-related
 direction: from-lport
 match: ip4.src == 10.10.0.2 && ip
 priority: 1002
 cidrBlock: 10.10.0.0/24
```

ACL

allowEWTraffic: true

### 3.1.6

ACL

ACL

Kube-OVN

Pod

CRD private true

allowSubnets allowSubnets

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: private
spec:
 protocol: IPv4
 default: false
 namespaces:
 - ns1
 - ns2
 cidrBlock: 10.69.0.0/16
 private: true
 allowSubnets:
 - 10.16.0.0/16
 - 10.18.0.0/16
```

### 3.1.7 Underlay

Underlay

- vlan: Underlay Subnet Vlan CR Underlay
- logicalGateway: Underlay OVN Underlay Overlay false

### 3.1.8

kube-ovn-cni Pod ICMP ARP

Underlay ICMP

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: disable-gw-check
spec:
 disableGatewayCheck: true
```

### 3.1.9 Multicast-Snoop

subnet Pod OVN

Pod subnet multicast snoop OVN South Database

Multicast\_Group

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: sample1
spec:
 enableMulticastSnoop: true
```

### 3.1.10 Subnet MTU

Subnet Pod MTU      Subnet Pod

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: sample1
spec:
 mtu: 1300
```

### 3.1.11

- IP
- VPC NAT
- QoS
- 
- DHCP
- 
- 
- IP

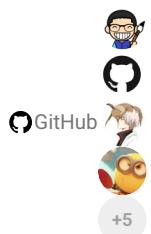
[!\[\]\(ab38aa34d7fd9612c4ea4d0692132800\_img.jpg\) PDF](#)

[!\[\]\(4754fc919b2e8116c30595fd4b918f00\_img.jpg\) Slack](#)

[!\[\]\(3d12067139d6a2b0989a839672a8beec\_img.jpg\) Support](#)

 2025 9 17

 2022 5 24



### 3.1.12

## 3.2

| Kube-OVN      | Pod     | Namespace | IP | Mac | Kube-OVN |
|---------------|---------|-----------|----|-----|----------|
| • Pod         | IP/Mac  |           |    |     |          |
| • Workload    | IP Pool |           |    |     |          |
| • StatefulSet |         |           |    |     |          |
| • KubeVirt VM |         |           |    |     |          |

### 3.2.1 Pod IP Mac

Pod annotation Pod IP/Mac, kube-ovn-controller

```
apiVersion: v1
kind: Pod
metadata:
 name: static-ip
 annotations:
 ovn.kubernetes.io/ip_address: 10.16.0.15 // 10.16.0.15,fd00:10:16::15
 ovn.kubernetes.io/mac_address: 00:00:00:53:6B:B6
spec:
 containers:
 - name: static-ip
 image: docker.io/library/nginx:alpine
```

annotation Pod IP/Mac

1. IP/Mac IP/Mac
2. IP CIDR
3. IP Mac

### 3.2.2 Workload IP Pool

| Kube-OVN   | annotation ovn.kubernetes.io/ip_pool | Workload Deployment/StatefulSet/DaemonSet/Job/CronJob | IP kube-ovn- |
|------------|--------------------------------------|-------------------------------------------------------|--------------|
| controller | ovn.kubernetes.io/ip_pool            | IP                                                    |              |
| IP Pool    | Annotation template annotation       | Kubernetes Workload                                   | Workload     |
| Deployment | IP                                   |                                                       |              |

#### Deployment IP

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: ippool
 labels:
 app: ippool
spec:
 replicas: 2
 selector:
 matchLabels:
 app: ippool
 template:
 metadata:
 labels:
 app: ippool
 annotations:
 ovn.kubernetes.io/ip_pool: 10.16.0.15,10.16.0.16,10.16.0.17 // 10.16.0.15,fd00:10:16::000E;10.16.0.16,fd00:10:16::000F;
 10.16.0.17,fd00:10:16::0010
 spec:
 containers:
 - name: ippool
 image: docker.io/library/nginx:alpine
```

| Workload                     | IP |          |      |     |  |          |  |                           |    |
|------------------------------|----|----------|------|-----|--|----------|--|---------------------------|----|
| 1. ovn.kubernetes.io/ip_pool | IP |          | CIDR |     |  |          |  |                           |    |
| 2. ovn.kubernetes.io/ip_pool | IP |          | IP   |     |  |          |  |                           |    |
| 3. ovn.kubernetes.io/ip_pool | IP | replicas |      | Pod |  | Workload |  | ovn.kubernetes.io/ip_pool | IP |

### 3.2.3 StatefulSet

| StatefulSet        | IP                        | Workload            | ovn.kubernetes.io/ip_pool | Pod | IP        |                           |             |       |                    |
|--------------------|---------------------------|---------------------|---------------------------|-----|-----------|---------------------------|-------------|-------|--------------------|
| StatefulSet        |                           | Kube-OVN            |                           |     |           |                           |             |       |                    |
| 1. Pod             | ovn.kubernetes.io/ip_pool | IP                  | StatefulSet               | web | web-0     | ovn.kubernetes.io/ip_pool | IP          | web-1 | IP                 |
| 2. StatefulSet Pod | OVN                       | logical_switch_port |                           | Pod | interface | Pod                       | IP/Mac      |       | StatefulSet Volume |
| 3. 2               | ovn.kubernetes.io/ip_pool | StatefulSet Pod     |                           |     | IP/Mac    |                           | StatefulSet |       |                    |

#### StatefulSet

```

apiVersion: apps/v1
kind: StatefulSet
metadata:
 name: web
spec:
 serviceName: "nginx"
 replicas: 2
 selector:
 matchLabels:
 app: nginx
 template:
 metadata:
 labels:
 app: nginx
 spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
 ports:
 - containerPort: 80
 name: web

```

StatefulSet Pod Pod IP

#### StatefulSet Pod IP

| StatefulSet        | IP | Pod Name    | Statefulset | ovn.kubernetes.io/ip_pool | Annotation | Pod         | IP |
|--------------------|----|-------------|-------------|---------------------------|------------|-------------|----|
| StatefulSet Pod IP |    | StatefulSet | scale       | 0                         | Annotation | StatefulSet |    |

### 3.2.4 KubeVirt VM

| KubeVirt | VM | kube-ovn-controller | StatefulSet Pod | IP | VM |
|----------|----|---------------------|-----------------|----|----|
| VM       | IP |                     |                 |    |    |

[PDF](#)
[Slack](#)
[Support](#)

⌚2025 9 10

⌚2022 5 20



⌚GitHub 🐱



3.2.5

---

### 3.3 IP

IP IPPool Subnet IPAM IP Namespace

#### 3.3.1

```
apiVersion: kubeovn.io/v1
kind: IPPool
metadata:
 name: pool-1
spec:
 subnet: ovn-default
 ips:
 - "10.16.0.201"
 - "10.16.0.210/30"
 - "10.16.0.220..10.16.0.230"
 namespaces:
 - ns-1
```

| subnet     |    |    |      |       |
|------------|----|----|------|-------|
| ips        | IP | .. | IPv6 |       |
| namespaces |    |    | Pod  | IP IP |

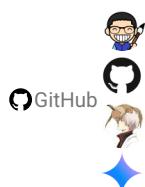
#### 3.3.2

1. Workload IP Pool IP IP
2. IP .spec.ips IP IP .spec.ips CIDR
3. IP IP IP
4. IP .spec.ips
5. IP IP IP IP IP
6. IP IP
7. IP Namespace



⌚2025 9 19

⌚2023 7 10



#### 3.3.3

## 3.4

### Pod Annotations

```
apiVersion: v1
kind: Pod
metadata:
 name: custom-routes
 annotations:
 ovn.kubernetes.io/routes: |
 [{ "dst": "192.168.0.101/24", "gw": "10.16.0.254" },
 { "gw": "10.16.0.254" }]
spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
```

dst

| Deployment | DaemonSet | StatefulSet | Annotation | .spec.template.metadata.annotations |
|------------|-----------|-------------|------------|-------------------------------------|
|------------|-----------|-------------|------------|-------------------------------------|

```
apiVersion: apps/v1
kind: Deployment
metadata:
 name: custom-routes
 labels:
 app: nginx
spec:
 replicas: 2
 selector:
 matchLabels:
 app: nginx
 template:
 metadata:
 labels:
 app: nginx
 annotations:
 ovn.kubernetes.io/routes: |
 [{ "dst": "192.168.0.101/24", "gw": "10.16.0.254" },
 { "gw": "10.16.0.254" }]
 spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
```

[PDF](#)
[Slack](#)
[Support](#)

⌚2023 9 26

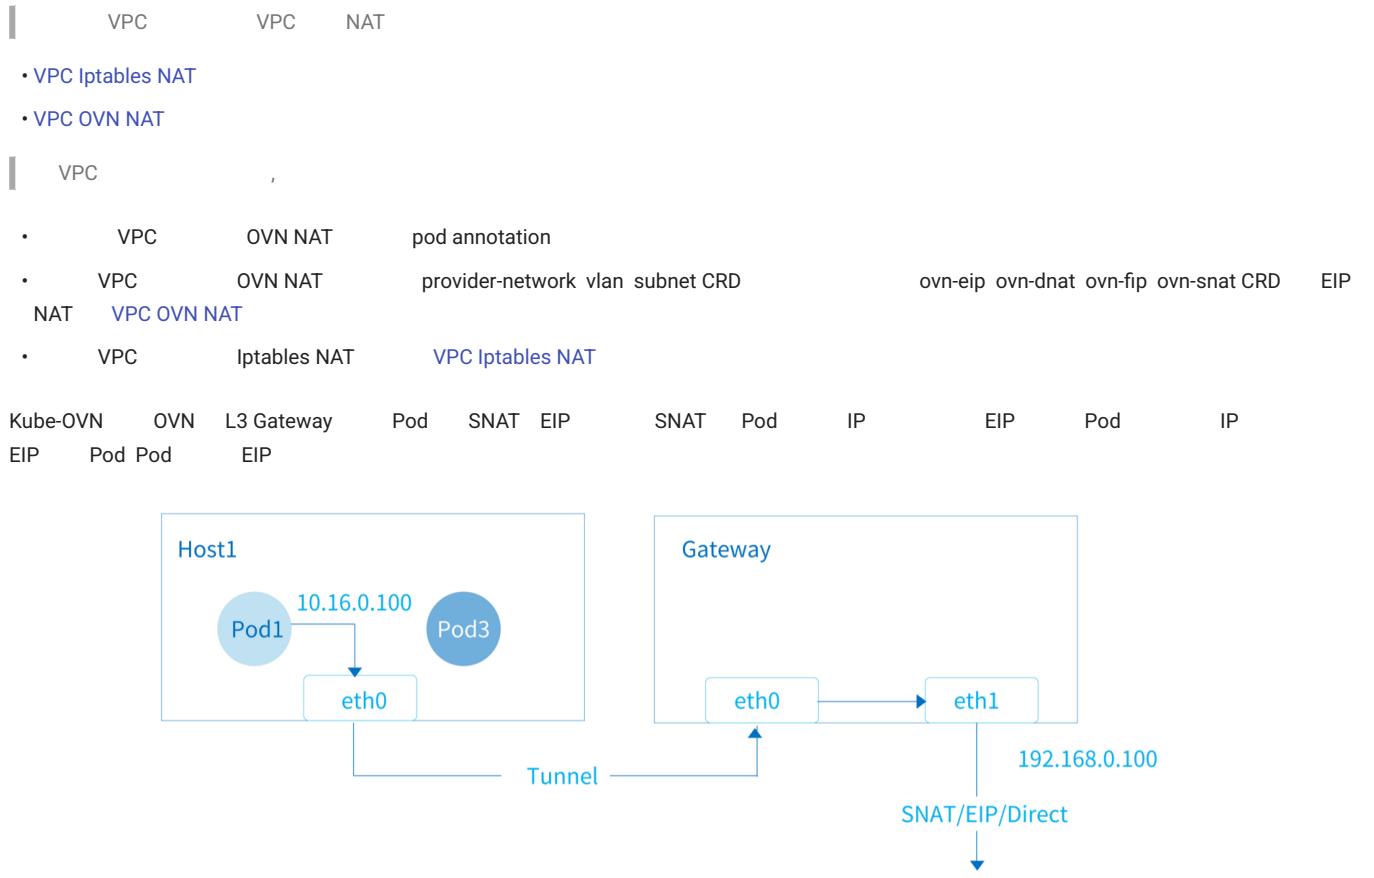
⌚2023 2 16



GitHub 

### 3.4.1

### 3.5 EIP SNAT



#### 3.5.1

- OVN L3 Gateway      OVS      Overlay Underlay
- NAT      Underlay
- EIP SNAT

#### 3.5.2

kube-system    ConfigMap ovn-external-gw-config

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-external-gw-config
 namespace: kube-system
data:
 enable-external-gw: "true"
 external-gw-nodes: "kube-ovn-worker"
 external-gw-nic: "eth1"
 external-gw-addr: "172.56.0.1/16"
```

```

nic-ip: "172.56.0.254/16"
nic-mac: "16:52:f3:13:6a:25"

• enable-external-gw: SNAT EIP
• type: centralized distributed centralized distributed
• external-gw-nodes: centralized
• external-gw-nic:
• external-gw-addr: IP
• nic-ip, nic-mac: IP Mac IP Mac

```

### 3.5.3 OVN OVS

|        |              |                          |         |
|--------|--------------|--------------------------|---------|
| OVN-NB | ovn-external | ovn-cluster-ovn-external | chassis |
|--------|--------------|--------------------------|---------|

```

kubectl get nbct1 show
switch 3de4cead7-1a71-43f3-8b62-435a57ef16a6 (external)
 port localnet.external
 type: localnet
 addresses: ["unknown"]
 port external-ovn-cluster
 type: router
 router-port: ovn-cluster-external
router e1eb83ad-34be-4ed5-9a02-fcc8b1d357c4 (ovn-cluster)
 port ovn-cluster-external
 mac: "ac:1f:6b:2d:33:f1"
 networks: ["172.56.0.100/16"]
 gateway chassis: [a5682814-2e2c-46dd-9c1c-6803ef0dab66]

```

|     |             |
|-----|-------------|
| OVS | br-external |
|-----|-------------|

```

kubectl get vsctl ${gateway node name} show
e7d81150-7743-4d6e-9e6f-5c688232e130
 Bridge br-external
 Port br-external
 Interface br-external
 type: internal
 Port eth1
 Interface eth1
 Port patch-localnet.external-to-br-int
 Interface patch-localnet.external-to-br-int
 type: patch
 options: {peer=patch-br-int-to-localnet.external}

```

### 3.5.4 Pod EIP SNAT

|     |                        |                                  |          |
|-----|------------------------|----------------------------------|----------|
| Pod | ovn.kubernetes.io/snat | ovn.kubernetes.io/eip annotation | SNAT EIP |
|-----|------------------------|----------------------------------|----------|

```

apiVersion: v1
kind: Pod
metadata:
 name: pod-snat
 annotations:
 ovn.kubernetes.io/snat: 172.56.0.200
spec:
 containers:
 - name: pod-snat
 image: docker.io/library/nginx:alpine

apiVersion: v1
kind: Pod
metadata:
 name: pod-eip
 annotations:
 ovn.kubernetes.io/eip: 172.56.0.233
spec:
 containers:
 - name: pod-eip
 image: docker.io/library/nginx:alpine

```

|         |     |          |                                     |
|---------|-----|----------|-------------------------------------|
| kubectl | Pod | EIP SNAT | ovn.kubernetes.io/routed annotation |
|---------|-----|----------|-------------------------------------|

```

kubectl annotate pod pod-gw ovn.kubernetes.io/eip=172.56.0.221 --overwrite
kubectl annotate pod pod-gw ovn.kubernetes.io/routed-

```

|          |                                     |
|----------|-------------------------------------|
| EIP SNAT | ovn.kubernetes.io/routed annotation |
|----------|-------------------------------------|

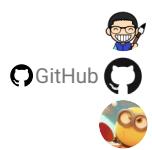
### 3.5.5

```
kube-ovn-controller SNAT EIP
• --external-gateway-config-ns: Configmap ovn-external-gw-config Namespace kube-system
• --external-gateway-net: external
• --external-gateway-vlanid: Vlan Tag 0 Vlan
```

[!\[\]\(05abdec45d3d9667a7f3c64e46754c68\_img.jpg\) PDF](#)[!\[\]\(b246cff486eb7c4dec085261afe05404\_img.jpg\) Slack](#)[!\[\]\(89b88375a21fd57206ff9ec1ea24d933\_img.jpg\) Support](#)

⌚2025 9 10

⌚2022 5 24



### 3.5.6

## 3.6 QoS



### 3.6.1 QoS

QoS      Pod annotation      Pod      Mbit/s

```

apiVersion: v1
kind: Pod
metadata:
 name: qos
 namespace: ls1
 annotations:
 ovn.kubernetes.io/ingress_rate: "3"
 ovn.kubernetes.io/egress_rate: "1"
spec:
 containers:
 - name: qos
 image: docker.io/library/nginx:alpine

```

annotation      QoS

```
kubectl annotate --overwrite pod nginx-74d5899f46-d7qkn ovn.kubernetes.io/ingress_rate=3
```

#### QoS

```

kind: DaemonSet
apiVersion: apps/v1
metadata:
 name: perf
 namespace: ls1
 labels:
 app: perf
spec:
 selector:
 matchLabels:
 app: perf
 template:
 metadata:
 labels:
 app: perf
 spec:
 containers:
 - name: nginx
 image: docker.io/kubeovn/perf

```

Pod      iperf3 server

```
kubectl exec -it perf-4n4gt -n ls1 sh
iperf3 -s

Server listening on 5201

```

Pod      Pod

```

kubectl exec -it perf-d4mqc -n ls1 sh
iperf3 -c 10.66.0.12
Connecting to host 10.66.0.12, port 5201
[4] local 10.66.0.14 port 51544 connected to 10.66.0.12 port 5201
[ID] Interval Transfer Bandwidth Retr Cwnd
[4] 0.00-1.00 sec 86.4 MBytes 725 Mbits/sec 3 350 KBytes
[4] 1.00-2.00 sec 89.9 MBytes 754 Mbits/sec 118 473 KBytes
[4] 2.00-3.00 sec 101 MBytes 848 Mbits/sec 184 586 KBytes
[4] 3.00-4.00 sec 104 MBytes 875 Mbits/sec 217 671 KBytes
[4] 4.00-5.00 sec 111 MBytes 935 Mbits/sec 175 772 KBytes

```

```
[4] 5.00-6.00 sec 100 MBytes 840 Mbits/sec 658 598 KBytes
[4] 6.00-7.00 sec 106 MBytes 890 Mbits/sec 742 668 KBytes
[4] 7.00-8.00 sec 102 MBytes 857 Mbits/sec 764 724 KBytes
[4] 8.00-9.00 sec 97.4 MBytes 817 Mbits/sec 1175 764 KBytes
[4] 9.00-10.00 sec 111 MBytes 934 Mbits/sec 1083 838 KBytes

[ID] Interval Transfer Bandwidth Retr
[4] 0.00-10.00 sec 1010 MBytes 848 Mbits/sec 5119 sender
[4] 0.00-10.00 sec 1008 MBytes 846 Mbits/sec receiver

iperf Done.
```

Pod QoS

```
kubectl annotate --overwrite pod perf-4n4gt -n ls1 ovn.kubernetes.io/ingress_rate=30
```

Pod Pod

```
iperf3 -c 10.66.0.12
Connecting to host 10.66.0.12, port 5201
[4] local 10.66.0.14 port 52372 connected to 10.66.0.12 port 5201
[ID] Interval Transfer Bandwidth Retr Cwnd
[4] 0.00-1.00 sec 3.66 MBytes 30.7 Mbits/sec 2 76.1 KBytes
[4] 1.00-2.00 sec 3.43 MBytes 28.8 Mbits/sec 0 104 KBytes
[4] 2.00-3.00 sec 3.50 MBytes 29.4 Mbits/sec 0 126 KBytes
[4] 3.00-4.00 sec 3.50 MBytes 29.3 Mbits/sec 0 144 KBytes
[4] 4.00-5.00 sec 3.43 MBytes 28.8 Mbits/sec 0 160 KBytes
[4] 5.00-6.00 sec 3.43 MBytes 28.8 Mbits/sec 0 175 KBytes
[4] 6.00-7.00 sec 3.50 MBytes 29.3 Mbits/sec 0 212 KBytes
[4] 7.00-8.00 sec 3.68 MBytes 30.9 Mbits/sec 0 294 KBytes
[4] 8.00-9.00 sec 3.74 MBytes 31.4 Mbits/sec 0 398 KBytes
[4] 9.00-10.00 sec 3.80 MBytes 31.9 Mbits/sec 0 526 KBytes

[ID] Interval Transfer Bandwidth Retr
[4] 0.00-10.00 sec 35.7 MBytes 29.9 Mbits/sec 2 sender
[4] 0.00-10.00 sec 34.5 MBytes 29.0 Mbits/sec receiver

iperf Done.
```

### 3.6.2 linux-netem QoS

RHEL netem yum install -y kernel-modules-extra && modprobe sch\_netem

|                                |                           |                         |                        |
|--------------------------------|---------------------------|-------------------------|------------------------|
| Pod annotation linux-netem QoS | ovn.kubernetes.io/latency | ovn.kubernetes.io/limit | ovn.kubernetes.io/loss |
| • ovn.kubernetes.io/latency    | Pod                       | ms                      |                        |
| • ovn.kubernetes.io/jitter     | Pod                       | ms                      |                        |
| • ovn.kubernetes.io/limit      | qdisc                     |                         | 1000                   |
| • ovn.kubernetes.io/loss       |                           | float                   | 20 20%                 |

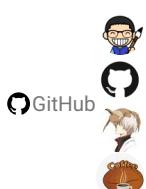
[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 5 24



### 3.6.3

## 3.7



### 3.7.1

CIDR    cidr=<IPv4 CIDR>,<IPv6 CIDR>    CIDR    IPv4    IPv6

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: ovn-test
spec:
 cidrBlock: 10.16.0.0/16,fd00:10:16::/64
 excludeIps:
 - 10.16.0.1
 - fd00:10:16::1
 gateway: 10.16.0.1,fd00:10:16::1

```

```

POD_CIDR="10.16.0.0/16,fd00:10:16::/64"
JOIN_CIDR="100.64.0.0/16,fd00:100:64::/64"

```

### 3.7.2 Pod

Pod    IPv4    IPv6    Pod annotation :

```

apiVersion: v1
kind: Pod
metadata:
 annotations:
 ovn.kubernetes.io/allocated: "true"
 ovn.kubernetes.io/cidr: 10.16.0.0/16,fd00:10:16::/64
 ovn.kubernetes.io/gateway: 10.16.0.1,fd00:10:16::1
 ovn.kubernetes.io/ip_address: 10.16.0.9,fd00:10:16::9
 ovn.kubernetes.io/logical_switch: ovn-default
 ovn.kubernetes.io/mac_address: 00:00:00:14:88:09
 ovn.kubernetes.io/network_types: geneve
 ovn.kubernetes.io/routed: "true"
...
podIP: 10.16.0.9
 podIPs:
 - ip: 10.16.0.9
 - ip: fd00:10:16::9

```

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 5 24



### 3.7.3

## 3.8 Webhook

```
Webhook Kube-OVN CRD Webhook IP Subnet CIDR
Webhook Subnet Pod Kube-OVN Webhook Pod
```

### 3.8.1 Cert-Manager

```
Webhook cert-manager Webhook cert-manager
```

cert-manager:

```
kubectl apply -f https://github.com/cert-manager/cert-manager/releases/download/v1.8.0/cert-manager.yaml
```

```
cert-manager cert-manager
```

### 3.8.2 Webhook

```
Webhook yaml :
```

```
kubectl apply -f https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/yamls/webhook.yaml
deployment.apps/kube-ovn-webhook created
service/kube-ovn-webhook created
validatingwebhookconfiguration.admissionregistration.k8s.io/kube-ovn-webhook created
certificate.cert-manager.io/kube-ovn-webhook-serving-cert created
issuer.cert-manager.io/kube-ovn-webhook-selfsigned-issuer created
```

### 3.8.3 Webhook

Pod Pod IP 10.16.0.15

```
kubectl get pod -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
static-7584848b74-fw9dm 1/1 Running 0 2d13h 10.16.0.15 kube-ovn-worker <none>
```

yaml IP Pod

```
apiVersion: v1
kind: Pod
metadata:
 annotations:
 ovn.kubernetes.io/ip_address: 10.16.0.15
 ovn.kubernetes.io/mac_address: 00:00:00:53:6B:B6
 labels:
 app: static
 managedFields:
 name: staticip-pod
 namespace: default
spec:
 containers:
 - image: docker.io/library/nginx:alpine
 imagePullPolicy: IfNotPresent
 name: qatest
```

yaml Pod IP

```
kubectl apply -f pod-static.yaml
Error from server (annotation ip address 10.16.0.15 is conflict with ip crd static-7584848b74-fw9dm.default 10.16.0.15): error when creating "pod-static.yaml": admission webhook "pod-ip-validation.kube-ovn.io" denied the request: annotation ip address 10.16.0.15 is conflict with ip crd static-7584848b74-fw9dm.default 10.16.0.15
```



PDF



Slack



Support

⌚2023 5 9

⌚2022 5 24

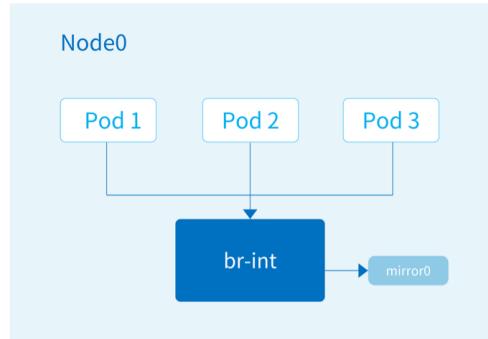
⌚GitHub 🧑‍💻

3.8.4

---

## 3.9

NPM  
CPU            5%~10%      CPU



### 3.9.1

kube-ovn-cni DaemonSet

```

• --enable-mirror=true
• --mirror-iface=mirror0:
 br-int Kube-OVN
 mirror0
tcpdump mirror0

tcpdump -ni mirror0

```

### 3.9.2 Pod

Pod            Pod    ovn.kubernetes.io/mirror annotation    Pod

```

apiVersion: v1
kind: Pod
metadata:
 name: mirror-pod
 namespace: ls1
 annotations:
 ovn.kubernetes.io/mirror: "true"
spec:
 containers:
 - name: mirror-pod
 image: docker.io/library/nginx:alpine

```

### 3.9.3

---

#### 1. Pod to Pod in the same Nodes

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 12.7 us     | 289 Mbits/sec  | 12.6 us     | (1.8%)        | 77.9 Mbits/sec |
| 128  | 15.5 us     | 517 Mbits/sec  | 12.7 us     | (0%)          | 155 Mbits/sec  |
| 512  | 12.2 us     | 1.64 Gbits/sec | 12.4 us     | (0%)          | 624 Mbits/sec  |
| 1k   | 13 us       | 2.96 Gbits/sec | 11.4 us     | (0.53%)       | 1.22 Gbits/sec |
| 4k   | 18 us       | 7.67 Gbits/sec | 25.7 us     | (0.41%)       | 1.50 Gbits/sec |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 11.9 us     | 324 Mbits/sec  | 12.2 us     | (0.22%)       | 102 Mbits/sec  |
| 128  | 10.5 us     | 582 Mbits/sec  | 9.5 us      | (0.21%)       | 198 Mbits/sec  |
| 512  | 11.6 us     | 1.84 Gbits/sec | 9.32 us     | (0.091%)      | 827 Mbits/sec  |
| 1k   | 10.5 us     | 3.44 Gbits/sec | 10 us       | (1.2%)        | 1.52 Gbits/sec |
| 4k   | 16.7 us     | 8.52 Gbits/sec | 18.2 us     | (1.3%)        | 2.42 Gbits/sec |

#### 2. Pod to Pod in the different Nodes

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 258 us      | 143 Mbits/sec  | 237 us      | (61%)         | 28.5 Mbits/sec |
| 128  | 240 us      | 252 Mbits/sec  | 231 us      | (64%)         | 54.9 Mbits/sec |
| 512  | 236 us      | 763 Mbits/sec  | 256 us      | (68%)         | 194 Mbits/sec  |
| 1k   | 242 us      | 969 Mbits/sec  | 225 us      | (62%)         | 449 Mbits/sec  |
| 4k   | 352 us      | 1.12 Gbits/sec | 382 us      | (0.71%)       | 21.4 Mbits/sec |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 278 us      | 140 Mbits/sec  | 227 us      | (24%)         | 59.6 Mbits/sec |
| 128  | 249 us      | 265 Mbits/sec  | 265 us      | (23%)         | 114 Mbits/sec  |
| 512  | 233 us      | 914 Mbits/sec  | 235 us      | (21%)         | 468 Mbits/sec  |
| 1k   | 238 us      | 1.14 Gbits/sec | 240 us      | (15%)         | 891 Mbits/sec  |
| 4k   | 370 us      | 1.25 Gbits/sec | 361 us      | (0.43%)       | 7.54 Mbits/sec |

### 3. Node to Node

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 205 us      | 162 Mbits/sec  | 183 us      | (11%)         | 74.2 Mbits/sec |
| 128  | 222 us      | 280 Mbits/sec  | 206 us      | (6.3%)        | 155 Mbits/sec  |
| 512  | 220 us      | 1.04 Gbits/sec | 177 us      | (20%)         | 503 Mbits/sec  |
| 1k   | 213 us      | 2.06 Gbits/sec | 201 us      | (8.6%)        | 1.14 Gbits/sec |
| 4k   | 280 us      | 5.01 Gbits/sec | 315 us      | (37%)         | 1.20 Gbits/sec |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 204 us      | 157 Mbits/sec  | 204 us      | (8.8%)        | 81.9 Mbits/sec |
| 128  | 213 us      | 262 Mbits/sec  | 225 us      | (19%)         | 136 Mbits/sec  |
| 512  | 220 us      | 1.02 Gbits/sec | 227 us      | (21%)         | 486 Mbits/sec  |
| 1k   | 217 us      | 1.79 Gbits/sec | 218 us      | (29%)         | 845 Mbits/sec  |
| 4k   | 275 us      | 5.27 Gbits/sec | 336 us      | (34%)         | 1.21 Gbits/sec |

### 4. Pod to the Node where the Pod is located

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 12.2 us     | 295 Mbits/sec  | 12.7 us     | (0.27%)       | 74.1 Mbits/sec |
| 128  | 14.1 us     | 549 Mbits/sec  | 10.6 us     | (0.41%)       | 153 Mbits/sec  |
| 512  | 13.5 us     | 1.83 Gbits/sec | 12.7 us     | (0.23%)       | 586 Mbits/sec  |
| 1k   | 12 us       | 2.69 Gbits/sec | 13 us       | (1%)          | 1.16 Gbits/sec |
| 4k   | 18.9 us     | 4.51 Gbits/sec | 21.8 us     | (0.42%)       | 1.81 Gbits/sec |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 10.4 us     | 335 Mbits/sec  | 12.2 us     | (0.75%)       | 95.4 Mbits/sec |
| 128  | 12.1 us     | 561 Mbits/sec  | 11.3 us     | (0.25%)       | 194 Mbits/sec  |
| 512  | 11.6 us     | 1.87 Gbits/sec | 10.7 us     | (0.66%)       | 745 Mbits/sec  |
| 1k   | 12.7 us     | 3.12 Gbits/sec | 10.9 us     | (1.2%)        | 1.46 Gbits/sec |
| 4k   | 16.5 us     | 8.23 Gbits/sec | 17.9 us     | (1.5%)        | 2.51 Gbits/sec |

## 5. Pod to the Node where the Pod is not located

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 234 us      | 153 Mbits/sec  | 232 us      | (63%)         | 29.4 Mbits/sec |
| 128  | 237 us      | 261 Mbits/sec  | 238 us      | (49%)         | 76.1 Mbits/sec |
| 512  | 231 us      | 701 Mbits/sec  | 238 us      | (57%)         | 279 Mbits/sec  |
| 1k   | 256 us      | 1.05 Gbits/sec | 228 us      | (56%)         | 524 Mbits/sec  |
| 4k   | 330 us      | 1.08 Gbits/sec | 359 us      | (1.5%)        | 35.7 Mbits/sec |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 283 us      | 141 Mbits/sec  | 230 us      | (26%)         | 55.8 Mbits/sec |
| 128  | 234 us      | 255 Mbits/sec  | 234 us      | (25%)         | 113 Mbits/sec  |
| 512  | 246 us      | 760 Mbits/sec  | 234 us      | (22%)         | 458 Mbits/sec  |
| 1k   | 268 us      | 1.23 Gbits/sec | 242 us      | (20%)         | 879 Mbits/sec  |
| 4k   | 326 us      | 1.20 Gbits/sec | 369 us      | (0.5%)        | 7.87 Mbits/sec |

## 6. Pod to the cluster ip service

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 237 us      | 133 Mbits/sec  | 213 us      | (65%)         | 25.5 Mbits/sec |
| 128  | 232 us      | 271 Mbits/sec  | 222 us      | (62%)         | 54.8 Mbits/sec |
| 512  | 266 us      | 800 Mbits/sec  | 234 us      | (60%)         | 232 Mbits/sec  |
| 1k   | 248 us      | 986 Mbits/sec  | 239 us      | (50%)         | 511 Mbits/sec  |
| 4k   | 314 us      | 1.03 Gbits/sec | 367 us      | (0.6%)        | 13.2 Mbits/sec |

| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 14305.17 | 0.87ms        | 1.48ms          | 24.46ms       |
| 100             | 29082.07 | 3.87ms        | 4.35ms          | 102.85ms      |

| Size | TCP Latency | TCP Bandwidth  | UDP Latency | UDP Lost Rate | UDP Bandwidth  |
|------|-------------|----------------|-------------|---------------|----------------|
| 64   | 241 us      | 145 Mbits/sec  | 225 us      | (19%)         | 60.2 Mbits/sec |
| 128  | 245 us      | 261 Mbits/sec  | 212 us      | (15%)         | 123 Mbits/sec  |
| 512  | 252 us      | 821 Mbits/sec  | 219 us      | (14%)         | 499 Mbits/sec  |
| 1k   | 253 us      | 1.08 Gbits/sec | 242 us      | (16%)         | 852 Mbits/sec  |
| 4k   | 320 us      | 1.32 Gbits/sec | 360 us      | (0.47%)       | 6.70 Mbits/sec |

| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 13634.07 | 0.96ms        | 1.72ms          | 30.07ms       |
| 100             | 30215.23 | 3.59ms        | 3.20ms          | 77.56ms       |

#### 7. Host to the Node port service where the Pod is not located on the target Node

| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 14802.73 | 0.88ms        | 1.66ms          | 31.49ms       |
| 100             | 29809.58 | 3.78ms        | 4.12ms          | 105.34ms      |

| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 14273.33 | 0.90ms        | 1.60ms          | 37.16ms       |
| 100             | 30757.81 | 3.62ms        | 3.41ms          | 59.78ms       |

#### 8. Host to the Node port service where the Pod is located on the target Node

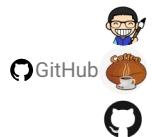
| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 15402.39 | 802.50us      | 1.42ms          | 30.91ms       |
| 100             | 29424.66 | 4.05ms        | 4.31ms          | 90.60ms       |

| TCP-Conn-Number | QPS      | Avg-Resp-Time | Stdev-Resp-Time | Max-Resp-Time |
|-----------------|----------|---------------|-----------------|---------------|
| 10              | 14649.21 | 0.91ms        | 1.72ms          | 43.92ms       |
| 100             | 32143.61 | 3.66ms        | 3.76ms          | 67.02ms       |



⌚2024 5 8

⌚2022 5 24



3.9.4

---

## 3.10 NetworkPolicy

|               |               |          |     |                    |               |               |
|---------------|---------------|----------|-----|--------------------|---------------|---------------|
| NetworkPolicy | Kubernetes    | Kube-OVN | OVN | ACL                | NetworkPolicy | NetworkPolicy |
| Kube-OVN      | NetworkPolicy |          |     | NetworkPolicy Drop |               |               |

### ⚠ Warning

|               |      |     |     |                               |
|---------------|------|-----|-----|-------------------------------|
| NetworkPolicy | Drop | CPU | OVN | <a href="#">ACL Log Meter</a> |
| ACL Log       |      |     |     |                               |

### 3.10.1 NetworkPolicy

NetworkPolicy annotation ovn.kubernetes.io/enable\_log

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
 name: default-deny-ingress
 namespace: kube-system
 annotations:
 ovn.kubernetes.io/enable_log: "true"
spec:
 podSelector: {}
 policyTypes:
 - Ingress
```

Pod /var/log/ovn/ovn-controller.log

```
tail -f /var/log/ovn/ovn-controller.log
2022-07-20T05:55:03.229Z |00394|acl_log(ovn_pinctrl0)|INFO|name="np/default-deny-ingress.kube-system/IPv4/0", verdict=drop, severity=warning, direction=to-lport: udp,vlan_tci=0x0000,d1_src=00:00:00:21:b7:d1,d1_dst=00:00:00:8d:0b:86,nw_src=10.16.0.10,nw_dst=10.
16.0.7,nw_tos=0,nw_ecn=0,nw_ttl=63,tp_src=54343,tp_dst=53
2022-07-20T05:55:06.229Z |00395|acl_log(ovn_pinctrl0)|INFO|name="np/default-deny-ingress.kube-system/IPv4/0", verdict=drop, severity=warning, direction=to-lport: udp,vlan_tci=0x0000,d1_src=00:00:00:21:b7:d1,d1_dst=00:00:00:8d:0b:86,nw_src=10.16.0.9,nw_dst=10.
16.0.7,nw_tos=0,nw_ecn=0,nw_ttl=63,tp_src=44187,tp_dst=53
2022-07-20T05:55:08.230Z |00396|acl_log(ovn_pinctrl0)|INFO|name="np/default-deny-ingress.kube-system/IPv4/0", verdict=drop, severity=warning, direction=to-lport: udp,vlan_tci=0x0000,d1_src=00:00:00:21:b7:d1,d1_dst=00:00:00:8d:0b:86,nw_src=10.16.0.10,nw_dst=10.
16.0.7,nw_tos=0,nw_ecn=0,nw_ttl=63,tp_src=54274,tp_dst=53
2022-07-20T05:55:11.231Z |00397|acl_log(ovn_pinctrl0)|INFO|name="np/default-deny-ingress.kube-system/IPv4/0", verdict=drop, severity=warning, direction=to-lport: udp,vlan_tci=0x0000,d1_src=00:00:00:21:b7:d1,d1_dst=00:00:00:8d:0b:86,nw_src=10.16.0.9,nw_dst=10.
16.0.7,nw_tos=0,nw_ecn=0,nw_ttl=63,tp_src=32778,tp_dst=53
2022-07-20T05:55:11.231Z |00398|acl_log(ovn_pinctrl0)|INFO|name="np/default-deny-ingress.kube-system/IPv4/0", verdict=drop, severity=warning, direction=to-lport: udp,vlan_tci=0x0000,d1_src=00:00:00:21:b7:d1,d1_dst=00:00:00:8d:0b:86,nw_src=10.16.0.9,nw_dst=10.
16.0.7,nw_tos=0,nw_ecn=0,nw_ttl=63,tp_src=34188,tp_dst=53
```

### 3.10.2 NetworkPolicy

|                                         |                                              |     |            |       |
|-----------------------------------------|----------------------------------------------|-----|------------|-------|
| ovn.kubernetes.io/enable_log annotation | Drop ACL                                     | ACL |            |       |
| Kube-OVN v1.13.0                        | annotation ovn.kubernetes.io/log_acl_actions | ACL | annotation | allow |
| NetworkPolicy                           | annotation ovn.kubernetes.io/log_acl_actions |     |            |       |

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
 name: default-deny-ingress
 namespace: kube-system
 annotations:
 ovn.kubernetes.io/enable_log: "true"
 ovn.kubernetes.io/log_acl_actions: "allow"
spec:
 podSelector: {}
 policyTypes:
 - Ingress
```

Pod Pod /var/log/ovn/ovn-controller.log Allow ACL

```
2024-08-14T09:27:49.590Z |00004|acl_log(ovn_pinctrl0)|INFO|name="np/test.default/ingress/IPv4/0", verdict=allow, severity=info, direction=to-lport: icmp,vlan_tci=0x0000,d1_src=96:7b:0:2fa0:1a,d1_dst=a6:e5:1b:c2:1b:f8,nw_src=10.16.0.7,nw_dst=10.
16.0.12,nw_tos=0,nw_ecn=0,nw_ttl=64,nw_frag=no,icmp_type=8,icmp_code=0
```

### 3.10.3 NetworkPolicy

```
NetworkPolicy annotation ovn.kubernetes.io/enable_log false NetworkPolicy
```

```
kubectl annotate networkpolicy -n kube-system default-deny-ingress ovn.kubernetes.io/enable_log=false --overwrite
```

### 3.10.4 AdminNetworkPolicy BaselineAdminNetworkPolicy

|         |          |                                    |                            |                    |                            |
|---------|----------|------------------------------------|----------------------------|--------------------|----------------------------|
| v1.13.0 | Kube-OVN | AdminNetworkPolicy                 | BaselineAdminNetworkPolicy | AdminNetworkPolicy | BaselineAdminNetworkPolicy |
|         |          | <a href="#">Network Policy API</a> |                            |                    |                            |

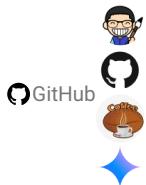
|                    |                                              |            |            |                   |
|--------------------|----------------------------------------------|------------|------------|-------------------|
| Network Policy API | ovn.kubernetes.io/log_acl_actions annotation | action ACL | annotation | allow, drop, pass |
|--------------------|----------------------------------------------|------------|------------|-------------------|

|            |            |            |                                   |
|------------|------------|------------|-----------------------------------|
| annotation | annotation | annotation | ovn.kubernetes.io/log_acl_actions |
| annotation |            |            |                                   |

[!\[\]\(974977f0ae7a9bd8eb2b52f26d2700fe\_img.jpg\) PDF](#)
[!\[\]\(646869a5c9c88531ae41e50ad8af3b92\_img.jpg\) Slack](#)
[!\[\]\(9500e4e0b06f494c3391c456117a0e37\_img.jpg\) Support](#)

 2025 9 17

 2022 7 20



### 3.10.5

## 3.11 LoadBalancer Service

```
Kube-OVN VPC VPC VPC
VPC VPC VPC LoadBalancer Service LoadBalancerIP VPC Service
```

1. multus-cni macvlan cni
2. LoadBalancer Service VPC vpc-nat-gw macvlan
3. VPC VPC LoadBalancer VPC VPC

### 3.11.1 VPC LoadBalancer Service

```
kube-system namespace deployment kube-ovn-controller args --enable-lb-svc=true false
```

```
containers:
- args:
 - /kube-ovn/start-controller.sh
 - --default-cidr=10.16.0.0/16
 - --default-gateway=10.16.0.1
 - --default-gateway-check=true
 - --enable-lb-svc=true
 // true
```

#### NetworkAttachmentDefinition CRD

```
yaml net-attach-def
```

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: lb-svc-attachment
 namespace: kube-system
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "macvlan",
 "master": "eth0",
 "mode": "bridge"
 }'
```

```
eth0 master
```

#### Subnet

```
Subnet LoadBalancer Service LoadBalancerIP Underlay Subnet
```

```
yaml
```

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: attach-subnet
spec:
 protocol: IPv4
 provider: lb-svc-attachment.kube-system # provider
 cidrBlock: 172.18.0.0/16
 gateway: 172.18.0.1
 excludeIps:
 - 172.18.0.0..172.18.0.10
```

```
Subnet provider ovn .ovn Kube-OVN logical switch
provider ovn .ovn Kube-OVN IPAM IP
```

## LoadBalancer Service

yaml LoadBalancer Service

```
apiVersion: v1
kind: Service
metadata:
 annotations:
 lb-svc-attachment.kube-system.kubernetes.io/logical_switch: attach-subnet #
 ovn.kubernetes.io/attachmentprovider: lb-svc-attachment.kube-system #
 labels:
 app: dynamic
 name: test-service
 namespace: default
spec:
 loadBalancerIP: 172.18.0.18
 ports:
 - name: test
 protocol: TCP
 port: 80
 targetPort: 80
 selector:
 app: dynamic
 sessionAffinity: None
 type: LoadBalancer
```

| yaml annotation ovn.kubernetes.io/attachmentprovider | net-attach-def | Name.Namespace | annotation | Pod |
|------------------------------------------------------|----------------|----------------|------------|-----|
| net-attach-def                                       |                |                |            |     |

| annotation     | annotation key      | net-attach-def | Name.Namespace.kubernetes.io/logical_switch |
|----------------|---------------------|----------------|---------------------------------------------|
| LoadBalancerIP |                     | LoadBalancerIP |                                             |
| LoadBalancerIP | spec.loadBalancerIP |                |                                             |

yaml Service Service Namespace Pod

```
kubectl get pod
NAME READY STATUS RESTARTS AGE
lb-svc-test-service-6869d98dd8-cjvll 1/1 Running 0 107m
kubectl get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
test-service LoadBalancer 10.109.201.193 172.18.0.18 80:30056/TCP 107m
```

| service.spec.loadBalancerIP | service external-ip |
|-----------------------------|---------------------|
|-----------------------------|---------------------|

Pod yaml

```
kubectl get pod -o yaml lb-svc-test-service-6869d98dd8-cjvll
apiVersion: v1
kind: Pod
metadata:
 annotations:
 k8s.v1.cni.cncf.io/network-status: |-
 [{
 "name": "kube-ovn",
 "ips": [
 "10.16.0.2"
],
 "default": true,
 "dns": {}
 }, {
 "name": "default/test-service",
 "interface": "net1",
 "mac": "ba:85:f7:02:9f:42",
 "dns": {}
 }]
 k8s.v1.cni.cncf.io/networks: default/test-service
 k8s.v1.cni.cncf.io/networks-status: |-
 [{
 "name": "kube-ovn",
 "ips": [
 "10.16.0.2"
],
 "default": true,
 "dns": {}
 }, {
 "name": "default/test-service",
 "interface": "net1",
 "mac": "ba:85:f7:02:9f:42",
 "dns": {}
 }]
 ovn.kubernetes.io/allocated: "true"
 ovn.kubernetes.io/cidr: 10.16.0.0/16
 ovn.kubernetes.io/gateway: 10.16.0.1
```

```

ovn.kubernetes.io/ip_address: 10.16.0.2
ovn.kubernetes.io/logical_router: ovn-cluster
ovn.kubernetes.io/logical_switch: ovn-default
ovn.kubernetes.io/mac_address: 00:00:00:45:F4:29
ovn.kubernetes.io/pod_nic_type: veth-pair
ovn.kubernetes.io/routed: "true"
test-service.default.kubernetes.io/allocated: "true"
test-service.default.kubernetes.io/cidr: 172.18.0.0/16
test-service.default.kubernetes.io/gateway: 172.18.0.1
test-service.default.kubernetes.io/ip_address: 172.18.0.18
test-service.default.kubernetes.io/logical_switch: attach-subnet
test-service.default.kubernetes.io/mac_address: 00:00:00:AF:AA:BF
test-service.default.kubernetes.io/pod_nic_type: veth-pair

```

## Service

```

kubectl get svc -o yaml test-service
apiVersion: v1
kind: Service
metadata:
 annotations:
 kubelet.kubernetes.io/last-applied-configuration: |
 {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{"test-service.default.kubernetes.io/logical_switch":"attach-subnet"},"labels":{"app":"dynamic","name":"test-service","namespace":"default"},"spec":{"ports":[{"name":"test","port":80,"protocol":"TCP","targetPort":80}],"selector":{"app":"dynamic","sessionAffinity":"None","type":"LoadBalancer"}}}
 ovn.kubernetes.io/vpc: ovn-cluster
 test-service.default.kubernetes.io/logical_switch: attach-subnet
 creationTimestamp: "2022-06-15T09:01:58Z"
 labels:
 app: dynamic
 name: test-service
 namespace: default
 resourceVersion: "38485"
 uid: 161edee1-7f6e-40f5-9e09-5a52c44267d0
spec:
 allocateLoadBalancerNodePorts: true
 clusterIP: 10.109.201.193
 clusterIPs:
 - 10.109.201.193
 externalTrafficPolicy: Cluster
 internalTrafficPolicy: Cluster
 ipFamilies:
 - IPv4
 ipFamilyPolicy: SingleStack
 ports:
 - name: test
 nodePort: 30056
 port: 80
 protocol: TCP
 targetPort: 80
 selector:
 app: dynamic
 sessionAffinity: None
 type: LoadBalancer
 status:
 loadBalancer:
 ingress:
 - ip: 172.18.0.18

```

## 3.11.2 LoadBalancerIP

yaml, Pod Service Endpoints

```

apiVersion: apps/v1
kind: Deployment
metadata:
 labels:
 app: dynamic
 name: dynamic
 namespace: default
spec:
 replicas: 2
 selector:
 matchLabels:
 app: dynamic
 strategy:
 rollingUpdate:
 maxSurge: 25%
 maxUnavailable: 25%
 type: RollingUpdate
 template:
 metadata:
 creationTimestamp: null
 labels:
 app: dynamic
 spec:
 containers:
 - image: docker.io/library/nginx:alpine

```

```
imagePullPolicy: IfNotPresent
name: nginx
dnsPolicy: ClusterFirst
restartPolicy: Always
```

#### Service LoadBalancerIP:Port

```
curl 172.18.0.18:80
<html>
<head>
 <title>Hello World!</title>
 <link href='//fonts.googleapis.com/css?family=Open+Sans:400,700' rel='stylesheet' type='text/css'>
 <style>
 body {
 background-color: white;
 text-align: center;
 padding: 50px;
 font-family: "Open Sans", "Helvetica Neue", Helvetica, Arial, sans-serif;
 }
 #logo {
 margin-bottom: 40px;
 }
 </style>
</head>
<body>
 <h1>Hello World!</h1>
 <h3>Links found</h3>
 <h3>I am on dynamic-7d8d7874f5-hsgc4</h3>
 <h3>Cookie =/>
 KUBERNETES listening in 443 available at tcp://10.96.0.1:443

 <h3>my name is hanhouchao!</h3>
 <h3> RequestURI='/'</h3>
</body>
</html>
```

#### Service Pod

```
ip a
4: net1@if62: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
 link/ether ba:85:f7:02:9f:42 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 172.18.0.18/16 scope global net1
 valid_lft forever preferred_lft forever
 inet6 fe80::ba85:f7ff:fe02:9f42/64 scope link
 valid_lft forever preferred_lft forever
36: eth0@if37: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP group default
 link/ether 00:00:00:45:f4:29 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 10.16.0.2/16 brd 10.16.255.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:fe45:f429/64 scope link
 valid_lft forever preferred_lft forever

ip rule
0: from all lookup local
32764: from all iif eth0 lookup 100
32765: from all iif net1 lookup 100
32766: from all lookup main
32767: from all lookup default

ip route show table 100
default via 172.18.0.1 dev net1
10.109.201.193 via 10.16.0.1 dev eth0
172.18.0.0/16 dev net1 scope link

iptables -t nat -L -n -v
Chain PREROUTING (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target prot opt in out source destination
 0 0 DNAT tcp -- * * 0.0.0.0/0 172.18.0.18 tcp dpt:80 to:10.109.201.193:80

Chain INPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target prot opt in out source destination

Chain OUTPUT (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target prot opt in out source destination

Chain POSTROUTING (policy ACCEPT 0 packets, 0 bytes)
 pkts bytes target prot opt in out source destination
 0 0 MASQUERADE all -- * * 0.0.0.0/0 10.109.201.193
```

#### lb service Pod nodeSelector

```
ovn-vpc-nat-config ConfigMap nodeSelector LoadBalancer service Pod
```

```
apiVersion: v1
data:
 image: docker.io/kubeovn/vpc-nat-gateway:v1.14.0
 nodeSelector: |
 kubernetes.io/hostname: kube-ovn-control-plane
```

```
kubernetes.io/os: linux
kind: ConfigMap
metadata:
 name: ovn-vpc-nat-config
 namespace: kube-system
```

[!\[\]\(8b6ba4e01d044f020536a6874a89c057\_img.jpg\) PDF](#)[!\[\]\(39fe6a447a07f214586d259b16ffba42\_img.jpg\) Slack](#)[!\[\]\(2a4282dc455b24a8719bbd3b8683d6a8\_img.jpg\) Support](#)

⌚2025 9 10

⌚2022 7 19



3.11.3

---

## 3.12



### 3.12.1 Prometheus Monitor

Kube-OVN    Prometheus Monitor CRD

```

#
kubectl apply -f https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/pinger-monitor.yaml
kube-ovn-controller
kubectl apply -f https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/controller-monitor.yaml
kube-ovn-cni
kubectl apply -f https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/cni-monitor.yaml
ovn
kubectl apply -f https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/ovn-monitor.yaml

```

Prometheus    15s    yaml    interval

### 3.12.2 Grafana

Kube-OVN    Grafana Dashboard

Dashboard

```

#
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/pinger-grafana.json
kube-ovn-controller
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/controller-grafana.json
kube-ovn-cni
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/cni-grafana.json
ovn
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/ovn-grafana.json
ovs
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/master/dist/monitoring/ovs-grafana.json

```

Grafana    Prometheus    Dashboard

kube-ovn-controller



## kube-ovn-pinger

**kube-ovn-cni**[PDF](#)[Slack](#)[Support](#)

⌚2025 9 17

⌚2022 5 23



3.12.3

---

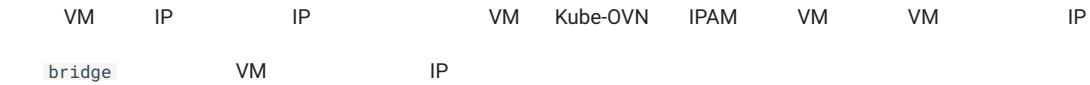
## 4. KubeVirt

---

### 4.1 VM IP



#### 4.1.1 IP VM



##### 1. VM

```

apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
 name: testvm
spec:
 runStrategy: Always
 template:
 metadata:
 labels:
 kubevirt.io/size: small
 kubevirt.io/domain: testvm
 annotations:
 kubevirt.io/allow-pod-bridge-network-live-migration: "true"
 spec:
 domain:
 devices:
 disks:
 - name: containerdisk
 disk:
 bus: virtio
 - name: cloudinitdisk
 disk:
 bus: virtio
 interfaces:
 - name: default
 bridge: {}
 resources:
 requests:
 memory: 64M
 networks:
 - name: default
 pod: {}
 volumes:
 - name: containerdisk
 containerDisk:
 image: quay.io/kubevirt/cirros-container-disk-demo
 - name: cloudinitdisk
 cloudInitNoCloud:
 userDataBase64: SGkuXG4=

```

##### 1. VM

```
kubectl get vmi testvm
```

### 1. VM

```
virtctl restart testvm
```

### 1. VM

```
virtctl migrate testvm
```

bridge    VM    IP

## 4.1.2 IP

VM	IP	VM	annotation	VM	IP	KubeVirt
----	----	----	------------	----	----	----------

```
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
 name: testvm
spec:
 runStrategy: Always
 template:
 metadata:
 labels:
 kubevirt.io/size: small
 kubevirt.io/domain: testvm
 annotations:
 ovn.kubernetes.io/ip_address: 10.16.0.15 #(1)
 kubevirt.io/allow-pod-bridge-network-live-migration: "true"
 spec:
 domain:
 devices:
 disks:
 - name: containerdisk
 disk:
 bus: virtio
 - name: cloudfinitdisk
 disk:
 bus: virtio
 interfaces:
 - name: default
 bridge: {}
 resources:
 requests:
 memory: 64M
 networks:
 - name: default
 pod: {}
 volumes:
 - name: containerdisk
 containerDisk:
 image: quay.io/kubevirt/cirros-container-disk-demo
 - name: cloudfinitdisk
 cloudInitNoCloud:
 user DataBase64: SGkuXG4=
```

1. 🖥 IP

## 4.1.3 VM IP

Kube-OVN	VM IP	IP	VM
----------	-------	----	----

VM IP

1. VM Annotation      IP

2. virtctl restart <vm name>    VM    IP

[PDF](#)

[Slack](#)

[Support](#)

Q2025 9 10

Q2025 3 3



4.1.4

---

## 4.2

KubeVirt	Clone	Label	Annotation	Kube-OVN	Annotation	IP	MAC
----------	-------	-------	------------	----------	------------	----	-----

### 4.2.1 Annotation

VirtualMachineClone	Kube-OVN	Annotation
---------------------	----------	------------

```

kind: VirtualMachineClone
apiVersion: "clone.kubevirt.io/v1beta1"
metadata:
 name: testclone
spec:
 source:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-source
 target:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-target
 template:
 annotationFilters:
 - "ovn.kubernetes.io/*"

```

### 4.2.2



#### Note

patches KubeVirt 1.6

IP

```

kind: VirtualMachineClone
apiVersion: "clone.kubevirt.io/v1beta1"
metadata:
 name: testclone
spec:
 source:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-source
 target:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-target
 patches:
 - {'op": "replace", "path": "/spec/template/metadata/annotations/ovn.kubernetes.io~1ip_address", "value": "10.16.0.15"}

```

```

kind: VirtualMachineClone
apiVersion: "clone.kubevirt.io/v1beta1"
metadata:
 name: testclone
spec:
 source:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-source
 target:
 apiGroup: kubevirt.io
 kind: VirtualMachine
 name: vm-target
 patches:
 - {"op": "remove", "path": "/spec/template/metadata/annotations/ovn.kubernetes.io~1ip_address"}

```

Annotation

KubeVirt Clone API



PDF



Slack



Support

⌚2025 9 10

⌚2025 8 13

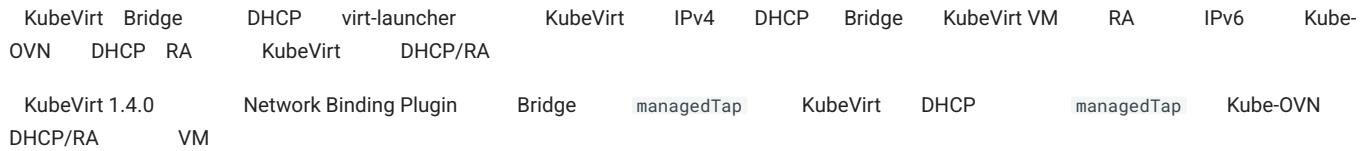


⌚GitHub⌚

4.2.3

---

4.3



### 4.3.1 DHCP

Kube-OVN Subnet DHCP IPv6 RA YAML

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: dual-stack-subnet
spec:
 cidrBlock: "10.244.0.0/16, fd00:10:244::/64"
 enableDHCP: true
 enableIPv6RA: true
```

#### 4.3.2 managedTap

KubeVirt      managedTap      Network Binding Plugin:

```
kubectl patch kubevirt -n kubevirt kubevirt --type=json -p=\
'[{"op": "add", "path": "/spec/configuration/network", "value": {
 "binding": {
 "managedtap": {
 "domainAttachmentType": "managedTap"
 }
 }
}}]'
```

### 4.3.3 managedTap

```
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
 name: dual-stack-vm
 namespace: default
spec:
 running: false
 template:
 spec:
 domain:
 devices:
 interfaces:
 - name: default
 binding:
 name: managedtap
 networks:
 - name: default
 pod: {}
```

VM    DHCP    IPv6 RA                    IPv4/IPv6



2025 9 10

Q+ 2025 3 3



#### 4.3.4

---

## 4.4

---

### KubeVirt

- KubeVirt      Bridge
  - KubeVirt
  - IP
  -
- Kube-OVN                                    0.5                            TCP

### 4.4.1

```
VM Spec kubevirt.io/allow-pod-bridge-network-live-migration: "true" annotation Kube-OVN
```

#### 1. VM

```
kubectl apply -f - <<EOF
apiVersion: kubevirt.io/v1
kind: VirtualMachine
metadata:
 name: testvm
spec:
 runStrategy: Always
 template:
 metadata:
 labels:
 kubevirt.io/size: small
 kubevirt.io/domain: testvm
 annotations:
 kubevirt.io/allow-pod-bridge-network-live-migration: "true"
 spec:
 domain:
 devices:
 disks:
 - name: containerdisk
 disk:
 bus: virtio
 - name: cloudinitdisk
 disk:
 bus: virtio
 interfaces:
 - name: default
 bridge: {}
 resources:
 requests:
 memory: 64M
 networks:
 - name: default
 pod: {}
 volumes:
 - name: containerdisk
 containerDisk:
 image: quay.io/kubevirt/cirros-container-disk-demo
 - name: cloudinitdisk
 cloudInitNoCloud:
 userDataBase64: SGkuXG4=
EOF
```

#### 1. SSH

```
password: gocubsgo
virtctl ssh cirros@testvm
ping 8.8.8.8
```

#### 1.

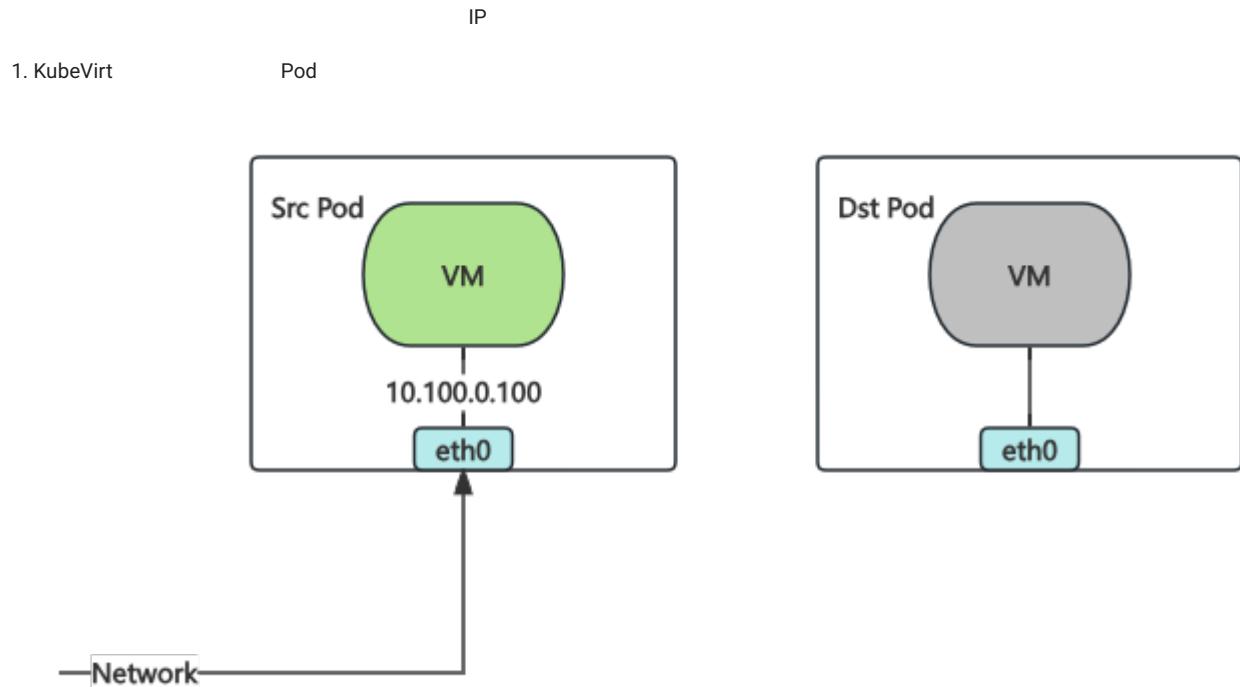
```
virtctl migrate testvm
```

VM	SSH	ping
----	-----	------

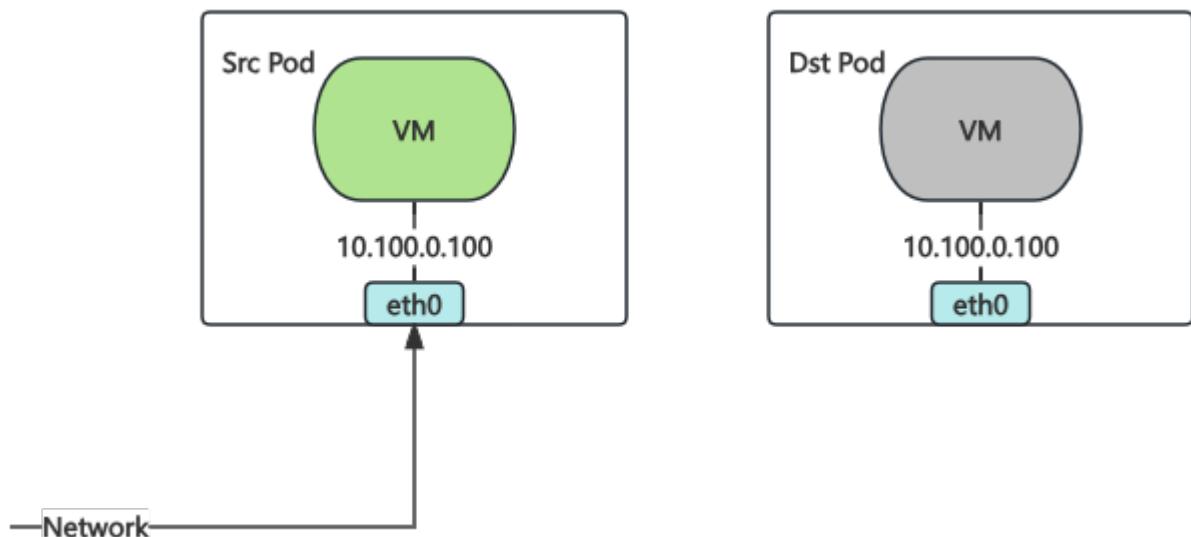
## 4.4.2

Kube-OVN

Live migration - Reducing downtime with multichassis port bindings



1. Kube-OVN Pod Pod Pod

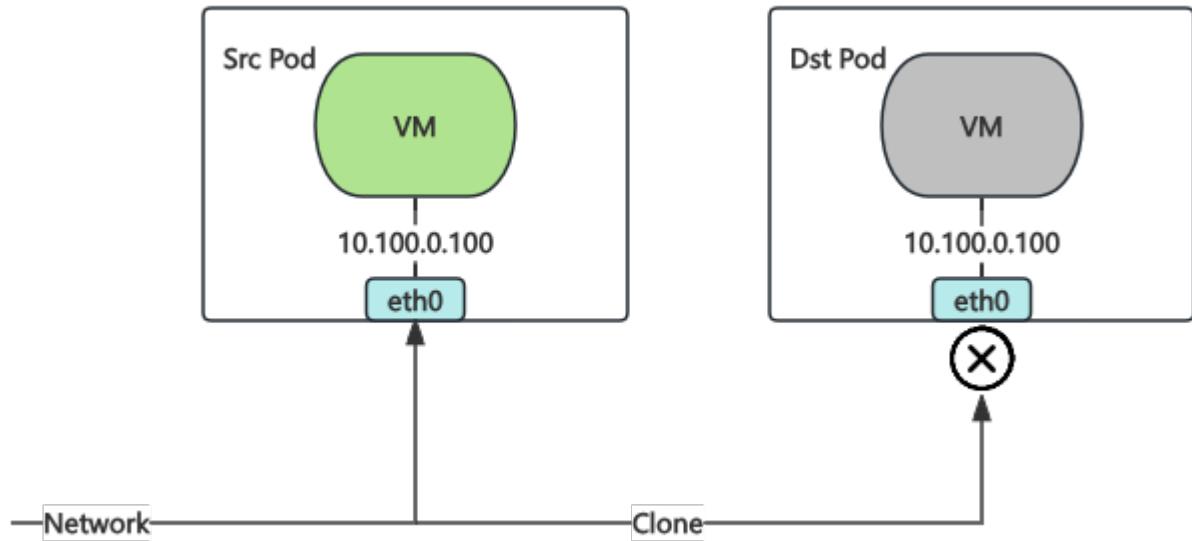


1. Kube-OVN

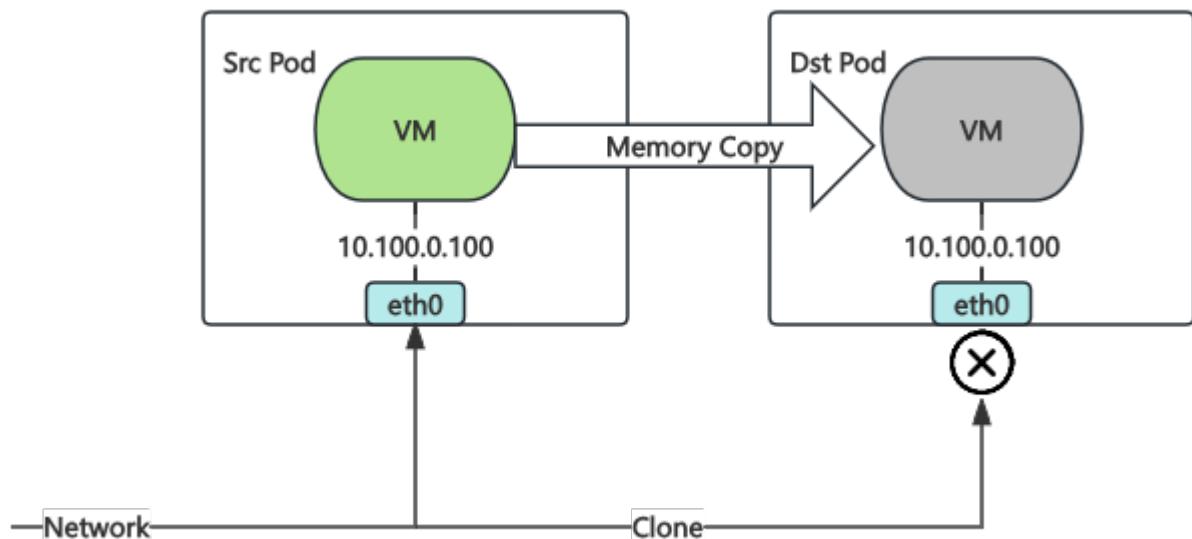
Pod Pod

Pod

Pod



1. KubeVirt VM

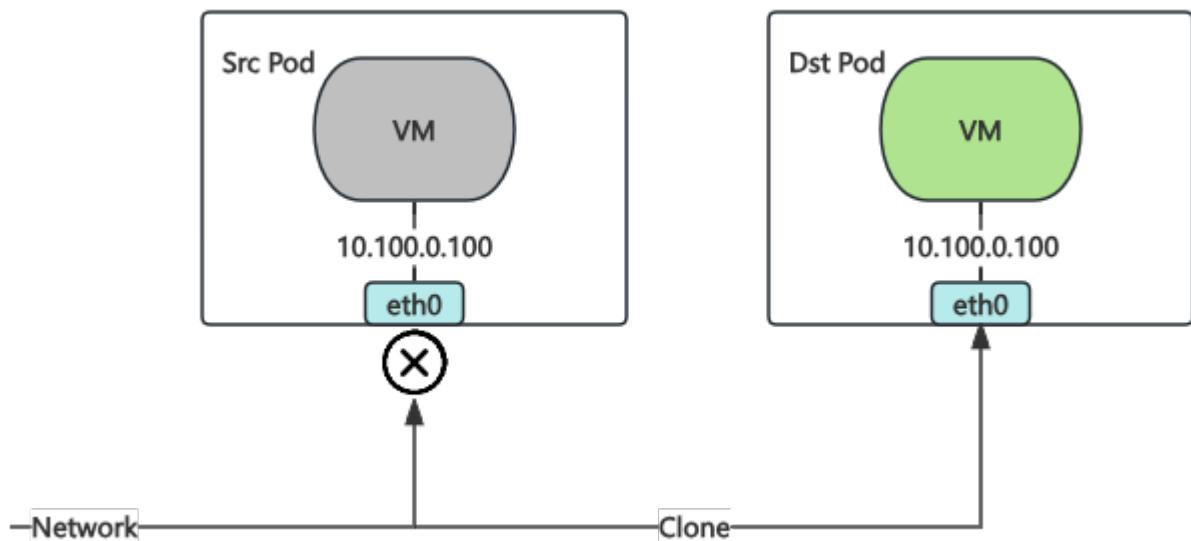


1. KubeVirt

Pod Pod



1. KubeVirt    Pod    libvirt    RARP    Pod    Pod



1. KubeVirt    Pod    Kube-OVN    Watch Migration CR



5 6

libvirt RARP

0.5 TCP



⌚2025 9 10

⌚2025 3 3



4.4.3

## 4.5 DHCP

managedTap	SR-IOV	DPDK	KubeVirt	DHCP	Kube-OVN	OVN	DHCP	DHCP	KubeVirt
DHCP	IP		Kube-OVN	DHCP	DHCPv6, IPv6RA, DNS	TFTP	DHCP	DHCP	

### ⚠ Warning

1. bridge      KubeVirt    DHCP    Kube-OVN    DHCP    Kube-OVN    DHCP    Kube-OVN    DHCP    managedTap  
bridge      managedTap      managedTap
2.            DHCP      Pod    DHCP

### DHCP

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: sn-dual
spec:
 cidrBlock: "10.0.0.0/24,240e::a00/120"
 default: false
 disableGatewayCheck: true
 disableInterConnection: false
 excludeIps:
 - 10.0.0.1
 - 240e::a01
 gateway: 10.0.0.1,240e::a01
 gatewayNode: ''
 gatewayType: distributed
 natOutgoing: false
 private: false
 protocol: Dual
 provider: ovn
 vpc: vpc-test
 enableDHCP: true
 dhcpV4Options: "lease_time=3600,router=10.0.0.1,server_id=169.254.0.254,server_mac=00:00:00:2E:2F:B8"
 dhcpV6Options: "server_id=00:00:00:2E:2F:C5"
 enableIPv6RA: true
 ipv6RAConfigs: "address_mode=dhcpv6_stateful,max_interval=30,min_interval=5,send_periodic=true"
```

- enableDHCP :      DHCP
- dhcpV4Options , dhcpV6Options :      ovn-nb    DHCP Options      "lease\_time=3600, router=\$ipv4\_gateway, server\_id=169.254.0.254, server\_mac=\$random\_mac"    server\_id=\$random\_mac
- enableIPv6RA :      DHCPv6
- ipv6RAConfigs :      ovn-nb    Logical\_Router\_Port      Logical Router Port      address\_mode=dhcpv6\_stateful, max\_interval=30, min\_interval=5, send\_periodic=true

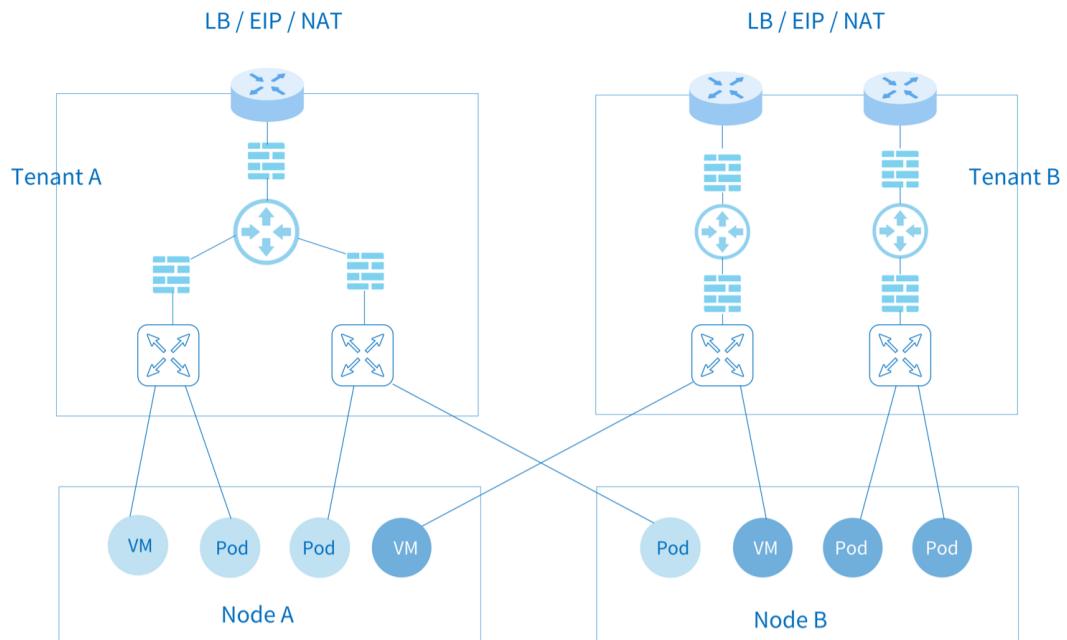
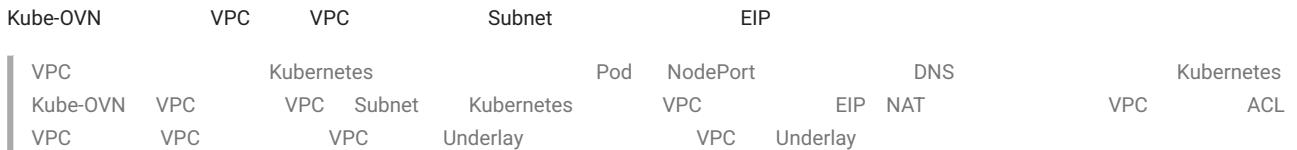
[PDF](#)
[Slack](#)
[Support](#)
 2025 9 10

 2022 5 24


4.5.1

## 5. VPC

### 5.1 VPC



#### 5.1.1

Kube-OVN	VPC	OVN	IP	VPC	IP	IP	VPC	OVN	Datapath ID
Datapath ID									

#### 5.1.2 VPC

##### VPC

```

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: test-vpc-1
spec:
 namespaces:
 - ns1

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: test-vpc-2
spec:
 namespaces:
 - ns2

```

- namespaces

Namespace

VPC

VPC CIDR:

```
kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: net1
spec:
 vpc: test-vpc-1
 cidrBlock: 10.0.1.0/24
 protocol: IPv4
 namespaces:
 - ns1

kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: net2
spec:
 vpc: test-vpc-2
 cidrBlock: 10.0.1.0/24
 protocol: IPv4
 namespaces:
 - ns2
```

Namespace Pod:

```
apiVersion: v1
kind: Pod
metadata:
 namespace: ns1
 name: vpc1-pod
spec:
 containers:
 - name: vpc1-pod
 image: docker.io/library/nginx:alpine

apiVersion: v1
kind: Pod
metadata:
 namespace: ns2
 name: vpc2-pod
spec:
 containers:
 - name: vpc2-pod
 image: docker.io/library/nginx:alpine
```

Pod CIDR VPC Pod

**VPC Pod livenessProbe readinessProbe**

VPC	Pod	kubelet	VPC	Pod	Kube-OVN	TProxy	kubelet	VPC	Pod	
			--enable-tproxy=true							

```
spec:
 template:
 spec:
 containers:
 - args:
 - --enable-tproxy=true
```

1. VPC Pod IP
2. tcpSocket httpGet

### 5.1.3 VPC

VPC	VPC	VPC	IP SNAT DNAT
VPC	Multus-CNI	multus-cni	

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: ovn-vpc-external-network
spec:
 protocol: IPv4
 provider: ovn-vpc-external-network.kube-system
 cidrBlock: 192.168.0.0/24
 gateway: 192.168.0.1 # IP address of the physical gateway
 excludeIps:
 - 192.168.0.1..192.168.0.10

apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: ovn-vpc-external-network
 namespace: kube-system
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "macvlan",
 "master": "eth1",
 "mode": "bridge",
 "ipam": {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "ovn-vpc-external-network.kube-system"
 }
}'

```

• Subnet	Macvlan	VPC	IP
• VPC	Macvlan	NetworkAttachmentDefinition	master
• name			
Macvlan	L2/L3	Vlan	
1. OpenStack VM PortSecurity			
2. VMware vSwitch	MAC Address Changes, Forged Transmits	Promiscuous Mode Operation	allow
3. Hyper-V MAC Address Spoofing			
4. AWS GCE Mac	Macvlan		
5. Macvlan	Macvlan	VpcNatGateway Pod	Pod
6. Trunk		Macvlan	

## VPC

VPC      kube-system      ovn-vpc-nat-gw-config      nodeSelector

```

kind: ConfigMap
apiVersion: v1
metadata:
 name: ovn-vpc-nat-config
 namespace: kube-system
data:
 image: 'docker.io/kubeovn/vpc-nat-gateway:v1.14.9'
 nodeSelector: |
 kubernetes.io/hostname: kube-ovn-control-plane

kind: ConfigMap
apiVersion: v1
metadata:
 name: ovn-vpc-nat-gw-config
 namespace: kube-system
data:
 enable-vpc-nat-gw: 'true'

```

- image: Pod
- enable-vpc-nat-gw: VPC

## VPC

```

kind: VpcNatGateway
apiVersion: kubeovn.io/v1
metadata:

```

```

 name: gw1
spec:
 vpc: test-vpc-1
 subnet: net1
 lanIp: 10.0.1.254
 selector:
 - "kubernetes.io/hostname: kube-ovn-worker"
 - "kubernetes.io/os: linux"
 externalSubnets:
 - ovn-vpc-external-network

```

- vpc VpcNatGateway VPC
- subnet VPC Subnet VPC Pod lanIp
- lanIp subnet IP VPC Pod IP VPC VpcNatGateway nextHopIP lanIp
- selector VpcNatGateway Pod Kubernetes NodeSelector
- externalSubnets VPC ovn-vpc-external-network

- tolerations VPC
- affinity VPC Pod

### VPC-NAT-GW

1. nat gw pod net1 arp ping eip arp ping

### EIP

EIP IP VPC DNAT SNAT IP  
EIP

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eip-random
spec:
 natGwDp: gw1

```

### EIP

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eip-static
spec:
 natGwDp: gw1
 v4ip: 192.168.0.100

```

### EIP

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eip-random
spec:
 natGwDp: gw1
 externalSubnet: ovn-vpc-external-network

```

- externalSubnet EIP ovn-vpc-external-network VPC externalSubnets

### DNAT

DNAT EIP VPC IP

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eipd01
spec:
 natGwDp: gw1

```

```

kind: IptablesDnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: dnat01
spec:
 eip: eipd01
 externalPort: '8888'
 internalIp: 10.0.1.10
 internalPort: '80'
 protocol: tcp
```

## SNAT

SNAT    VPC    Pod                  EIP    SNAT

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eips01
spec:
 natGwDp: gw1

kind: IptablesSnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: snat01
spec:
 eip: eips01
 internalCIDR: 10.0.1.0/24
```

## IP

IP    VPC    IP    EIP                  EIP    VPC    IP VPC    IP                  SNAT    EIP

```

kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eipf01
spec:
 natGwDp: gw1

kind: IptablesFIPRule
apiVersion: kubeovn.io/v1
metadata:
 name: fip01
spec:
 eip: eipf01
 internalIp: 10.0.1.5
```

## 5.1.4

VPC                                      Kube-OVN

```
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: test-vpc-1
spec:
 staticRoutes:
 - cidr: 0.0.0.0/0
 nextHopIP: 10.0.1.254
 policy: policyDst
 - cidr: 172.31.0.0/24
 nextHopIP: 10.0.1.253
 policy: policySrc
 routeTable: "rtb1"
```

- policy:                policyDst                policySrc
- CIDR
- routeTable:

OVN

Logical Router Policy

```

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: test-vpc-1
spec:
 policyRoutes:
 - action: drop
 match: ip4.src==10.0.1.0/24 && ip4.dst==10.0.1.250
 priority: 11
 - action: reroute
 match: ip4.src==10.0.1.0/24
 nextHopIP: 10.0.1.252
 priority: 10

```

### 5.1.5

Kubernetes	Service	Kubernetes	Service	IP	VPC	VPC
------------	---------	------------	---------	----	-----	-----

Kubernetes	Service
------------	---------

Kube-OVN      SwitchLBRule

SwitchLBRule

```

apiVersion: kubeovn.io/v1
kind: SwitchLBRule
metadata:
 name: cjh-slr-nginx
spec:
 vip: 1.1.1.1
 sessionAffinity: ClientIP
 namespace: default
 selector:
 - app:nginx
 ports:
 - name: dns
 port: 8888
 targetPort: 80
 protocol: TCP

```

- vip
- namespace      Pod      Namespace
- sessionAffinity    Service    sessionAffinity
- selector    Service    selector
- ports    Service    port

```

kubectl get slr
NAME VIP PORT(S) SERVICE AGE
vpc-dns-test-cjh2 10.96.0.3 53/UDP,53/TCP,9153/TCP kube-system/slr-vpc-dns-test-cjh2 88m

```

### 5.1.6 vpc-dns

VPC	VPC	VPC	Pod	coredns	VPC	CoreDNS	Service	Kube-OVN	vpc-dns
-----	-----	-----	-----	---------	-----	---------	---------	----------	---------

```

apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: ovn-nad
 namespace: default
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "kube-ovn",

```

```

 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "ovn-nad.default.ovn"
}'

```

## ovn-default provider

```
ovn-default provider nad provider ovn-nad.default.ovn
```

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: ovn-default
spec:
 cidrBlock: 10.16.0.0/16
 default: true
 disableGatewayCheck: false
 disableInterConnection: false
 enableDHCP: false
 enableIPv6RA: false
 excludeIps:
 - 10.16.0.1
 gateway: 10.16.0.1
 gatewayType: distributed
 logicalGateway: false
 natOutgoing: true
 private: false
 protocol: IPv4
 provider: ovn-nad.default.ovn
 vpc: ovn-cluster

```

## vpc-dns ConfigMap

```
kube-system configmap vpc-dns vpc-dns
```

```

apiVersion: v1
kind: ConfigMap
metadata:
 name: vpc-dns-config
 namespace: kube-system
data:
 coredns-vip: 10.96.0.3
 enable-vpc-dns: "true"
 nad-name: ovn-nad
 nad-provider: ovn-nad.default.ovn

```

- enable-vpc-dns true false true
- coredns-image dns coredns
- coredns-template dns URL yamls/coredns-template.yaml
- coredns-vip coredns lb vip
- nad-name network-attachment-definitions
- nad-provider provider
- k8s-service-host coredns k8s apiserver ip
- k8s-service-port coredns k8s apiserver port

## vpc-dns

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
 labels:
 kubernetes.io/bootstrapping: rbac-defaults
 name: system:vpc-dns
rules:
- apiGroups:
 - ""
 resources:
 - endpoints
 - services
 - pods
 - namespaces
 verbs:
 - list
 - watch
- apiGroups:

```

```

 - discovery.k8s.io
 resources:
 - endpointslices
 verbs:
 - list
 - watch

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
 annotations:
 rbac.authorization.kubernetes.io/autoupdate: "true"
 labels:
 kubernetes.io/bootstrapping: rbac-defaults
 name: vpc-dns
roleRef:
 apiGroup: rbac.authorization.k8s.io
 kind: ClusterRole
 name: system:vpc-dns
subjects:
- kind: ServiceAccount
 name: vpc-dns
 namespace: kube-system

apiVersion: v1
kind: ServiceAccount
metadata:
 name: vpc-dns
 namespace: kube-system

apiVersion: v1
kind: ConfigMap
metadata:
 name: vpc-dns-corefile
 namespace: kube-system
data:
 Corefile: |
 .:53 {
 errors
 health {
 lameduck 5s
 }
 ready
 kubernetes cluster.local in-addr.arpa ip6.arpa {
 pods insecure
 fallthrough in-addr.arpa ip6.arpa
 }
 prometheus :9153
 forward . /etc/resolv.conf {
 prefer_udp
 }
 cache 30
 loop
 reload
 loadbalance
 }
}

```

## vpc-dns

```

kind: VpcDns
apiVersion: kubeovn.io/v1
metadata:
 name: test-cjh1
spec:
 vpc: cjh-vpc-1
 subnet: cjh-subnet-1

```

- vpc dns vpc
- subnet dns

```
[root@hci-dev-mst-1 kubeovn]# kubectl get vpc-dns
NAME ACTIVE VPC SUBNET
test-cjh1 false cjh-vpc-1 cjh-subnet-1
test-cjh2 true cjh-vpc-1 cjh-subnet-2
```

- ACTIVE : true dns false

- VPC DNS

- VPC      vpc-dns      VPC      subnet      vpc-dns      true      false
- true    vpc-dns      false    vpc-dns

### 5.1.7

VPC      VPC

```
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: test-vpc-1
spec:
 namespaces:
 - ns1
 defaultSubnet: test
```

- defaultSubnet VPC

Namespace      ovn.kubernetes.io/logical\_switch      ovn.kubernetes.io/logical\_switch      Pod

[!\[\]\(5089d549bd5af430aa197ded06da9e4e\_img.jpg\) PDF](#)

[!\[\]\(d447a39d466d76c338e0c282e090b13b\_img.jpg\) Slack](#)

[!\[\]\(4924db801212cbdeeeb8d4ec3fb4cef8\_img.jpg\) Support](#)

 2025 9 19

 2022 5 24



### 5.1.8

## 5.2 VPC Egress Gateway

VPC Egress Gateway    VPC    VPC    Pod

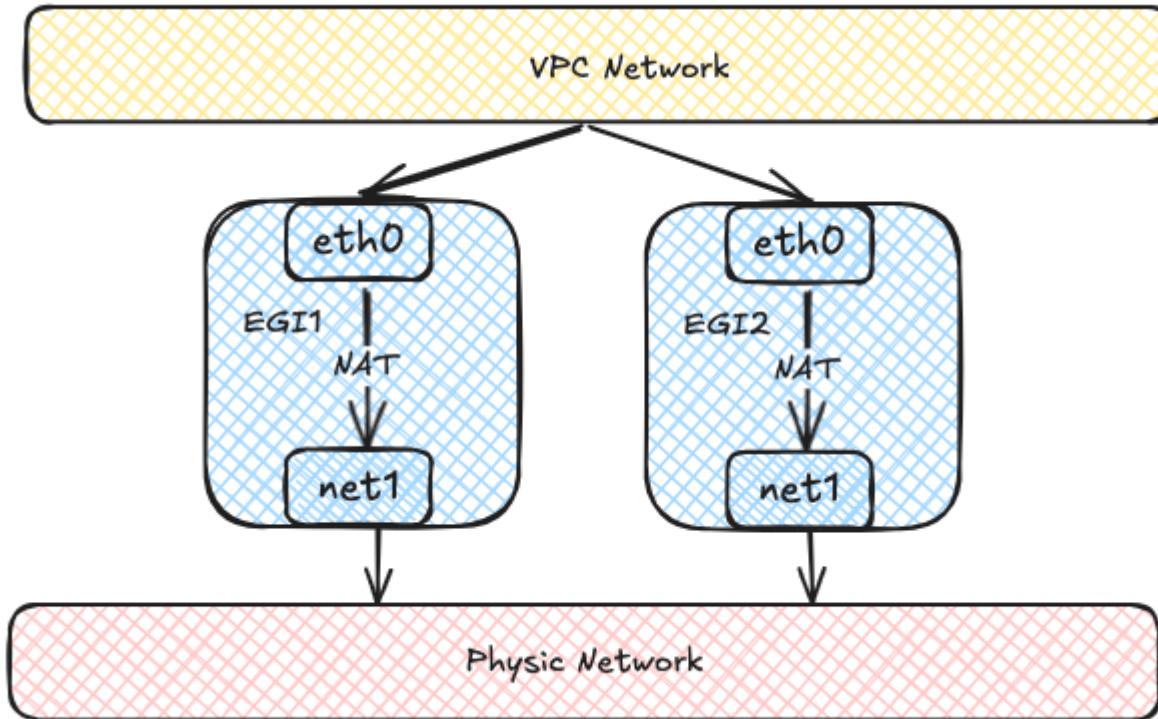
- ECMP    Active-Active
- BFD    <1s
- IPv6
- Namespace    Pod
- Node    Egress Gateway

VPC Egress Gateway

- Macvlan                      Underlay
- Gateway                      Egress IP
- SNAT    EIP    DNAT
- 

### 5.2.1

Egress Gateway    Pod    Pod    VPC    Macvlan    Egress Gateway    NAT



Egress Gateway

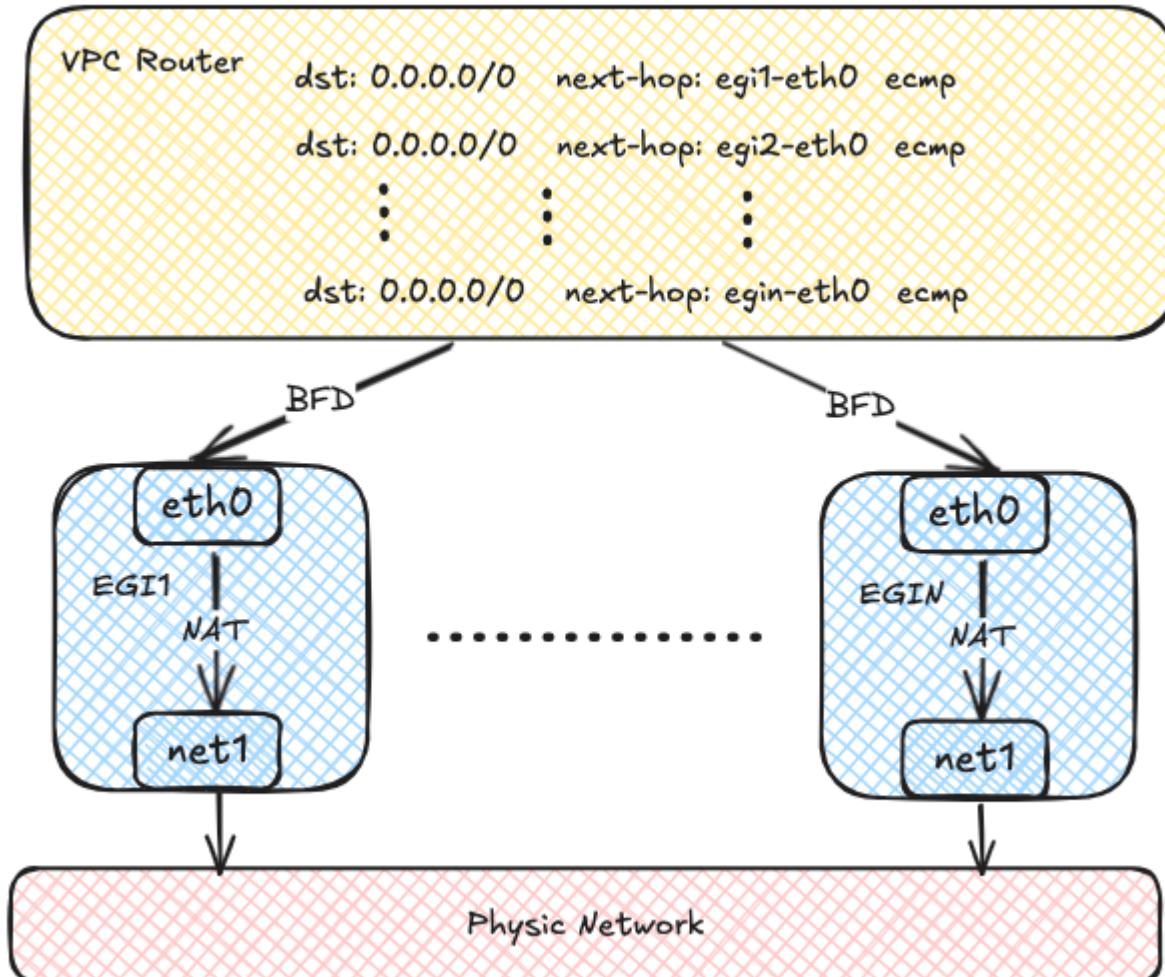
OVN

VPC    Pod

OVN

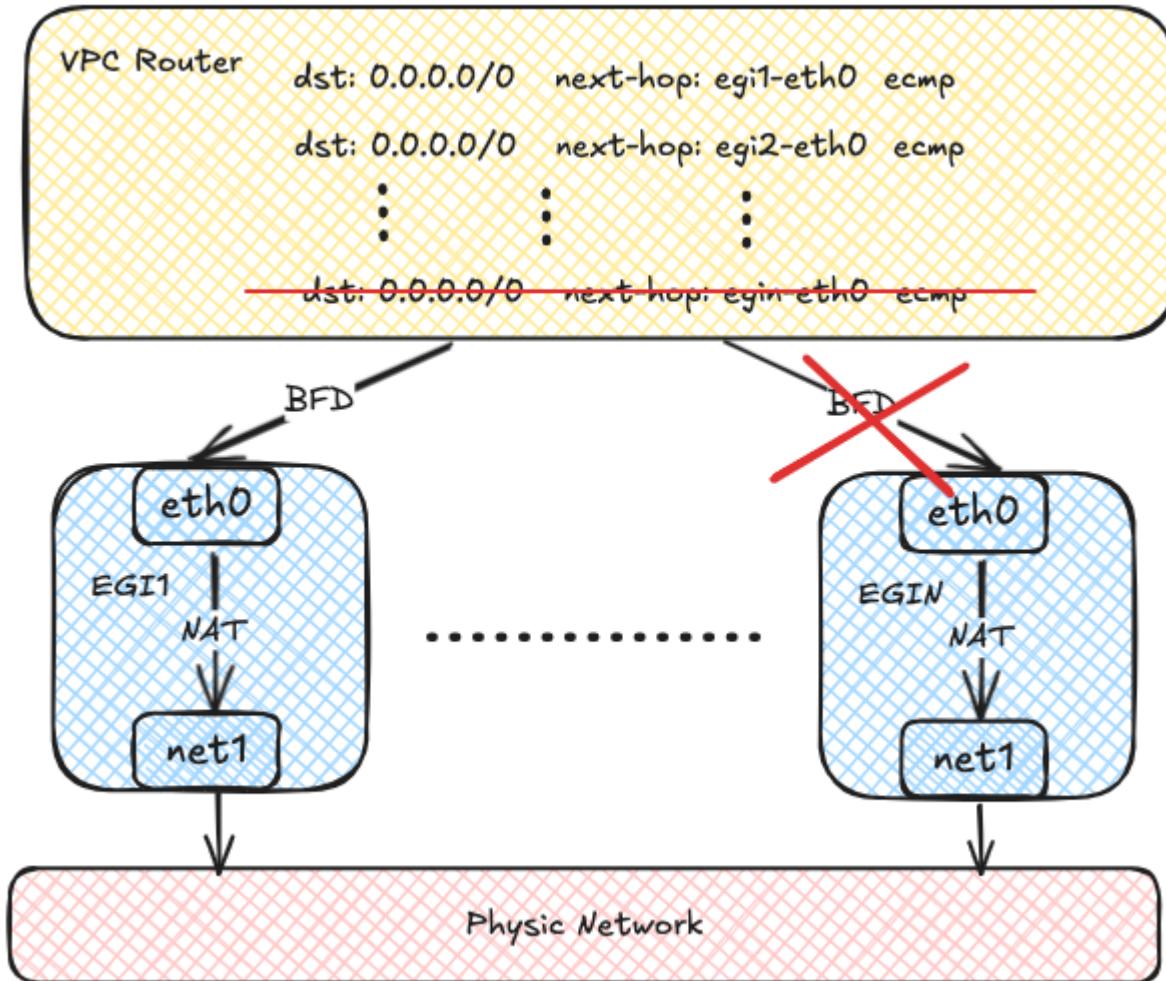
Egress Gateway

Egress Gateway



OVN   BFD      Egress Gateway

Egress Gateway      OVN



### 5.2.2

VPC Egress Gateway	VPC NAT Gateway	Multus-CNI
VPC Egress Gateway	ConfigMap	

### 5.2.3

#### NetworkAttachmentDefinition

VPC Egress Gateway	VPC	NetworkAttachmentDefinition	macvlan	Kube-OVN	IPAM
--------------------	-----	-----------------------------	---------	----------	------

```

apiVersion: k8s.cni.cncf.io/v1
kind: NetworkAttachmentDefinition
metadata:
 name: eth1
 namespace: default
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "macvlan",
 "master": "eth1",
 "mode": "bridge",
 "ipam": {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "eth1.default"
 }
 }'

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: macvlan1

```

```
spec:
 protocol: IPv4
 provider: eth1.default
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 excludeIps:
 - 172.17.0.0..172.17.0.10
```

CNI NetworkAttachmentDefinition VPC Egress Gateway

## VPC Egress Gateway

### VPC Egress Gateway

```
apiVersion: kubeovn.io/v1
kind: VpcEgressGateway
metadata:
 name: gateway1
 namespace: default
spec:
 vpc: ovn-cluster
 replicas: 1
 externalSubnet: macvlan1
 policies:
 - snat: true
 subnets:
 - ovn-default
```

default	VPC	ovn-cluster	gateway1	VPC Egress Gateway	ovn-cluster	ovn-default	10.16.0.0/16	Pod
macvlan1								SNAT

### VPC Egress Gateway

```
$ kubectl get veg gateway1
NAME VPC REPLICAS BFD ENABLED EXTERNAL SUBNET PHASE READY AGE
gateway1 ovn-cluster 1 false macvlan1 Completed true 13s
```

```
kubectl get veg gateway1 -o wide
NAME VPC REPLICAS BFD ENABLED EXTERNAL SUBNET PHASE READY INTERNAL IPS EXTERNAL IPS WORKING NODES AGE
gateway1 ovn-cluster 1 false macvlan1 Completed true ["10.16.0.12"] ["172.17.0.11"] ["kube-ovn-worker"] 82s
```

```
$ kubectl get deployment -l ovn.kubernetes.io/vpc-egress-gateway=gateway1
NAME READY UP-TO-DATE AVAILABLE AGE
gateway1 1/1 1 1 4m40s
```

```
$ kubectl get pod -l ovn.kubernetes.io/vpc-egress-gateway=gateway1 -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
gateway1-b9f8b4448-761hm 1/1 Running 0 4m48s 10.16.0.12 kube-ovn-worker <none> <none>
```

### Pod IP iptables

```
$ kubectl exec gateway1-b9f8b4448-761hm -c gateway -- ip address show
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::1/128 scope host
 valid_lft forever preferred_lft forever
2: net1@if13: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default qlen 1000
 link/ether 62:d8:71:90:7b:86 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 172.17.0.11/16 brd 172.17.255.255 scope global net1
 valid_lft forever preferred_lft forever
 inet6 fe80::60d8:71ff:fe90:7b86/64 scope link
 valid_lft forever preferred_lft forever
17: eth0@if18: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UP group default
 link/ether 36:7c:6b:c7:82:6b brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 10.16.0.12/16 brd 10.16.255.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet6 fe80::347c:6bff:fec7:826b/64 scope link
 valid_lft forever preferred_lft forever
```

```
$ kubectl exec gateway1-b9f8b4448-761hm -c gateway -- ip route show
default via 172.17.0.1 dev net1
10.16.0.0/16 dev eth0 proto kernel scope link src 10.16.0.12
```

```
172.17.0.0/16 dev net1 proto kernel scope link src 172.17.0.11

$ kubectl exec gateway1-b9f8b4448-76lhm -c gateway -- iptables -t nat -S
-P PREROUTING ACCEPT
-P INPUT ACCEPT
-P OUTPUT ACCEPT
-P POSTROUTING ACCEPT
-A POSTROUTING -s 10.16.0.0/16 -j MASQUERADE --random-fu
```

## Gateway Pod

```
$ kubectl exec -ti gateway1-b9f8b4448-76lhm -c gateway -- bash
nobody@gateway1-b9f8b4448-76lhm:/kube-ovn$ tcpdump -i any -nnve icmp and host 172.17.0.1
tcpdump: data link type LINUX_SLL2
tcpdump: listening on any, link-type LINUX_SLL2 (Linux cooked v2), snapshot length 262144 bytes
06:50:58.936528 eth0 In ifindex 17 92:26:b8:9e:f2:1c ethertype IPv4 (0x0800), length 104: (tos 0x0, ttl 63, id 30481, offset 0, flags [DF], proto ICMP (1), length 84)
 10.16.0.9 > 172.17.0.1: ICMP echo request, id 37989, seq 0, length 64
06:50:58.936574 net1 Out ifindex 2 62:d8:71:90:7b:86 ethertype IPv4 (0x0800), length 104: (tos 0x0, ttl 62, id 30481, offset 0, flags [DF], proto ICMP (1), length 84)
 172.17.0.11 > 172.17.0.1: ICMP echo request, id 39449, seq 0, length 64
06:50:58.936613 net1 In ifindex 2 02:42:39:79:7f:08 ethertype IPv4 (0x0800), length 104: (tos 0x0, ttl 64, id 26701, offset 0, flags [none], proto ICMP (1), length 84)
 172.17.0.1 > 172.17.0.11: ICMP echo reply, id 39449, seq 0, length 64
06:50:58.936621 eth0 Out ifindex 17 36:7c:6b:c7:82:6b ethertype IPv4 (0x0800), length 104: (tos 0x0, ttl 63, id 26701, offset 0, flags [none], proto ICMP (1), length 84)
 172.17.0.1 > 10.16.0.9: ICMP echo reply, id 37989, seq 0, length 64
```

## OVN Logical Router

### VPC

```
$ kubectl ko nbctl lr-policy-list ovn-cluster
Routing Policies
 31000 ip4.dst == 10.16.0.0/16 allow
 31000 ip4.dst == 100.64.0.0/16 allow
 30000 ip4.dst == 172.18.0.2 reroute 100.64.0.3
 30000 ip4.dst == 172.18.0.3 reroute 100.64.0.2
 30000 ip4.dst == 172.18.0.4 reroute 100.64.0.4
 29100 ip4.src == 10.16.0.0/16 reroute 10.16.0.12
 29000 ip4.src == $ovn.default.kube.ovn.control_plane_ip4 reroute 100.64.0.2
 29000 ip4.src == $ovn.default.kube.ovn.worker2_ip4 reroute 100.64.0.4
 29000 ip4.src == $ovn.default.kube.ovn.worker_ip4 reroute 100.64.0.3
```

```
.spec.replicas
```

```
$ kubectl scale veg gateway1 --replicas=2
vpceregressgateway.kubeovn.io/gateway1 scaled

$ kubectl get veg gateway1
NAME VPC REPLICAS BFD ENABLED EXTERNAL SUBNET PHASE READY AGE
gateway1 ovn-cluster 2 false macvlan Completed true 39m

$ kubectl get pod -l ovn.kubernetes.io/vpc-egress-gateway=gateway1 -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
gateway1-b9f8b4448-76lhm 1/1 Running 0 40m 10.16.0.12 kube-ovn-worker <none> <none>
gateway1-b9f8b4448-zd4dl 1/1 Running 0 64s 10.16.0.13 kube-ovn-worker2 <none> <none>

$ kubectl ko nbctl lr-policy-list ovn-cluster
Routing Policies
 31000 ip4.dst == 10.16.0.0/16 allow
 31000 ip4.dst == 100.64.0.0/16 allow
 30000 ip4.dst == 172.18.0.2 reroute 100.64.0.3
 30000 ip4.dst == 172.18.0.3 reroute 100.64.0.2
 30000 ip4.dst == 172.18.0.4 reroute 100.64.0.4
 29100 ip4.src == 10.16.0.0/16 reroute 10.16.0.12, 10.16.0.13
 29000 ip4.src == $ovn.default.kube.ovn.control_plane_ip4 reroute 100.64.0.2
 29000 ip4.src == $ovn.default.kube.ovn.worker2_ip4 reroute 100.64.0.4
 29000 ip4.src == $ovn.default.kube.ovn.worker_ip4 reroute 100.64.0.3
```

## BFD

BFD	VPC	BFD LRP	VPC	BFD Port

```
apiVersion: kubeovn.io/v1
kind: Vpc
metadata:
 name: vpc
spec:
 bfdPort:
 enabled: true
 ip: 10.255.255.255

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: subnet
```

```
spec:
 vpc: vpc1
 protocol: IPv4
 cidrBlock: 192.168.0.0/24
```

## BFD Port      OVN LR

```
$ kubectl ko nbctl show vpc1
router 0c1d1e8f-4c86-4d96-88b2-c4171c7ff824 (vpc1)
 port bfd@vpc1
 mac: "8e:51:4b:16:3c:90"
 networks: ["10.255.255.255"]
 port vpc1-subnet1
 mac: "de:c9:5c:38:7a:61"
 networks: ["192.168.0.1/24"]
```

## VPC Egress Gateway .spec.bfd.enabled true

```
apiVersion: kubeovn.io/v1
kind: VpcEgressGateway
metadata:
 name: gateway2
 namespace: default
spec:
 vpc: vpc1
 replicas: 2
 internalSubnet: subnet1
 externalSubnet: macvlan
 bfd:
 enabled: true
 policies:
 - snat: true
 ipBlocks:
 - 192.168.0.0/24
```

## VPC Egress Gateway

```
$ kubectl get veg gateway2 -o wide
NAME VPC REPLICAS BFD ENABLED EXTERNAL SUBNET PHASE READY INTERNAL IPS EXTERNAL IPs WORKING
NODES AGE
gateway2 vpc1 2 true macvlan Completed true ["192.168.0.2", "192.168.0.3"] ["172.17.0.13", "172.17.0.14"] ["kube-ovn-worker", "kube-ovn-worker2"]
 58s

$ kubectl get pod -l ovn.kubernetes.io/vpc-egress-gateway=gateway2 -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
gateway2-fcc6b8b87-8lgvx 1/1 Running 0 2m18s 192.168.0.3 kube-ovn-worker2 <none> <none>
gateway2-fcc6b8b87-wmww6 1/1 Running 0 2m18s 192.168.0.2 kube-ovn-worker <none> <none>

$ kubectl ko nbctl lr-route-list vpc1
IPv4 Routes
Route Table <main>:
 192.168.0.0/24 192.168.0.2 src-ip ecmp ecmp-symmetric-reply bfd
 192.168.0.0/24 192.168.0.3 src-ip ecmp ecmp-symmetric-reply bfd

$ kubectl ko nbctl list bfd
_uuid : 223ede10-9169-4c7d-9524-a546e24bfab5
detect_mult : 3
dst_ip : "192.168.0.2"
external_ids : {af="4", vendor=kube-ovn, vpc-egress-gateway="default/gateway2"}
logical_port : "bfd@vpc1"
min_rx : 1000
min_tx : 1000
options : {}
status : up

_uuid : b050c75e-2462-470b-b89c-7bd38889b758
detect_mult : 3
dst_ip : "192.168.0.3"
external_ids : {af="4", vendor=kube-ovn, vpc-egress-gateway="default/gateway2"}
logical_port : "bfd@vpc1"
min_rx : 1000
min_tx : 1000
options : {}
status : up
```

## Pod BFD

```
$ kubectl exec gateway2-fcc6b8b87-8lgvx -c bfdd -- bfdd-control status
There are 1 sessions:
Session 1
id=1 local=192.168.0.3 (p) remote=10.255.255.255 state=Up

$ kubectl exec gateway2-fcc6b8b87-wmww6 -c bfdd -- bfdd-control status
There are 1 sessions:
Session 1
id=1 local=192.168.0.2 (p) remote=10.255.255.255 state=Up
```

## VPC BFD PORT

<code>enabled</code>	<code>boolean</code>	<code>false</code>	<code>BFD Port</code>	<code>true</code>
<code>ip</code>	<code>string</code>	-	<code>BFD Port</code>	<code>IP</code> <code>IPv6</code> 169.255.255.2 169.255.255.2
<code>nodeSelector</code>	<code>object</code>	-	<code>BFD Port</code>	-
			<code>OVN HA Chassis</code> <code>Group</code> Active/ Backup Active <code>nodeSelector</code> Kube-OVN	
			<code>kubectl ko nbctl</code> <code>list</code> <code>ha_chassis_group</code> OVN HA Chassis Group	
<code>nodeSelector.matchLabels</code>	<code>dict/map</code>	-		-
<code>nodeSelector.matchExpressions</code>	<code>object array</code>	-		-

## VPC EGRESS GATEWAY

## Spec

vpc	string	VPC	ovn-cluster	VPC	vpc1
replicas	integer/int32	1			2
prefix	string	-		Deployment	veg-
image	string	-		Deployment	docker.io/kubeovn/kube-ovn
internalSubnet	string	VPC		VPC	subnet1
externalSubnet	string	-			ext1
internalIPs	string array	-	VPC IPv6	IP IP	10.16.0.101 / fd00::11 / 1
			<replicas> + 1 Pod		
externalIPs	string array	-	IPv6	IP IP	10.16.0.101 / fd00::11 / 1
			<replicas> + 1 Pod		
bfd	object	-		BFD	-
policies	object array	-		Egress selectors	-
selectors	object array	-		Namespace Selector Pod Selector Egress Pod SNAT/ MASQUERADE policies	-
nodeSelector	object array	-		Deployment/Pod	-
trafficPolicy	string	Cluster		Cluster / Local <b>BFD</b> Local Egress VPC	Local
				Egress Gateway VPC Egress Gateway Egress	

**BFD**

<code>enabled</code>	<code>boolean</code>	<code>false</code>	<code>BFD</code>	<code>true</code>
<code>minRX</code>	<code>integer/int32</code>	<code>1000</code>	<code>BFD minRX</code>	<code>500</code>
<code>minTX</code>	<code>integer/int32</code>	<code>1000</code>	<code>BFD minTX</code>	<code>500</code>
<code>multiplier</code>	<code>integer/int32</code>	<code>3</code>	<code>BFD multiplier</code>	<code>1</code>

**Egress**

<code>snat</code>	<code>boolean</code>	<code>false</code>	<code>SNAT/MASQUERADE</code>	<code>true</code>
<code>ipBlocks</code>	<code>string array</code>	<code>-</code>	<code>Gateway IP</code>	<code>192.168.0.1 / 192.168.0.0/24</code>
<code>subnets</code>	<code>string array</code>	<code>-</code>	<code>Gateway VPC</code>	<code>subnet1 IPv6</code>

**Selectors**

<code>namespaceSelector</code>	<code>object</code>	<code>-</code>	<code>Namespace</code>	<code>-</code>
<code>podSelector</code>	<code>object</code>	<code>-</code>	<code>Pod</code>	<code>-</code>
<code>podSelector.matchLabels</code>	<code>dict/map</code>	<code>-</code>	<code>Pod</code>	<code>-</code>
<code>podSelector.matchExpressions</code>	<code>object array</code>	<code>-</code>	<code>Pod</code>	<code>-</code>
<code>namespaceSelector.matchLabels</code>	<code>dict/map</code>	<code>-</code>	<code>Namespace</code>	<code>-</code>
<code>namespaceSelector.matchExpressions</code>	<code>object array</code>	<code>-</code>	<code>Namespace</code>	<code>-</code>

<code>matchLabels</code>	<code>dict/map</code>	<code>-</code>	<code>-</code>
<code>matchExpressions</code>	<code>object array</code>	<code>-</code>	<code>-</code>
<code>matchFields</code>	<code>object array</code>	<code>-</code>	<code>-</code>

## Status

ready	boolean	Gateway	true
phase	string	Gateway	Pending / Processing / Completed
internalIPs	string array	VPC	IP
externalIPs	string array	IP	-
workload	object		-
workload.apiVersion	string	API	apps/v1
workload.kind	string		Deployment
workload.name	string		gateway1
workload.nodes	string array		-
conditions	object array	-	-

[PDF](#)[Slack](#)[Support](#)

⌚2025 7 29

⌚2024 12 12



GitHub



5.2.4

## 5.3 VPC QoS

Kube-OVN    QoS Policy CRD    VPC

### 5.3.1 EIP QoS

EIP	1Mbps	1	shared=false	QoS Policy	EIP	QoS Policy	QoS
-----	-------	---	--------------	------------	-----	------------	-----

QoS Policy

```
apiVersion: kubeovn.io/v1
kind: QoS Policy
metadata:
 name: qos-eip-example
spec:
 shared: false
 bindingType: EIP
 bandwidthLimitRules:
 - name: eip-ingress
 rateMax: "1" # Mbps
 burstMax: "1" # Mbps
 priority: 1
 direction: ingress
 - name: eip-egress
 rateMax: "1" # Mbps
 burstMax: "1" # Mbps
 priority: 1
 direction: egress
```

Iptables EIP

```
kind: IptablesEIP
apiVersion: kubeovn.io/v1
metadata:
 name: eip-1
spec:
 natGwDp: gw1
 qosPolicy: qos-eip-example
```

.spec.qosPolicy

### 5.3.2 QoS EIP

label	qos	eip
-------	-----	-----

```
kubectl get eip -l ovn.kubernetes.io/qos=qos-eip-example
NAME IP MAC NAT NATGWDP READY
eip-1 172.18.11.24 00:00:00:34:41:0B fip gw1 true
```

### 5.3.3 VPC NATGW net1 QoS

VPC NATGW	net1	10Mbps	3	shared=true	QoS Policy	QoS Policy
-----------	------	--------	---	-------------	------------	------------

QoS Policy

```
apiVersion: kubeovn.io/v1
kind: QoS Policy
metadata:
 name: qos-natgw-example
spec:
 shared: true
 bindingType: NATGW
 bandwidthLimitRules:
 - name: net1-ingress
 interface: net1
 rateMax: "10" # Mbps
 burstMax: "10" # Mbps
 priority: 3
 direction: ingress
 - name: net1-egress
 interface: net1
 rateMax: "10" # Mbps
 burstMax: "10" # Mbps
```

```
priority: 3
direction: egress
```

#### VpcNatGateway

```
kind: VpcNatGateway
apiVersion: kubeovn.io/v1
metadata:
 name: gw1
spec:
 vpc: test-vpc-1
 subnet: net1
 lanIp: 10.0.1.254
 qosPolicy: qos-natgw-example
 selector:
 - "kubernetes.io/hostname: kube-ovn-worker"
 - "kubernetes.io/os: linux"
```

```
.spec.qosPolicy
```

### 5.3.4 net1 QoS

net1	5Mbps	2	shared=true	QoS Policy	QoS Policy
------	-------	---	-------------	------------	------------

#### QoS Policy

```
apiVersion: kubeovn.io/v1
kind: QoS Policy
metadata:
 name: qos-natgw-example
spec:
 shared: true
 bindingType: NATGW
 bandwidthLimitRules:
 - name: net1-extip-ingress
 interface: net1
 rateMax: "5" # Mbps
 burstMax: "5" # Mbps
 priority: 2
 direction: ingress
 matchType: ip
 matchValue: src 172.18.11.22/32
 - name: net1-extip-egress
 interface: net1
 rateMax: "5" # Mbps
 burstMax: "5" # Mbps
 priority: 2
 direction: egress
 matchType: ip
 matchValue: dst 172.18.11.23/32
```

#### VpcNatGateway

```
kind: VpcNatGateway
apiVersion: kubeovn.io/v1
metadata:
 name: gw1
spec:
 vpc: test-vpc-1
 subnet: net1
 lanIp: 10.0.1.254
 qosPolicy: qos-natgw-example
 selector:
 - "kubernetes.io/hostname: kube-ovn-worker"
 - "kubernetes.io/os: linux"
```

### 5.3.5 QoS NATGW

label	qos	eip
-------	-----	-----

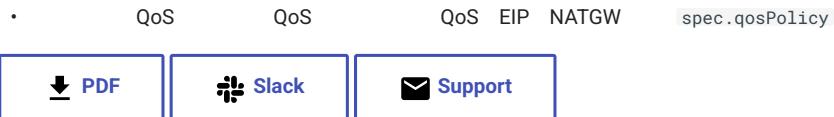
```
kubectl get vpc-nat-gw -l ovn.kubernetes.io/qos=qos-natgw-example
NAME VPC SUBNET LANIP
gw1 test-vpc-1 net1 10.0.1.254
```

### 5.3.6 qos

```
kubectl get qos -A
NAME SHARED BINDINGTYPE
```

qos-eip-example	false	EIP
qos-natgw-example	true	NATGW

### 5.3.7



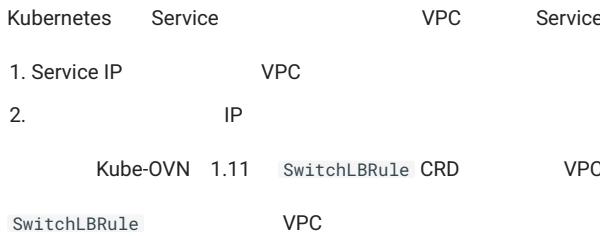
⌚2025 3 3

⌚2023 5 9



### 5.3.8

## 5.4 VPC



### 5.4.1 Selector

```

 selector label pod
 SwitchLBRule

```

```

apiVersion: kubeovn.io/v1
kind: SwitchLBRule
metadata:
 name: cjh-slr-nginx
spec:
 vip: 1.1.1.1
 sessionAffinity: ClientIP
 namespace: default
 selector:
 - app:nginx
 ports:
 - name: dns
 port: 8888
 targetPort: 80
 protocol: TCP

```

- **selector**, **sessionAffinity**, **port**      Kubernetes Service
- **vip**      IP
- **namespace** **selector**      Pod

```

Kube-OVN SwitchLBRule Pod Pod VPC L2 LB

```

### 5.4.2 Endpoints

```

 endpoints selector kubevirt vm
 SwitchLBRule

```

```

apiVersion: kubeovn.io/v1
kind: SwitchLBRule
metadata:
 name: cjh-slr-nginx
spec:
 vip: 1.1.1.1
 sessionAffinity: ClientIP
 namespace: default
 endpoints:
 - 192.168.0.101
 - 192.168.0.102
 - 192.168.0.103
 ports:
 - name: dns
 port: 8888
 targetPort: 80
 protocol: TCP

```

- **sessionAffinity**, **port**      Kubernetes Service
- **vip**      IP
- **namespace** **selector**      Pod
- **endpoints**      IP

```
 selector endpoints , selector
```

### 5.4.3

OVN IPv4

[Health Checks](<https://www.ovn.org/support/dist-docs/ovn-nb.5.html>)

ovn	SwitchLBRule	SwitchLBRule	VPC	subnet	vip	ip_port_mappings
load_balancer_health_check						
•	vip	subnet	subnet		SwitchLBRule	
•	Selector					

```
root@server:~# kubectl get po -o wide -n vulpecula
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
nginx-78d9578975-f4qn4 1/1 Running 3 4d16h 10.16.0.4 worker <none> <none>
nginx-78d9578975-t8tm5 1/1 Running 3 4d16h 10.16.0.6 worker <none> <none>

slr
root@server:~# cat << END > slr.yaml
apiVersion: kubeovn.io/v1
kind: SwitchLBRule
metadata:
 name: nginx
 namespace: vulpecula
spec:
 vip: 1.1.1.1
 sessionAffinity: ClientIP
 namespace: default
 selector:
 - app:nginx
 ports:
 - name: dns
 port: 8888
 targetPort: 80
 protocol: TCP
END
root@server:~# kubectl apply -f slr.yaml
root@server:~# kubectl get slr
NAME VIP PORT(S) SERVICE AGE
vulpecula-nginx 1.1.1.1 8888/TCP default/slr-vulpecula-nginx 3d21h
```

subnet vip

```
vip

root@server:~# kubectl get vip
NAME NS V4IP MAC V6IP PMAC SUBNET READY TYPE
vulpecula-subnet 10.16.0.2 00:00:00:39:95:C1 <nil> vulpecula-subnet true
```

Load\_Balancer\_Health\_Check Service\_Monitor

```
root@server:~# kubectl ko nbctl list Load_Balancer
_uuid : 3ccb6d43-44aa-4028-962f-30d2dba9f0b8
external_ids : {}
health_check : [5bee3f12-6b54-411c-9cc8-c9def8f67356]
ip_port_mappings : {"10.16.0.4":"nginx-78d9578975-f4qn4.default:10.16.0.2", "10.16.0.6":"nginx-78d9578975-t8tm5.default:10.16.0.2"}
name : cluster-tcp-session-loadbalancer
options : {affinity_timeout="10800"}
protocol : tcp
selection_fields : [ip_src]
vips : {"1.1.1.1:8888":"10.16.0.4:80,10.16.0.6:80"}

root@server:~# kubectl ko nbctl list Load_Balancer_Health_Check
_uuid : 5bee3f12-6b54-411c-9cc8-c9def8f67356
external_ids : {switch_lb_subnet=vulpecula-subnet}
options : {failure_count="3", interval="5", success_count="3", timeout="20"}
vip : "1.1.1.1:8888"

root@server:~# kubectl ko sbctl list Service_Monitor
_uuid : 1bddc541-cc49-44ea-9935-a4208f627a91
external_ids : {}
ip : "10.16.0.4"
logical_port : nginx-78d9578975-f4qn4.default
options : {failure_count="3", interval="5", success_count="3", timeout="20"}
port : 80
protocol : tcp
```

```

src_ip : "10.16.0.2"
src_mac : "c6:d4:b8:08:54:e7"
status : online

_uuid : 84dd24c5-e1b4-4e97-9daa-13687ed59785
external_ids: {}
ip : "10.16.0.6"
logical_port: nginx-78d9578975-t8tm5.default
options : {"failure_count="3", interval="5", success_count="3", timeout="20"}
port : 80
protocol : tcp
src_ip : "10.16.0.2"
src_mac : "c6:d4:b8:08:54:e7"
status : online

```

## vip

```

root@server:~# kubectl exec -it -n vulpecula nginx-78d9578975-t8tm5 -- curl 1.1.1.1:8888
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<p>Thank you for using nginx.</p>
</body>
</html>

```

## pod

```

kubectl delete po nginx-78d9578975-f4qn4
kubectl get po -o wide -n vulpecula
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
nginx-78d9578975-lxmvh 1/1 Running 0 31s 10.16.0.8 worker <none> <none>
nginx-78d9578975-t8tm5 1/1 Running 3 4d16h 10.16.0.6 worker <none> <none>

```

## Load\_Balancer\_Health\_Check Service\_Monitor

```

root@server:~# kubectl ko nbctl list Load_Balancer
_uuid : 3ccb6d43-44aa-4028-962f-30d2dba9f0b8
external_ids : {}
health_check : [5bee3f12-6b54-411c-9cc8-c9def8f67356]
ip_port_mappings: {"10.16.0.4": "nginx-78d9578975-f4qn4.default:10.16.0.2", "10.16.0.6": "nginx-78d9578975-t8tm5.default:10.16.0.2", "10.16.0.8": "nginx-78d9578975-lxmvh.default:10.16.0.2"}
name : cluster-tcp-session-loadbalancer
options : {"affinity_timeout": "10800"}
protocol : tcp
selection_fields: [ip_src]
vips : {"1.1.1.1:8888": "10.16.0.6:80,10.16.0.8:80"}

root@server:~# kubectl ko nbctl list Load_Balancer_Health_Check
_uuid : 5bee3f12-6b54-411c-9cc8-c9def8f67356
external_ids : {switch_lb_subnet=vulpecula-subnet}
options : {"failure_count": "3", interval="5", success_count="3", timeout="20"}
vip : "1.1.1.1:8888"

root@server:~# kubectl ko sbctl list Service_Monitor
_uuid : 84dd24c5-e1b4-4e97-9daa-13687ed59785
external_ids : {}
ip : "10.16.0.6"
logical_port : nginx-78d9578975-t8tm5.default
options : {"failure_count": "3", interval="5", success_count="3", timeout="20"}
port : 80
protocol : tcp
src_ip : "10.16.0.2"
src_mac : "c6:d4:b8:08:54:e7"
status : online

_uuid : 5917b7b7-a999-49f2-a42d-da81f1eeb28f
external_ids : {}
ip : "10.16.0.8"
logical_port : nginx-78d9578975-lxmvh.default
options : {"failure_count": "3", interval="5", success_count="3", timeout="20"}
port : 80
protocol : tcp
src_ip : "10.16.0.2"
src_mac : "c6:d4:b8:08:54:e7"
status : online

```

## SwitchLBRule Load\_Balancer\_Health\_Check Service\_Monitor vip

```

root@server:~# kubectl delete -f slr.yaml
switchlrule.kubeovn.io "vulpecula-nginx" deleted
root@server:~# kubectl get vip
No resources found

```

```
root@server:~# kubectl ko sbctl list Service_Monitor
root@server:~#
root@server:~# kubectl ko nbctl list Load_Balancer_Health_Check
root@server:~#
```

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2022 12 15



5.4.4

---

## 5.5 VPC DNS



### 5.5.1 vpc-dns

```

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
 labels:
 kubernetes.io/bootstrapping: rbac-defaults
 name: system:vpc-dns
rules:
- apiGroups:
 - ""
 resources:
 - endpoints
 - services
 - pods
 - namespaces
 verbs:
 - list
 - watch
- apiGroups:
 - discovery.k8s.io
 resources:
 - endpointslices
 verbs:
 - list
 - watch

apiVersion: rbac.authorization.k8s.io/v1
kind: ClusterRoleBinding
metadata:
 annotations:
 rbac.authorization.kubernetes.io/autoupdate: "true"
 labels:
 kubernetes.io/bootstrapping: rbac-defaults
 name: vpc-dns
roleRef:
 apiGroup: rbac.authorization.k8s.io
 kind: ClusterRole
 name: system:vpc-dns
subjects:
- kind: ServiceAccount
 name: vpc-dns
 namespace: kube-system

apiVersion: v1
kind: ServiceAccount
metadata:
 name: vpc-dns
 namespace: kube-system

apiVersion: v1
kind: ConfigMap
metadata:
 name: vpc-dns-corefile
 namespace: kube-system
data:
 Corefile: |
 .:53 {
 errors
 health {
 lameduck 5s
 }
 ready
 kubernetes cluster.local in-addr.arpa ip6.arpa {
 pods insecure
 fallthrough in-addr.arpa ip6.arpa
 }
 prometheus :9153
 forward . /etc/resolv.conf {
 prefer_udp
 }
 }
 cache 30

```

```

 loop
 reload
 loadbalance
}

```

nat-gw-pod

### 5.5.2

```

apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: ovn-nad
 namespace: default
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "ovn-nad.default.ovn"
}'

```

### 5.5.3 vpc-dns Configmap

kube-system configmap vpc-dns vpc-dns

```

apiVersion: v1
kind: ConfigMap
metadata:
 name: vpc-dns-config
 namespace: kube-system
data:
 coredns-vip: 10.96.0.3
 enable-vpc-dns: "true"
 nad-name: ovn-nad
 nad-provider: ovn-nad.default.ovn

```

- enable-vpc-dns true
- coredns-image dns coredns
- coredns-vip coredns lb vip
- coredns-template coredns URL ovn coredns-template.yaml <https://raw.githubusercontent.com/kubeovn/kube-ovn/> /yamls/coredns-template.yaml
- nad-name network-attachment-definitions
- nad-provider provider
- k8s-service-host coredns k8s apiserver ip apiserver
- k8s-service-port coredns k8s apiserver port apiserver

### 5.5.4 vpc-dns

vpc-dns yaml

```

kind: VpcDns
apiVersion: kubeovn.io/v1
metadata:
 name: test-cjh1
spec:
 vpc: cjh-vpc-1
 subnet: cjh-subnet-1
 replicas: 2

```

- vpc dns vpc
- subnet dns
- replicas: vpc dns deployment replicas

```
kubectl get vpc-dns
NAME ACTIVE VPC SUBNET
test-cjh1 false cjh-vpc-1 cjh-subnet-1
test-cjh2 true cjh-vpc-1 cjh-subnet-2
```

```
ACTIVE : true dns false
 VPC DNS
• VPC vpc-dns VPC subnet vpc-dns true false
• true vpc-dns false vpc-dns
```

### 5.5.5

vpc-dns Pod      label app=vpc-dns      vpc-dns Pod

```
kubectl -n kube-system get pods -l app=vpc-dns
NAME READY STATUS RESTARTS AGE
vpc-dns-test-cjh1-7b878d96b4-g5979 1/1 Running 0 28s
vpc-dns-test-cjh1-7b878d96b4-ltmf9 1/1 Running 0 28s
```

slr

```
kubectl -n kube-system get slr
NAME VIP PORT(S) SERVICE AGE
vpc-dns-test-cjh1 10.96.0.3 53/UDP, 53/TCP kube-system/vpc-dns-test-cjh1 113s
```

VPC   Pod   dns :

```
nslookup kubernetes.default.svc.cluster.local 10.96.0.3
```

VPC   switch lb rule      VPC      Pod

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 12 26



### 5.5.6

## 5.6 SecurityGroup

Kube-OVN	Pod	Pod	IPAM	L2/L3	MAC	IP	
<b>⚠ Warning</b>							
Kube-OVN	<a href="#">NetworkPolicy</a>	<a href="#">Network Policy API</a>	<a href="#">Subnet ACL</a>	<a href="#">Security Group</a>	<a href="#">OVN ACL</a>	<a href="#">NetworkPolicy</a>	<a href="#">NetworkPolicy API</a>

### 5.6.1

```
apiVersion: kubeovn.io/v1
kind: SecurityGroup
metadata:
 name: sg-example
spec:
 allowSameGroupTraffic: true
 egressRules:
 - ipVersion: ipv4
 policy: allow
 priority: 1
 protocol: all
 remoteAddress: 10.16.0.13 # 10.16.0.0/16
 remoteType: address
 ingressRules:
 - ipVersion: ipv4
 policy: deny
 priority: 1
 protocol: icmp
 remoteAddress: 10.16.0.14
 remoteType: address
```

#### Kube-OVN

Pod	annotation	annotation	Kube-OVN	IPAM	L2/L3	Pod	L2/L3	Pod
• port_security:			Kube-OVN	IPAM	L2/L3	Pod	L2/L3	Pod
• security_groups:		ACL						
annotation								
ovn.kubernetes.io/port_security: "true" ovn.kubernetes.io/security_groups: sg-example								

### 5.6.2

•	ACL	OVN	ACL	ACL				
•	priority	1-200	ACL	ACL	ACL	= 2300 -		ACL
•		Kube-OVN	CNI	Pod	Pod	Pod	ContainerCreating	Running

### 5.6.3

YAML Pod annotation

```
apiVersion: v1
kind: Pod
metadata:
 labels:
 app: static
 annotations:
 ovn.kubernetes.io/port_security: 'true'
 ovn.kubernetes.io/security_groups: 'sg-example'
 name: sg-test-pod
 namespace: default
spec:
 nodeName: kube-ovn-worker
 containers:
```

```
- image: docker.io/library/nginx:alpine
 imagePullPolicy: IfNotPresent
 name: qatest
```

```
kubectl get pod -o wide
NAME READY STATUS RESTARTS AGE IP NODE
sg-test-pod 0/1 ContainerCreating 0 5h32m <none> kube-ovn-worker
test-99fffff86-52h9r 1/1 Running 0 5h41m 10.16.0.14 kube-ovn-control-plane
test-99fffff86-qcgjw 1/1 Running 0 5h43m 10.16.0.13 kube-ovn-worker
```

kubectl describe pod Pod

```
kubectl describe pod sg-test-pod
Name: sg-test-pod
Namespace: default
Priority: 0
Node: kube-ovn-worker/172.18.0.2
Start Time: Tue, 28 Feb 2023 10:29:36 +0800
Labels: app=static
Annotations: ovn.kubernetes.io/allocated: true
 ovn.kubernetes.io/cidr: 10.16.0.0/16
 ovn.kubernetes.io/gateway: 10.16.0.1
 ovn.kubernetes.io/ip_address: 10.16.0.15
 ovn.kubernetes.io/logical_router: ovn-cluster
 ovn.kubernetes.io/logical_switch: ovn-default
 ovn.kubernetes.io/mac_address: 00:00:00:FA:17:97
 ovn.kubernetes.io/pod_nic_type: veth-pair
 ovn.kubernetes.io/port_security: true
 ovn.kubernetes.io/routed: true
 ovn.kubernetes.io/security_groups: sg-allow-reject
Status: Pending
IP:
IPs: <none>
.
.
Events:
 Type Reason Age From Message
 ---- ---- -- ---- -----
 Warning FailedCreatePodSandBox 5m3s (x70 over 4h59m) kubelet (combined from similar events): Failed to create pod sandbox: rpc error: code = Unknown desc = failed to setup network for sandbox "40636e0c7f1ade5500fa958486163d74f2e2300051a71522a9af7ba0538afb6": plugin type="kube-ovn" failed (add): RPC failed; request ip return 500 configure nic failed 10.16.0.15 network not ready after 200 ping 10.16.0.1
```

```
apiVersion: kubeovn.io/v1
kind: SecurityGroup
metadata:
 name: sg-gw-both
spec:
 allowSameGroupTraffic: true
 egressRules:
 - ipVersion: ipv4
 policy: allow
 priority: 2
 protocol: all
 remoteAddress: 10.16.0.13
 remoteType: address
 - ipVersion: ipv4
 policy: allow
 priority: 1
 protocol: all
 remoteAddress: 10.16.0.1
 remoteType: address
 ingressRules:
 - ipVersion: ipv4
 policy: deny
 priority: 2
 protocol: icmp
 remoteAddress: 10.16.0.14
 remoteType: address
 - ipVersion: ipv4
 policy: allow
 priority: 1
 protocol: icmp
 remoteAddress: 10.16.0.1
 remoteType: address
```

yaml      Pod      Pod

```
apiVersion: v1
kind: Pod
```

```

metadata:
 labels:
 app: static
 annotations:
 ovn.kubernetes.io/port_security: 'true'
 ovn.kubernetes.io/security_groups: 'sg-gw-both'
 name: sg-gw-both
 namespace: default
spec:
 nodeName: kube-ovn-worker
 containers:
 - image: docker.io/library/nginx:alpine
 imagePullPolicy: IfNotPresent
 name: qatest

```

## Pod

# kubectl get pod -o wide									
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	NODE	READINESS GATES
sg-test-pod	0/1	ContainerCreating	0	5h41m	<none>	kube-ovn-worker	<none>	<none>	<none>
sg-gw-both	1/1	Running	0	5h37m	10.16.0.19	kube-ovn-worker	<none>	<none>	<none>

[PDF](#)
[Slack](#)
[Support](#)

⌚2025 9 10

⌚2023 2 28



5.6.4

## 5.7 OVN EIP FIP SNAT DNAT

VPC OVN NAT

provider-network vlan (external) subnet

VPC EIP/SNAT

### 5.7.1

- kube-ovn-controller kube-ovn-cni ovn-external-gw-config VPC spec enableExternal
- CRD provider-network vlan subnet VPC spec extraExternalSubnets ovn-eip ovn-dnat ovn-fip ovn-snats CRD

```
graph LR
pod --> subnet --> vpc --> lrp --bind--> gw-chassis --> snat --> lsp --> external-subnet
lrp -. peer -. -lsp
```

Pod SNAT Pod Fip

```
graph LR
pod --> subnet --> vpc --> lrp --bind--> local-chassis --> snat --> lsp --> external-subnet
lrp -. peer -. -lsp
```

Pod FIP (dnat\_and\_snat)

- CRD iptables nat gw
- ovn eip: ip underlay provider network vlan subnet
  - ovn fip dnat snat VPC ip vip
  - ovn snat VPC ip snat
  - ovn dnat router lb , ip + VPC endpoints

### 5.7.2 1.

OpenStack Neutron ovn

provider network

VPC EIP/SNAT

vlan vlan 0 vlan id

```
#
1. kube-ovn-controller
- --external-gateway-vlanid=204
- --external-gateway-switch=external204

2. kube-ovn-cni :
- --external-gateway-switch=external204

vlan id underlay
```

- provider network vlan subnet
- VPC enable\_eip\_snat vlan subnet ip ipam
- VPC enable\_eip\_snat , pod annotation fip snat
- VPC enable\_eip\_snat vlan subnet VPC eip snat

## 1.1 underlay

```
provider-network vlan subnet
cat 01-provider-network.yaml

apiVersion: kubeovn.io/v1
kind: ProviderNetwork
metadata:
 name: external204
spec:
 defaultInterface: vlan

cat 02-vlan.yaml

apiVersion: kubeovn.io/v1
kind: Vlan
metadata:
 name: vlan204
spec:
 id: 204
 provider: external204

cat 03-vlan-subnet.yaml

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: external204
spec:
 protocol: IPv4
 cidrBlock: 10.5.204.0/24
 gateway: 10.5.204.254
 vlan: vlan204
 excludeIps:
 - 10.5.204.1..10.5.204.100
```

## 1.2 VPC eip\_snat

```
VPC underlay provider subnet
cat 00-centralized-external-gw-no-ip.yaml

apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-external-gw-config
 namespace: kube-system
data:
 enable-external-gw: "true"
 external-gw-nodes: "pc-node-1,pc-node-2,pc-node-3"
 type: "centralized"
 external-gw-nic: "vlan" # ovs
 external-gw-addr: "10.5.204.254/24" # underlay ip
```

logical router port (lrp)	ip	mac	underlay	lrp	ovn eip
ip	lrp	ovn-eip		lrp	ovn eip

## 1.3 VPC eip snat fip

node

```
external-gw-nodes
kubectl label nodes pc-node-1 pc-node-2 pc-node-3 ovn.kubernetes.io/external-gw=true

cat 00-ns.yaml

apiVersion: v1
kind: Namespace
metadata:
 name: vpc1

cat 01-vpc-ecmp-enable-external-bfd.yaml

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc1
spec:
 namespaces:
 - vpc1
 enableExternal: true
 staticRoutes:
 - cidr: 0.0.0.0/0
 nextHopIP: 10.5.204.254
```

```

policy: policyDst

VPC enableExternal lrp

cat 02-subnet.yml

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: vpc1-subnet1
spec:
 cidrBlock: 192.168.0.0/24
 default: false
 disableGatewayCheck: false
 disableInterConnection: true
 enableEcmp: true
 gatewayNode: ""
 gatewayType: distributed
 #gatewayType: centralized
 natOutgoing: false
 private: false
 protocol: IPv4
 provider: ovn
 vpc: vpc1
 namespaces:
 - vpc1

#
subnet

```

```

kubectl ko nbctl show vpc1

router 87ad06fd-71d5-4ff8-a1f0-54fa3bba1a7f (vpc1)
 port vpc1-vpc1-subnet1
 mac: "00:00:00:ED:8E:C7"
 networks: ["192.168.0.1/24"]
 port vpc1-external294
 mac: "00:00:00:EF:05:C7"
 networks: ["10.5.204.105/24"]
 gateway chassis: [7cedd14f-265b-42e5-ac17-e03e7a1f2342 276baccb-fe9c-4476-b41d-05872a94976d fd9f140c-c45d-43db-a6c0-0d4f8ea298dd]
 nat 21d853b0-f7b4-40bd-9a53-31d2e2745739
 external ip: "10.5.204.115"
 logical ip: "192.168.0.0/24"
 type: "snat"

kubectl ko nbctl lr-route-list vpc1

IPv4 Routes
Route Table <main>:
 0.0.0.0/0 10.5.204.254 dst-ip
VPC CRD

```

enableExternal

VPC CRD

## 1.4

### 1.4.1 UNDERLAY

eip snat fip 1 eip snat fip

```

provider-network vlan subnet
cat 01-extra-provider-network.yaml
apiVersion: kubeovn.io/v1
kind: ProviderNetwork
metadata:
 name: extra
spec:
 defaultInterface: vlan
cat 02-extra-vlan.yaml
apiVersion: kubeovn.io/v1
kind: Vlan
metadata:
 name: vlan0
spec:
 id: 0
 provider: extra
cat 03-extra-vlan-subnet.yaml
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: extra
spec:
 protocol: IPv4
 cidrBlock: 10.10.204.0/24
 gateway: 10.10.204.254
 vlan: vlan0

```

```
excludeIps:
- 10.10.204.1..10.10.204.100
```

#### 1.4.2 VPC

```
apiVersion: kubeovn.io/v1
kind: Vpc
metadata:
 name: vpc1
spec:
 namespaces:
 - vpc1
 enableExternal: true # enableExternal VPC external ls
 extraExternalSubnets: # extraExternalSubnets
 - extra
```

```
kubectl ko nbctl show vpc1
router 87ad06fd-71d5-4ff8-a1f0-54fa3bba1a7f (vpc1)
 port vpc1-vpc1-subnet
 mac: "00:00:00:ED:8E:C7"
 networks: ["192.168.0.1/24"]
 port vpc1-external204
 mac: "00:00:00:EF:05:C7"
 networks: ["10.5.204.105/24"]
 gateway chassis: [7cedd14f-265b-42e5-ac17-e03e7a1f2342 276baccb-fe9c-4476-b41d-05872a94976d fd9f140c-c45d-43db-a6c0-0d4f8ea298dd]
 port vpc1-extra
 mac: "00:00:00:EF:6A:C7"
 networks: ["10.10.204.105/24"]
 gateway chassis: [7cedd14f-265b-42e5-ac17-e03e7a1f2342 276baccb-fe9c-4476-b41d-05872a94976d fd9f140c-c45d-43db-a6c0-0d4f8ea298dd]
```

### 5.7.3 2. ovn-eip

	iptables-eip	ovn-eip	type
• nat:	ovn dnat fip, snat	nat	
• lrp:	underlay	lrp ip	dnat snat
• lsp:	ovn bfd ecmp	ovs internal port	ecmp

```

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-static
spec:
 externalSubnet: external204
 type: nat
eip fip
```

	externalSubnet	external204
	externalSubnet	
		extra

#### 2.1 ovn-fip pod fip

```
kubectl get po -o wide -n vpc1 vpc-1-busybox01
NAME READY STATUS RESTARTS AGE IP NODE
vpc-1-busybox01 1/1 Running 0 3d15h 192.168.0.2 pc-node-2

kubectl get ip vpc-1-busybox01.vpc1
NAME V4IP V6IP MAC NODE SUBNET
vpc-1-busybox01.vpc1 192.168.0.2 00:00:00:0A:DD:27 pc-node-2 vpc1-subnet1

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-static
spec:
 externalSubnet: external204
 type: nat

kind: OvnFip
apiVersion: kubeovn.io/v1
metadata:
```

```

 name: eip-static
spec:
 ovnEip: eip-static
 ipName: vpc-1-busybox01.vpc1 # ip crd
 type: "centralized" # centralized distributed

--
VPC ip

kind: OvnFip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-static
spec:
 ovnEip: eip-static
 vpc: vpc1
 v4Ip: 192.168.0.2
 type: "centralized" # centralized distributed

kubectl get ofip
NAME VPC V4EIP V4IP READY IPTYPE IPNAME
eip-for-vip vpc1 10.5.204.106 192.168.0.3 true vip test-fip-vip
eip-static vpc1 10.5.204.101 192.168.0.2 true vpc vpc-1-busybox01.vpc1
kubectl get ofip eip-static
NAME VPC V4EIP V4IP READY IPTYPE IPNAME
eip-static vpc1 10.5.204.101 192.168.0.2 true vpc vpc-1-busybox01.vpc1

[root@pc-node-1 03-cust-vpc]# ping 10.5.204.101
PING 10.5.204.101 (10.5.204.101) 56(84) bytes of data.
64 bytes from 10.5.204.101: icmp_seq=2 ttl=62 time=1.21 ms
64 bytes from 10.5.204.101: icmp_seq=3 ttl=62 time=0.624 ms
64 bytes from 10.5.204.101: icmp_seq=4 ttl=62 time=0.368 ms
^C
--- 10.5.204.101 ping statistics ---
4 packets transmitted, 3 received, 25% packet loss, time 3049ms
rtt min/avg/max/mdev = 0.368/0.734/1.210/0.352 ms
[root@pc-node-1 03-cust-vpc]#

node ping VPC pod ip

ip
kubectl get pods --selector=vpc=vpc1
router 87ad6fd-71d5-4ff8-a1f0-54fa3bba1a7f (vpc1)
 port vpc1-vpc1-subnet1
 mac: "00:00:00:ED:8E:C7"
 networks: ["192.168.0.1/24"]
 port vpc1-external204
 mac: "00:00:00:EF:05:C7"
 networks: ["10.5.204.105/24"]
 gateway chassis: [7cedd14f-265b-42e5-ac17-e03e7a1f2342 276baccb-fe9c-4476-b41d-05872a94976d fd9f140c-c45d-43db-a6c0-0d4f8ea298dd]
 nat 813523e7-c68c-408f-bd8c-cba30cb2e4f4
 external ip: "10.5.204.101"
 logical ip: "192.168.0.2"
 type: "dnat_and_snat"

```

## 2.2 ovn-fip vip fip

```

vip kubevirt vip keepalived kube-vip

fip VPC vip vip

```

```

vip eip eip vip
cat vip.yaml

apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: test-fip-vip
spec:
 subnet: vpc1-subnet1

cat 04-fip.yaml

kind: OvnFip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-for-vip
spec:
 externalSubnet: external204
 type: nat

kind: OvnFip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-for-vip
spec:

```

```

ovnEip: eip-for-vip
ipType: vip # fip pod ip vip
ipName: test-fip-vip

VPC ip

kind: OvnFip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-for-vip
spec:
 ovnEip: eip-for-vip
 ipType: vip # fip pod ip vip
 vpc: vpc1
 v4Ip: 192.168.0.3

kubectl get ofip
NAME VPC V4EIP V4IP READY IPTYPE IPNAME
eip-for-vip vpc1 10.5.204.106 192.168.0.3 true vip test-fip-vip

[root@pc-node-1 fip-vip]# ping 10.5.204.106
PING 10.5.204.106 (10.5.204.106) 56(84) bytes of data.
64 bytes from 10.5.204.106: icmp_seq=1 ttl=62 time=0.694 ms
64 bytes from 10.5.204.106: icmp_seq=2 ttl=62 time=0.436 ms

node ping

pod ip

[root@pc-node-1 fip-vip]# kubectl -n vpc1 exec -it vpc-1-busybox03 -- bash
[root@vpc-1-busybox03 /]#
[root@vpc-1-busybox03 /]#
[root@vpc-1-busybox03 /]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::1/128 scope host
 valid_lft forever preferred_lft forever
1568: eth0@if1569: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
 link/ether 00:00:00:56:40:e5 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 192.168.0.5/24 brd 192.168.0.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet 192.168.0.3/24 scope global secondary eth0 #
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:fe56:40e5/64 scope link
 valid_lft forever preferred_lft forever

[root@vpc-1-busybox03 /]# tcpdump -i eth0 host 192.168.0.3 -netvv
tcpdump: listening on eth0, link-type EN10MB (Ethernet), capture size 262144 bytes
00:00:00:8e:c7 > 00:00:00:56:40:e5, ethertype IPv4 (0x0800), length 98: (tos 0x0, ttl 62, id 44830, offset 0, flags [DF], proto ICMP (1), length 84)
 10.5.32.51 > 192.168.0.3: ICMP echo request, id 177, seq 1, length 64
00:00:00:56:40:e5 > 00:00:00:8e:c7, ethertype IPv4 (0x0800), length 98: (tos 0x0, ttl 64, id 43962, offset 0, flags [none], proto ICMP (1), length 84)
 192.168.0.3 > 10.5.32.51: ICMP echo reply, id 177, seq 1, length 64

pod fip icmp

```

## 5.7.4 3. ovn-snats

### 3.1 ovn-snats subnet cidr

#### iptables-snats

```

cat 03-subnet-snats.yaml

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: snat-for-subnet-in-vpc
spec:
 externalSubnet: external204
 type: nat

kind: OvnNatRule
apiVersion: kubeovn.io/v1
metadata:
 name: snat-for-subnet-in-vpc
spec:
 ovnEip: snat-for-subnet-in-vpc
 vpcSubnet: vpc1-subnet1 # eip

VPC subnet cidr

```

```
kind: OvnSnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: snat-for-subnet-in-vpc
spec:
 ovnIip: snat-for-subnet-in-vpc
 vpc: vpc1
 v4IpCidr: 192.168.0.0/24 # cidr ip
```

externalSubnet extra

## 3.2 ovn-snat pod ip

## iptables-snat

```
cat 03-pod-snats.yaml

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: snat-for-pod-vpc-ip
spec:
 externalSubnet: external1204
 type: nat

kind: OvnSnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: snat01
spec:
 ovnEip: snat-for-pod-vpc-ip
 ipName: vpc-1-busybox02.vpc1 # eip pod ip

VPC ip

kind: OvnSnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: snat-for-subnet-in-vpc
spec:
 ovnEip: snat-for-subnet-in-vpc
 vpc: vpc1
 v4Ipcidr: 192.168.0.4
```

externalSubnet extra

snat

```
kubectl get vpc1
router 87ad06fd-71d5-4ff8-a1f0-54fa3bba1a7f (vpc1)
 port vpc1-vpc1-subnet1
 mac: "00:00:00:ED:8E:C7"
 networks: ["192.168.0.1/24"]
 port vpc1-external204
 mac: "00:00:00:EF:05:C7"
 networks: ["10.5.204.105/24"]
 gateway chassis: [7cddd14f-265b-42e5-ac17-e03e7a1f2342 276baccb-fe9c-4476-b41d-05872a94976d fd9f140c-c45d-43db-a6c0-0d4f8ea298dd]
nat 21d853b0-f7b4-40bd-9a53-31de2745739
 external ip: "10.5.204.115"
 logical ip: "192.168.0.0/24"
 type: "snat"
nat da77a11f-c523-439c-b1d1-72c664196a0f
 external ip: "10.5.204.116"
 logical ip: "192.168.0.4"
 type: "snat"
```

```
[root@pc-node-1 03-cust-vpc]# kubectl get po -A -o wide | grep busy
vpc1 vpc-1-busybox01 1/1 Running 0 3d15h 192.168.0.2 pc-node-2 <none> <none>
vpc1 vpc-1-busybox02 1/1 Running 0 17h 192.168.0.4 pc-node-1 <none> <none>
vpc1 vpc-1-busybox03 1/1 Running 0 17h 192.168.0.5 pc-node-1 <none> <none>
vpc1 vpc-1-busybox04 1/1 Running 0 17h 192.168.0.6 pc-node-3 <none> <none>
vpc1 vpc-1-busybox05 1/1 Running 0 17h 192.168.0.7 pc-node-1 <none> <none>

kubectl exec -it -n vpc1 vpc-1-busybox04 bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
[root@vpc-1-busybox04 ~]#
[root@vpc-1-busybox04 ~]#
[root@vpc-1-busybox04 ~]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::/128 scope host

```

```

valid_lft forever preferred_lft forever
17095: eth0@if17096: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
 link/ether 00:00:00:76:94:55 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 192.168.0.6/24 brd 192.168.0.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:fe76:9455/64 scope link
 valid_lft forever preferred_lft forever
[root@vpc-1-busybox04/]# ping 223.5.5.5
PING 223.5.5.5 (223.5.5.5) 56(84) bytes of data.
64 bytes from 223.5.5.5: icmp_seq=1 ttl=114 time=22.2 ms
64 bytes from 223.5.5.5: icmp_seq=2 ttl=114 time=21.8 ms

[root@pc-node-1 03-cust-vpc]# kubectl exec -it -n vpc1 vpc-1-busybox02 bash
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND] instead.
[root@vpc-1-busybox02/]#
[root@vpc-1-busybox02/]#
[root@vpc-1-busybox02/]#
[root@vpc-1-busybox02/]# ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::1/128 scope host
 valid_lft forever preferred_lft forever
1566: eth0@if1567: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
 link/ether 00:00:00:0b:e9:d0 brd ff:ff:ff:ff:ff:ff link-netnsid 0
 inet 192.168.0.4/24 brd 192.168.0.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:fe0b:e9d0/64 scope link
 valid_lft forever preferred_lft forever
[root@vpc-1-busybox02/]# ping 223.5.5.5
PING 223.5.5.5 (223.5.5.5) 56(84) bytes of data.
64 bytes from 223.5.5.5: icmp_seq=2 ttl=114 time=22.7 ms
64 bytes from 223.5.5.5: icmp_seq=3 ttl=114 time=22.6 ms
64 bytes from 223.5.5.5: icmp_seq=4 ttl=114 time=22.1 ms
^C
--- 223.5.5.5 ping statistics ---
4 packets transmitted, 3 received, 25% packet loss, time 3064ms
rtt min/avg/max/mdev = 22.126/22.518/22.741/0.278 ms

pod snat

```

## 5.7.5 4. ovn-dnat

### 4.1 ovn-dnat pod dnat

```

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: eip-dnat
spec:
 externalSubnet: underlay

kind: OvnDnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: eip-dnat
spec:
 ovnEip: eip-dnat
 ipName: vpc-1-busybox01.vpc1 # pod ip crd
 protocol: tcp
 internalPort: "22"
 externalPort: "22"

VPC ip

kind: OvnDnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: eip-dnat
spec:
 ovnEip: eip-dnat
 protocol: tcp
 internalPort: "22"
 externalPort: "22"
 vpc: vpc1
 v4Ip: 192.168.0.3

```

externalSubnet	extra
----------------	-------

OvnDnatRule      IptablesDnatRule

```

kubectl get oeip eip-dnat
NAME V4IP V6IP MAC TYPE READY
eip-dnat 10.5.49.4 00:00:00:4D:CE:49 dnat true

```

# kubectl get odnat	NAME	EIP	PROTOCOL	V4EIP	V4IP	INTERNALPORT	EXTERNALPORT	IPNAME	READY
	eip-dnat	eip-dnat	tcp	10.5.49.4	192.168.0.3	22	22	vpc-1-busybox01.vpc1	true

## 4.2 ovn-dnat vip dnat

```
kind: OvnDnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: eip-dnat
spec:
 ipType: vip # dnat pod ip vip
 ovnEip: eip-dnat
 ipName: test-dnat-vip
 protocol: tcp
 internalPort: "22"
 externalPort: "22"

VPC ip

kind: OvnDnatRule
apiVersion: kubeovn.io/v1
metadata:
 name: eip-dnat
spec:
 ipType: vip # dnat pod ip vip
 ovnEip: eip-dnat
 ipName: test-dnat-vip
 protocol: tcp
 internalPort: "22"
 externalPort: "22"
 vpc: vpc1
 v4Ip: 192.168.0.4
```

## OvnDnatRule IptablesDnatRule

```
kubectl get vip test-dnat-vip
NAME V4IP PV4IP MAC PMAC V6IP PV6IP SUBNET READY
test-dnat-vip 192.168.0.4 00:00:00:D0:C0:B5 vpc1-subnet1 true

kubectl get oeip eip-dnat
NAME V4IP V6IP MAC TYPE READY
eip-dnat 10.5.49.4 00:00:00:4D:CE:49 dnat true

kubectl get odnat eip-dnat
NAME EIP PROTOCOL V4EIP V4IP INTERNALPORT EXTERNALPORT IPNAME READY
eip-dnat eip-dnat tcp 10.5.49.4 192.168.0.4 22 22 test-dnat-vip true
```

[PDF](#)
[Slack](#)
[Support](#)

⌚2025 9 10

⌚2023 3 3



5.7.6

## 5.8 OVN SNAT ECMP BFD L3 HA

---

VPC OVN SNAT ECMP Gateway Node ovnnext0

- bfd
- hash

```
graph LR
pod --> vpc-subnet --> vpc --> snat --> ecmp --> external-subnet --> gw-node1-ovnnext0 --> node1-external-switch
external-subnet --> gw-node2-ovnnext0 --> node2-external-switch
external-subnet --> gw-node3-ovnnext0 --> node3-external-switch
```

[ovn-eip-fip-snat.md](#) install.sh provider-network vlan subnet

lsp ovn-eip vpc enable\_bfd bfd ecmp

### 5.8.1 1.

#### 1.1 underlay

#### 1.2 vpc eip\_snat

#### 1.3 VPC eip snat fip

[ovn-eip-fip-snat.md](#) VPC ecmp bfd

VPC 2 ovn-eip

```
cat gw-node-eip.yaml

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: pc-node-1
spec:
 externalSubnet: external204
 type: lsp

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: pc-node-2
spec:
 externalSubnet: external204
 type: lsp

kind: OvnEip
apiVersion: kubeovn.io/v1
metadata:
 name: pc-node-3
spec:
 externalSubnet: external204
 type: lsp
```

vpc ecmp vpc bfd enable\_bfd lrp ovn\_eip bfd

### 5.8.2 2. vpc ecmp bfd L3 HA

```
cat 01-vpc-ecmp-enable-external-bfd.yml
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc1
spec:
 namespaces:
 - vpc1
 enableExternal: true
 enableBfd: true # bfd
 #enableBfd: false
```

```
cat 02-subnet.yml
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: vpc1-subnet1
spec:
 cidrBlock: 192.168.0.0/24
 default: false
 disableGatewayCheck: false
 disableInterConnection: true
 enableEcmp: true # ecmp
 gatewayNode: ""
 gatewayType: distributed
 #gatewayType: centralized
 natOutgoing: false
 private: false
 protocol: IPv4
 provider: ovn
 vpc: vpc1
 namespaces:
 - vpc1
```

:

1. vpc ecmp ecmp bfd vpc enableBfd subnet enableEcmp ecmp bfd
- 2.
3. VPC VPC VPC snat
4. vpc subnet enableEcmp gatewayType
5. EnableExternal vpc
6. EnableExternal EnableBfd

```
ovn
vpc
k get vpc
NAME ENABLEEXTERNAL ENABLEBFD STANDBY SUBNETS NAMESPACES
ovn-cluster true true true ["external204","join","ovn-default"]
vpc1 true true true ["vpc1-subnet1"] ["vpc1"]

vpc ENABLEBFD
vpc

1. bfd
k ko nbctl list bfd
_uuid : be7df545-2c4c-4751-878f-b3507987f050
detect_mult : 3
dst_ip : "10.5.204.121"
external_ids : {}
logical_port : vpc1-external204
min_rx : 100
min_tx : 100
options : {}
status : up

_uuid : 684c4489-5b59-4693-8d8c-3beab93f8093
detect_mult : 3
dst_ip : "10.5.204.109"
external_ids : {}
logical_port : vpc1-external204
min_rx : 100
min_tx : 100
options : {}
status : up

_uuid : f0f62077-2ae9-4e79-b4f8-a446ec6e784c
detect_mult : 3
dst_ip : "10.5.204.108"
external_ids : {}
logical_port : vpc1-external204
min_rx : 100
min_tx : 100
options : {}
status : up

status up

2. bfd
k ko nbctl lr-route-list vpc1
IPv4 Routes
Route Table <main>:
 192.168.0.0/24 10.5.204.108 src-ip ecmp ecmp-symmetric-reply bfd
 192.168.0.0/24 10.5.204.109 src-ip ecmp ecmp-symmetric-reply bfd
 192.168.0.0/24 10.5.204.121 src-ip ecmp ecmp-symmetric-reply bfd
```

```

3.

kubectl find Logical_Router_Static_Route policy=src-ip options=ecmp_symmetric_reply="true"
_uuid : 3aacb384-d5ee-4b14-aebf-59e8c11717ba
bfd : 684c4489-5b59-4693-8d8c-3beab93f8093
external_ids : {}
ip_prefix : "192.168.0.0/24"
nexthop : "10.5.204.109"
options : {ecmp_symmetric_reply="true"}
output_port : []
policy : src-ip
route_table : ""

_uuid : 18bcc585-bc05-430b-925b-ef673c8e1aef
bfd : f0f62077-2ae9-4e79-b4f8-a446ec6e784c
external_ids : {}
ip_prefix : "192.168.0.0/24"
nexthop : "10.5.204.108"
options : {ecmp_symmetric_reply="true"}
output_port : []
policy : src-ip
route_table : ""

_uuid : 7d0a4e6b-cde0-4110-8176-fbaf19738498
bfd : be7df545-2c4c-4751-878f-b3507987f050
external_ids : {}
ip_prefix : "192.168.0.0/24"
nexthop : "10.5.204.121"
options : {ecmp_symmetric_reply="true"}
output_port : []
policy : src-ip
route_table : ""

[root@pc-node-1 ~]# ip netns exec ovnnext bash ip a
/usr/sbin/ip: /usr/sbin/ip: cannot execute binary file
[root@pc-node-1 ~]#
[root@pc-node-1 ~]# ip netns exec ovnnext ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default qlen 1000
 link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
 inet 127.0.0.1/8 scope host lo
 valid_lft forever preferred_lft forever
 inet6 ::1/128 scope host
 valid_lft forever preferred_lft forever
1541: ovnnext0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1400 qdisc noqueue state UNKNOWN group default qlen 1000
 link/ether 00:00:00:ab:bd:87 brd ff:ff:ff:ff:ff:ff
 inet 10.5.204.108/24 brd 10.5.204.255 scope global ovnnext0
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:feab:bd87/64 scope link
 valid_lft forever preferred_lft forever
[root@pc-node-1 ~]#
[root@pc-node-1 ~]# ip netns exec ovnnext route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.5.204.254 0.0.0.0 UG 0 0 0 ovnnext0
10.5.204.0 0.0.0.0 255.255.255.0 U 0 0 0 ovnnext0

internal port unerlay pod ns

[root@pc-node-1 ~]# ip netns exec ovnnext bfdd-control status
There are 1 sessions:
Session 1
 id=1 local=10.5.204.108 (p) remote=10.5.204.122 state=Up

lrp bfd lrp ecmp

[root@pc-node-1 ~]# ip netns exec ovnnext ping -c1 223.5.5.5
PING 223.5.5.5 (223.5.5.5) 56(84) bytes of data.
64 bytes from 223.5.5.5: icmp_seq=1 ttl=115 time=21.6 ms

#

```

### ovnnext ns

```

tcpdump -i ovnnext0 host 223.5.5.5 -netvv
dropped privs to tcpdump
tcpdump: listening on ovnnext0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C
0 packets captured
0 packets received by filter
0 packets dropped by kernel
[root@pc-node-1 ~]# exit
[root@pc-node-1 ~]# ssh pc-node-2
Last login: Thu Feb 23 09:21:08 2023 from 10.5.32.51
[root@pc-node-2 ~]# ip netns exec ovnnext bash
[root@pc-node-2 ~]# tcpdump -i ovnnext0 host 223.5.5.5 -netvv
dropped privs to tcpdump
tcpdump: listening on ovnnext0, link-type EN10MB (Ethernet), capture size 262144 bytes
^C
0 packets captured
0 packets received by filter

```

```

0 packets dropped by kernel
[root@pc-node-2 ~]# exit
[root@pc-node-2 ~]# logout
Connection to pc-node-2 closed.
[root@pc-node-1 ~]# ssh pc-node-3
Last login: Thu Feb 23 08:32:41 2023 from 10.5.32.51
[root@pc-node-3 ~]# ip netns exec ovnnext bash
[root@pc-node-3 ~]# tcpdump -i ovnnext0 host 223.5.5.5 -netvv
dropped privs to tcpdump
tcpdump: listening on ovnnext0, link-type EN10MB (Ethernet), capture size 262144 bytes
00:00:00:2d:f8:ce > 00:00:00:fd:b2:a4, ethertype IPv4 (0x0800), length 98: (tos 0x0, ttl 63, id 57978, offset 0, flags [DF], proto ICMP (1), length 84)
 10.5.204.102 > 223.5.5.5: ICMP echo request, id 22, seq 71, length 64
00:00:00:fd:b2:a4 > dc:ef:80:5a:44:1a, ethertype IPv4 (0x0800), length 98: (tos 0x0, ttl 62, id 57978, offset 0, flags [DF], proto ICMP (1), length 84)
 10.5.204.102 > 223.5.5.5: ICMP echo request, id 22, seq 71, length 64
^C
2 packets captured
2 packets received by filter
0 packets dropped by kernel
[root@pc-node-3 ~]#
down pod
3

```

### 5.8.3 3. bfd

#### vpc enable\_eip\_snat

```

cat 01-vpc-ecmp-enable-external-bfd.yml
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc2
spec:
 namespaces:
 - vpc2
 enableExternal: true
 #enableBfd: true
 enableBfd: false

bfd

k ko nbctl lr-route-list vpc2
IPv4 Routes
Route Table <main>:
 0.0.0.0/0 10.5.204.254 dst-ip

#
nbctl list bfd lrp bfd
ovnnext ns bfd
vpc subnet ping ()
()

```



⌚2025 9 10

⌚2023 3 3



### 5.8.4

## 5.9 VPC

---

VPC	VPC	VPC	NAT
-----	-----	-----	-----

### 5.9.1

1. VPC
2. VPC CIDR
3. VPC VPC

### 5.9.2

VPC	VPC	Subnet	CIDR
-----	-----	--------	------

```
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc-1
spec: {}

kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: net1
spec:
 vpc: vpc-1
 cidrBlock: 10.0.0.0/16

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc-2
spec: {}

kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: net2
spec:
 vpc: vpc-2
 cidrBlock: 172.31.0.0/16
```

VPC	vpcPeerings		
-----	-------------	--	--

```
kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc-1
spec:
 vpcPeerings:
 - remoteVpc: vpc-2
 localConnectIP: 169.254.0.1/30
 staticRoutes:
 - cidr: 172.31.0.0/16
 nextHopIP: 169.254.0.2
 policy: policyDst

kind: Vpc
apiVersion: kubeovn.io/v1
metadata:
 name: vpc-2
spec:
 vpcPeerings:
 - remoteVpc: vpc-1
 localConnectIP: 169.254.0.2/30
 staticRoutes:
 - cidr: 10.0.0.0/16
 nextHopIP: 169.254.0.1
 policy: policyDst
```

- remoteVpc: VPC
- localConnectIP: IP CIDR IP CIDR
- cidr Subnet CIDR
- nextHopIP VPC localConnectIP

## Subnet Pod

```
apiVersion: v1
kind: Pod
metadata:
 annotations:
 ovn.kubernetes.io/logical_switch: net1
 name: vpc-1-pod
spec:
 containers:
 - name: vpc-1-pod
 image: docker.io/library/nginx:alpine

apiVersion: v1
kind: Pod
metadata:
 annotations:
 ovn.kubernetes.io/logical_switch: net2
 name: vpc-2-pod
spec:
 containers:
 - name: vpc-2-pod
 image: docker.io/library/nginx:alpine
```

```
kubectl exec -it vpc-1-pod -- ping $(kubectl get pod vpc-2-pod -o jsonpath='{.status.podIP}')
PING 172.31.0.2 (172.31.0.2): 56 data bytes
64 bytes from 172.31.0.2: seq=0 ttl=62 time=0.655 ms
64 bytes from 172.31.0.2: seq=1 ttl=62 time=0.886 ms
64 bytes from 172.31.0.2: seq=2 ttl=62 time=0.098 ms
^C
--- 172.31.0.2 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.086/0.279/0.655 ms
kubectl exec -it vpc-2-pod -- ping $(kubectl get pod vpc-1-pod -o jsonpath='{.status.podIP}')
PING 10.0.0.2 (10.0.0.2): 56 data bytes
64 bytes from 10.0.0.2: seq=0 ttl=62 time=0.594 ms
64 bytes from 10.0.0.2: seq=1 ttl=62 time=0.093 ms
64 bytes from 10.0.0.2: seq=2 ttl=62 time=0.088 ms
^C
--- 10.0.0.2 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.088/0.258/0.594 ms
```

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 7 2

⌚2022 5 24



5.9.3

## 6.

---

### 6.1 kubectl

Kube-OVN    kubectl    OVN    OVN    OVS    tcpdump

#### 6.1.1

Kube-OVN    kubectl

kubectl-ko

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/kubectl-ko
```

\$PATH

```
mv kubectl-ko /usr/local/bin/kubectl-ko
```

```
chmod +x /usr/local/bin/kubectl-ko
```

```
kubectl plugin list
The following compatible plugins are available:
/usr/local/bin/kubectl-ko
```

#### 6.1.2

kubectl ko

```
kubectl ko
kubectl ko {subcommand} [option...]
Available Subcommands:
[nb|sb] [status|kick|backup|dbstatus|restore] ovn-db operations show cluster status, kick stale server, backup database, get db consistency status or
restore ovn nb db when met 'inconsistent data' error
nbctl [ovn-nbctl options ...] invoke ovn-nbctl
sbctl [ovn-sbctl options ...] invoke ovn-sbctl
vsctl {node_name} [ovs-vsctl options ...] invoke ovs-vsctl on the specified node
ofctl {node_name} [ovs-ofctl options ...] invoke ovs-ofctl on the specified node
dpctl {node_name} [ovs-dpctl options ...] invoke ovs-dpctl on the specified node
appctl {node_name} [ovs-appctl options ...] invoke ovs-appctl on the specified node
tcpdump {namespace/podname} [tcpdump options ...] capture pod traffic
{trace|ovn-trace} ... trace ovn microflow of specific packet
 {trace|ovn-trace} {namespace/podname} {target ip address} [target mac address] {icmp|tcp|udp} [target tcp/udp port] trace ICMP/TCP/UDP
 {trace|ovn-trace} {namespace/podname} {target ip address} [target mac address] arp {request|reply} trace ARP request/reply
 {trace|ovn-trace} {node//node_name} {target ip address} [target mac address] {icmp|tcp|udp} [target tcp/udp port] trace ICMP/TCP/UDP
 {trace|ovn-trace} {node//node_name} {target ip address} [target mac address] arp {request|reply} trace ARP request/reply
diagnose {all|node|subnet|IPPorts} [{node_name|subnet_name}|{proto1}-{IP1}-{Port1},{proto2}-{IP2}-{Port2}] diagnose connectivity of all nodes or a specific
node or specify subnet's ds pods or IPPorts like 'tcp-172.18.0.2-53,udp-172.18.0.3-53'
env-check check the environment configuration
reload restart all kube-ovn components
log {kube-ovn|ovn|ovs|linux|all} save log to ./kubectl-ko-log/
perf [image] performance test default image is docker.io/kubeovn/test:v1.13.0
icnbctl [ovn-nbctl options ...] invoke ovn-ic-nbctl
icsbctl [ovn-sbctl options ...] invoke ovn-ic-sbctl
```

**[nb | sb] [status | kick | backup | dbstatus | restore]**

OVN

OVN leader ovs-appctl cluster/status :

```
kubectl ko nb status
306b
Name: OVN_Northbound
Cluster ID: 9a87 (9a872522-3e7d-47ca-83a3-d74333e1a7ca)
Server ID: 306b (306b256b-b5e1-4eb0-be91-4ca96adf6bad)
Address: tcp:[172.18.0.2]:6643
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 280309 ms ago, reason: timeout
Last Election won: 280309 ms ago
Election timer: 5000
Log: [139, 139]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-8723 ->8723 <-85d6 ->85d6
Disconnections: 0
Servers:
 85d6 (85d6 at tcp:[172.18.0.4]:6643) next_index=139 match_index=138 last msg 763 ms ago
 8723 (8723 at tcp:[172.18.0.3]:6643) next_index=139 match_index=138 last msg 763 ms ago
 306b (306b at tcp:[172.18.0.2]:6643) (self) next_index=2 match_index=138
status: ok
```

Server match\_index last msg Server

OVN 172.18.0.3 :

```
kubectl ko nb kick 8723
started removal
```

```
kubectl ko nb status
306b
Name: OVN_Northbound
Cluster ID: 9a87 (9a872522-3e7d-47ca-83a3-d74333e1a7ca)
Server ID: 306b (306b256b-b5e1-4eb0-be91-4ca96adf6bad)
Address: tcp:[172.18.0.2]:6643
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 324356 ms ago, reason: timeout
Last Election won: 324356 ms ago
Election timer: 5000
Log: [140, 140]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-85d6 ->85d6
Disconnections: 2
Servers:
 85d6 (85d6 at tcp:[172.18.0.4]:6643) next_index=140 match_index=139 last msg 848 ms ago
 306b (306b at tcp:[172.18.0.2]:6643) (self) next_index=2 match_index=139
status: ok
```

OVN

```
kubectl ko nb backup
tar: Removing leading '/' from member names
backup ovn-nb db to /root/ovnnb_db.060223191654183154.backup
```

```
kubectl ko nb dbstatus
status: ok
```

inconsistent data

### inconsistent data

```
kubectl ko nb restore
deployment.apps/ovn-central scaled
ovn-central original replicas is 3
first nodeIP is 172.18.0.5
ovs-ovn pod on node 172.18.0.5 is ovs-ovn-8jxv9
ovs-ovn pod on node 172.18.0.3 is ovs-ovn-sjzb6
ovs-ovn pod on node 172.18.0.4 is ovs-ovn-t87zk
backup nb db file
restore nb db file, operate in pod ovs-ovn-8jxv9
deployment.apps/ovn-central scaled
finish restore nb db file and ovn-central replicas
recreate ovs-ovn pods
pod "ovs-ovn-8jxv9" deleted
pod "ovs-ovn-sjzb6" deleted
pod "ovs-ovn-t87zk" deleted
```

### [nbctl | sbctl] [options ...]

OVN

leader

ovn-nbctl ovn-sbctl

OVN

ovn-nbctl(8) ovn-sbctl(8)

```
kubectl ko nbctl show
switch c7cd17e8-ceee-4a91-9bb3-e5a313fe1ece (snat)
 port snat-ovn-cluster
 type: router
 router-port: ovn-cluster-snat
switch 20e0c6d0-023a-4756-aec5-200e0c60f95d (join)
 port node-liumengxin-ovn3-192.168.137.178
 addresses: ["00:00:00:64:FF:A8 100.64.0.4"]
 port node-liumengxin-ovn1-192.168.137.176
 addresses: ["00:00:00:AF:98:62 100.64.0.2"]
 port node-liumengxin-ovn2-192.168.137.177
 addresses: ["00:00:00:D9:58:B8 100.64.0.3"]
 port join-ovn-cluster
 type: router
 router-port: ovn-cluster-join
switch 0191705c-f827-427b-9de3-3c3b7d971ba5 (central)
 port central-ovn-cluster
 type: router
 router-port: ovn-cluster-central
switch 2a45ff05-388d-4f85-9daf-e6fcfd5833dc (ovn-default)
 port alertmanager-main-0.monitoring
 addresses: ["00:00:00:6C:DF:A3 10.16.0.19"]
 port kube-state-metrics-5d6885d89-4nf8h.monitoring
 addresses: ["00:00:00:6F:02:1C 10.16.0.15"]
 port fake-kubelet-67c55dfdf89-pv86k.kube-system
 addresses: ["00:00:00:5C:12:E8 10.16.19.177"]
 port ovn-default-ovn-cluster
 type: router
 router-port: ovn-cluster-ovn-default
router 212f73dd-d63d-4d72-864b-a537e9afbee1 (ovn-cluster)
 port ovn-cluster-snat
 mac: "00:00:00:7A:82:8F"
 networks: ["172.22.0.1/16"]
 port ovn-cluster-join
 mac: "00:00:00:F8:18:5A"
 networks: ["100.64.0.1/16"]
 port ovn-cluster-central
 mac: "00:00:00:4D:8C:F5"
 networks: ["192.168.0.1/16"]
 port ovn-cluster-ovn-default
 mac: "00:00:00:A3:F8:18"
 networks: ["10.16.0.1/16"]
```

### vsctl {nodeName} [options ...]

nodeName ovs-ovn

ovs-vsctl

vswitchd

OVS

ovs-vsctl(8)

```
kubectl ko vsctl kube-ovn-01 show
0d4c4675-c9cc-440a-8c1a-878e17f81b88
Bridge br-int
 fail_mode: secure
 datapath_type: system
 Port a2c1a8a8b83a_h
 Interface a2c1a8a8b83a_h
 Port "4fa5c4cbb1a5_h"
 Interface "4fa5c4cbb1a5_h"
 Port ovn-eef07d-0
 Interface ovn-eef07d-0
 type: stt
 options: {csum="true", key=flow, remote_ip="192.168.137.178"}
 Port ovn0
```

```

Interface ovn0
 type: internal
Port mirror0
 Interface mirror0
 type: internal
Port ovn-efa253-0
 Interface ovn-efa253-0
 type: stt
 options: {csum="true", key=flow, remote_ip="192.168.137.177"}
Port br-int
 Interface br-int
 type: internal
ovs_version: "2.17.2"

```

**ofctl {nodeName} [options ...]**

nodeName	ovs-ovn	ovs-ofctl	OpenFlow	OVS	ovs-ofctl(8)
----------	---------	-----------	----------	-----	--------------

```

kubectl ko ofctl kube-ovn-01 dump-flows br-int
NXST_FLOW reply (xid=0x4): flags=[more]
cookie=0xcf3429e6, duration=671791.432s, table=0, n_packets=0, n_bytes=0, idle_age=65534, hard_age=65534, priority=100, in_port=2 actions=load:0x4->NXM_NX_REG13[], load:0x9->NXM_NX_REG11[], load:0xb->NXM_NX_REG12[], load:0x4->0XM_OF_METADATA[], load:0x1->NXM_NX_REG14[], resubmit(.8)
cookie=0xc91413c6, duration=671791.431s, table=0, n_packets=9978275, n_bytes=99978275, idle_age=0, hard_age=65534, priority=100, in_port=7 actions=load:0x1->NXM_NX_REG13[], load:0x9->NXM_NX_REG11[], load:0xb->NXM_NX_REG12[], load:0x4->0XM_OF_METADATA[], load:0x4->NXM_NX_REG14[], resubmit(.8)
cookie=0xf180459, duration=671791.431s, table=0, n_packets=17348582, n_bytes=2667811214, idle_age=0, hard_age=65534, priority=100, in_port=6317 actions=load:0xa->NXM_NX_REG13[], load:0x9->NXM_NX_REG11[], load:0xb->NXM_NX_REG12[], load:0x4->0XM_OF_METADATA[], load:0x9->NXM_NX_REG14[], resubmit(.8)
cookie=0x7806dd90, duration=671791.431s, table=0, n_packets=3235428, n_bytes=833821312, idle_age=0, hard_age=65534, priority=100, in_port=1 actions=load:0xd->NXM_NX_REG13[], load:0x9->NXM_NX_REG11[], load:0xb->NXM_NX_REG12[], load:0x4->0XM_OF_METADATA[], load:0x3->NXM_NX_REG14[], resubmit(.8)
...

```

**dpctl {nodeName} [options ...]**

nodeName	ovs-ovn	ovs-dpctl	OVS datapath	OVS	ovs-dpctl(8)
----------	---------	-----------	--------------	-----	--------------

```

kubectl ko dpctl kube-ovn-01 show
system@ovs-system:
lookups: hit:35080505 missed:21983648 lost:7:
flows: 105
masks: hit:1970748791 total:22 hit/pkt:5.29
port 0: ovs-system (internal)
port 1: ovn0 (internal)
port 2: mirror0 (internal)
port 3: br-int (internal)
port 4: stt_sys_7471 (stt: packet_type=ptap)
port 5: eeb4d9e51b5d_h
port 6: a2c1a8a8b83a_h
port 7: 4fa5c4cbb1a5_h

```

**appctl {nodeName} [options ...]**

nodeName	ovs-ovn	ovs-appctl	daemon	OVS	ovs-appctl(8)
----------	---------	------------	--------	-----	---------------

```

kubectl ko appctl kube-ovn-01 vlog/list
 console syslog file
 ----- -----
backtrace OFF ERR INFO
bfd OFF ERR INFO
bond OFF ERR INFO
bridge OFF ERR INFO
bundle OFF ERR INFO
bundles OFF ERR INFO
...

```

**tcpdump {namespace/podname} [tcpdump options ...]**

namespace/podname	kube-ovn-cni	tcpdump	veth
-------------------	--------------	---------	------

```

kubectl ko tcpdump default/ds1-l6n7p icmp
+ kubectl exec -it kube-ovn-cni-wlg4s -n kube-ovn -- tcpdump -nn -i d7176fe7b4e0_h icmp
tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
listening on d7176fe7b4e0_h, link-type EN10MB (Ethernet), capture size 262144 bytes
06:52:36.619688 IP 10.64.0.3 > 10.16.0.4: ICMP echo request, id 2, seq 1, length 64
06:52:36.619746 IP 10.16.0.4 > 100.64.0.3: ICMP echo reply, id 2, seq 1, length 64
06:52:37.619588 IP 100.64.0.3 > 10.16.0.4: ICMP echo request, id 2, seq 2, length 64
06:52:37.619630 IP 10.16.0.4 > 100.64.0.3: ICMP echo reply, id 2, seq 2, length 64
06:52:38.619933 IP 100.64.0.3 > 10.16.0.4: ICMP echo request, id 2, seq 3, length 64
06:52:38.619973 IP 10.16.0.4 > 100.64.0.3: ICMP echo reply, id 2, seq 3, length 64

```

## trace [arguments ...]

Pod	OVN	Openflow
-----	-----	----------

```
kubectl ko trace {namespace/podname} {target ip address} [target mac address] {icmp|tcp|udp} [target tcp/udp port]
kubectl ko trace {namespace/podname} {target ip address} [target mac address] arp {request|reply}
kubectl ko trace {node/nodename} {target ip address} [target mac address] {icmp|tcp|udp} [target tcp/udp port]
kubectl ko trace {node/nodename} {target ip address} [target mac address] arp {request|reply}
```

```
kubectl ko trace default/ds1-16n7p 8.8.8.8 icmp
+ kubectl exec ovn-central-5bc494cb5-n kube-ovn -- ovn-trace --ct=new ovn-default 'inport == "ds1-16n7p.default" && ip.ttl == 64 && icmp && eth.src == 0a:00:00:10:00:05 && ip4.src == 10.16.0.4 && eth.dst == 00:00:00:B8:CA:43 && ip4.dst == 8.8.8.8'
icmp,reg14=0xf,vlan_tci=0x0000,d1_src=0a:00:00:10:00:05,d1_dst=00:00:00:B8:CA:43,nw_src=10.16.0.4,nw_dst=8.8.8.8,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=0,icmp_code=0

ingress(dp="ovn-default", inport="ds1-16n7p.default")

0. ls_in_port_sec_l2 (ovn-northd.c:4143): inport == "ds1-16n7p.default" && eth.src == {0a:00:00:10:00:05}, priority 50, uuid 39453393
 next;
1. ls_in_port_sec_ip (ovn-northd.c:2898): inport == "ds1-16n7p.default" && eth.src == 0a:00:00:10:00:05 && ip4.src == {10.16.0.4}, priority 90, uuid 81bcd485
 next;
3. ls_in_pre_acl (ovn-northd.c:3269): ip, priority 100, uuid 7b4f4971
 reg0[0] = 1;
 next;
5. ls_in_pre_stateful (ovn-northd.c:3396): reg0[0] == 1, priority 100, uuid 36cd577
 ct_next;

ct_next(ct_state=new|trk)

6. ls_in_acl (ovn-northd.c:3759): ip && (!ct.est || (ct.est && ct_label.blocked == 1)), priority 1, uuid 7608af5b
 reg0[1] = 1;
 next;
10. ls_in_stateful (ovn-northd.c:3995): reg0[1] == 1, priority 100, uuid 2aba1b90
 ct_commit(ct_label=0x01);
 next;
16. ls_in_l2_1kup (ovn-northd.c:4470): eth.dst == 00:00:00:B8:CA:43, priority 50, uuid 5c9c3c9f
 output = "ovn-default-ovn-cluster";
 output;
...
...
```

trace	Underlay	Mac
-------	----------	-----

```
kubectl ko trace default/virt-handler-7lvm1 8.8.8.8 82:7c:9f:83:8c:01 icmp
```

## diagnose {all|node|subnet|IPPorts} [nodename|subnetName|{proto1}-{IP1}-{Port1},{proto2}-{IP2}-{Port2}]

kube-ovn-pinger

```
kubectl ko diagnose all
switch c7cd17e8-ceee-4a91-9bb3-e5a313fe1ece (snat)
 port snat-ovn-cluster
 type: router
 router-port: ovn-cluster-snat
switch 20e0c6d0-023a-4756-aec5-200e0c60f95d (join)
 port node-liumengxin-ovn3-192.168.137.178
 addresses: ["00:00:00:64:FF:A8 100.64.0.4"]
 port node-liumengxin-ovn1-192.168.137.176
 addresses: ["00:00:00:AF:98:62 100.64.0.2"]
 port join-ovn-cluster
 type: router
 router-port: ovn-cluster-join
switch 0191705c-f827-427b-9de3-3c3b7d971ba5 (central)
 port central-ovn-cluster
 type: router
 router-port: ovn-cluster-central
switch 2a45ff05-388d-4f85-9daf-e6fccd5833dc (ovn-default)
 port ovn-default-ovn-cluster
 type: router
 router-port: ovn-cluster-ovn-default
 port prometheus-k8s-1.monitoring
 addresses: ["00:00:00:AA:37:DF 10.16.0.23"]
router 212f73dd-d63d-4d72-864b-a537e9afbee1 (ovn-cluster)
 port ovn-cluster-snats
 mac: "00:00:00:7A:82:8F"
 networks: ["172.22.0.1/16"]
 port ovn-cluster-join
 mac: "00:00:00:F8:18:5A"
 networks: ["100.64.0.1/16"]
 port ovn-cluster-central
```

```

mac: "00:00:00:4D:8C:F5"
networks: ["192.101.0.1/16"]
port ovn-cluster-ovn-default
 mac: "00:00:00:A3:F8:18"
 networks: ["10.16.0.1/16"]

Routing Policies
 31000 ip4.dst == 10.16.0.0/16 allow
 31000 ip4.dst == 100.64.0.0/16 allow
 30000 ip4.dst == 192.168.137.177 reroute 100.64.0.3
 30000 ip4.dst == 192.168.137.178 reroute 100.64.0.4
 29000 ip4.src == $ovn.default.fake.6_ip4 reroute 100.64.0.22
 29000 ip4.src == $ovn.default.fake.7_ip4 reroute 100.64.0.21
 29000 ip4.src == $ovn.default.fake.8_ip4 reroute 100.64.0.23
 29000 ip4.src == $ovn.default.liumengxin.ovn3.192.168.137.178_ip4 reroute 100.64.0.4
 20000 ip4.src == $ovn.default.liumengxin.ovn1.192.168.137.176_ip4 & ip4.dst != $ovn.cluster.overlay.subnets.IPv4 reroute 100.
64.0.2
 20000 ip4.src == $ovn.default.liumengxin.ovn2.192.168.137.177_ip4 & ip4.dst != $ovn.cluster.overlay.subnets.IPv4 reroute 100.
64.0.3
 20000 ip4.src == $ovn.default.liumengxin.ovn3.192.168.137.178_ip4 & ip4.dst != $ovn.cluster.overlay.subnets.IPv4 reroute 100.
64.0.4
IPv4 Routes
Route Table <main>:
 0.0.0.0/0 100.64.0.1 dst-ip
UUID LB PROTO VIP IPs
e9bcfd9d-793e-4431-9073-6dec96b75d71 cluster-tcp-load tcp 10.100.209.132:10660 192.168.137.176:10660
 tcp 10.101.239.192:6641 192.168.137.177:6641
 tcp 10.101.240.101:3000 10.16.0.7:3000
 tcp 10.103.184.186:6642 192.168.137.177:6642
35d2b7a5-e3a7-485a-a4b7-b4970eb0e63b cluster-tcp-sess tcp 10.100.158.128:8080 10.16.0.10:8080,10.16.0.5:8080,10.16.63.30:8080
 tcp 10.107.26.215:8080 10.16.0.19:8080,10.16.0.20:8080,10.16.0.21:8080
 tcp 10.107.26.215:9093 10.16.0.19:9093,10.16.0.20:9093,10.16.0.21:9093
 tcp 10.98.187.99:8080 10.16.0.22:8080,10.16.0.23:8080
 tcp 10.98.187.99:9090 10.16.0.22:9090,10.16.0.23:9090
f43303e4-89aa-4d3e-a3dc-278a552fe27b cluster-udp-load udp 10.96.0.10:53 10.16.0.4:53,10.16.0.9:53
_uuid : 06776304-5a96-43ed-90c4-c4854c251699
addresses : []
external_ids : {vendor=kube-ovn}
name : node_liumengxin_ovn2_192.168.137.177_underlay_v6

_uuid : 62690625-87d5-491c-8675-9fd83b1f433c
addresses : []
external_ids : {vendor=kube-ovn}
name : node_liumengxin_ovn1_192.168.137.176_underlay_v6

_uuid : b03a9bae-94d5-4562-b34c-b5f6198e180b
addresses : ["10.16.0.0/16", "100.64.0.0/16", "172.22.0.0/16", "192.101.0.0/16"]
external_ids : {vendor=kube-ovn}
name : ovn.cluster.overlay.subnets.IPv4

_uuid : e1056f3a-24cc-4666-8a91-75ee6c3c2426
addresses : []
external_ids : {vendor=kube-ovn}
name : ovn.cluster.overlay.subnets.IPv6

_uuid : 3e5dffff-e670-47b2-a2f5-a39f4698a8c5
addresses : []
external_ids : {vendor=kube-ovn}
name : node_liumengxin_ovn3_192.168.137.178_underlay_v6
_uuid : 2d85dbdc-d0db-4abe-b19e-cc886d32b492
action : drop
direction : from-lport
external_ids : {}
label : 0
log : false
match : "inport==@ovn.sg.kubeovn_deny_all && ip"
meter : []
name : []
options : {}
priority : 2003
severity : []

_uuid : de790cc8-f155-405f-bb32-5a51f30c545f
action : drop
direction : to-lport
external_ids : {}
label : 0
log : false
match : "outport==@ovn.sg.kubeovn_deny_all && ip"
meter : []
name : []
options : {}
priority : 2003
severity : []

Chassis "e15ed4d4-1780-4d50-b09e-ea8372ed48b8"
 hostname: liumengxin-ovn1-192.168.137.176
 Encap stt
 ip: "192.168.137.176"
 options: {csum="true"}
 Port_Binding node-liumengxin-ovn1-192.168.137.176
 Port_Binding perf-6vxkn.default
 Port_Binding kube-state-metrics-5d6885d89-4nf8h.monitoring
 Port_Binding alertmanager-main-0.monitoring
 Port_Binding kube-ovn-pinger-6ftdf.kube-system
 Port_Binding fake-kubelet-67c55dfd89-pv86k.kube-system

```

```

Port_Binding prometheus-k8s-0.monitoring
Chassis "eef07da1-f8ad-4775-b14d-bd6a3b4eb0d5"
 hostname: liumengxin-ovn3-192.168.137.178
 Encap stt
 ip: "192.168.137.178"
 options: {csum="true"}
Port_Binding kube-ovn-pinger-7twb4.kube-system
Port_Binding prometheus-adapter-86df476d87-r188g.monitoring
Port_Binding prometheus-k8s-1.monitoring
Port_Binding node-liumengxin-ovn3-192.168.137.178
Port_Binding perf-ff475.default
Port_Binding alertmanager-main-1.monitoring
Port_Binding blackbox-exporter-676d976865-tvsjd.monitoring
Chassis "efa253c9-494d-4719-83ae-b48ab0f11c03"
 hostname: liumengxin-ovn2-192.168.137.177
 Encap stt
 ip: "192.168.137.177"
 options: {csum="true"}
Port_Binding grafana-6c4c6b8fb7-pzd2c.monitoring
Port_Binding node-liumengxin-ovn2-192.168.137.177
Port_Binding alertmanager-main-2.monitoring
Port_Binding coredns-6789c94dd8-9jqsz.kube-system
Port_Binding coredns-6789c94dd8-25d4r.kube-system
Port_Binding prometheus-operator-7bbc99fc8b-wgjm4.monitoring
Port_Binding prometheus-adapter-86df476d87-gdxmc.monitoring
Port_Binding perf-fjnws.default
Port_Binding kube-ovn-pinger-vh2xg.kube-system
ds kube-proxy ready
kube-proxy ready
deployment ovn-central ready
deployment kube-ovn-controller ready
ds kube-ovn-cni ready
ds ovs-ovn ready
deployment coredns ready
ovn-nb leader check ok
ovn-sb leader check ok
ovn-northd leader check ok
kube-ovn-controller recent log

start to diagnose node liumengxin-ovn1-192.168.137.176
ovn-controller log:
2022-06-03T00:56:44.897Z|16722|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:06:44.912Z|16723|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:16:44.925Z|16724|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:26:44.936Z|16725|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:36:44.959Z|16726|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:46:44.974Z|16727|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T01:56:44.988Z|16728|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T02:06:45.001Z|16729|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T02:16:45.025Z|16730|inc_proc_eng|INFO|User triggered force recompute.
2022-06-03T02:26:45.040Z|16731|inc_proc_eng|INFO|User triggered force recompute.

ovs-vsctl log:
2022-06-02T23:03:00.137Z|00079|dpif(handler)|WARN|system@ovs-system: execute ct(commit,zone=14,label=0/0x1,nat(src)),8 failed (Invalid argument) on packet
 icmp,vlan_tci=0x0000,dl_src=00:00:00:f8:07:c8,dl_dst=00:00:00:fa:1e:50,nw_src=10.16.0.5,nw_dst=10.16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0
 icmp_csum=f0d1
 with metadata skb_priority(0),tunnel(tun_id=0x160017000004,src=192.168.137.177,dst=192.168.137.176,ttl=64,tp_src=38881,tp_dst=7471,flags(csum|key)),skb_mark(0),ct_state(0x21),ct_zone(0xe),ct_tuple4(src=10.16.0.5,dst=10.16.0.10,proto=1,tp_src=8,tp_dst=0),in_port(4) mtu 0
2022-06-02T23:23:31.840Z|00080|dpif(handler)|WARN|system@ovs-system: execute ct(commit,zone=14,label=0/0x1,nat(src)),8 failed (Invalid argument) on packet
 icmp,vlan_tci=0x0000,dl_src=00:00:00:f8:07:c8,dl_dst=00:00:00:fa:1e:50,nw_src=10.16.0.5,nw_dst=10.16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0
 icmp_csum=15b2
 with metadata skb_priority(0),tunnel(tun_id=0x160017000004,src=192.168.137.177,dst=192.168.137.176,ttl=64,tp_src=38881,tp_dst=7471,flags(csum|key)),skb_mark(0),ct_state(0x21),ct_zone(0xe),ct_tuple4(src=10.16.0.5,dst=10.16.0.10,proto=1,tp_src=8,tp_dst=0),in_port(4) mtu 0
2022-06-03T00:09:15.659Z|00081|dpif(handler)|WARN|system@ovs-system: execute ct(commit,zone=14,label=0/0x1,nat(src)),8 failed (Invalid argument) on packet
 icmp,vlan_tci=0x0000,dl_src=00:00:00:dc:a3:63,dl_dst=00:00:00:fa:1e:50,nw_src=10.16.63.30,nw_dst=10.16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0
 16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0 icmp_csum=e5a5
 with metadata skb_priority(0),tunnel(tun_id=0x150017000004,src=192.168.137.178,dst=192.168.137.176,ttl=64,tp_src=9239,tp_dst=7471,flags(csum|key)),skb_mark(0),ct_state(0x21),ct_zone(0xe),ct_tuple4(src=10.16.0.5,dst=10.16.0.10,proto=1,tp_src=8,tp_dst=0),in_port(4) mtu 0
2022-06-03T00:30:13.409Z|00064|dpif(handler2)|WARN|system@ovs-system: execute ct(commit,zone=14,label=0/0x1,nat(src)),8 failed (Invalid argument) on packet
 icmp,vlan_tci=0x0000,dl_src=00:00:00:f8:07:c8,dl_dst=00:00:00:fa:1e:50,nw_src=10.16.0.5,nw_dst=10.16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0
 icmp_csum=6b4a
 with metadata skb_priority(0),tunnel(tun_id=0x160017000004,src=192.168.137.177,dst=192.168.137.176,ttl=64,tp_src=38881,tp_dst=7471,flags(csum|key)),skb_mark(0),ct_state(0x21),ct_zone(0xe),ct_tuple4(src=10.16.0.5,dst=10.16.0.10,proto=1,tp_src=8,tp_dst=0),in_port(4) mtu 0
2022-06-03T02:03:23.832Z|00082|dpif(handler)|WARN|system@ovs-system: execute ct(commit,zone=14,label=0/0x1,nat(src)),8 failed (Invalid argument) on packet
 icmp,vlan_tci=0x0000,dl_src=00:00:00:f8:07:c8,dl_dst=00:00:00:fa:1e:50,nw_src=10.16.0.5,nw_dst=10.16.0.10,nw_tos=0,nw_ecn=0,nw_ttl=64,icmp_type=8,icmp_code=0
 icmp_csum=a819
 with metadata skb_priority(0),tunnel(tun_id=0x160017000004,src=192.168.137.177,dst=192.168.137.176,ttl=64,tp_src=38881,tp_dst=7471,flags(csum|key)),skb_mark(0),ct_state(0x21),ct_zone(0xe),ct_tuple4(src=10.16.0.5,dst=10.16.0.10,proto=1,tp_src=8,tp_dst=0),in_port(4) mtu 0

ovs-vsctl show results:
0d4c4675-c9cc-448a-8c1a-878e17f81b88
Bridge br-int
 fail_mode: secure
 datapath_type: system
 Port a2c1a8a8b83a_h
 Interface a2c1a8a8b83a_h
 "4fa5c4ccb1a5_h"
 Interface "4fa5c4ccb1a5_h"
 Port ovn-eef07d-0
 Interface ovn-eef07d-0
 type: stt
 options: {csum="true", key=flow, remote_ip="192.168.137.178"}
 Port ovn0
 Interface ovn0

```

```

 type: internal
Port "04d03360e9a0_h"
 Interface "04d03360e9a0_h"
Port eeb4d9e51b5d_h
 Interface eeb4d9e51b5d_h
Port mirror0
 Interface mirror0
 type: internal
Port "8e5d887ccdb0_h"
 Interface "8e5d887ccdb0_h"
Port ovn-efa253-0
 Interface ovn-efa253-0
 type: stt
 options: {csum="true", key=flow, remote_ip="192.168.137.177"}
Port "17512d5be1f1_h"
 Interface "17512d5be1f1_h"
Port br-int
 Interface br-int
 type: internal
ovs_version: "2.17.2"

pinger diagnose results:
I0603 10:35:04.349404 17619 pinger.go:19]

Kube-OVN:
Version: v1.14.9
Build: 2022-04-24_08:02:50
Commit: git-73f9d15
Go Version: go1.17.8
Arch: amd64

I0603 10:35:04.376797 17619 config.go:166] pinger config is &{KubeConfigFile: KubeClient:0xc000493380 Port:8080 DaemonSetNameSpace:kube-system DaemonSetName:kube-ovn-pinger Interval:5 Mode:job ExitCode:0 InternalDNS:kubernetes.default ExternalDNS: NodeName:liumengxin-ovn1-192.168.137.176 HostIP:192.168.137.176 PodName:kube-ovn-pinger-6ftdf PodIP:10.16.0.10 PodProtocols:[IPv4] ExternalAddress: NetworkMode:kube-ovn PollTimeout:2 PollInterval:15 SystemRunDir:/var/run/openvswitch DatabaseVswitchName:Open_vSwitch DatabaseVswitchSocketRemote:unix:/var/run/openvswitch/db.sock DatabaseVswitchFilePath:/etc/openvswitch/conf.db DatabaseVswitchFileLogPath:/var/log/openvswitch/ovsdb-server.log DatabaseVswitchFilePidPath:/var/run/openvswitch/ovsdb-server.pid DatabaseVswitchFileSystemIDPath:/etc/openvswitch/system-id.conf ServiceVswitchchFileLogPath:/var/log/openvswitch/ovs-vswitchd.log ServiceVswitchchFilePidPath:/var/run/openvswitch/ovs-vswitchd.pid ServiceOvnControllerFileLogPath:/var/log/ovn/ovn-controller.log ServiceOvnControllerFilePidPath:/var/run/ovn-controller.pid}
I0603 10:35:04.449166 17619 exporter.go:75] liumengxin-ovn1-192.168.137.176: exporter connect successfully
I0603 10:35:04.554011 17619 ovn.go:21] ovs-vswitchd and ovsvdb are up
I0603 10:35:04.651293 17619 ovn.go:33] ovn_controller is up
I0603 10:35:04.651342 17619 ovn.go:39] start to check port binding
I0603 10:35:04.749613 17619 ovn.go:135] chassis id is 1d7f3d6c-eec5-4b3c-adca-2969d9cdfd80
I0603 10:35:04.763487 17619 ovn.go:49] port in sb is [node-liumengxin-ovn1-192.168.137.176 perf-6vxkn.default kube-state-metrics-5d6805d89-4nf8h.monitoring alertmanager-main-0.monitoring kube-ovn-pinger-6ftdf.kube-system fake-kubelet-67c55dfd89-pv86k.kube-system prometheus-k8s-0.monitoring]
I0603 10:35:04.763583 17619 ovn.go:61] ovs and ovn-sb binding check passed
I0603 10:35:05.049309 17619 ping.go:259] start to check apiserver connectivity
I0603 10:35:05.053666 17619 ping.go:268] connect to apiserver success in 4.27ms
I0603 10:35:05.053786 17619 ping.go:129] start to check pod connectivity
I0603 10:35:05.249590 17619 ping.go:159] ping pod: kube-ovn-pinger-6ftdf 10.16.0.10, count: 3, loss count 0, average rtt 16.30ms
I0603 10:35:05.354135 17619 ping.go:159] ping pod: kube-ovn-pinger-7wb4 10.16.63.30, count: 3, loss count 0, average rtt 1.81ms
I0603 10:35:05.458460 17619 ping.go:159] ping pod: kube-ovn-pinger-vh2xg 10.16.0.5, count: 3, loss count 0, average rtt 1.92ms
I0603 10:35:05.458523 17619 ping.go:83] start to check node connectivity

```

diagnose	subnet	subnet	daemonset	kube-ovn-pinger	daemonset	pod	daemonset
diagnose	IPPorts			kube-ovn-pinger		IP Port	

## reload

### Kube-OVN

```

kubectl ko reload
pod "ovn-central-8684dd94bd-vzgr" deleted
Waiting for deployment "ovn-central" rollout to finish: 0 of 1 updated replicas are available...
deployment "ovn-central" successfully rolled out
pod "ovs-ovn-bsnvz" deleted
pod "ovs-ovn-m9b98" deleted
pod "kube-ovn-controller-8459db5ff4-64c62" deleted
Waiting for deployment "kube-ovn-controller" rollout to finish: 0 of 1 updated replicas are available...
deployment "kube-ovn-controller" successfully rolled out
pod "kube-ovn-cni-2klhn" deleted
pod "kube-ovn-cni-t2jz4" deleted
Waiting for daemon set "kube-ovn-cni" rollout to finish: 0 of 2 updated pods are available...
Waiting for daemon set "kube-ovn-cni" rollout to finish: 1 of 2 updated pods are available...
daemon set "kube-ovn-cni" successfully rolled out
pod "kube-ovn-pinger-ln72z" deleted
pod "kube-ovn-pinger-w8lrk" deleted
Waiting for daemon set "kube-ovn-pinger" rollout to finish: 0 of 2 updated pods are available...
Waiting for daemon set "kube-ovn-pinger" rollout to finish: 1 of 2 updated pods are available...
daemon set "kube-ovn-pinger" successfully rolled out
pod "kube-ovn-monitor-7fb67d5488-7q6zb" deleted
Waiting for deployment "kube-ovn-monitor" rollout to finish: 0 of 1 updated replicas are available...
deployment "kube-ovn-monitor" successfully rolled out

```

**log**

kube-ovn      Kube-OVN OVN Open vSwitch log linux debug

```
kubectl ko log all
Collecting kube-ovn logging files
Collecting ovn logging files
Collecting openvswitch logging files
Collecting linux dmesg files
Collecting linux iptables-legacy files
Collecting linux iptables-nft files
Collecting linux route files
Collecting linux link files
Collecting linux neigh files
Collecting linux memory files
Collecting linux top files
Collecting linux sysctl files
Collecting linux netstat files
Collecting linux addr files
Collecting linux ipset files
Collecting linux tcp files
Collected files have been saved in the directory /root/kubectl-ko-log
```

```
tree kubectl-ko-log/
kubectl-ko-log/
|-- kube-ovn-control-plane
| |-- kube-ovn
| | |-- kube-ovn-cni.log
| | |-- kube-ovn-monitor.log
| | '-- kube-ovn-pinger.log
| '-- linux
| |-- addr.log
| |-- dmesg.log
| |-- ipset.log
| |-- iptables-legacy.log
| |-- iptables-nft.log
| |-- link.log
| |-- memory.log
| |-- neigh.log
| |-- netstat.log
| |-- route.log
| |-- sysctl.log
| |-- tcp.log
| '-- top.log
| '-- openvswitch
| |-- ovs-vswitchd.log
| '-- ovsdb-server.log
| '-- ovn
| |-- ovn-controller.log
| |-- ovn-northd.log
| '-- ovsdb-server-nb.log
| '-- ovsdb-server-sb.log
```

**perf [image]**

Kube-OVN

- 1.
2. Hostnetwork
- 3.
4. OVN-NB, OVN-SB, OVN-Northd leader

image      Pod      kubeovn/test:v1.12.0

```
kubectl ko perf
===== Preparing Performance Test Resources =====
pod/test-client created
pod/test-host-client created
pod/test-server created
pod/test-host-server created
service/test-server created
pod/test-client condition met
pod/test-host-client condition met
pod/test-host-server condition met
pod/test-server condition met
=====
===== Start Pod Network Unicast Performance Test =====
Size TCP Latency TCP Bandwidth UDP Latency UDP Lost Rate UDP Bandwidth
64 82.8 us 97.7 Mbits/sec 67.6 us (0%) 8.42 Mbits/sec
128 85.4 us 167 Mbits/sec 67.2 us (0%) 17.2 Mbits/sec
```

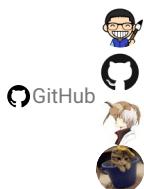
```

512 85.8 us 440 Mbits/sec 68.7 us (0%) 68.4 Mbits/sec
1k 85.1 us 567 Mbits/sec 68.7 us (0%) 134 Mbits/sec
4k 138 us 826 Mbits/sec 78.1 us (1.4%) 503 Mbits/sec
=====
===== Start Host Network Performance Test =====
Size TCP Latency TCP Bandwidth UDP Latency UDP Lost Rate UDP Bandwidth
64 49.7 us 120 Mbits/sec 37.9 us (0%) 18.6 Mbits/sec
128 49.7 us 200 Mbits/sec 38.1 us (0%) 35.5 Mbits/sec
512 51.9 us 588 Mbits/sec 38.9 us (0%) 142 Mbits/sec
1k 51.7 us 944 Mbits/sec 37.2 us (0%) 279 Mbits/sec
4k 74.9 us 1.66 Gbits/sec 39.9 us (0%) 1.20 Gbits/sec
=====
===== Start Service Network Performance Test =====
Size TCP Latency TCP Bandwidth UDP Latency UDP Lost Rate UDP Bandwidth
64 111 us 96.3 Mbits/sec 88.4 us (0%) 7.59 Mbits/sec
128 83.7 us 150 Mbits/sec 69.2 us (0%) 16.9 Mbits/sec
512 87.4 us 374 Mbits/sec 75.8 us (0%) 60.9 Mbits/sec
1k 88.2 us 521 Mbits/sec 73.1 us (0%) 123 Mbits/sec
4k 148 us 813 Mbits/sec 77.6 us (0.0044%) 451 Mbits/sec
=====
===== Start Pod Multicast Network Performance Test =====
Size UDP Latency UDP Lost Rate UDP Bandwidth
64 0.014 ms (0.17%) 5.80 Mbits/sec
128 0.012 ms (0%) 11.4 Mbits/sec
512 0.016 ms (0%) 46.1 Mbits/sec
1k 0.023 ms (0.073%) 89.8 Mbits/sec
4k 0.035 ms (1.3%) 126 Mbits/sec
=====
===== Start Host Multicast Network Performance =====
Size UDP Latency UDP Lost Rate UDP Bandwidth
64 0.007 ms (0%) 9.95 Mbits/sec
128 0.005 ms (0%) 21.8 Mbits/sec
512 0.008 ms (0%) 86.8 Mbits/sec
1k 0.013 ms (0.045%) 168 Mbits/sec
4k 0.010 ms (0.31%) 242 Mbits/sec
=====
===== Start Leader Recover Time Test =====
Delete ovn central nb pod
pod "ovn-central-5cb9c67d75-tlz9w" deleted
Waiting for ovn central nb pod running
===== OVN nb Recovery takes 3.305236803 s =====
Delete ovn central sb pod
pod "ovn-central-5cb9c67d75-szx4c" deleted
Waiting for ovn central sb pod running
===== OVN sb Recovery takes 3.462698535 s =====
Delete ovn central northd pod
pod "ovn-central-5cb9c67d75-zqmqv" deleted
Waiting for ovn central northd pod running
===== OVN northd Recovery takes 2.691291403 s =====
=====
===== Remove Performance Test Resource =====
rm -f unicast-test-client.log
rm -f unicast-test-host-client.log
rm -f unicast-test-client.log
kubectl ko nbctl lb-del test-server
rm -f multicast-test-server.log
kubectl exec ovs-ovn-gxdrf -n kube-system -- ip maddr del 01:00:5e:00:00:64 dev eth0
kubectl exec ovs-ovn-h57bf -n kube-system -- ip maddr del 01:00:5e:00:00:64 dev eth0
rm -f multicast-test-host-server.log
pod "test-client" deleted
pod "test-host-client" deleted
pod "test-host-server" deleted
pod "test-server" deleted
service "test-server" deleted
=====
```

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2022 5 24



6.1.3

## 6.2



### 6.2.1

```

kubectl drain kube-ovn-worker --ignore-daemonsets --force
node/kube-ovn-worker cordoned
WARNING: ignoring DaemonSet-managed Pods: kube-system/kube-ovn-cni-zt74b, kube-system/kube-ovn-pinger-5rxf5, kube-system/kube-proxy-jpmnm, kube-system/ovs-ovn-v2k11
evicting pod kube-system/coredns-64897985d-qsgpt
evicting pod local-path-storage/local-path-provisioner-5ddd94ff66-llss6
evicting pod kube-system/kube-ovn-controller-8459db5ff4-941xb
pod/kube-ovn-controller-8459db5ff4-941xb evicted
pod/coredns-64897985d-qsgpt evicted
pod/local-path-provisioner-5ddd94ff66-llss6 evicted
node/kube-ovn-worker drained

```

### 6.2.2 kubelet docker

ovs-ovn      ovn-central

```

systemctl stop kubelet
systemctl stop docker

```

CRI containerd      ovs-ovn

```

crictl rm -f $(crictl ps | grep openvswitch | awk '{print $1}')

```

### 6.2.3 Node

```

rm -rf /var/run/openvswitch
rm -rf /var/run/ovn
rm -rf /etc/origin/openvswitch/
rm -rf /etc/origin/ovn/
rm -rf /etc/cni/net.d/00-kube-ovn.conflist
rm -rf /etc/cni/net.d/01-kube-ovn.conflist
rm -rf /var/log/openvswitch
rm -rf /var/log/ovn

```

### 6.2.4 kubectl

```

kubectl delete no kube-ovn-01

```

### 6.2.5 ovn-sb

kube-ovn-worker

```

kubectl ko sbctl show
Chassis "b0564934-5a0d-4804-a4c0-476c93596a17"
 hostname: kube-ovn-worker
 Encap geneve
 ip: "172.18.0.2"
 options: {csum="true"}
 Port_Binding kube-ovn-pinger-5rxf5.kube-system
Chassis "6a29de7e-d731-4eaf-bacd-2f239ee52b28"
 hostname: kube-ovn-control-plane
 Encap geneve
 ip: "172.18.0.3"
 options: {csum="true"}
 Port_Binding coredns-64897985d-nbfln.kube-system
 Port_Binding node-kube-ovn-control-plane
 Port_Binding local-path-provisioner-5ddd94ff66-h4tn9.local-path-storage
 Port_Binding kube-ovn-pinger-hf2p6.kube-system
 Port_Binding coredns-64897985d-fhwlw.kube-system

```

## 6.2.6 chassis

uuid      Chassis    id

```
kubectl ko sbctl chassis-del b0564934-5a0d-4804-a4c0-476c93596a17
kubectl ko sbctl show
Chassis "6a29de7e-d731-4eaf-bacd-2f239ee52b28"
 hostname: kube-ovn-control-plane
 Encap geneve
 ip: "172.18.0.3"
 options: {csum="true"}
 Port_Binding coredns-64897985d-nbf1n.kube-system
 Port_Binding node-kube-ovn-control-plane
 Port_Binding local-path-provisioner-5dd94ff66-h4tn9.local-path-storage
 Port_Binding kube-ovn-pinger-hf2p6.kube-system
 Port_Binding coredns-64897985d-fhw1w.kube-system
```

[PDF](#)[Slack](#)[Support](#)

⌚2023 4 25

⌚2022 6 3

GitHub 

## 6.2.7

## 6.3 ovn-central

```
ovn-central ovn-nb ovn-sb etcd Raft ovn-central
```

### 6.3.1 ovn-central

```
kube-ovn-control-plane2 ovn-central
```

```
kubectl -n kube-system get pod -o wide | grep central
ovn-central-6bf58cbc97-2cdhg 1/1 Running 0 21m 172.18.0.3 kube-ovn-control-plane <none> <none>
ovn-central-6bf58cbc97-crmfp 1/1 Running 0 21m 172.18.0.5 kube-ovn-control-plane2 <none> <none>
ovn-central-6bf58cbc97-lxmpl 1/1 Running 0 21m 172.18.0.4 kube-ovn-control-plane3 <none> <none>
```

#### ovn-nb

ID

```
kubectl ko nb status
1b9a
Name: OVN_Northbound
Cluster ID: 32ca (32ca07fb-739b-4257-b510-12fa18e7cce8)
Server ID: 1b9a (1b9a5d76-e69b-410c-8085-39943d0cd38c)
Address: tcp:[172.18.0.3]:6643
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 2135194 ms ago, reason: timeout
Last Election won: 2135188 ms ago
Election timer: 5000
Log: [135, 135]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-d64b ->d64b <-4984 ->4984
Disconnections: 0
Servers:
 4984 (4984 at tcp:[172.18.0.4]:6643) next_index=135 match_index=134 last msg 1084 ms ago
 1b9a (1b9a at tcp:[172.18.0.3]:6643) (self) next_index=2 match_index=134
 d64b (d64b at tcp:[172.18.0.5]:6643) next_index=135 match_index=134 last msg 1084 ms ago
status: ok
```

```
kube-ovn-control-plane2 IP 172.18.0.5 ID d64b ovn-nb
```

```
kubectl ko nb kick d64b
started removal
```

```
kubectl ko nb status
1b9a
Name: OVN_Northbound
Cluster ID: 32ca (32ca07fb-739b-4257-b510-12fa18e7cce8)
Server ID: 1b9a (1b9a5d76-e69b-410c-8085-39943d0cd38c)
Address: tcp:[172.18.0.3]:6643
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 2297649 ms ago, reason: timeout
Last Election won: 2297643 ms ago
Election timer: 5000
Log: [136, 136]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-4984 ->4984
Disconnections: 2
Servers:
 4984 (4984 at tcp:[172.18.0.4]:6643) next_index=136 match_index=135 last msg 1270 ms ago
 1b9a (1b9a at tcp:[172.18.0.3]:6643) (self) next_index=2 match_index=135
status: ok
```

**ovn-sb**

ovn-sb ID

```
kubectl ko sb status
3722
Name: OVN_Southbound
Cluster ID: d4bd (d4bd37a4-0400-499f-b4df-b4fd389780f0)
Server ID: 3722 (3722d5ae-2ced-4820-a6b2-8b744d11fb3e)
Address: tcp:[172.18.0.3]:6644
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 2395317 ms ago, reason: timeout
Last Election won: 2395316 ms ago
Election timer: 5000
Log: [130, 130]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-e9f7 ->e9f7 <-6e84 ->6e84
Disconnections: 0
Servers:
 e9f7 (e9f7 at tcp:[172.18.0.5]:6644) next_index=130 match_index=129 last msg 1006 ms ago
 6e84 (6e84 at tcp:[172.18.0.4]:6644) next_index=130 match_index=129 last msg 1004 ms ago
 3722 (3722 at tcp:[172.18.0.3]:6644) (self) next_index=2 match_index=129
status: ok
```

kube-ovn-control-plane2 IP 172.18.0.5 ID e9f7 ovn-sb

```
kubectl ko sb kick e9f7
started removal
```

```
kubectl ko sb status
3722
Name: OVN_Southbound
Cluster ID: d4bd (d4bd37a4-0400-499f-b4df-b4fd389780f0)
Server ID: 3722 (3722d5ae-2ced-4820-a6b2-8b744d11fb3e)
Address: tcp:[172.18.0.3]:6644
Status: cluster member
Role: leader
Term: 1
Leader: self
Vote: self

Last Election started 2481636 ms ago, reason: timeout
Last Election won: 2481635 ms ago
Election timer: 5000
Log: [131, 131]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: <-6e84 ->6e84
Disconnections: 2
Servers:
 6e84 (6e84 at tcp:[172.18.0.4]:6644) next_index=131 match_index=130 last msg 642 ms ago
 3722 (3722 at tcp:[172.18.0.3]:6644) (self) next_index=2 match_index=130
status: ok
```

**ovn-central**

ovn-central NODE\_IPS

```
kubectl label node kube-ovn-control-plane2 kube-ovn/role-
kubectl scale deployment -n kube-system ovn-central --replicas=2
kubectl set env deployment/ovn-central -n kube-system NODE_IPS="172.18.0.3,172.18.0.4"
kubectl rollout status deployment/ovn-central -n kube-system
```

**ovn-central**

ovs-ovn

```
kubectl set env daemonset/ovs-ovn -n kube-system OVN_DB_IPS="172.18.0.3,172.18.0.4"
daemonset.apps/ovs-ovn env updated
kubectl delete pod -n kube-system -lapp=ovs
pod "ovs-ovn-4f6jc" deleted
```

```

pod "ovs-ovn-csn2w" deleted
pod "ovs-ovn-mpbmb" deleted

kube-ovn-controller

kubectl set env deployment/kube-ovn-controller -n kube-system OVN_DB_IPS="172.18.0.3,172.18.0.4"
deployment.apps/kube-ovn-controller env updated

kubectl rollout status deployment/kube-ovn-controller -n kube-system
Waiting for deployment "kube-ovn-controller" rollout to finish: 1 of 3 updated replicas are available...
Waiting for deployment "kube-ovn-controller" rollout to finish: 2 of 3 updated replicas are available...
deployment "kube-ovn-controller" successfully rolled out

```

## kube-ovn-control-plane2

```
rm -rf /etc/origin/ovn
```

Kubernetes

## 6.3.2 ovn-central

Kubernetes ovn-central

```
/etc/origin/ovn ovnnb_db.db ovnsb_db.db
```

```
rm -rf /etc/origin/ovn
```

**ovn-central**

## ovn-central

```

kubectl ko nb status
1b9a
Name: OVN_Northbound
Cluster ID: 32ca (32ca07fb-739b-4257-b510-12fa18e7cce8)
Server ID: 1b9a (1b9a5d76-e69b-410c-8085-39943d0cd38c)
Address: tcp:[172.18.0.3]:6643
Status: cluster member
Role: leader
Term: 44
Leader: self
Vote: self

Last Election started 1855739 ms ago, reason: timeout
Last Election won: 1855729 ms ago
Election timer: 5000
Log: [147, 147]
Entries not yet committed: 0
Entries not yet applied: 0
Connections: ->4984 <-4984
Disconnections: 0
Servers:
 4984 (4984 at tcp:[172.18.0.4]:6643) next_index=147 match_index=146 last msg 367 ms ago
 1b9a (1b9a at tcp:[172.18.0.3]:6643) (self) next_index=140 match_index=146
status: ok

kubectl ko sb status
3722
Name: OVN_Southbound
Cluster ID: d4bd (d4bd37a4-0400-499f-b4df-b4fd389780f0)
Server ID: 3722 (3722d5ae-2ced-4820-a6b2-8b744d11fb3e)
Address: tcp:[172.18.0.3]:6644
Status: cluster member
Role: leader
Term: 33
Leader: self
Vote: self

Last Election started 1868589 ms ago, reason: timeout
Last Election won: 1868579 ms ago
Election timer: 5000
Log: [142, 142]
Entries not yet committed: 0

```

```
Entries not yet applied: 0
Connections: ->6e84 <-6e84
Disconnections: 0
Servers:
 6e84 (6e84 at tcp:[172.18.0.4]:6644) next_index=142 match_index=141 last msg 728 ms ago
 3722 (3722 at tcp:[172.18.0.3]:6644) (self) next_index=134 match_index=141
status: ok
```

**ovn-central**    **NODE\_IPS**

```
kubectl label node kube-ovn-control-plane2 kube-ovn/role=master
kubectl scale deployment -n kube-system ovn-central --replicas=3
kubectl set env deployment/ovn-central -n kube-system NODE_IPS="172.18.0.3,172.18.0.4,172.18.0.5"
kubectl rollout status deployment/ovn-central -n kube-system
```

**ovn-central****ovs-ovn**

```
kubectl set env daemonset/ovs-ovn -n kube-system OVN_DB_IPS="172.18.0.3,172.18.0.4,172.18.0.5"
daemonset.apps/ovs-ovn env updated
kubectl delete pod -n kube-system -lapp=ovs
pod "ovs-ovn-4f6jc" deleted
pod "ovs-ovn-csn2w" deleted
pod "ovs-ovn-mpbmb" deleted
```

**kube-ovn-controller**

```
kubectl set env deployment/kube-ovn-controller -n kube-system OVN_DB_IPS="172.18.0.3,172.18.0.4,172.18.0.5"
deployment.apps/kube-ovn-controller env updated

kubectl rollout status deployment/kube-ovn-controller -n kube-system
Waiting for deployment "kube-ovn-controller" rollout to finish: 1 of 3 updated replicas are available...
Waiting for deployment "kube-ovn-controller" rollout to finish: 2 of 3 updated replicas are available...
deployment "kube-ovn-controller" successfully rolled out
```

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2022 5 24



6.3.3

## 6.4 OVN

---

### 6.4.1

```
kubectl backup
```

```
kubectl ko nb backup
tar: Removing leading `/' from member names
backup ovn-nb db to /root/ovnnb_db.060223191654183154.backup

kubectl ko sb backup
tar: Removing leading `/' from member names
backup ovn-sb db to /root/ovnsb_db.060223191654183154.backup
```

### 6.4.2

```
/var/log/ovn/ovn-northd.log
```

```
* ovn-northd is not running
ovsdb-server: ovsdb error: error reading record 2739 from OVN_Northbound log: record 2739 advances commit index to 6308 but last log index is 6307
* Starting ovsdb-nb
```

OVN_Northbound	OVN_Southbound	OVN_Northbound	ovn-nb
----------------	----------------	----------------	--------

```
kubectl ko nb status
9182
Name: OVN_Northbound
Cluster ID: e75f (e75fa340-49ed-45ab-990e-26cb865ebc85)
Server ID: 9182 (9182e8dd-b5b0-4dd8-8518-598cc1e374f3)
Address: tcp:[10.0.128.61]:6643
Status: cluster member
Role: leader
Term: 1454
Leader: self
Vote: self

Last Election started 1732603 ms ago, reason: timeout
Last Election won: 1732587 ms ago
Election timer: 1000
Log: [7332, 12512]
Entries not yet committed: 1
Entries not yet applied: 1
Connections: ->f080 <-f080 <-e631 ->e631
Disconnections: 1
Servers:
 f080 (f080 at tcp:[10.0.129.139]:6643) next_index=12512 match_index=12510 last msg 63 ms ago
 9182 (9182 at tcp:[10.0.128.61]:6643) () next_index=10394 match_index=12510
 e631 (e631 at tcp:[10.0.131.173]:6643) next_index=12512 match_index=0
```

```
kubectl ko nb kick e631
```

```
mv /etc/origin/ovn/ovnnb_db.db /tmp
```

```
ovn-central Pod
```

```
kubectl delete pod -n kube-system ovn-central-xxxx
```

### 6.4.3

leader

#### ovn-central

ovn-central ovn-central

```
kubectl scale deployment -n kube-system ovn-central --replicas=0
```

ovsdb-tool cluster-to-standalone

ovn-central NODE\_IPS

/etc/origin/ovn

```
docker run -it -v /etc/origin/ovn:/etc/ovn kubeovn/kube-ovn:v1.14.9 bash
cd /etc/ovn/
ovsdb-tool cluster-to-standalone ovnnb_db_standalone.db ovnnb_db.db
ovsdb-tool cluster-to-standalone ovnsb_db_standalone.db ovnsb_db.db
```

#### ovn-central

```
mv /etc/origin/ovn/ovnnb_db.db /tmp
mv /etc/origin/ovn/ovnsb_db.db /tmp
```

ovnnb\_db.db ovnsb\_db.db ovn-central NODE\_IPS /etc/origin/ovn/

```
mv /etc/origin/ovn/ovnnb_db_standalone.db /etc/origin/ovn/ovnnb_db.db
mv /etc/origin/ovn/ovnsb_db_standalone.db /etc/origin/ovn/ovnsb_db.db
```

ovn-central

```
kubectl scale deployment -n kube-system ovn-central --replicas=3
kubectl rollout status deployment/ovn-central -n kube-system
```

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 5 24



### 6.4.4

## 6.5 CIDR

---

CIDR

CIDR Pod

CIDR

Join CIDR

Join CIDR

### 6.5.1

```
kubectl edit cidrBlock gateway excludeIps
```

```
kubectl edit subnet test-subnet
```

### 6.5.2 Namespace Pod

test Namespace

```
for pod in $(kubectl get pod --no-headers -n "$ns" --field-selector spec.restartPolicy=Always -o custom-columns=NAME:.metadata.name,HOST:spec.hostNetwork | awk '{if ($2!="true") print $1}'); do
 kubectl delete pod "$pod" -n test --ignore-not-found
done
```

host Pod

```
for ns in $(kubectl get ns --no-headers -o custom-columns=NAME:.metadata.name); do
 for pod in $(kubectl get pod --no-headers -n "$ns" --field-selector spec.restartPolicy=Always -o custom-columns=NAME:.metadata.name,HOST:spec.hostNetwork | awk '{if ($2!="true") print $1}'); do
 kubectl delete pod "$pod" -n "$ns" --ignore-not-found
 done
done
```

### 6.5.3

CIDR kube-ovn-controller Deployment

```
args:
- --default-cidr=10.17.0.0/16
- --default-gateway=10.17.0.1
- --default-exclude-ipss=10.17.0.1
```

 PDF

 Slack

 Support

 2025 9 10

 2022 5 24



### 6.5.4

## 6.6 Join CIDR

Join CIDR

Join CIDR Pod

### 6.6.1 Join

```
kubectl patch subnet join --type='json' -p '[{"op": "replace", "path": "/metadata/finalizers", "value": []}]'
kubectl delete subnet join
```

### 6.6.2

```
kubectl annotate node ovn.kubernetes.io/allocated=false --all --overwrite
```

### 6.6.3 Join

`kube-ovn-controller` Join

```
kubectl edit deployment -n kube-system kube-ovn-controller
```

```
args:
- --node-switch-cidr=100.51.0.0/16
```

`kube-ovn-controller` join

```
kubectl delete pod -n kube-system -lapp=kube-ovn-controller
```

Join

```
kubectl get subnet
NAME PROVIDER VPC PROTOCOL CIDR PRIVATE NAT DEFAULT GATEWAYTYPE V4USED V4AVAILABLE V6USED V6AVAILABLE
EXCLUDEIPS
join ovn ovn-cluster IPv4 100.51.0.0/16 false false false distributed 2 65531 0 0
["100.51.0.1"]
ovn-default ovn ovn-cluster IPv4 10.17.0.0/16 true true true distributed 5 65528 0 0
["10.17.0.1"]
```

### 6.6.4 ovn0

ovn0 kube-ovn-cni

```
kubectl delete pod -n kube-system -l app=kube-ovn-cni
```

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2022 5 24



### 6.6.5

## 6.7

kube-ovn.yaml

```
vi kube-ovn.yaml
...
- name: kube-ovn-controller
 image: "docker.io/kubeovn/kube-ovn:v1.14.9"
 imagePullPolicy: IfNotPresent
 args:
 - /kube-ovn/start-controller.sh
 - --v=3
...
#
```

[PDF](#)

[Slack](#)

[Support](#)

⌚2025 9 10

⌚2023 4 4



⌚ GitHub



### 6.7.1

## 6.8

---

### Kube-OVN

#### 6.8.1

1.

- Pod Pod
- Pod
- Pod Service
- Pod
- Pod
- kube-ovn-pinger

2.

- Pod
- kube-ovn-cni CNI
- ovs-ovn OVS
- kube-ovn-controller
- ovn-central OVN
- dmesg
- netstat -s

3. CPU, IO

```
kubectl ko logs
```

#### 6.8.2 Pod IP

```
Pod Running kubectl describe Pod duplicate IPv4 address <ip> found on logical switch port <port>
```

1. Pod IP Pod IP IP IP IP Pod
2. kube-ovn-controller IP
3. kube-ovn-controller IP
4. kubectl ko nbctl show OVN IP
5. OVN Kubernetes IP kubectl ko nbctl del-port <port>

#### 6.8.3 Pod ping gateway failed

```
Pod Running kubectl describe Pod network <ip> with gateway <gw ip> is not ready for interface eth0 after 30 checks
```

1. kubectl ko sbctl show Pod
2. ovn-central ovs-ovn ovn-central ovs-ovn

3. Pod
4. Underlay [Underlay](#)

#### 6.8.4 Pod

---

VPC Pod

1. kubectl ko trace OVN ACL
2. ACL ACL
3. ACL Subnet stats
4. Subnet Spec
5. kube-ovn-controller

#### 6.8.5 Pod IP CIDR

---

Pod IP CIDR

1. /etc/cni/net.d/ Kube-OVN CNI
- 2.
3. kubelet Pod

#### 6.8.6 Debug Pod

---

kubectl debug Pod ContainerCreating Pod Event network not ready no address allocated

debug Pod yaml yaml Annotation

```
ovn.kubernetes.io/ip_address
ovn.kubernetes.io/mac_address
ovn.kubernetes.io/allocated
ovn.kubernetes.io/routed
```

debug Pod yaml debug Pod

#### 6.8.7 ARM

---

ARM Offload

netstat

```
netstat -us
IcmpMsg:
 InType0: 22
 InType3: 24
 InType8: 117852
 OutType0: 117852
```

```

OutType3: 29
OutType8: 22
Udp:
 3040636 packets received
 0 packets to unknown port received.
 4 packet receive errors
 602 packets sent
 0 receive buffer errors
 0 send buffer errors
 InCsumErrors: 4
UdpLite:
IpExt:
 InBcastPkts: 10244
 InOctets: 4446320361
 OutOctets: 1496815600
 InBcastOctets: 3095950
 InNoECTPkts: 7683903

```

InCsumErrors

tx offload      TCP

ethtool -K eth0 tx off

CentOS 7

4.19.90-25.16.v2101

## 6.8.8 Pod Service

Pod      Service      dmesg

```

netlink: Unknown conntrack attr (type=6, max=5)
openvswitch: netlink: Flow actions may not be safe on all matching packets.

```

OVS      NAT

1.      OVS

2.      Overlay      kube-ovn-controller      --enable-lb=false      OVN LB      kube-proxy      Service

## 6.8.9 ovn-central

v1.11.x      1      Pod      OVN NB      SB      Kube-OVN      ovsdb-server/compact

ovn-central      compact

```

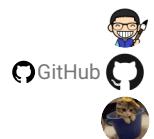
- name: ENABLE_COMPACT
 value: "false"

```

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2022 5 24



6.8.10

---

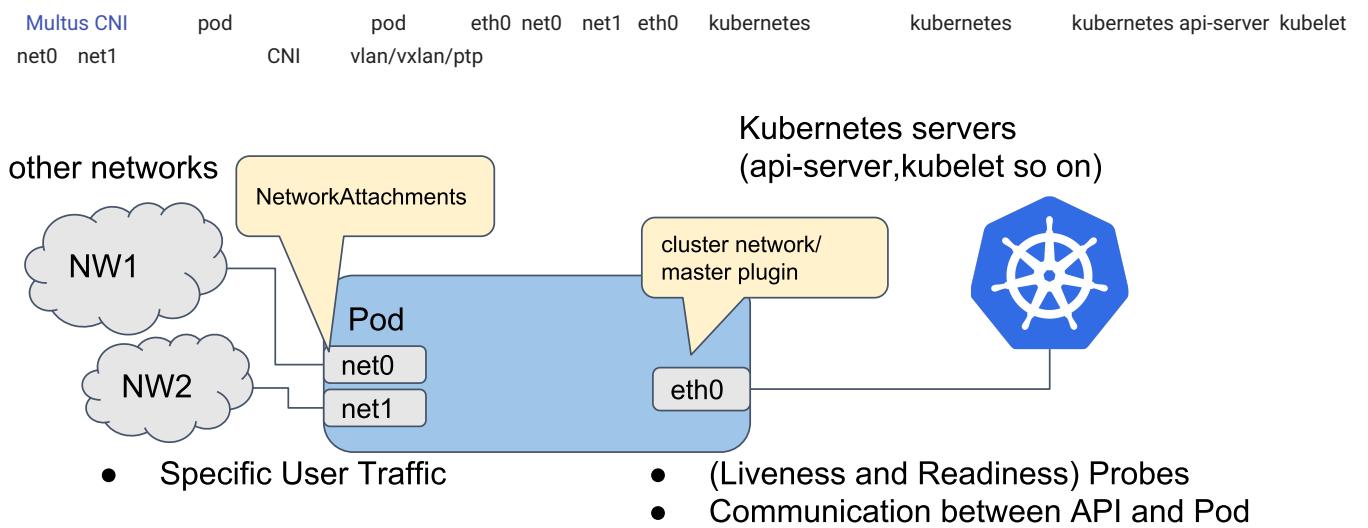
## 7.

---

### 7.1

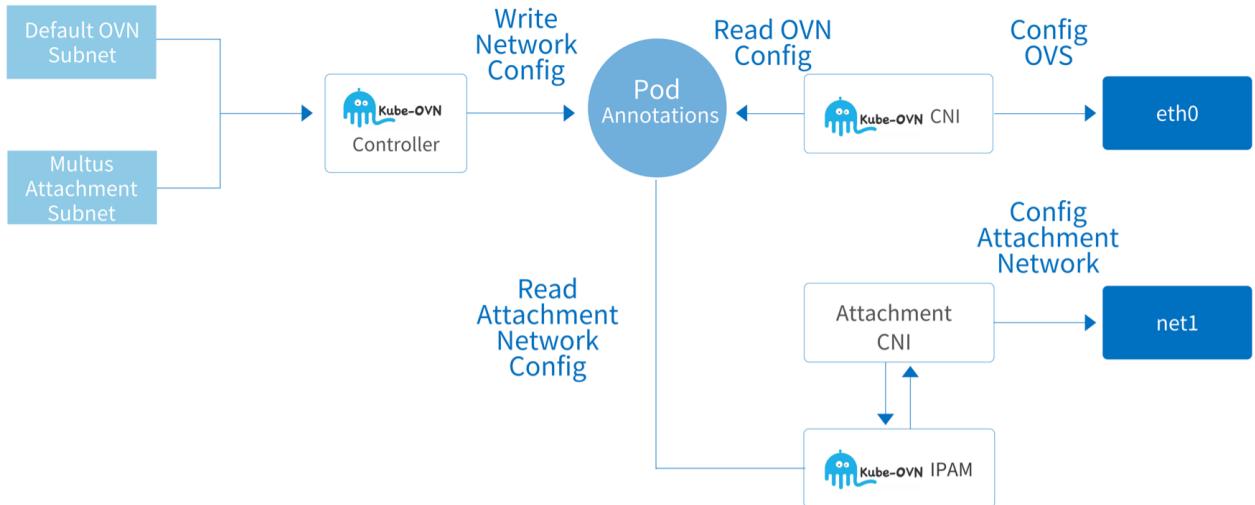
Kube-OVN      CNI      macvlan vlan host-device      IPAM      Kube-OVN      IP  
 Kube-OVN      Kube-OVN

#### 7.1.1



### IPAM

Multus CNI,      Pod      IP  
 Kube-OVN      IPAM      Subnet      IP      CRD      IP      IP



	Kube-OVN	IP	eth0	OVN	net1	CNI	net1	multus-cni	NetworkAttachmentDefinition
Pod	kube-ovn-controller	Pod	Pod	annotation	Subnet	IP	Pod	Pod	Pod annotation
CNI		kube-ovn-cni	ipam	, kube-ovn-cni	Pod annotation	CNI		CNI	CNI

## 7.1.2

NetworkAttachmentDefinition spec multus defaultConfDir CNI

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: macvlan-conf-2
```

kube-ovn-controller NetworkAttachmentDefinition provider spec Kube-OVN IPAM

## 7.1.3

### Kube-OVN Multus

Kube-OVN Multus how to use Kube-OVN Multus-CNI

### CNI IPAM

Kube-OVN CNI

NETWORKATTACHMENTDEFINITION

macvlan ipam kube-ovn

```
macvlan
sudo modprobe macvlan
```

```

apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: macvlan
 namespace: default
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "macvlan",
 "master": "eth0",
 "mode": "bridge",
 "ipam": {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "macvlan.default"
 }
}'

```

- spec.config.ipam.type: kube-ovn kube-ovn
- server\_socket: Kube-OVN socket /run/openvswitch/kube-ovn-daemon.sock
- provider: NetworkAttachmentDefinition <name>. <namespace>, Kube-OVN Subnet
- master:

#### KUBE-OVN SUBNET

Kube-OVN Subnet,	cidrBlock	exclude_ips	provider	NetworkAttachmentDefinition	<name>. <namespace>	macvlan
Subnet						

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: macvlan
spec:
 protocol: IPv4
 provider: macvlan.default
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 excludeIps:
 - 172.17.0.0..172.17.0.10

```

gateway, private, nat provider ovn attachment network

#### Pod

Pod	annotation k8s.v1.cni.cncf.io/networks,	NetworkAttachmentDefinition	<namespace>/<name>
-----	-----------------------------------------	-----------------------------	--------------------

```

apiVersion: v1
kind: Pod
metadata:
 name: samplepod
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/networks: default/macvlan
spec:
 containers:
 - name: samplepod
 command: ["/bin/ash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge

```

#### IP Pod

IP Pod	<networkAttachmentName>. <networkAttachmentNamespace>.kubernetes.io/ip_address annotation
--------	-------------------------------------------------------------------------------------------

```

apiVersion: v1
kind: Pod
metadata:
 name: static-ip
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/networks: default/macvlan
 ovn.kubernetes.io/ip_address: 10.16.0.15
 ovn.kubernetes.io/mac_address: 00:00:00:53:6B:B6
 macvlan.default.kubernetes.io/ip_address: 172.17.0.100
 macvlan.default.kubernetes.io/mac_address: 00:00:00:53:6B:BB
spec:
 containers:
 - name: static-ip
 image: docker.io/library/nginx:alpine

```

IP

```
ippool , <networkAttachmentName>.<networkAttachmentNamespace>.kubernetes.io/ip_pool annotations:
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
 namespace: default
 name: static-workload
 labels:
 app: static-workload
spec:
 replicas: 2
 selector:
 matchLabels:
 app: static-workload
 template:
 metadata:
 labels:
 app: static-workload
 annotations:
 k8s.v1.cni.cncf.io/networks: default/macvlan
 ovn.kubernetes.io/ip_pool: 10.16.0.15,10.16.0.16,10.16.0.17
 macvlan.default.kubernetes.io/ip_pool: 172.17.0.200,172.17.0.201,172.17.0.202
 spec:
 containers:
 - name: static-workload
 image: docker.io/library/nginx:alpine
```

macvlan Pod

macvlan	Pod	Pod	annotation default-route
---------	-----	-----	--------------------------

```
apiVersion: v1
kind: Pod
metadata:
 name: samplepod-route
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/networks: '[{
 "name": "macvlan",
 "namespace": "default",
 "default-route": ["172.17.0.1"]
 }]'
spec:
 containers:
 - name: samplepod-route
 command: ["/bin/ash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge
```

macvlan Pod

macvlan	Pod	annotation v1.multus-cni.io/default-network ,	NetworkAttachmentDefinition <namespace>/<name>
---------	-----	-----------------------------------------------	------------------------------------------------

```
apiVersion: v1
kind: Pod
metadata:
 name: samplepod-macvlan
 namespace: default
 annotations:
 v1.multus-cni.io/default-network: default/macvlan
spec:
 containers:
 - name: samplepod-macvlan
 command: ["/bin/ash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge
```

**KUBE-OVN SUBNET PROVIDER OVN**

Kube-OVN Subnet	cidrBlock	exclude_ips	provider ovn	Subnet
-----------------	-----------	-------------	--------------	--------

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: macvlan
spec:
 protocol: IPv4
 provider: ovn
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 excludeIps:
 - 172.17.0.0..172.17.0.10
```

**Pod**

```
provider ovn subnet IP Pod annotation k8s.v1.cni.cncf.io/networks
<networkAttachmentName>.<networkAttachmentNamespace>.kubernetes.io/logical_switch
```

```
apiVersion: v1
kind: Pod
metadata:
 name: samplepod
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/networks: default/macvlan
 macvlan.default.kubernetes.io/logical_switch: macvlan
spec:
 containers:
 - name: samplepod
 command: ["/bin/ash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge
```

- k8s.v1.cni.cncf.io/networks : NetworkAttachmentDefinition <namespace>/<name>

- macvlan.default.kubernetes.io/logical\_switch :

<networkAttachmentName>.<networkAttachmentNamespace>.kubernetes.io/logical_switch	provider	ovn	subnet
ipam          IP Pod        IP           macvlan Pod       macvlan Pod			

**Kube-OVN****Kube-OVN****NETWORKATTACHMENTDEFINITION**

```
provider ovn
```

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: attachnet
 namespace: default
spec:
 config: '{
 "cniVersion": "0.3.0",
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "attachnet.default.ovn"
 }'
```

- spec.config.type : kube-ovn CNI Kube-OVN

- server\_socket : Kube-OVN socket /run/openvswitch/kube-ovn-daemon.sock

- provider : NetworkAttachmentDefinition <name>.<namespace>.ovn , Kube-OVN Subnet ovn

**KUBE-OVN SUBNET**

Kube-OVN	provider	NetworkAttachmentDefinition <name>.<namespace>.ovn	ovn	Kube-OVN	Subnet
----------	----------	----------------------------------------------------	-----	----------	--------

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: attachnet
spec:
 protocol: IPv4
 provider: attachnet.default.ovn
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 excludeIps:
 - 172.17.0.0..172.17.0.10
```

**Pod**

Pod	annotation k8s.v1.cni.cncf.io/networks ,	NetworkAttachmentDefinition <namespace>/<name>
-----	------------------------------------------	------------------------------------------------

```
apiVersion: v1
kind: Pod
metadata:
 name: samplepod
```

```

namespace: default
annotations:
 k8s.v1.cni.cncf.io/networks: default/attachnet
spec:
 containers:
 - name: samplepod
 command: ["/bin/bash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge

```

## KUBE-OVN SUBNET PROVIDER OVN

Kube-OVN Subnet      cidrBlock    exclude\_ips , provider    ovn    Subnet

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: attachnet
spec:
 protocol: IPv4
 provider: ovn
 cidrBlock: 172.17.0.0/16
 gateway: 172.17.0.1
 excludeIps:
 - 172.17.0.0..172.17.0.10

```

## Pod

provider	ovn	subnet	IP	Pod	annotation k8s.v1.cni.cncf.io/networks
<networkAttachmentName>.<networkAttachmentNamespace>.					ovn.kubernetes.io/logical_switch

```

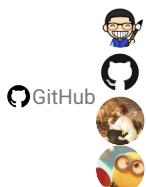
apiVersion: v1
kind: Pod
metadata:
 name: samplepod
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/networks: default/attachnet
 attachnet.default.ovn.kubernetes.io/logical_switch: attachnet
spec:
 containers:
 - name: samplepod
 command: ["/bin/bash", "-c", "trap : TERM INT; sleep infinity & wait"]
 image: docker.io/library/alpine:edge

```

- k8s.v1.cni.cncf.io/networks : NetworkAttachmentDefinition <namespace>/<name>
- attachnet.default.ovn.kubernetes.io/logical\_switch :

: <networkAttachmentName>.<networkAttachmentNamespace>.	kubernetes.io/logical_switch	provider	Kube-OVN
Pod      IP      Pod      IP	macvlan      Pod	Kube-OVN      Pod	

[!\[\]\(1cff35817f9b26f1ebd3caafd81eb6a4\_img.jpg\) PDF](#)
[!\[\]\(a7af5d58cbd8aa4306c2a25713e8d664\_img.jpg\) Slack](#)
[!\[\]\(76145ff3cab98c5aa8507684243c45a5\_img.jpg\) Support](#)
 2025 9 10

 2022 5 20



---

 7.1.4

## 7.2

---

Kube-OVN

Kube-OVN

### 7.2.1

---

1.	CNI	CNI
2.	10ns	20ns
3.	CPU	CPU
4.	Macvlan	SR-IOV

### 7.2.2

---

#### Overlay

- Kubernetes: 1.22.0
- OS: CentOS 7
- Kube-OVN: 1.8.0 *Overlay*
- CPU: Intel(R) Xeon(R) E-2278G
- Network: 2\*10Gbps, xmit\_hash\_policy=layer3+4

```
qperf -t 60 <server ip> -ub -oo msg_size:1 -vu tcp_lat tcp_bw udp_lat udp_bw 1 tcp/udp
```

Type	tcp_lat (us)	udp_lat (us)	tcp_bw (Mb/s)	udp_bw(Mb/s)
Kube-OVN Default	25.7	22.9	27.1	1.59
Kube-OVN Optimized	13.9	12.9	27.6	5.57
HOST Network	13.1	12.4	28.2	6.02

#### Overlay Underlay

Kube-OVN      Overlay    Underlay

*Environment:*

- Kubernetes: 1.22.0
- OS: CentOS 7
- Kube-OVN: 1.8.0
- CPU: AMD EPYC 7402P 24-Core Processor
- Network: Intel Corporation Ethernet Controller XXV710 for 25GbE SFP28

```
qperf -t 60 <server ip> -ub -oo msg_size:1 -vu tcp_lat tcp_bw udp_lat udp_bw
```

Type	tcp_lat (us)	udp_lat (us)	tcp_bw (Mb/s)	udp_bw(Mb/s)
Kube-OVN Overlay	15.2	14.6	23.6	2.65
Kube-OVN Underlay	14.3	13.8	24.2	3.46
HOST Network	16.6	15.4	24.8	2.64

```
qperf -t 60 <server ip> -ub -oo msg_size:1K -vu tcp_lat tcp_bw udp_lat udp_bw
```

Type	tcp_lat (us)	udp_lat (us)	tcp_bw (Gb/s)	udp_bw(Gb/s)
Kube-OVN Overlay	16.5	15.8	10.2	2.77
Kube-OVN Underlay	15.9	14.5	9.6	3.22
HOST Network	18.1	16.6	9.32	2.66

```
qperf -t 60 <server ip> -ub -oo msg_size:4K -vu tcp_lat tcp_bw udp_lat udp_bw
```

Type	tcp_lat (us)	udp_lat (us)	tcp_bw (Gb/s)	udp_bw(Gb/s)
Kube-OVN Overlay	34.7	41.6	16.0	9.23
Kube-OVN Underlay	32.6	44	15.1	6.71
HOST Network	35.9	45.9	14.6	5.59

netfilter      kube-proxy      netfilter

## 7.2.3

### CPU

CPU

CPU

```
cpupower frequency-set -g performance
```

```
ethtool -g eno1
Ring parameters for eno1:
Pre-set maximums:
RX: 4096
RX Mini: 0
RX Jumbo: 0
TX: 4096
Current hardware settings:
RX: 255
RX Mini: 0
RX Jumbo: 0
TX: 255
```

```
ethtool -G eno1 rx 4096
ethtool -G eno1 tx 4096
```

**tuned**

tuned profile

tuned-adm profile network-latency

tuned-adm profile network-throughput

irqbalance CPU CPU

**OVN LB**

OVN L2 LB	conntrack	recirculate	CPU	20%	CPU	Overlay	kube-proxy	Service
Pod-to-Pod	kube-ovn-controller							

```
command:
- ./kube-ovn/start-controller.sh
args:
...
- --enable-lb=false
...
```

Underlay kube-proxy iptables ipvs LB Service

**FastPath**

network ns	netfilter	20% CPU	netfilter	FastPath	netfilter CPU
------------	-----------	---------	-----------	----------	---------------

netfilter iptables ipvs nftables

**FastPath**

insmod kube\_ovn\_fastpath.ko FastPath dmesg

```
dmesg
...
[619631.323788] init_module,kube_ovn_fastpath_local_out
[619631.323798] init_module,kube_ovn_fastpath_post_routing
[619631.323800] init_module,kube_ovn_fastpath_pre_routing
[619631.323801] init_module,kube_ovn_fastpath_local_in
...
```

**OVS**

OVS flow	10% CPU	x86 CPU	popcnt sse4.2	flow CPU
5%				

**FastPath**

CPU

```
cat /proc/cpuinfo | grep popcnt
cat /proc/cpuinfo | grep sse4_2
```

## CENTOS

```
yum install -y gcc kernel-devel-$(uname -r) python3 autoconf automake libtool rpm-build openssl-devel
```

OVS      RPM :

```
git clone -b branch-3.5 --depth=1 https://github.com/openvswitch/ovs.git
cd ovs
curl -s https://github.com/kubeovn/ovs/commit/2d2c83c26d4217446918f39d5cd5838e9ac27b32.patch | git apply
./boot.sh
./configure --with-linux=/lib/modules/$(uname -r)/build CFLAGS="-g -O2 -mpopcnt -msse4.2"
make rpm-fedora-kmod
cd rpm/rpmbuild/RPMS/x86_64/
```

RPM

```
rpm -i openvswitch-kmod-3.5.1-1.el7.x86_64.rpm
```

Kube-OVN    OVS

## UBUNTU

```
apt install -y autoconf automake libtool gcc build-essential libssl-dev
```

OVS

```
apt install -y autoconf automake libtool gcc build-essential libssl-dev

git clone -b branch-3.5 --depth=1 https://github.com/openvswitch/ovs.git
cd ovs
curl -s https://github.com/kubeovn/ovs/commit/2d2c83c26d4217446918f39d5cd5838e9ac27b32.patch | git apply
./boot.sh
./configure --prefix=/usr/ --localstatedir=/var --enable-ssl --with-linux=/lib/modules/$(uname -r)/build
make -j `nproc`
make install
make modules_install

cat > /etc/modprobe.d/openvswitch.conf << EOF
override openvswitch * extra
override vport-* * extra
EOF

depmod -a
cp debian/openvswitch-switch.init /etc/init.d/openvswitch-switch
/etc/init.d/openvswitch-switch force-reload-kmod
```

Kube-OVN    OVS

## STT



**Warning**

OpenVswitch    3.6      STT    Tunnel

CPU	Geneve	Vxlan	UDP	UDP	TCP	TCP	offload	TCP
STT		TCP		TCP		TCP		TCP

STT                        OVS                        OVS

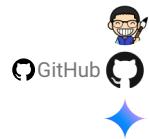
STT

```
kubectl set env daemonset/ovs-ovn -n kube-system TUNNEL_TYPE=stt
kubectl delete pod -n kube-system -lapp=ovs
```



⌚2025 9 10

⌚2022 5 24



---

7.2.4

## 7.3 FastPath

Profile Netfilter

20% CPU FastPath

Netfilter CPU

FastPath

### 7.3.1

```
git clone --depth=1 https://github.com/kubeovn/kube-ovn.git
```

### 7.3.2

CentOS

```
yum install -y kernel-devel-$(uname -r) gcc elfutils-libelf-devel
```

### 7.3.3

3.x

```
cd kube-ovn/fastpath
make all
```

4.x

```
cd kube-ovn/fastpath/4.18
cp ../Makefile .
make all
```

### 7.3.4

```
kube_ovn_fastpath.ko /tmp kube-ovn-cni
```

dmesg

```
dmesg
[619631.323788] init_module,kube_ovn_fastpath_local_out
[619631.323798] init_module,kube_ovn_fastpath_post_routing
[619631.323800] init_module,kube_ovn_fastpath_pre_routing
[619631.323801] init_module,kube_ovn_fastpath_local_in
```

```
/tmp kube-ovn-cni
```



PDF



Slack



Support

⌚2025 9 10

⌚2022 5 24



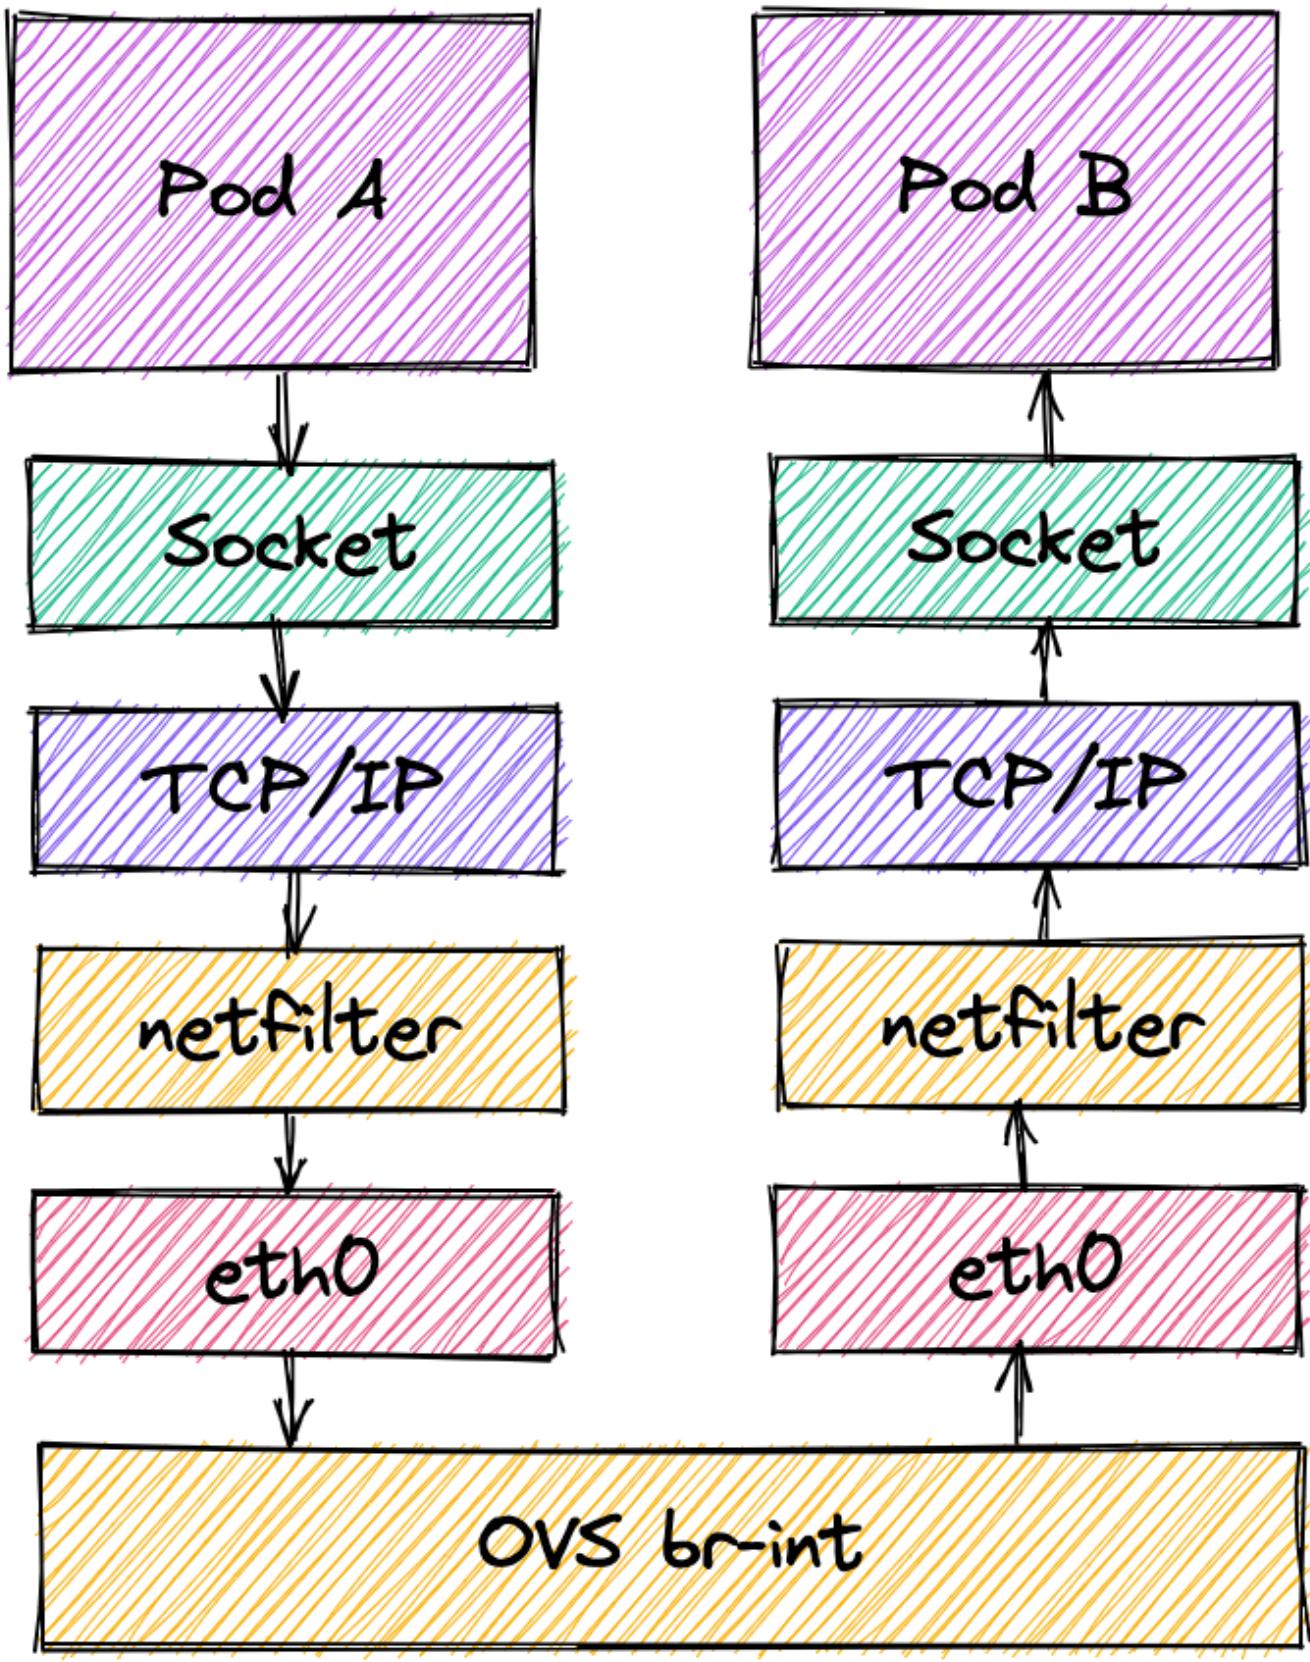
### 7.3.5

## 7.4 eBPF TCP

5G Pod TCP Intel **istio-tcpip-bypass** Pod eBPF TCP/IP socket

### 7.4.1

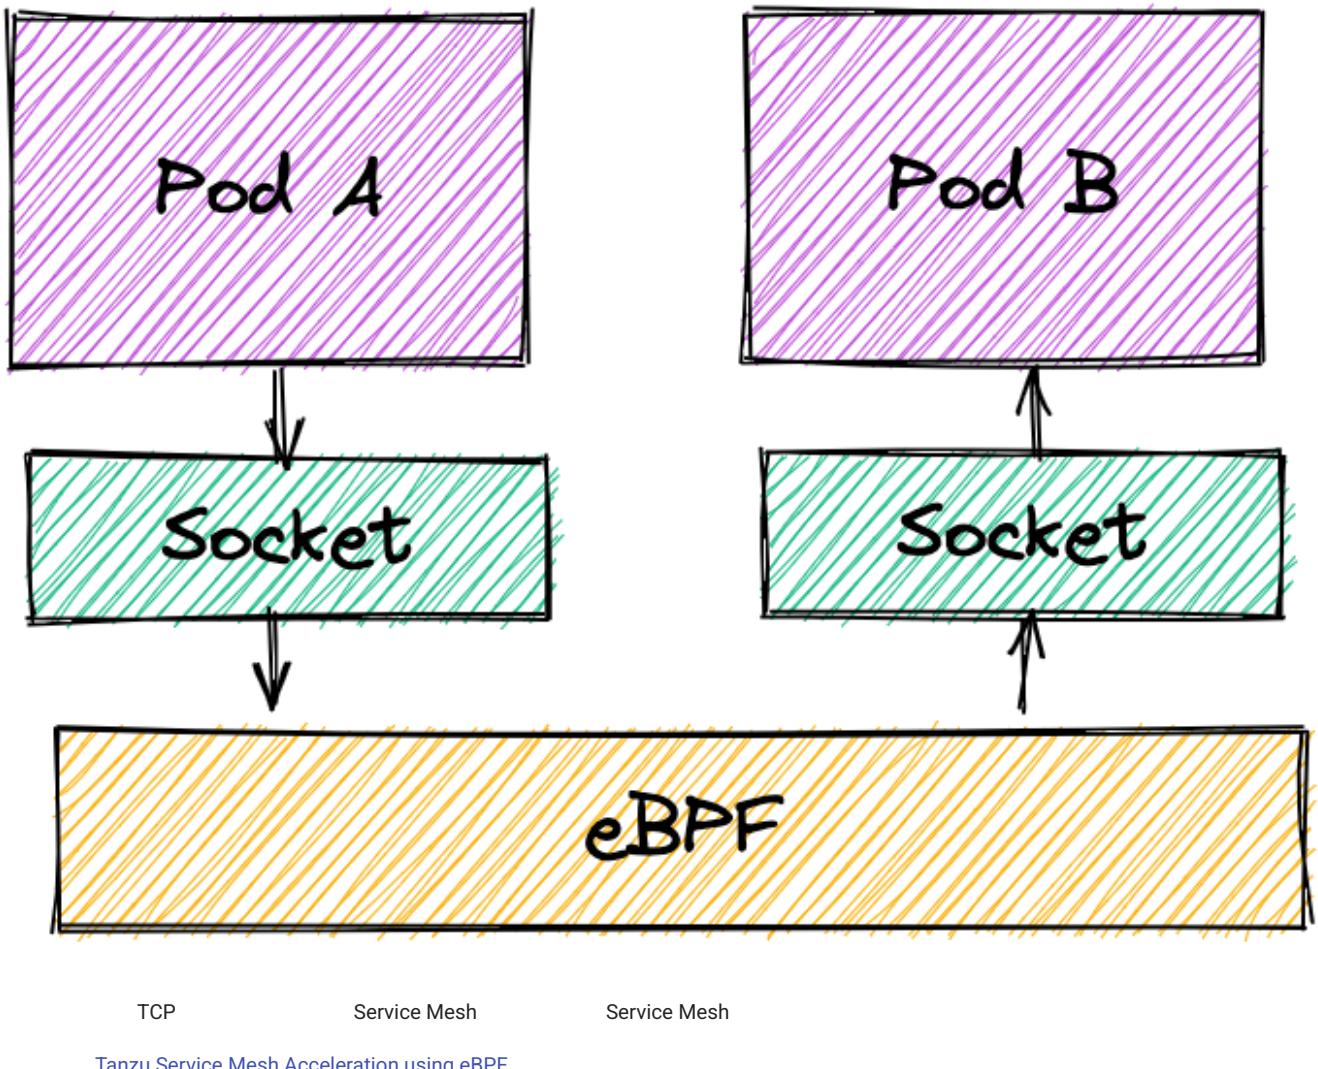
Pod TCP TCP/IP, netfilter, OVS



istio-tcpip-bypass

TCP

socket



## 7.4.2

eBPF      Ubuntu 20.04    Linux 5.4.0-74-generic

## 7.4.3

Pod                  nodeSelector

```
kubectl create deployment perf --image=kubeovn/perf:dev --replicas=2
deployment.apps/perf created
kubectl get pod -o wide
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES
perf-7697bc6ddf-b2cpv 1/1 Running 0 28s 100.64.0.3 sealos <none> <none>
perf-7697bc6ddf-p2xpt 1/1 Running 0 28s 100.64.0.2 sealos <none> <none>
```

Pod qperf server    Pod qperf client

```
kubectl exec -it perf-7697bc6ddf-b2cpv sh
/ # qperf

kubectl exec -it perf-7697bc6ddf-p2xpt sh
/ # qperf -t 60 100.64.0.3 -ub -oo msg_size:1:16K:*4 -vu tcp_lat tcp_bw
```

istio-tcpip-bypass

```
kubectl apply -f https://raw.githubusercontent.com/intel/istio-tcpip-bypass/main/bypass-tcpip-daemonset.yaml
```

### perf client

```
kubectl exec -it perf-7697bc6ddf-p2xpt sh
/ # qperf -t 60 100.64.0.3 -ub -oo msg_size:1:16K:*4 -vu tcp_lat tcp_bw
```

#### 7.4.4

TCP            40% ~ 60%            1024            40% ~ 80%

Packet Size (byte)	eBPF tcp_lat (us)	Default tcp_lat (us)	eBPF tcp_bw (Mb/s)	Default tcp_bw(Mb/s)
1	20.2	44.5	1.36	4.27
4	20.2	48.7	5.48	16.7
16	19.6	41.6	21.7	63.5
64	18.8	41.3	96.8	201
256	19.2	36	395	539
1024	18.3	42.4	1360	846
4096	16.5	62.6	4460	2430
16384	20.2	58.8	9600	6900

512            eBPF            TCP            eBPF            eBPF TCP

#### 7.4.5

1. [istio-tcpip-bypass](#)
2. [Deep Dive TCP/IP Bypass with eBPF in Service Mesh](#)
3. [Tanzu Service Mesh Acceleration using eBPF](#)

[!\[\]\(84348fb9a6f94c21b366d9a45d06d619\_img.jpg\) PDF](#)

[!\[\]\(e515b70648c87fb3e6c87966e4dbe96b\_img.jpg\) Slack](#)

[!\[\]\(9651bdf58a2c7ea87962f5ada9f6bdfa\_img.jpg\) Support](#)

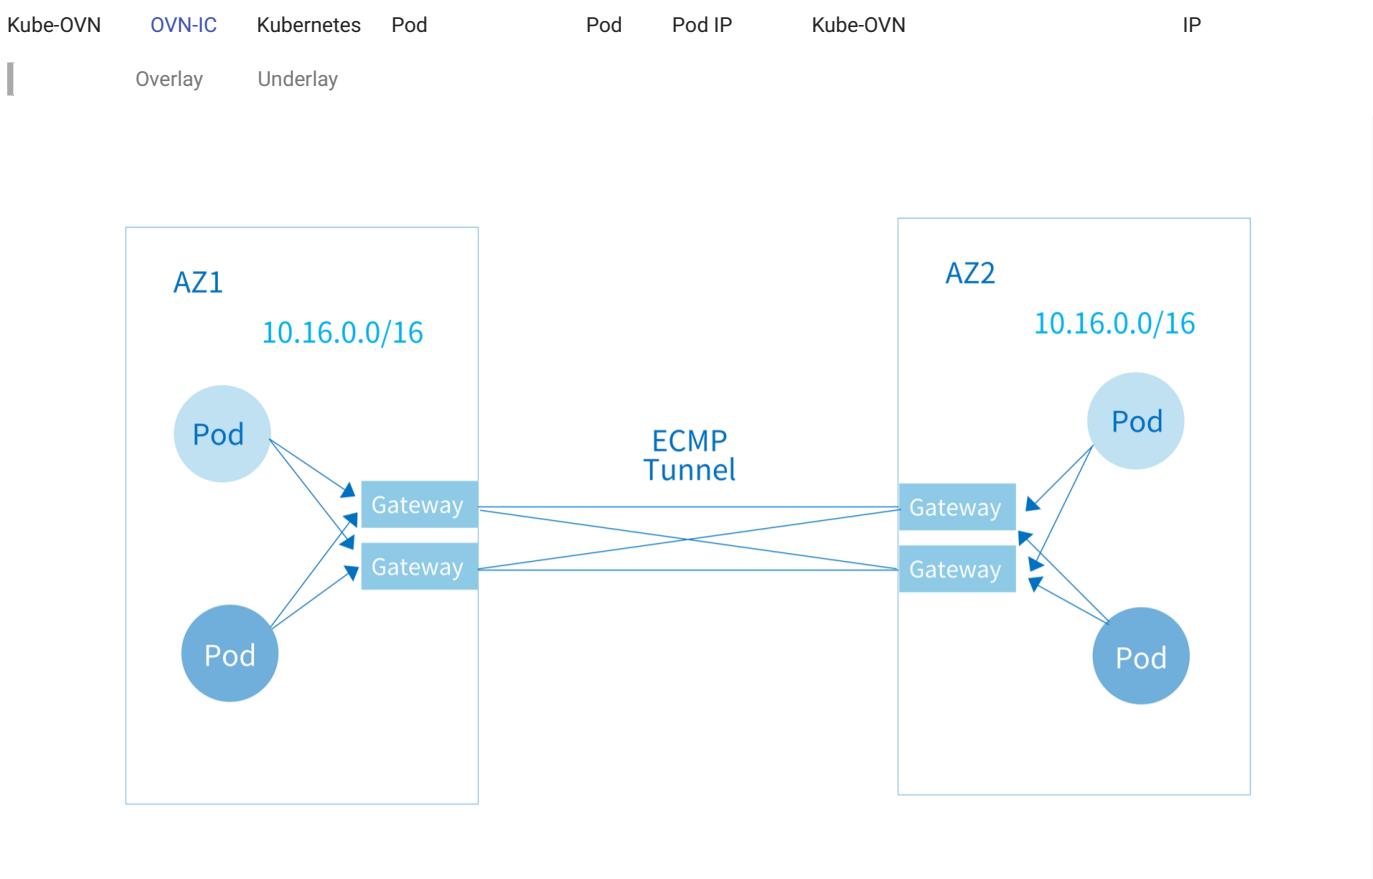
 2025 9 10

 2022 9 27



#### 7.4.6

## 7.5 OVN-IC



## Limitation

OVN-IC	Pod IP	Service	DNS	NetworkPolicy	Istio
--------	--------	---------	-----	---------------	-------

### 7.5.1

1. 1.11.16 install.sh

ENABLE\_IC=true

## deployment ovn-ic-controller

## 2. CIDR

### 3. kube-ovn-controller IP

4. IP

5. VPC VPC

## 7.5.2 OVN-IC

1

1 Kube-OVN v1.11.16

" " " " Deployment master master 1 master

```
install-ovn-ic.sh
```

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install-ic-server.sh
```

TS_NUM	ECMP Path
--------	-----------

```
sed 's/VERSION=.*/VERSION=v1.14.9/' dist/images/install-ic-server.sh | TS_NUM=3 bash
```

```
deployment.apps/ovn-ic-server created
Waiting for deployment spec update to be observed...
Waiting for deployment "ovn-ic-server" rollout to finish: 0 out of 3 new replicas have been updated...
Waiting for deployment "ovn-ic-server" rollout to finish: 0 of 3 updated replicas are available...
Waiting for deployment "ovn-ic-server" rollout to finish: 1 of 3 updated replicas are available...
Waiting for deployment "ovn-ic-server" rollout to finish: 2 of 3 updated replicas are available...
Waiting for deployment "ovn-ic-server" rollout to finish: 3 of 3 updated replicas are available...
deployment "ovn-ic-server" successfully rolled out
OVN IC Server installed Successfully
```

```
kubectl ko icsbctl show
```

```
kubectl ko icsbctl show
availability-zone az0
gateway 059b5c54-c540-4d77-b009-02d65f181a02
 hostname: kube-ovn-worker
 type: geneve
 ip: 172.18.0.3
 port ts-az0
 transit switch: ts
 address: ["00:00:00:B4:8E:BE 169.254.100.97/24"]
gateway 74ee4b9a-ba48-4a07-861e-1a8e4b9f905f
 hostname: kube-ovn-worker2
 type: geneve
 ip: 172.18.0.2
 port ts1-az0
 transit switch: ts1
 address: ["00:00:00:19:2E:F7 169.254.101.90/24"]
gateway 7e2428b6-344c-4dd5-a0d5-972c1cc581
 hostname: kube-ovn-control-plane
 type: geneve
 ip: 172.18.0.4
 port ts2-az0
 transit switch: ts2
 address: ["00:00:00:EA:32:BA 169.254.102.103/24"]
availability-zone az1
gateway 034da7cb-3826-4318-81ce-6a877a9bf285
 hostname: kube-ovn1-worker
 type: geneve
 ip: 172.18.0.6
 port ts-az1
 transit switch: ts
 address: ["00:00:00:25:3A:B9 169.254.100.51/24"]
gateway 2531a683-283e-4fb8-a619-bdbcb33539b8
 hostname: kube-ovn1-worker2
 type: geneve
 ip: 172.18.0.5
 port ts1-az1
 transit switch: ts1
 address: ["00:00:00:52:87:F4 169.254.101.118/24"]
gateway b0efb0be-e5a7-4323-ad4b-317637a757c4
 hostname: kube-ovn1-control-plane
 type: geneve
 ip: 172.18.0.8
 port ts2-az1
 transit switch: ts2
 address: ["00:00:00:F6:93:1A 169.254.102.17/24"]
```

## 2

kube-ovn-controller	IP	OVN-IC
---------------------	----	--------

docker		OVN-IC
--------	--	--------

```
docker run --name=ovn-ic-db -d --env "ENABLE_OVN_LEADER_CHECK=false" --network=host --privileged -v /etc/ovn/:/etc/ovn -v /var/run/ovn:/var/run/ovn -v /var/log/ovn:/var/log/ovn kubeovn/kube-ovn:v1.14.9 bash start-ic-db.sh
```

containerd	docker
------------	--------

```
ctr -n k8s.io run -d --env "ENABLE_OVN_LEADER_CHECK=false" --net-host --privileged --mount="type=bind,src=/etc/ovn/,dst=/etc/ovn,options=rbind:rw" --mount="type=bind,src=/var/run/ovn,dst=/var/run/ovn,options=rbind:rw" --mount="type=bind,src=/var/log/ovn,dst=/var/log/ovn,options=rbind:rw" docker.io/kubeovn/kube-ovn:v1.14.9 ovn-ic-db bash start-ic-db.sh
```

## 7.5.3

VPC	Subnet	CIDR	OVN-IC	Subnet CIDR
-----	--------	------	--------	-------------

kube-system Namespace    ovn-ic-config ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-ic-config
 namespace: kube-system
data:
 enable-ic: "true"
 az-name: "az1"
 ic-db-host: "192.168.65.3"
 ic-nb-port: "6645"
 ic-sb-port: "6646"
 gw-nodes: "az1-gw"
 auto-route: "true"
```

- enable-ic:
- az-name:
- ic-db-host: OVN-IC
- ic-nb-port: OVN-IC              6645
- ic-sb-port: OVN-IC              6646
- gw-nodes:
- auto-route:

ovn-ic-config	ConfigMap	ConfigMap	ConfigMap
---------------	-----------	-----------	-----------

ovn-ic              ts

```
ovn-ic-sbctl show
availability-zone az1
 gateway deee03e0-af16-4f45-91e9-b50c3960f809
 hostname: az1-gw
 type: geneve
 ip: 192.168.42.145
 port ts-az1
 transit switch: ts
 address: ["00:00:00:50:AC:8C 169.254.100.45/24"]
availability-zone az2
 gateway e94cc831-8143-40e3-a478-90352773327b
 hostname: az2-gw
 type: geneve
 ip: 192.168.42.149
 port ts-az2
 transit switch: ts
 address: ["00:00:00:07:4A:59 169.254.100.63/24"]
```

```
kubectl ko nbctl lr-route-list ovn-cluster
IPv4 Routes
 10.42.1.1 169.254.100.45 dst-ip (learned)
 10.42.1.3 100.64.0.2 dst-ip
 10.16.0.2 100.64.0.2 src-ip
 10.16.0.3 100.64.0.2 src-ip
 10.16.0.4 100.64.0.2 src-ip
 10.16.0.6 100.64.0.2 src-ip
 10.17.0.0/16 169.254.100.45 dst-ip (learned)
 100.65.0.0/16 169.254.100.45 dst-ip (learned)
```

1    Pod    ping    2    Pod IP

Subnet    disableInterConnection

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
```

```

 name: no-advertise
spec:
 cidrBlock: 10.199.0.0/16
 disableInterConnection: true

```

## 7.5.4

### CIDR

```
kube-system Namespace ovn-ic-config ConfigMap auto-route false
```

```

apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-ic-config
 namespace: kube-system
data:
 enable-ic: "true"
 az-name: "az1"
 ic-db-host: "192.168.65.3"
 ic-nb-port: "6645"
 ic-sb-port: "6646"
 gw-nodes: "az1-gw"
 auto-route: "false"

```

```

[root@az1 ~]# kubectl ko nbctl show
switch a391d3a1-14a0-4841-9836-4bd930c447fb (ts)
 port ts-az1
 type: router
 router-port: az1-ts
 port ts-az2
 type: remote
 addresses: ["00:00:00:4B:E2:9F 169.254.100.31/24"]

[root@az2 ~]# kubectl ko nbctl show
switch d46138b8-de81-4908-abf9-b2224ec4edf3 (ts)
 port ts-az2
 type: router
 router-port: az2-ts
 port ts-az1
 type: remote
 addresses: ["00:00:00:FB:2A:F7 169.254.100.79/24"]

```

```

az1 az2 169.254.100.31 az2 az1 169.254.100.79
 CIDR 10.16.0.0/24 az2 CIDR 10.17.0.0/24

```

```
az1 az2
```

```
kubectl ko nbctl lr-route-add ovn-cluster 10.17.0.0/24 169.254.100.31
```

```
az2 az1
```

```
kubectl ko nbctl lr-route-add ovn-cluster 10.16.0.0/24 169.254.100.79
```

## 7.5.5 OVN-IC

**1**

1 Kube-OVN v1.11.16

**1**

**2**

OVN-IC	Raft	3
OVN-IC	leader	
docker		

```
docker run --name=ovn-ic-db -d --env "ENABLE_OVN_LEADER_CHECK=false" --network=host --privileged -v /etc/ovn:/etc/ovn -v /var/run/ovn:/var/run/ovn -v /var/log/ovn:/var/log/ovn -e LOCAL_IP="192.168.65.3" -e NODE_IPS="192.168.65.3,192.168.65.2,192.168.65.1" kubeovn/kube-ovn:v1.14.9 bash start-ic-db.sh
```

containerd

```
ctr -n k8s.io run -d --env "ENABLE_OVN_LEADER_CHECK=false" --net-host --privileged --mount="type=bind,src=/etc/ovn/,dst=/etc/ovn,options=rbind:rw" --mount="type=bind,src=/var/run/ovn,dst=/var/run/ovn,options=rbind:rw" --mount="type=bind,src=/var/log/ovn,dst=/var/log/ovn,options=rbind:rw" --env="NODE_IPS=192.168.65.3,192.168.65.2,192.168.65.1" --env="LOCAL_IP=192.168.65.3" docker.io/kubeovn/kube-ovn:v1.14.9 ovn-ic-db bash start-ic-db.sh
```

- LOCAL\_IP IP
- NODE\_IPS OVN-IC IP
- OVN-IC follower

docker

```
docker run --name=ovn-ic-db -d --network=host --privileged -v /etc/ovn:/etc/ovn -v /var/run/ovn:/var/run/ovn -v /var/log/ovn:/var/log/ovn -e LOCAL_IP="192.168.65.2" -e NODE_IPS="192.168.65.3,192.168.65.2,192.168.65.1" -e LEADER_IP="192.168.65.3" kubeovn/kube-ovn:v1.14.9 bash start-ic-db.sh
```

containerd

```
ctr -n k8s.io run -d --net-host --privileged --mount="type=bind,src=/etc/ovn/,dst=/etc/ovn,options=rbind:rw" --mount="type=bind,src=/var/run/ovn,dst=/var/run/ovn,options=rbind:rw" --mount="type=bind,src=/var/log/ovn,dst=/var/log/ovn,options=rbind:rw" --env="NODE_IPS=192.168.65.3,192.168.65.2,192.168.65.1" --env="LOCAL_IP=192.168.65.2" --env="LEADER_IP=192.168.65.3" docker.io/kubeovn/kube-ovn:v1.14.9 ovn-ic-db bash start-ic-db.sh
```

- LOCAL\_IP IP
- NODE\_IPS OVN-IC IP
- LEADER\_IP: OVN-IC leader IP

ovn-ic-config OVN-IC

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-ic-config
 namespace: kube-system
data:
 enable-ic: "true"
 az-name: "az1"
 ic-db-host: "192.168.65.3,192.168.65.2,192.168.65.1"
 ic-nb-port: "6645"
 ic-sb-port: "6646"
 gw-nodes: "az1-gw"
 auto-route: "true"
```

## 7.5.6 ECMP

1

ECMP ECMP path 3 ECMP path

```
kubectl edit deployment ovn-ic-server -n kube-system
```

'TS\_NUM' TS\_NUM ECMP Path

## 7.5.7

ovn-ic-config Configmap

```
kubectl -n kube-system delete cm ovn-ic-config
```

ts

```
kubectl ko nbctl ls-del ts
```

### 7.5.8 az-name

```
kubectl edit ovn-ic-config configmap az-name ovn-cni pod 10
```

```
ovn-appctl -t ovn-controller inc-engine/recompute
```

### 7.5.9

ovn-ic-config Configmap

```
kubectl -n kube-system delete cm ovn-ic-config
```

ts

```
kubectl ko nbctl ls-del ts
```

OVN-IC

docker

```
docker stop ovn-ic-db
docker rm ovn-ic-db
```

containerd

```
ctr -n k8s.io task kill ovn-ic-db
ctr -n k8s.io containers rm ovn-ic-db
```

deployment ovn-ic-server

```
kubectl delete deployment ovn-ic-server -n kube-system
```

master

DB

```
rm -f /etc/origin/ovn/ovn_ic_nb_db.db
rm -f /etc/origin/ovn/ovn_ic_sb_db.db
```

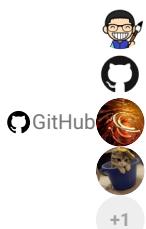
[PDF](#)

[Slack](#)

[Support](#)

⌚2025 6 19

⌚2022 5 24



### 7.5.10

## 7.6 Submariner



### 7.6.1

- Service CIDR CIDR

### 7.6.2 Submariner

```

subctl

curl -Ls https://get.submariner.io | bash
export PATH=$PATH:~/local/bin
echo export PATH=\$PATH:~/local/bin >> ~/.profile

kubeconfig submariner-broker

subctl deploy-broker

cluster0 CIDR 10.16.0.0/16 join CIDR 100.64.0.0/16 cluster1 CIDR 11.16.0.0/16 join CIDR 100.68.0.0/16

kubeconfig cluster0 broker :

subctl join broker-info.subm --clusterid cluster0 --clustercidr 100.64.0.0/16,10.16.0.0/16 --natt=false --cable-driver vxlan --health-check=false
kubectl label nodes cluster0 submariner.io/gateway=true

kubeconfig cluster1 broker :

subctl join broker-info.subm --clusterid cluster1 --clustercidr 100.68.0.0/16,11.16.0.0/16 --natt=false --cable-driver vxlan --health-check=false
kubectl label nodes cluster1 submariner.io/gateway=true

join gateway, routeagent pod , submariner-operator clusterrole :

- apiGroups:
 - "apps"
 resources:
 - daemonsets
 verbs:
 - create
 - get
 - list
 - watch
 - update

subnet ovn-default centralized submariner gateway subnet

Pod IP

subctl

subctl show all
subctl diagnose all

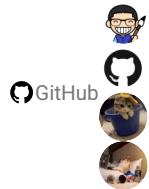
Submariner Submariner

```

[PDF](#)  [Slack](#)  [Support](#)

⌚2025 9 10

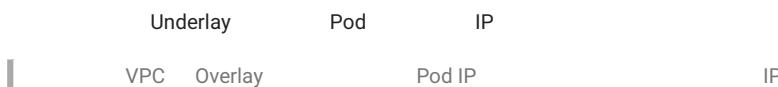
⌚2022 5 24



7.6.3

---

## 7.7 Overlay



### 7.7.1

- ip\_forward
- iptables forward Drop
- ct INVALID

### 7.7.2

natOutgoing false nat Pod IP

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: routed
spec:
 protocol: IPv4
 cidrBlock: 10.166.0.0/16
 default: false
 excludeIps:
 - 10.166.0.1
 gateway: 10.166.0.1
 gatewayType: distributed
 natOutgoing: false

```

Pod

Kubernetes

```
ip route add 10.166.0.0/16 via 192.168.2.10 dev eth0
```

10.166.0.0/16 192.168.2.10 Kubernetes

IP	Keepalived	VIP	VIP	
Subnet	gatewayType	centralized	gatewayNode	IP

```

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: routed
spec:
 protocol: IPv4
 cidrBlock: 10.166.0.0/16
 default: false
 excludeIps:
 - 10.166.0.1
 gateway: 10.166.0.1
 gatewayType: centralized
 gatewayNode: "node1"
 natOutgoing: false

```

nat VPC NAT

[PDF](#)

[Slack](#)

[Support](#)

⌚2023 12 25

⌚2022 7 6

⌚GitHub 🧑‍💻

7.7.3

---

## 7.8 BGP



### 7.8.1 kube-ovn-speaker

kube-ovn-speaker    GoBGP

`kube-ovn-speaker`

```
kubectl label nodes speaker-node-1 ovn.kubernetes.io/bgp=true
kubectl label nodes speaker-node-2 ovn.kubernetes.io/bgp=true
```

`kube-ovn-speaker`

ECMP

`yaml:`

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/yamls/speaker.yaml
```

`yaml`

```
- --neighbor-address=10.32.32.254
- --neighbor-ipv6-address=2409:AB00:AB00:2000::AFB:8AFE
- --neighbor-as=65030
- --cluster-as=65000
```

```
- --neighbor-address=10.32.32.252,10.32.32.253
- --neighbor-ipv6-address=2409:AB00:AB00:2000::AFB:8AFC,2409:AB00:AB00:2000::AFB:8AFD
- --neighbor-as=65030
- --cluster-as=65000
```

- `neighbor-address`: BGP Peer
- `neighbor-as`: BGP Peer AS
- `cluster-as`: AS

`yaml:`

```
kubectl apply -f speaker.yaml
```

### 7.8.2 Pod/Subnet

BGP              Subnet    natOutgoing    false    Pod IP

`annotation`

```
kubectl annotate pod sample ovn.kubernetes.io/bgp=true
kubectl annotate subnet ovn-default ovn.kubernetes.io/bgp=true
```

`annotation`

```
kubectl annotate pod sample ovn.kubernetes.io/bgp-
kubectl annotate subnet ovn-default ovn.kubernetes.io/bgp-
```

`BGP`

### 7.8.3 ClusterIP Service

```
Service ClusterIP kube-ovn-speaker --announce-cluster-ip true BGP
```

annotation

```
kubectl annotate service sample ovn.kubernetes.io/bgp=true
```

annotation

```
kubectl annotate service sample ovn.kubernetes.io/bgp-
```

### 7.8.4 EIPs

EIPs	VPC NAT Gateway	VpcNatGateway	BGP	BGP Sidecar			
VPC NAT Gateway	BGP	BGP Speaker Sidecar	NetworkAttachmentDefinition	NAD	VPC	Subnet	Sidecar
Kubernetes API	EIPs	VPC	CoreDNS	NAD			

NetworkAttachmentDefinition Subnet provider {nadName}.{nadNamespace}.ovn

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: api-ovn-nad
 namespace: default
spec:
 config: '{'
 "cniVersion": "0.3.0",
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "api-ovn-nad.default.ovn"
 '}'

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: vpc-apiserver-subnet
spec:
 protocol: IPv4
 cidrBlock: 100.100.100.0/24
 provider: api-ovn-nad.default.ovn
```

ovn-vpc-nat-config ConfigMap apiNadProvider BGP Speaker :

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-vpc-nat-config
 namespace: kube-system
data:
 apiNadProvider: api-ovn-nad.default.ovn # What NetworkAttachmentDefinition provider to use so that the sidecar
 # can access the K8S API, as it can't by default due to VPC segmentation
 bgpSpeakerImage: docker.io/kubeovn/kube-ovn:v1.13.0 # Sets the BGP speaker image used
 image: docker.io/vpc-nat-gateway:v1.13.0
```

ovn-default provider

```
provider: api-ovn-nad.default.ovn
```

VPC NAT Gateway BGP

```
kind: VpcNatGateway
apiVersion: kubeovn.io/v1
metadata:
 name: vpc-natgw
spec:
 vpc: vpc1
 subnet: net1
 lanIp: 10.0.1.10
 bgpSpeaker:
 enabled: true
 asn: 65500
 remoteAsn: 65000
 neighbors:
 - 100.127.4.161
 - fd:01::1
 enableGracefulRestart: true # Optional
```

```

routerId: 1.1.1.1 # Optional
holdTime: 1m # Optional
password: "password123" # Optional
extraArgs: # Optional, passed directly to the BGP speaker
 - -v5 # Enables verbose debugging of the BGP speaker sidecar
selector:
 - "kubernetes.io/os: linux"
externalSubnets:
 - ovn-vpc-external-network # Network on which we'll speak BGP and receive/send traffic to the outside world
 # BGP neighbors need to be on that network

```

BGP EIP

```
kubectl annotate eip sample ovn.kubernetes.io/bgp=true
```

## 7.8.5

kube-ovn-speaker											
• Cluster:		speaker	Pod IPs/Subnet CIDRs		IP	CIDR	Pod	speaker	Pod		
		Pod	Subnet								
• Local:		Pod IPs	Pod		Cluster		Pod				
	:	Local		kube-ovn-speaker	Pod		speaker				
		Cluster	Pod/Subnet	annotation ovn.kubernetes.io/bgp							
• ovn.kubernetes.io/bgp=cluster		ovn.kubernetes.io/bgp=yes		Cluster							
• ovn.kubernetes.io/bgp=local		Local									
		Service	kube-proxy	ClusterIP	Service		Cluster				

## 7.8.6 BGP

kube-ovn-speaker BGP		
• announce-cluster-ip:	Service	false
• auth-password:	BGP peer	
• holdtime:	BGP	90
• graceful-restart:	BGP Graceful Restart	
• graceful-restart-time:	BGP Graceful restart time	RFC4724 3
• graceful-restart-deferral-time:	BGP Graceful restart deferral time	RFC4724 4.1
• passivemode:	Speaker passive	peer
• ebgp-multipath:	ebgp ttl	1

## 7.8.7 BGP routes debug

```

show peer neighbor
gobgp neighbor

show announced routes to one peer
gobgp neighbor 10.32.32.254 adj-out

```

[PDF](#)
[Slack](#)
[Support](#)

⌚2025 9 10

⌚2022 6 14



7.8.8

---

## 7.9 MetalLB Kube-OVN Underlay

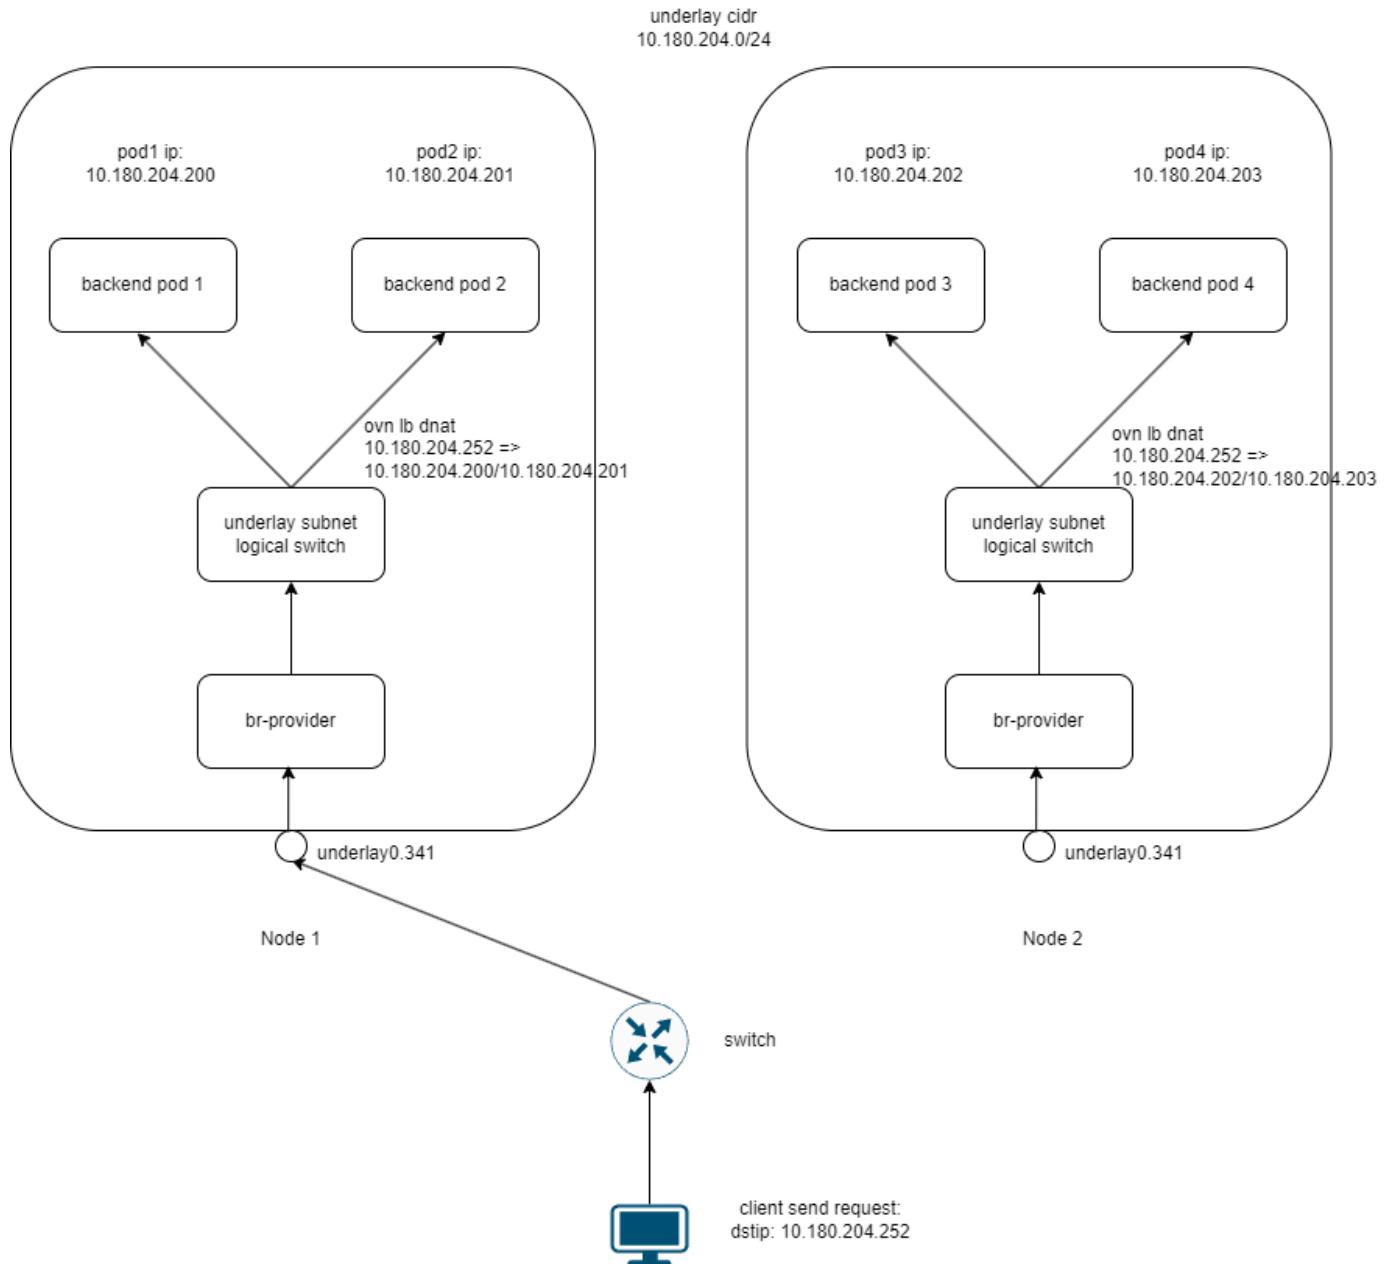


### 7.9.1

Kube-OVN 1.14.0 MetalLB Underlay

- MetalLB IP
- Pod MetalLB VIP Underlay
- IP SNAT

### 7.9.2



### 1 MetalLB VIP Kube-OVN Underlay

MetalLB Kube-OVN Underlay

```

1. VIP 10.180.204.252 IP MetalLB L2 Node1 metallb VIP
2. VIP underlay0.341
3. br-provider Underlay
4. br-provider OpenFlow OVN
5. underlay subnet OVN ovn lb dnat
6. OVN Pod
10.180.204.0/24 VIP Pod IP

```

### 7.9.3

- Kube-OVN --enable-ovn-lb-prefer-local=true
- Underlay enableExternalLBAddress=true
- Underlay excludeIps MetalLB IP

### 7.9.4

#### 1. Kube-OVN

```
Kube-OVN Kube-OVN --enable-ovn-lb-prefer-local=true --ls-ct-skip-dst-lport-ips=false
```

```
kube-ovn-controller Deployment
kubectl edit deployment -n kube-system kube-ovn-controller
```

```
--enable-ovn-lb-prefer-local=true
--ls-ct-skip-dst-lport-ips=false
```

#### 2. Underlay

```
Underlay LoadBalancer excludeIps MetalLB IP
```

```
apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: underlay-subnet
spec:
 protocol: IPv4
 provider: ovn
 cidrBlock: 10.180.204.0/24 #
 gateway: 10.180.204.1
 excludeIps:
 - 10.180.204.250
 - 10.180.204.251
 - 10.180.204.252 # MetalLB
 natOutgoing: false
 enableExternalLBAddress: true # subnet cidr ip metallb vip
```

#### 3. MetalLB

```
MetalLB MetalLB
```

```
kubectl apply -f https://raw.githubusercontent.com/metallb/metallb/v0.13.7/config/manifests/metallb-native.yaml
```

```
MetalLB L2
```

```
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
```

```

 name: underlay-pool
 namespace: metallb-system
spec:
 addresses:
 - 10.180.204.250-10.180.204.254 # VIP 10.180.204.252

apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
 name: l2-advert
 namespace: metallb-system
spec:
 ipAddressPools:
 - underlay-pool

```

#### 4. LoadBalancer Service

LoadBalancer    Service    Underlay    Pod

```

apiVersion: apps/v1
kind: Deployment
metadata:
 labels:
 app: nginx
 name: deploy-16940264
spec:
 replicas: 3
 selector:
 matchLabels:
 app: nginx
 template:
 metadata:
 annotations:
 ovn.kubernetes.io/logical_switch: underlay-subnet
 labels:
 app: nginx
 spec:
 containers:
 - args:
 - netexec
 - --http-port
 - "80"
 image: kubeovn/agnhost:2.47
 imagePullPolicy: IfNotPresent
 name: nginx

apiVersion: v1
kind: Service
metadata:
 name: nginx-lb
spec:
 externalTrafficPolicy: Local
 ipFamilies:
 - IPv4
 ipFamilyPolicy: PreferDualStack
 ports:
 - port: 80
 protocol: TCP
 targetPort: 80
 selector:
 app: nginx
 type: LoadBalancer

```

### 7.9.5

1. Service    MetalLB    IP

```
kubectl get svc nginx-lb
```

EXTERNAL-IP    IP    10.180.204.252

1.    Service    IP

```
curl http://10.180.204.252
```

1.              Pod

Service    endpoints    Pod

```
Service endpoints
kubectl get endpoints nginx-lb
```

```
Pod
kubectl get pods -l app=nginx -o wide
```

## 1. IP

nginx Pod IP IP SNAT IP

```
kubectl exec -it $(kubectl get pods -l app=nginx -o name | head -n1) -- cat /var/log/nginx/access.log
```

## 7.9.6

IP

MetalLB	Underlay	CIDR	Underlay	excludeIps	IP
---------	----------	------	----------	------------	----

MetalLB	Kube-OVN Underlay	underlay0.341	VLAN	VLAN	ARP	MetalLB VIP
---------	-------------------	---------------	------	------	-----	-------------

<code>- Kube-OVN</code>	<code>--enable-ovn-lb-prefer-local=true</code>	<code>- Service</code>	<code>externalTrafficPolicy: Local</code>
-------------------------	------------------------------------------------	------------------------	-------------------------------------------

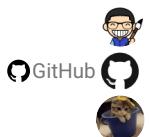
PDF

Slack

Support

2025 9 10

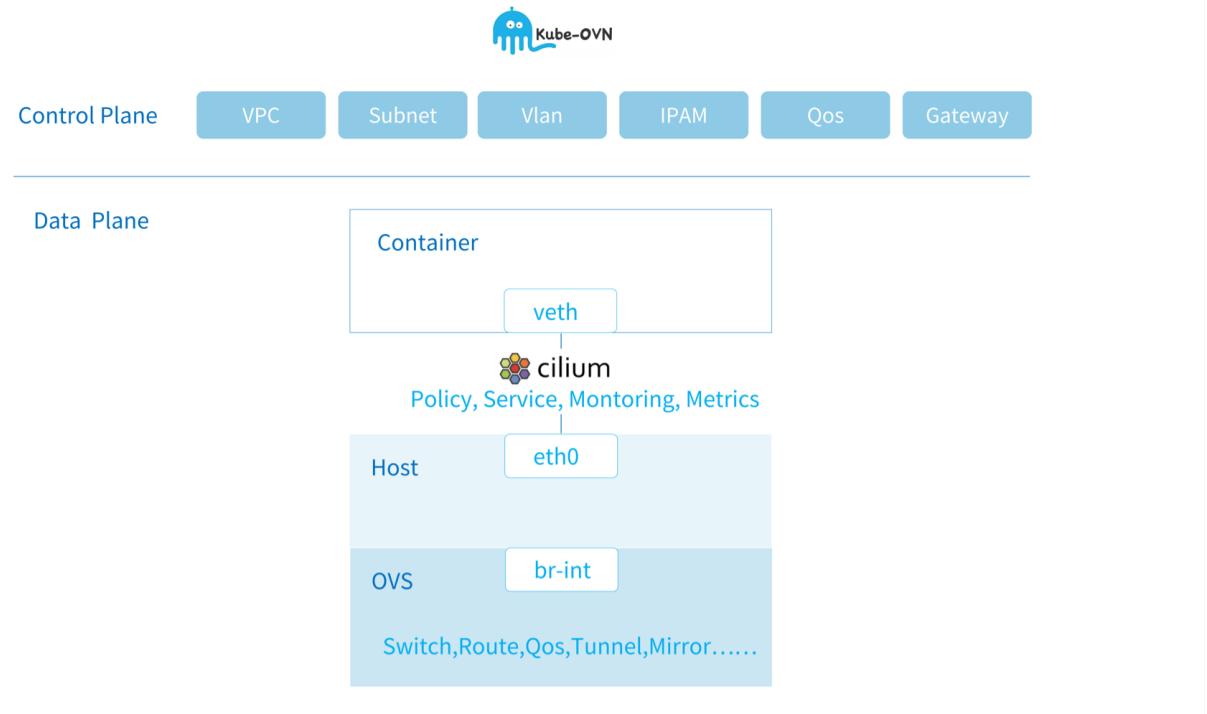
2025 4 2



## 7.9.7

## 7.10 Cilium

- Cilium eBPF Kube-OVN CNI Chaining Kube-OVN eBPF
- Cilium Kube-OVN
- Hubble



### 7.10.1

1. Linux 4.19 eBPF
2. Helm Cilium Helm [Installing Helm](#)

### 7.10.2 Kube-OVN

Cilium Kube-OVN networkpolicy CNI

```
install.sh
```

```
ENABLE_NP=false
CNI_CONFIG_PRIORITY=10
```

```
kube-ovn-controller networkpolicy
```

```
args:
- --enable-np=false
```

```
kube-ovn-cni CNI
```

```
args:
- --cni-conf-name=10-kube-ovn.conflist
```

Kube-OVN

Cilium

```
mv /etc/cni/net.d/01-kube-ovn.conflist /etc/cni/net.d/10-kube-ovn.conflist
```

### 7.10.3 Cilium

chaining.yaml Cilium generic-veth

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: cni-configuration
 namespace: kube-system
data:
 cni-config: |-
 {
 "name": "generic-veth",
 "cniVersion": "0.3.1",
 "plugins": [
 {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "ipam": {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock"
 }
 },
 {
 "type": "portmap",
 "snat": true,
 "capabilities": {"portMappings": true}
 },
 {
 "type": "cilium-cni"
 }
]
 }
```

```
kubectl apply -f chaining.yaml
```

Helm Cilium

```
helm repo add cilium https://helm.cilium.io/
helm install cilium cilium/cilium --version 1.11.6 \
 --namespace kube-system \
 --set cni.chainingMode=generic-veth \
 --set cni.customConf=true \
 --set cni.configMap=cni-configuration \
 --set tunnel=disabled \
 --set enableIPv4Masquerade=false \
 --set devices="eth+ ovn0 genev_sys_6081 vxlan_sys_4789" \
 --set enableIdentityMark=false
```

Cilium

```
cilium status
 /--\
 /--\ /--\ Cilium: OK
 \--/ \--/ Operator: OK
 /--\ /--\ Hubble: disabled
 \--/ \--\ ClusterMesh: disabled
 \--/

DaemonSet cilium Desired: 2, Ready: 2/2, Available: 2/2
Deployment cilium-operator Desired: 2, Ready: 2/2, Available: 2/2
Containers: cilium Running: 2
 cilium-operator Running: 2
Cluster Pods: 8/11 managed by Cilium
Image versions cilium quay.io/cilium/cilium:v1.10.5@sha256:0612218e28288db360c63677c09fafafa2d17edda4f13867bcabf87056046b33bb: 2
 cilium-operator quay.io/cilium/operator-generic:v1.10.5@sha256:2d2f730f219d489ff0702923bf24c0002cd93eb4b47ba344375566202f56d972: 2
```



PDF



Slack



Support

⌚2024 10 8

⌚2022 5 24



7.10.4

---

## 7.11 Cilium NetworkPolicy

Kube-OVN	Cilium	Cilium
Cilium	Cilium	Cilium L3 L4

### 7.11.1

#### Pod

namespace test yaml test namespace label app=test Pod Pod

```
apiVersion: apps/v1
kind: Deployment
metadata:
 labels:
 app: test
 name: test
 namespace: test
spec:
 replicas: 1
 selector:
 matchLabels:
 app: test
 strategy:
 rollingUpdate:
 maxSurge: 25%
 maxUnavailable: 25%
 type: RollingUpdate
 template:
 metadata:
 labels:
 app: test
 spec:
 containers:
 - image: docker.io/library/nginx:alpine
 imagePullPolicy: IfNotPresent
 name: nginx
```

yaml default namespace label app=dynamic Pod Pod

```
apiVersion: apps/v1
kind: Deployment
metadata:
 labels:
 app: dynamic
 name: dynamic
 namespace: default
spec:
 replicas: 2
 selector:
 matchLabels:
 app: dynamic
 strategy:
 rollingUpdate:
 maxSurge: 25%
 maxUnavailable: 25%
 type: RollingUpdate
 template:
 metadata:
 creationTimestamp: null
 labels:
 app: dynamic
 spec:
 containers:
 - image: docker.io/library/nginx:alpine
 imagePullPolicy: IfNotPresent
 name: nginx
```

Pod Label :

```
kubectl get pod -o wide --show-labels
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES LABELS
dynamic-7d8d7874f5-9v5c4 1/1 Running 0 28h 10.16.0.35 kube-ovn-worker <none> <none>
template-hash=7d8d7874f5
dynamic-7d8d7874f5-s822n 1/1 Running 0 28h 10.16.0.36 kube-ovn-control-plane <none> <none>
template-hash=7d8d7874f5
kubectl get pod -o wide -n test --show-labels
NAME READY STATUS RESTARTS AGE IP NODE NOMINATED NODE READINESS GATES LABELS
dynamic-7d8d7874f5-6ds96 1/1 Running 0 7h20m 10.16.0.2 kube-ovn-control-plane <none> <none>
template-hash=7d8d7874f5
```

```

dynamic-7d8d7874f5-tjgtp 1/1 Running 0 7h46m 10.16.0.42 kube-ovn-worker <none> <none> app=dynamic,pod-
template-hash=7d8d7874f5
label-test1-77b6764857-sq4k 1/1 Running 0 3h43m 10.16.0.12 kube-ovn-worker <none> <none> app=test1,pod-
template-hash=77b6764857

// Pod
test-54c98bc466-mft5s 1/1 Running 0 8h 10.16.0.41 kube-ovn-worker <none> <none> app=test,pod-
template-hash=54c98bc466

```

**L3**

`yaml CiliumNetworkPolicy :`

```

apiVersion: "cilium.io/v2"
kind: CiliumNetworkPolicy
metadata:
 name: "l3-rule"
 namespace: test
spec:
 endpointSelector:
 matchLabels:
 app: test
 ingress:
 - fromEndpoints:
 - matchLabels:
 app: dynamic

```

default namespace	Pod	Pod	test namespace	Pod
-------------------	-----	-----	----------------	-----

`default namespace :`

```

kubectl exec -it dynamic-7d8d7874f5-9v5c4 -- bash
bash-5.0# ping -c 3 10.16.0.41
PING 10.16.0.41 (10.16.0.41): 56 data bytes
--- 10.16.0.41 ping statistics ---
3 packets transmitted, 0 packets received, 100% packet loss

```

`test namespace Pod :`

```

kubectl exec -it -n test dynamic-7d8d7874f5-6dsg6 -- bash
bash-5.0# ping -c 3 10.16.0.41
PING 10.16.0.41 (10.16.0.41): 56 data bytes
64 bytes from 10.16.0.41: seq=0 ttl=64 time=2.558 ms
64 bytes from 10.16.0.41: seq=1 ttl=64 time=0.223 ms
64 bytes from 10.16.0.41: seq=2 ttl=64 time=0.304 ms

--- 10.16.0.41 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.223/1.028/2.558 ms

```

Cilium	CiliumNetworkPolicy	Namespace	Cilium
--------	---------------------	-----------	--------

Namespace	Pod	Namespace	Pod
-----------	-----	-----------	-----

Namespace	Namespace
-----------	-----------

`CiliumNetworkPolicy namespace :`

```

ingress:
- fromEndpoints:
 - matchLabels:
 app: dynamic
 k8s.io.kubernetes.pod.namespace: default // Namespace Pod

```

`CiliumNetworkPolicy :`

```

kubectl get cnp -n test -o yaml l3-rule
apiVersion: cilium.io/v2
kind: CiliumNetworkPolicy
metadata:
 name: l3-rule
 namespace: test
spec:
 endpointSelector:
 matchLabels:
 app: test
 ingress:
 - fromEndpoints:
 - matchLabels:

```

```
app: dynamic
- matchLabels:
 app: dynamic
k8s:io.kubernetes.pod.namespace: default
```

default namespace Pod Pod :

```
kubectl exec -it dynamic-7d8d7874f5-9v5c4 -n test -- bash
bash-5.0# ping -c 3 10.16.0.41
PING 10.16.0.41 (10.16.0.41): 56 data bytes
64 bytes from 10.16.0.41: seq=0 ttl=64 time=2.383 ms
64 bytes from 10.16.0.41: seq=1 ttl=64 time=0.115 ms
64 bytes from 10.16.0.41: seq=2 ttl=64 time=0.142 ms

--- 10.16.0.41 ping statistics ---
3 packets transmitted, 3 packets received, 0% packet loss
round-trip min/avg/max = 0.115/0.880/2.383 ms
```

Kubernetes	<a href="#">networkpolicy</a>	Cilium	Namespace	Namespace	Pod	Namespace	Namespace	Pod
Kube-OVN	Kube-OVN	k8s	Namespace	<b>Pod</b>	Pod	Namespace	Namespace	Pod
Pod								

#### L4

yaml L4 :

```
apiVersion: "cilium.io/v2"
kind: CiliumNetworkPolicy
metadata:
 name: "14-rule"
 namespace: test
spec:
 endpointSelector:
 matchLabels:
 app: test
 ingress:
 - fromEndpoints:
 - matchLabels:
 app: dynamic
 toPorts:
 - ports:
 - port: "80"
 protocol: TCP
```

Namespace Pod

```
kubectl exec -it -n test dynamic-7d8d7874f5-6dsg6 -- bash
bash-5.0# ping -c 3 10.16.0.41
PING 10.16.0.41 (10.16.0.41): 56 data bytes

--- 10.16.0.41 ping statistics ---
3 packets transmitted, 0 packets received, 100% packet loss
bash-5.0#
bash-5.0# curl 10.16.0.41:80
<html>
<head>
 <title>Hello World!</title>
 <link href='//fonts.googleapis.com/css?family=Open+Sans:400,700' rel='stylesheet' type='text/css'>
 <style>
 body {
 background-color: white;
 text-align: center;
 padding: 50px;
 font-family: "Open Sans", "Helvetica Neue", Helvetica, Arial, sans-serif;
 }
 #logo {
 margin-bottom: 40px;
 }
 </style>
</head>
<body>
 <h1>Hello World!</h1>
 <h3>Links found</h3>
 <h3>I am on test-54c98bc466-mft5s</h3>
 <h3>Cookie
 KUBERNETES listening in 443 available at tcp://10.96.0.1:443

 <h3>my name is hanhouchao!</h3>
 <h3> RequestURI='/'</h3>
 </body>
</html>
```

Namespace Pod

```
kubectl exec -it -n test label-test1-77b6764857-sq4k4 -- bash
bash-5.0# ping -c 3 10.16.0.41
PING 10.16.0.41 (10.16.0.41): 56 data bytes

--- 10.16.0.41 ping statistics ---
3 packets transmitted, 0 packets received, 100% packet loss
bash-5.0#
bash-5.0# curl -v 10.16.0.41:80 --connect-timeout 10
* Trying 10.16.0.41:80...
* After 1000ms connect time, move on!
* connect to 10.16.0.41 port 80 failed: Operation timed out
* Connection timeout after 10001 ms
* Closing connection 0
curl: (28) Connection timeout after 10001 ms
```

Namespace	L3
L4	TCP
ping	
ICMP	4

L7

## chaining L7 Cilium Generic Veth Chaining

issue 12454



2025 9 10

Q 2022 8 2



7.11.2

## 7.12 Cilium

```
Kube-OVN Cilium Cilium
Cilium Hubble Hubble
```

### 7.12.1 Hubble

```
Cilium Hubble Hubble
```

```
helm Hubble
```

```
helm upgrade cilium cilium/cilium --version 1.11.6 \
--namespace kube-system \
--reuse-values \
--set hubble.relay.enabled=true \
--set hubble.ui.enabled=true
```

```
Hubble cilium status Hubble
```

```
cilium status
 /--\
 /--\ /--\ Cilium: OK
 \--\ /--\ Operator: OK
 /--\ /--\ Hubble: OK
 \--\ /--\ ClusterMesh: disabled
 \--\

Deployment hubble-relay Desired: 1, Ready: 1/1, Available: 1/1
Deployment cilium-operator Desired: 2, Ready: 2/2, Available: 2/2
DaemonSet cilium Desired: 2, Ready: 2/2, Available: 2/2
Deployment hubble-ui Desired: 1, Ready: 1/1, Available: 1/1
Containers: cilium Running: 2
 hubble-ui Running: 1
 hubble-relay Running: 1
 cilium-operator Running: 2
Cluster Pods: 16/17 managed by Cilium
Image versions hubble-relay quay.io/cilium/hubble-relay:v1.11.6@sha256:fd9034a2d04d5b973f1e8ed44f230ea195b89c37955ff32e34e5aa68f3ed675a: 1
 cilium-operator quay.io/cilium/operator-generic:v1.11.6@sha256:9f6063c7bcaede801a39315ec7c166389f6a6783e98665f6693939cf1701bc17: 2
 cilium quay.io/cilium/cilium:v1.11.6@sha256:f7f93c26739b6641a3a3d76b1e1605b15989f25d06625260099e01c8243f54c: 2
 hubble-ui quay.io/cilium/hubble-ui:v0.9.0@sha256:0ef04e9a29212925da6bdfdb0a5b581765e41a01fcc30563cef9b30b457fea0: 1
 hubble-ui quay.io/cilium/hubble-ui-backend:v0.9.0@sha256:0000f6b76719f607a9edefb9af94fdf1811a6f1b6a8a9c537cba90bf12df474b: 1
apple@bogon cilium %
```

```
Hubble Hubble CLI :
```

```
curl -L --fail --remote-name-all https://github.com/cilium/hubble/releases/download/v0.10.0/hubble-linux-amd64.tar.gz
sudo tar xzvfC hubble-linux-amd64.tar.gz /usr/local/bin
```

### 7.12.2

```
Cilium
```

```
cilium connectivity test Cilium cilium-test Namespace cilium-test
```

```
cilium-test namespace
```

```
kubectl get all -n cilium-test
NAME READY STATUS RESTARTS AGE
pod/client-7df6cfbf7b-z5t2j 1/1 Running 0 21s
pod/client2-547996d7d8-nvgxg 1/1 Running 0 21s
pod/echo-other-node-d79544ccf-h14gg 2/2 Running 0 21s
pod/echo-same-node-5d466d5444-m17tc 2/2 Running 0 21s

NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
service/echo-other-node NodePort 10.109.58.126 <none> 8080:32269/TCP 21s
service/echo-same-node NodePort 10.188.70.32 <none> 8080:32490/TCP 21s

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/client 1/1 1 1 21s
deployment.apps/client2 1/1 1 1 21s
deployment.apps/echo-other-node 1/1 1 1 21s
deployment.apps/echo-same-node 1/1 1 1 21s

NAME DESIRED CURRENT READY AGE

```

```
replicaset.apps/client-7df6cfbf7b 1 1 1 21s
replicaset.apps/client2-547996d7d8 1 1 1 21s
replicaset.apps/echo-other-node-d79544ccf 1 1 1 21s
replicaset.apps/echo-same-node-5d466d5444 1 1 1 21s
```

### 7.12.3

Cilium	kube-system namespace	Cilium	pod	hubble observe				
# kubectl get pod -n kube-system -o wide								
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
cilium-d6h56	1/1	Running	0	2d20h	172.18.0.2	kube-ovn-worker	<none>	<none>
cilium-operator-5887f78bbb-c7sb2	1/1	Running	0	2d20h	172.18.0.2	kube-ovn-worker	<none>	<none>
cilium-operator-5887f78bbb-wj8gt	1/1	Running	0	2d20h	172.18.0.3	kube-ovn-control-plane	<none>	<none>
cilium-tq5xb	1/1	Running	0	2d20h	172.18.0.3	kube-ovn-control-plane	<none>	<none>
kube-ovn-pinger-7lgk8	1/1	Running	0	21h	10.16.0.19	kube-ovn-control-plane	<none>	<none>
kube-ovn-pinger-msvcn	1/1	Running	0	21h	10.16.0.18	kube-ovn-worker	<none>	<none>
# kubectl exec -it -n kube-system cilium-d6h56 -- bash								
root@kube-ovn-worker:/home/cilium# hubble observe --from-namespace kube-system								
Jul 29 03:24:25.551: kube-system/kube-ovn-pinger-msvcn:35576 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: ACK, PSH)								
Jul 29 03:24:25.561: kube-system/kube-ovn-pinger-msvcn:35576 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: RST)								
Jul 29 03:24:25.561: kube-system/kube-ovn-pinger-msvcn:35576 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: ACK, RST)								
Jul 29 03:24:25.572: kube-system/kube-ovn-pinger-msvcn:35578 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: SYN)								
Jul 29 03:24:25.572: kube-system/kube-ovn-pinger-msvcn:35578 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:25.651: kube-system/kube-ovn-pinger-msvcn:35578 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: ACK, PSH)								
Jul 29 03:24:25.661: kube-system/kube-ovn-pinger-msvcn:35578 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: RST)								
Jul 29 03:24:25.661: kube-system/kube-ovn-pinger-msvcn:35578 -> 172.18.0.3:6642 to-stack FORWARDED (TCP Flags: ACK, RST)								
Jul 29 03:24:25.761: kube-system/kube-ovn-pinger-msvcn:52004 -> 172.18.0.3:6443 to-stack FORWARDED (TCP Flags: ACK, PSH)								
Jul 29 03:24:25.779: kube-system/kube-ovn-pinger-msvcn -> kube-system/kube-ovn-pinger-7lgk8 to-stack FORWARDED (ICMPv4 EchoRequest)								
Jul 29 03:24:25.779: kube-system/kube-ovn-pinger-msvcn -> kube-system/kube-ovn-pinger-7lgk8 to-endpoint FORWARDED (ICMPv4 EchoReply)								
Jul 29 03:24:25.866: kube-system/hubble-ui-7596f7ffff-7j6f2:55836 -> kube-system/hubble-relay-959988db5-zc5vv:4245 to-stack FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:25.866: kube-system/hubble-ui-7596f7ffff-7j6f2:55836 -> kube-system/hubble-relay-959988db5-zc5vv:80 to-endpoint FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:25.866: kube-system/hubble-ui-7596f7ffff-7j6f2:55836 -> kube-system/hubble-relay-959988db5-zc5vv:4245 to-stack FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:25.866: kube-system/hubble-ui-7596f7ffff-7j6f2:55836 -> kube-system/hubble-relay-959988db5-zc5vv:4245 to-endpoint FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:25.975: kube-system/kube-ovn-pinger-7lgk8 -> kube-system/kube-ovn-pinger-msvcn to-endpoint FORWARDED (ICMPv4 EchoRequest)								
Jul 29 03:24:25.975: kube-system/kube-ovn-pinger-7lgk8 -> kube-system/kube-ovn-pinger-msvcn to-stack FORWARDED (ICMPv4 EchoReply)								
Jul 29 03:24:25.979: kube-system/kube-ovn-pinger-msvcn -> 172.18.0.3 to-stack FORWARDED (ICMPv4 EchoRequest)								
Jul 29 03:24:26.037: kube-system/coredns-6d4b75cb6d-lbgjg:36430 -> 172.18.0.3:6443 to-stack FORWARDED (TCP Flags: ACK)								
Jul 29 03:24:26.282: kube-system/kube-ovn-pinger-msvcn -> 172.18.0.2 to-stack FORWARDED (ICMPv4 EchoRequest)								

### Hubble Relay Hubble

```
Hubble API Hubble Service kubectl port-forward deployment/hubble-relay -n kube-system 4245:4245
```

```
kubectl port-forward
```

```
hubble status
```

```
hubble status
Healthcheck (via localhost:4245): Ok
Current/Max Flows: 8,190/8,190 (100.00%)
Flows/s: 22.86
Connected Nodes: 2/2
```

**hubble observe** **hubble help observe** **Hubble CLI**

#### 7.12.4 UI

```
cilium status Hubble UI Hubble UI
cilium hubble ui hubble-ui service Hubble UI
http://localhost:12000 UI
```

cilium-test namespace Cilium

localhost:12000/cilium-test

Filter by: label key=val, ip=1.1.1.1, dns=google.com, identity=42, pod=frontend

Any verdict Visual 11.5 flows/s • 2/2 nodes

Source Service	Destination Service	Destination Port	L7 info	Verdict	Timestamp
client cilium-test	echo-other-node cilium-test	8080	—	dropped	less than 5 seconds
client cilium-test	echo-other-node cilium-test	8080	—	dropped	less than 5 seconds
client cilium-test	echo-other-node cilium-test	8080	—	forwarded	less than 5 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds
client2 cilium-test	echo-same-node cilium-test	8080	—	forwarded	less than 20 seconds

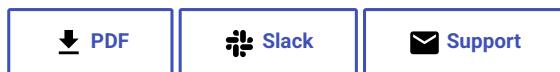
## 7.12.5 Hubble

Hubble Pod Hubble

```
hubble.metrics.enabled :
```

```
helm upgrade cilium cilium/cilium --version 1.11.6 \
--namespace kube-system \
--reuse-values \
--set hubble.relay.enabled=true \
--set hubble.ui.enabled=true \
--set hubble.metrics.enabled="{dns,drop,tcp,flow,icmp,http}"
```

kube-system namespace	hubble-metrics	Endpoints	Hubble	:
# curl 172.18.0.2:9091/metrics				
# HELP hubble_drop_total Number of drops				
# TYPE hubble_drop_total counter				
hubble_drop_total{protocol="ICMPv6",reason="Unsupported L3 protocol"} 2				
# HELP hubble_flows_processed_total Total number of flows processed				
# TYPE hubble_flows_processed_total counter				
hubble_flows_processed_total{protocol="ICMPv4",subtype="to-endpoint",type="Trace",verdict="FORWARDED"} 335				
hubble_flows_processed_total{protocol="ICMPv4",subtype="to-stack",type="Trace",verdict="FORWARDED"} 335				
hubble_flows_processed_total{protocol="ICMPv6",subtype="",type="Drop",verdict="DROPPED"} 2				
hubble_flows_processed_total{protocol="TCP",subtype="to-endpoint",type="Trace",verdict="FORWARDED"} 8282				
hubble_flows_processed_total{protocol="TCP",subtype="to-stack",type="Trace",verdict="FORWARDED"} 6767				
hubble_flows_processed_total{protocol="UDP",subtype="to-endpoint",type="Trace",verdict="FORWARDED"} 1642				
hubble_flows_processed_total{protocol="UDP",subtype="to-stack",type="Trace",verdict="FORWARDED"} 1642				
# HELP hubble_icmp_total Number of ICMP messages				
# TYPE hubble_icmp_total counter				
hubble_icmp_total{family="IPv4",type="EchoReply"} 335				
hubble_icmp_total{family="IPv4",type="EchoRequest"} 335				
hubble_icmp_total{family="IPv4",type="RouterSolicitation"} 2				
# HELP hubble_tcp_flags_total TCP flag occurrences				
# TYPE hubble_tcp_flags_total counter				
hubble_tcp_flags_total{family="IPv4",flag="FIN"} 2043				
hubble_tcp_flags_total{family="IPv4",flag="RST"} 301				
hubble_tcp_flags_total{family="IPv4",flag="SYN"} 1169				
hubble_tcp_flags_total{family="IPv4",flag="SYN-ACK"} 1169				



⌚2023 3 26

⌚2022 8 2



7.12.6

---

## 7.13

### Kube-OVN

#### 7.13.1

```
kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: external
spec:
 cidrBlock: 172.31.0.0/16
 gatewayType: centralized
 natOutgoing: false
 externalEgressGateway: 192.168.0.1
 policyRoutingTableID: 1000
 policyRoutingPriority: 1500

 • natOutgoing: false
 • externalEgressGateway
 • policyRoutingTableID TableID
 • policyRoutingPriority
```

[!\[\]\(68a75be2bf574996694613cb93c6fa0e\_img.jpg\) PDF](#)[!\[\]\(47692797c073287681b8e0dcfc4ebe20\_img.jpg\) Slack](#)[!\[\]\(b0eefcb1e9f23e5a14f31baf71970556\_img.jpg\) Support](#)

⌚2022 6 8

⌚2022 5 24

 GitHub 

#### 7.13.2

## 7.14 VIP IP

```

VIP IP IP VIP kube-ovn IP POD IP IP VIP Openstack neutron Allowed-Address-Pairs
AAP Openstack octavia POD IP aliyun terway neutron IP VIP VIP
POD IP VIP IP OVN Switch LB IP LB VIP VIP OVN
Switch LB Rule VIP

```

- Allowed-Address-Pairs VIP
- Switch LB rule VIP
- Pod VIP IP

### 7.14.1 1. Allowed-Address-Pairs VIP

```

IP Pod
• Kubernetes Kubernetes Kubernetes Underlay Subnet
• LB Subnet IP Pod
VIP Allowed-Address-Pairs IP IP
• Keepalived IP

```

#### 1.1 VIP

IP IP yaml

```

apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: vip-dynamic-01
spec:
 subnet: ovn-default
 type: ""

```

- subnet: Subnet IP
- type: ipam ip switch\_lb\_vip vip switch lb vip ip

VIP

```

kubectl get vip
NAME V4IP PV4IP MAC PMAC V6IP PV6IP SUBNET READY
vip-dynamic-01 10.16.0.12 00:00:00:F0:DB:25

```

VIP 10.16.0.12 IP

#### 1.2 VIP

VIP IP yaml

```

apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: static-vip01
spec:
 subnet: ovn-default
 v4ip: "10.16.0.121"

```

- subnet: Subnet IP
- v4ip: IP subnet CIDR

VIP

```
kubectl get vip
NAME V4IP PV4IP MAC PMAC V6IP PV6IP SUBNET READY
static-vip01 10.16.0.121 00:00:00:F0:DB:26 ovn-default true
```

VIP IP

### 1.3 Pod VIP AAP

Pod	annotation	VIP	AAP	labels	VIP
Pod annotation		VIP			ovn.kubernetes.io/aaps: vip-aap,vip-aap2,vip-aap3
AAP		Pod	AAP	Pod	VIP subnet Port

#### 1.3.1 VIP AAP

```
apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: vip-aap
spec:
 subnet: ovn-default
 namespace: default
 selector:
 - "app: aap1"
```

VIP

- namespace : AAP      VIP      VIP      AAP
- selector : AAP      VIP      Pod      Kubernetes      NodeSelector

VIP      Port

```
kubectl ko nbctl show ovn-default
switch e32e1d3b-c539-45f4-ab19-be4e33a061f6 (ovn-default)
 port aap-vip
 type: virtual
```

```
apiVersion: v1
kind: Pod
metadata:
 name: busybox
 annotations:
 ovn.kubernetes.io/aaps: vip-aap
 labels:
 app: aap1
spec:
 containers:
 - name: busybox
 image: busybox
 command: ["sleep", "3600"]
 securityContext:
 capabilities:
 add:
 - NET_ADMIN
```

AAP

```
kubectl ko nbctl list logical_switch_port aap-vip
_uuid : cd930750-0533-4f06-a6c0-217ddac73272
addresses : []
dhcpv4_options : []
dhcpv6_options : []
dynamic_addresses : []
enabled : []
external_ids : {ls=ovn-default, vendor=kube-ovn}
ha_chassis_group : []
mirror_rules : []
name : aap-vip
options : {virtual-ip="10.16.0.100", virtual-parents="busybox.default"}
parent_name : []
port_security: []
tag : []
tag_request : []
type : virtual
up : false
```

virtual-ip      VIP      IP virtual-parents      AAP      Pod      Port

## Pod

```
kubectl exec -it busybox -- ip addr add 10.16.0.100/16 dev eth0
kubectl exec -it busybox01 -- ip addr show eth0
35: eth0@if36: <BROADCAST,MULTICAST,UP,LOWER_UP,M-DOWN> mtu 1400 qdisc noqueue
 link/ether 00:00:00:e2:ab:0c brd ff:ff:ff:ff:ff:ff
 inet 10.16.0.7/16 brd 10.16.255.255 scope global eth0
 valid_lft forever preferred_lft forever
 inet 10.16.0.100/16 scope global secondary eth0
 valid_lft forever preferred_lft forever
 inet6 fe80::200:ff:fee2:ab0c/64 scope link
 valid_lft forever preferred_lft forever
```

Pod	IP	VIP	IP	subnet	Pod	IP
-----	----	-----	----	--------	-----	----

## 7.14.2 2. Switch LB rule vip

```
apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: slr-01
spec:
 subnet: ovn-default
 type: switch_lb_vip

 • subnet: Subnet IP
 • type: ipam ip switch_lb_vip vip switch lb vip ip
```

## 7.14.3 3. Pod VIP IP

v1.12

IP

```
apiVersion: kubeovn.io/v1
kind: Vip
metadata:
 name: pod-use-vip
spec:
 subnet: ovn-default
 type: ""
```

annotation	VIP	Pod
------------	-----	-----

```
apiVersion: v1
kind: Pod
metadata:
 name: static-ip
 annotations:
 ovn.kubernetes.io/vip: pod-use-vip # vip
 namespace: default
spec:
 containers:
 - name: static-ip
 image: docker.io/library/nginx:alpine
```

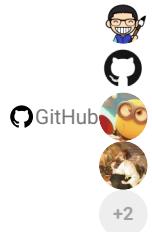
## 3.1 StatefulSet Kubevirt VM VIP

StatefulSet	VM	Pod	VIP
VM	VIP	kube-ovn-controller	keep-vm-ip true Kubevirt VM

[PDF](#)[Slack](#)[Support](#)

2025 9 17

2022 5 24

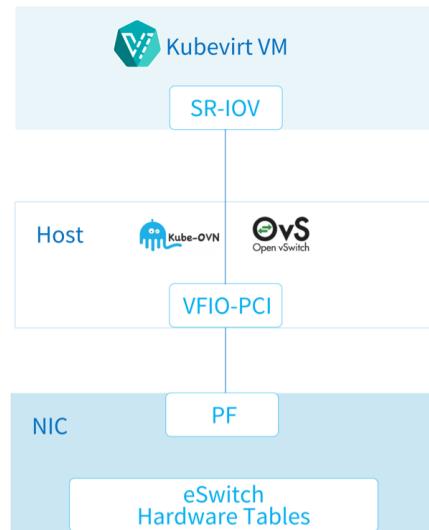


7.14.4

---

## 7.15 Mellanox Offload

Kube-OVN	OVS	CPU	CPU	Mellanox	Accelerated Switching And Packet
Processing (ASAP <sup>2</sup> )	OVS	eSwitch	OVS	CPU	



Note

2022

### 7.15.1

- Mellanox CX5/CX6/CX7/BlueField      ASAP<sup>2</sup>
- CentOS 8 Stream      Linux 5.7
- dp\_hash      hash      OVN LB
- bond

### 7.15.2 SR-IOV Device Plugin

Mellanox	offload	SR-IOV Device Plugin	<a href="#">srivio-network-operator</a>
----------	---------	----------------------	-----------------------------------------

#### SR-IOV Device Plugin

SR-IOV

ID      84:00.0      84.00.1

```
lspci -nn | grep ConnectX-5
84:00.0 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5] [15b3:1017]
84:00.1 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5] [15b3:1017]
```

## ID

```
ls -l /sys/class/net/ | grep 84:00.0
lrwxrwxrwx 1 root root 0 Feb 4 16:16 enp132s0f0np0 -> ../../devices/pci0000:80/0000:80:08.0/0000:84:00.0/net/enp132s0f0np0
ls -l /sys/class/net/ | grep 84:00.1
lrwxrwxrwx 1 root root 0 Feb 4 16:16 enp132s0f1np1 -> ../../devices/pci0000:80/0000:80:08.0/0000:84:00.1/net/enp132s0f1np1
```

## bond

```
enp132s0f0np0 enp132s0f1np1 bond1
```

```
ip link show enp132s0f0np0 | grep bond
160: enp132s0f0np0: <BROADCAST,MULTICAST,SLAVE,UP,LOWER_UP> mtu 1500 qdisc mq master bond1 state UP mode DEFAULT group default qlen 1000
ip link show enp132s0f1np1 | grep bond
169: enp132s0f1np1: <BROADCAST,MULTICAST,SLAVE,UP,LOWER_UP> mtu 1500 qdisc mq master bond1 state UP mode DEFAULT group default qlen 1000
```

## bond VF

```
ifenslave -d bond1 enp132s0f0np0
ifenslave -d bond1 enp132s0f1np1
echo 0 > /sys/class/net/enp132s0f0np0/device/sriov_numvfs
echo 0 > /sys/class/net/enp132s0f1np1/device/sriov_numvfs
ip link set enp132s0f0np0 down
ip link set enp132s0f1np1 down
```

## OVS

## SMFS DMFS

- SMFS (software-managed flow steering)
- DMFS (device-managed flow steering)

## sysfs devlink API

```
sysfs
echo <smfs|dmfs> > /sys/class/net/enp132s0f0np0/compat/devlink/steering_mode
echo <smfs|dmfs> > /sys/class/net/enp132s0f1np1/compat/devlink/steering_mode
devlink
devlink dev param set pci/84:00.0 name flow_steering_mode value smfs cmode runtime
devlink dev param set pci/84:00.1 name flow_steering_mode value smfs cmode runtime
```

## VF

```
cat /sys/class/net/enp132s0f0np0/device/sriov_totalvfs
127
cat /sys/class/net/enp132s0f1np1/device/sriov_totalvfs
127
```

## VF

```
echo '4' > /sys/class/net/enp132s0f0np0/device/sriov_numvfs
echo '4' > /sys/class/net/enp132s0f1np1/device/sriov_numvfs
ip link show enp132s0f0np0
160: enp132s0f0np0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq state DOWN mode DEFAULT group default qlen 1000
 link/ether 00:c0:eb:74:c3:4b brd ff:ff:ff:ff:ff:ff
 vf 0 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 1 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 2 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 3 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
ip link show enp132s0f1np1
169: enp132s0f1np1: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc mq state DOWN mode DEFAULT group default qlen 1000
 link/ether 00:c0:eb:74:c3:4b brd ff:ff:ff:ff:ff:ff
 vf 0 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 1 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 2 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
 vf 3 link/ether 00:00:00:00:00:00 brd ff:ff:ff:ff:ff:ff, spoof checking off, link-state disable, trust off, query_rss off
ip link set enp132s0f0np0 up
ip link set enp132s0f1np1 up
```

## VF ID

```
lspci -nn | grep ConnectX-5 | grep Virtual
84:00.2 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:00.3 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:00.4 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
```

```
84:00.5 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:00.6 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:00.7 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:01.0 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
84:01.1 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
```

VF

```
echo 0000:84:00.2 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:00.3 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:00.4 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:00.5 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:00.6 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:00.7 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:01.0 > /sys/bus/pci/drivers/mlx5_core/unbind
echo 0000:84:01.1 > /sys/bus/pci/drivers/mlx5_core/unbind
```

eSwitch

```
devlink dev eswitch set pci/0000:84:00.0 mode switchdev
devlink dev eswitch set pci/0000:84:00.1 mode switchdev
ethtool -K enp132s0f0np0 hw-tc-offload on
ethtool -K enp132s0f1np1 hw-tc-offload on
```

SR-IOV VF

- Active-backup
  - XOR
  - LACP



LACP

```
modprobe bonding mode=802.3ad
ip link set enp132s0f0np0 master bond1
ip link set enp132s0f1np1 master bond1
ip link set enp132s0f0np0 up
ip link set enp132s0f1np1 up
ip link set bond1 up
```



VF

```
echo 0000:84:00.02 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:00.3 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:00.4 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:00.5 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:00.6 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:00.7 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:01.0 > /sys/bus/pci/drivers/mlx5_core/bind
echo 0000:84:01.1 > /sys/bus/pci/drivers/mlx5_core/bind
```

NetworkManager

NetworkManager

```
systemctl stop NetworkManager
systemctl disable NetworkManager
```

## DEVICE PLUGIN

VF

Pod VF

## SR-IOV Device Plugin

## SR-IOV Configmap

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: sriovdp-config
 namespace: kube-system
data:
 config.json: |
```

```
{
 "resourceList": [
 {
 "resourcePrefix": "mellanox.com",
 "resourceName": "cx5_sriov_switchdev",
 "selectors": {
 "vendors": ["15b3"],
 "devices": ["1018"],
 "drivers": ["mlx5_core"]
 }
 }
]
}
```

SR-IOV Device Plugin      ConfigMap      ConfigMap name sriovdp-config

- **selectors:VF**
- **vendors:**
- **devices:**
- **drivers:**

selectors    pciAddresses acpiIndexes    VF      SR-IOV ConfigMap

## SR-IOV

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/sriov-network-device-plugin/v3.6.2/deployments/sriovdp-daemonset.yaml
```

SR-IOV      Kubernetes Node

```
kubectl describe node kube-ovn-01 | grep mellanox

mellanox.com/cx5_sriov_switchdev: 8
mellanox.com/cx5_sriov_switchdev: 8
mellanox.com/cx5_sriov_switchdev 0 0
```

## sriov-network-operator    SR-IOV Device Plugin

### node-feature-discovery

```
kubectl apply -k https://github.com/kubernetes-sigs/node-feature-discovery/deployment/overlays/default?ref=v0.11.3
```

offload      annotation:

```
kubectl label nodes [offloadNicNode] feature.node.kubernetes.io/network-sriov.capable=true
```

### Operator

```
git clone --depth=1 https://github.com/kubeovn/sriov-network-operator.git
kubectl apply -k sriov-network-operator/deploy
```

### Operator

```
kubectl get -n kube-system all | grep sriov
NAME READY STATUS RESTARTS AGE
pod/sriov-network-config-daemon-bf9nt 1/1 Running 0 8s
pod/sriov-network-operator-54d7545f65-296gb 1/1 Running 0 10s

NAME DESIRED CURRENT READY UP-TO-DATE AVAILABLE NODE
SELECTOR AGE
daemonset.apps/sriov-network-config-daemon 1 1 1 1 1 beta.kubernetes.io/os=linux,feature.node.kubernetes.io/network-sriov.capable=true 8s

NAME READY UP-TO-DATE AVAILABLE AGE
deployment.apps/sriov-network-operator 1/1 1 1 10s

NAME DESIRED CURRENT READY AGE
replicaset.apps/sriov-network-operator-54d7545f65 1 1 1 10s
```

SriovNetworkNodeState    node1      Mellanox

```
kubectl get sriovnetworknodestates.sriovnetwork.openshift.io -n kube-system node1 -o yaml
apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetworkNodeState
spec: ...
```

```

status:
 interfaces:
 - deviceID: "1017"
 driver: mlx5_core
 mtu: 1500
 pciAddress: "0000:5f:00.0"
 totalvfs: 8
 vendor: "15b3"
 linkSeed: 25000Mb/s
 linkType: ETH
 mac: 08:c0:eb:f4:85:bb
 name: ens41f0np0
 - deviceID: "1017"
 driver: mlx5_core
 mtu: 1500
 pciAddress: "0000:5f:00.1"
 totalvfs: 8
 vendor: "15b3"
 linkSeed: 25000Mb/s
 linkType: ETH
 mac: 08:c0:eb:f4:85:bb
 name: ens41f1np1

```

#### SriovNetworkNodePolicy      nicSelector

```

apiVersion: sriovnetwork.openshift.io/v1
kind: SriovNetworkNodePolicy
metadata:
 name: policy
 namespace: kube-system
spec:
 nodeSelector:
 feature.node.kubernetes.io/network-sriov.capable: "true"
 eSwitchMode: switchdev
 numVfs: 3
 nicSelector:
 pfNames:
 - ens41f0np0
 - ens41f1np1
 resourceName: cx_sriov_switchdev

```

#### SriovNetworkNodeState      status

```

kubectl get sriovnetworknodestates.sriovnetwork.openshift.io -n kube-system node1 -o yaml

...
spec:
 interfaces:
 - eSwitchMode: switchdev
 name: ens41f0np0
 numVfs: 3
 pciAddress: 0000:5f:00.0
 vfGroups:
 - policyName: policy
 vfRange: 0-2
 resourceName: cx_sriov_switchdev
 - eSwitchMode: switchdev
 name: ens41f1np1
 numVfs: 3
 pciAddress: 0000:5f:00.1
 vfGroups:
 - policyName: policy
 vfRange: 0-2
 resourceName: cx_sriov_switchdev
 status:
 interfaces
 - Vfs:
 - deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.2
 vendor: "15b3"
 - deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.3
 vendor: "15b3"
 - deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.4
 vendor: "15b3"
 deviceID: "1017"
 driver: mlx5_core
 linkSeed: 25000Mb/s
 linkType: ETH
 mac: 08:c0:eb:f4:85:ab
 mtu: 1500
 name: ens41f0np0
 numVfs: 3
 pciAddress: 0000:5f:00.0
 totalvfs: 3
 vendor: "15b3"
 - Vfs:

```

```

- deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.5
 vendor: "15b3"
- deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.6
 vendor: "15b3"
- deviceID: 1018
 driver: mlx5_core
 pciAddress: 0000:5f:00.7
 vendor: "15b3"
deviceID: "1017"
driver: mlx5_core
linkSeed: 25000Mb/s
linkType: ETH
mac: 08:c0:eb:f4:85:bb
mtu: 1500
name: ens41f1np1
numVfs: 3
pciAddress: 0000:5f:00.1
totalvfs: 3
vendor: "15b3"

```

## VF

```

lspci -nn | grep ConnectX
5f:00.0 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5] [15b3:1017]
5f:00.1 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5] [15b3:1017]
5f:00.2 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
5f:00.3 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
5f:00.4 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
5f:00.5 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
5f:00.6 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]
5f:00.7 Ethernet controller [0200]: Mellanox Technologies MT27800 Family [ConnectX-5 Virtual Function] [15b3:1018]

```

## PF

```

cat /sys/class/net/ens41f0np0/compat/devlink/mode
switchdev

```

## 7.15.3 Multus-CNI

SR-IOV Device Plugin	ID	Multus-CNI	Kube-OVN	Multus-CNI
----------------------	----	------------	----------	------------

### Multus-CNI

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/multus-cni/v4.0.2/deployments/multus-daemonset-thick.yml
```

multus	Thin	Thick	SR-IOV	Thick
--------	------	-------	--------	-------

### NetworkAttachmentDefinition

```

apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: sriov
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/resourceName: mellanox.com/cx5_sriov_switchdev
spec:
 config: '{
 "cniVersion": "0.3.1",
 "name": "kube-ovn",
 "plugins": [
 {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "sriov.default.ovn"
 },
 {
 "type": "portmap",
 "capabilities": {
 "portMappings": true
 }
 }
]
}'

```

- provider: NetworkAttachmentDefinition {name}.{namespace}.ovn

## 7.15.4 Overlay

### Kube-OVN

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install.sh
```

IFACE	IP
-------	----

```
ENABLE_MIRROR=${ENABLE_MIRROR:-false}
HW_OFFLOAD=${HW_OFFLOAD:-true}
ENABLE_LB=${ENABLE_LB:-false}
IFACE="bond1"
SR-IoV Device Plugin bond IFACE bond1 bond IFACE enp132s0f0np0 enp132s0f1np1
```

### Kube-OVN

```
bash install.sh
```

### VF Pod

yaml	VF	Pod:
------	----	------

```
apiVersion: v1
kind: Pod
metadata:
 name: nginx-overlay
 annotations:
 v1.multus-cni.io/default-network: default/sriov
 sriov.default.ovn.kubernetes.io/logical_switch: ovn-default
spec:
 containers:
 - name: nginx-overlay
 image: docker.io/library/nginx:alpine
 resources:
 requests:
 mellanox.com/cx5_sriov_switchdev: '1'
 limits:
 mellanox.com/cx5_sriov_switchdev: '1'
```

- v1.multus-cni.io/default-network: NetworkAttachmentDefinition {namespace}/{name}
- sriov.default.ovn.kubernetes.io/logical\_switch: Pod Pod

## 7.15.5 Underlay

### Kube-OVN

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install.sh
```

IFACE	IP
-------	----

```
ENABLE_MIRROR=${ENABLE_MIRROR:-false}
HW_OFFLOAD=${HW_OFFLOAD:-true}
ENABLE_LB=${ENABLE_LB:-false}
IFACE=""
Underlay IFACE PF IFACE K8s PF
```

### Kube-OVN

```
bash install.sh
```

### VF Pod

yaml	VF	Pod:
------	----	------

```

apiVersion: kubeovn.io/v1
kind: ProviderNetwork
metadata:
 name: underlay-offload
spec:
 defaultInterface: bond1

apiVersion: kubeovn.io/v1
kind: Vlan
metadata:
 name: vlan0
spec:
 id: 0
 provider: underlay-offload

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: vlan0
spec:
 protocol: IPv4
 provider: ovn
 cidrBlock: 10.10.204.0/24
 gateway: 10.10.204.254
 vlan: vlan0
 excludeIps:
 - 10.10.204.1..10.10.204.100

apiVersion: v1
kind: Pod
metadata:
 name: nginx-underlay
 annotations:
 k8s.v1.cni.cncf.io/networks: '[{
 "name": "sriov",
 "namespace": "default",
 "default-route": ["10.10.204.254"]
 }]'
 ovn.default.ovn.kubernetes.io/logical_switch: vlan0
spec:
 containers:
 - name: nginx-underlay
 image: docker.io/library/nginx:alpine
 resources:
 requests:
 mellanox.com/cx5_sriov_switchdev: '1'
 limits:
 mellanox.com/cx5_sriov_switchdev: '1'

```

• v1.multus-cni.io/default-network: NetworkAttachmentDefinition {namespace}/{name}

	multus	VF	Pod	VF	Pod	multus
--	--------	----	-----	----	-----	--------

yaml VF Pod:

```

apiVersion: v1
kind: Pod
metadata:
 name: nginx-underlay-noVF
 annotations:
 ovn.kubernetes.io/logical_switch: vlan0
spec:
 containers:
 - name: nginx-underlay-noVF
 image: docker.io/library/nginx:alpine

```

	VF	Pod	ovs-kernel	e-switch
--	----	-----	------------	----------

## 7.15.6

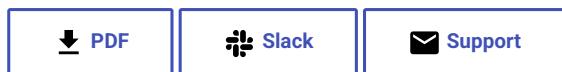
Pod ovs-ovn

```

ovs-appctl dptcl/dump-flows -m type=offloaded
ufid:91cc45de-e7e9-4935-8f82-1890430b0f66, skb_priority(0/0),skb_mark(0/0),ct_state(0/0x23),ct_zone(0/0),ct_mark(0/0),ct_label(0/0x1),recirc_id(0),dp_hash(0/0),in_port(5b45c61b307e_h),packet_type(ns=0/0,id=0/0),eth(src=00:00:00:c5:6d:4e,dst=00:00:00:e7:16:ce),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0.0/0.0.0.0,proto=0/0,tos=0/0,ttl=0/0,frag=no), packets:941539, bytes:62142230, used:0.260s, offloaded:yes, dp:tc, actions:54235e5753b8_h
ufid:e00768d7-e652-4d79-8182-3291d852b791, skb_priority(0/0),skb_mark(0/0),ct_state(0/0x23),ct_zone(0/0),ct_mark(0/0),ct_label(0/0x1),recirc_id(0),dp_hash(0/0),in_port(54235e5753b8_h),packet_type(ns=0/0,id=0/0),eth(src=00:00:00:e7:16:ce,dst=00:00:00:c5:6d:4e),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0/0.0.0.0,proto=0/0,tos=0/0,ttl=0/0,frag=no), packets:82386659, bytes:115944854173, used:0.260s, offloaded:yes, dp:tc, actions:5b45c61b307e_h

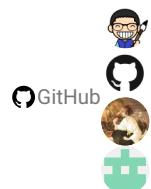
```

offloaded:yes, dp:tc



⌚2025 9 10

⌚2022 5 24

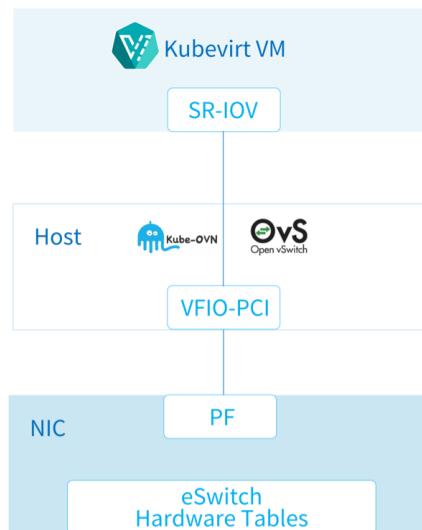


7.15.7

---

## 7.16 Offload

Kube-OVN	OVS	CPU	CPU	Agilio CX	OVS
OVS					



### Note

2022

### 7.16.1

- Agilio CX
- CentOS 8 Stream    Linux 5.7
- dp\_hash    hash              OVN LB

### 7.16.2 SR-IOV

[Agilio Open vSwitch TC User Guide](#)

```

#!/bin/bash
DEVICE=${1}
DEFAULT_ASSY=scan
ASSY=${2:-$DEFAULT_ASSY}
APP=${3:-flower}

if ["x${DEVICE}" = "x" -o ! -e /sys/class/net/${DEVICE}]; then
 echo Syntax: ${0} device [ASSY] [APP]
 echo
 echo This script associates the TC Offload firmware
 echo with a Netronome SmartNIC.
 echo
 echo device: is the network device associated with the SmartNIC

```

```

echo ASSY: defaults to ${DEFAULT_ASSY}
echo APP: defaults to flower. flower-next is supported if updated
echo firmware has been installed.
exit 1
fi

It is recommended that the assembly be determined by inspection
The following code determines the value via the debug interface
if ["${ASSY}" = "scanx"]; then
 ethtool -w ${DEVICE} 0
 DEBUG=$(ethtool -w ${DEVICE} data /dev/stdout | strings)
 SERIAL=$(echo "$DEBUG" | grep '^SN:')
 ASSY=$(echo ${SERIAL} | grep -oE AMDA[0-9]{4})
fi

PCIADDR=$(basename $(readlink -e /sys/class/net/${DEVICE}/device))
FWDIR="/lib/firmware/netronome"

AMDA0081 and AMDA0097 uses the same firmware
if ["${ASSY}" = "AMDA0081"]; then
 if [! -e ${FWDIR}/${APP}/nic_AMDA0081.nffw]; then
 ln -sf nic_AMDA0097.nffw ${FWDIR}/${APP}/nic_AMDA0081.nffw
 fi
fi

FW="${FWDIR}/pci-${PCIADDR}.nffw"
ln -sf "${APP}/nic_${ASSY}.nffw" "${FW}"

insert distro-specific initramfs section here...

```

```

./agilio-tc-fw-select.sh ens47np0 scan
rmmod nfp
modprobe nfp

```

VF        VF

```

cat /sys/class/net/ens3/device/sriov_totalvfs
65

echo 4 > /sys/class/net/ens47/device/sriov_numvfs

```

### 7.16.3 SR-IOV Device Plugin

VF	Pod	VF	SR-IOV Device Plugin
----	-----	----	----------------------

SR-IOV    Configmap

```

apiVersion: v1
kind: ConfigMap
metadata:
 name: sriovdp-config
 namespace: kube-system
data:
 config.json: |
 {
 "resourceList": [
 {
 "resourcePrefix": "coragine.com",
 "resourceName": "agilio_sriov",
 "selectors": {
 "vendors": ["19ee"],
 "devices": ["6003"],
 "drivers": ["nfp_netvf"]
 }
 }
]
 }

```

SR-IOV    :

```
kubectl apply -f https://raw.githubusercontent.com/intel/sriov-network-device-plugin/master/deployments/k8s-v1.16/sriovdp-daemonset.yaml
```

SR-IOV    Kubernetes Node

```

kubectl describe no containerserver | grep coragine

coragine.com/agilio_sriov: 4
coragine.com/agilio_sriov: 4
coragine.com/agilio_sriov 0 0

```

## 7.16.4 Multus-CNI

SR-IOV Device Plugin      ID      Multus-CNI      Kube-OVN      Multus-CNI

### Multus-CNI

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/multus-cni/master/deployments/multus-daemonset.yaml
```

#### NetworkAttachmentDefinition

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: default
 namespace: default
 annotations:
 k8s.v1.cni.cncf.io/resourceName: corigine.com/agilio_sriov
spec:
 config: '{
 "cniVersion": "0.3.1",
 "name": "kube-ovn",
 "plugins": [
 {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "default.default.ovn"
 },
 {
 "type": "portmap",
 "capabilities": {
 "portMappings": true
 }
 }
]
}'

```

- provider: NetworkAttachmentDefinition {name}.{namespace}.ovn

## 7.16.5 Kube-OVN

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install.sh
```

IFACE      IP

```
ENABLE_MIRROR=${ENABLE_MIRROR:-false}
HW_OFFLOAD=${HW_OFFLOAD:-true}
ENABLE_LB=${ENABLE_LB:-false}
IFACE="ensp01"
```

### Kube-OVN

```
bash install.sh
```

## 7.16.6 VF      Pod

yaml      VF      Pod:

```
apiVersion: v1
kind: Pod
metadata:
 name: nginx
 namespace: default
 annotations:
 v1.multus-cni.io/default-network: default/default
spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
 resources:
 requests:
 corigine.com/agilio_sriov: '1'
```

```

limits:
 coragine.com/agilio_sriov: '1'

• v1.multus-cni.io/default-network: NetworkAttachmentDefinition {namespace}/{name}

Pod ovs-ovn

ovs-appctl dpctl/dump-flows -m type=offloaded
ufid:91cc45de-e7e9-4935-8f82-1890430b0f66, skb_priority(0/0),skb_mark(0/0),ct_state(0/0x23),ct_zone(0/0),ct_mark(0/0),ct_label(0/0x1),recirc_id(0),dp_hash(0/0),in_port(5b45c61b307e_h),packet_type(ns=0/0,id=0/0),eth(src=00:00:00:c5:6d:4e,dst=00:00:00:e7:16:ce),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0/0.0.0.0,proto=0/0,tos=0/0,ttl=0/0,frag=no), packets:941539, bytes:62142230, used:0.260s, offloaded:yes, dp:tc, actions:54235e5753b8_h
ufid:e00768d7-e652-4d79-8182-3291d852b791, skb_priority(0/0),skb_mark(0/0),ct_state(0/0x23),ct_zone(0/0),ct_mark(0/0),ct_label(0/0x1),recirc_id(0),dp_hash(0/0),in_port(54235e5753b8_h),packet_type(ns=0/0,id=0/0),eth(src=00:00:00:e7:16:ce,dst=00:00:00:c5:6d:4e),eth_type(0x0800),ipv4(src=0.0.0.0/0.0.0.0,dst=0.0.0/0.0.0.0,proto=0/0,tos=0/0,ttl=0/0,frag=no), packets:82386659, bytes:115944854173, used:0.260s, offloaded:yes, dp:tc, actions:5b45c61b307e_h

offloaded:yes, dp:tc

```

[PDF](#)[Slack](#)[Support](#)

⌚2025 7 21

⌚2022 5 24

[GitHub](#)

7.16.7

## 7.17 Offload



### Note

1. 2024
2. 1.11 Kube-OVN

### 7.17.1

- metaScale
- MCR
- BIOS SR-IOV VT-d

### 7.17.2

#### hw-offload Kube-OVN

1.

```
wget https://github.com/yunsilicon/kube-ovn/blob/release-1.11/dist/images/install.sh
```

1.

```
/opt/ovs-config/ovs-dpdk-config
```

```
specify log level for ovs dpdk, the value is info or dbg, default is info
VLOG=info
specify nic offload, the value is true or false, default is true
HW_OFFLOAD=true
specify cpu mask for ovs dpdk, not specified by default
CPU_MASK=0x02
specify socket memory, not specified by default
SOCKET_MEM="2048,2048"
specify encap IP
ENCAP_IP=6.6.208/24
specify pci device
DPDK_DEV=0000:b3:00.0
specify mtu, default is 1500
PF_MTU=1500
specify bond name if bond enabled, not specified by default
BR_PHY_BOND_NAME=bond0
```

#### 1. Kube-OVN

```
bash install.sh
```

#### SR-IOV

1. metaScale ID b3:00.0:

```
[root@k8s-master ~]# lspci -d 1f67:
b3:00.0 Ethernet controller: Device 1f67:1111 (rev 02)
b3:00.1 Ethernet controller: Device 1f67:1111 (rev 02)
```

1. ID p3p1

```
ls -l /sys/class/net/ | grep b3:00.0
lrwxrwxrwx 1 root root 0 May 7 16:30 p3p1 -> ../../devices/pci0000:b2/0000:b2:00.0/0000:b3:00.0/net/p3p1
```

## 1. VF

```
cat /sys/class/net/p3p1/device/sriov_totalvfs
512
```

## 1. VF

```
echo '10' > /sys/class/net/p3p1/device/sriov_numvfs
```

## 1. VF

```
lspci -d 1f67:
b3:00.0 Ethernet controller: Device 1f67:1111 (rev 02)
b3:00.1 Ethernet controller: Device 1f67:1111 (rev 02)
b3:00.2 Ethernet controller: Device 1f67:1112
b3:00.3 Ethernet controller: Device 1f67:1112
b3:00.4 Ethernet controller: Device 1f67:1112
b3:00.5 Ethernet controller: Device 1f67:1112
b3:00.6 Ethernet controller: Device 1f67:1112
b3:00.7 Ethernet controller: Device 1f67:1112
b3:01.0 Ethernet controller: Device 1f67:1112
b3:01.1 Ethernet controller: Device 1f67:1112
b3:01.2 Ethernet controller: Device 1f67:1112
b3:01.3 Ethernet controller: Device 1f67:1112
```

## 1. switchdev

```
devlink dev eswitch set pci/0000:b3:00.0 mode switchdev
```

## SR-IOV Device Plugin

### 1. SR-IOV ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: sriovdp-config
 namespace: kube-system
data:
 config.json: |
 {
 "resourceList": [
 {
 "resourceName": "xsc_sriov",
 "resourcePrefix": "yunsilicon.com",
 "selectors": {
 "vendors": ["1f67"],
 "devices": ["1012", "1112"]
 }
 }
]
 }
```

### 1. SR-IOV Device Plugin DevicePlugin

### 2. SR-IOV

```
kubectl describe node <node name> | grep yunsilicon.com/xsc_sriov
yunsilicon.com/xsc_sriov: 10
yunsilicon.com/xsc_sriov: 10
yunsilicon.com/xsc_sriov 0 0
```

## Multus-CNI

### 1. Multus-CNI Multus-CNI

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/multus-cni/master/deployments/multus-daemonset.yaml
```

### 1. NetworkAttachmentDefinition

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
```

```

metadata:
 name: sriov-net1
 namespace: default
 annotations:
 k8s.v1.cncf.io/resourceName: yunsilicon.com/xsc_sriov
spec:
 config: '{
 "cniVersion": "0.3.1",
 "name": "kube-ovn",
 "plugins": [
 {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "sriov-net1.default.ovn"
 },
 {
 "type": "portmap",
 "capabilities": {
 "portMappings": true
 }
 }
]
}'

```

## SR-IOV Pod

```

apiVersion: v1
kind: Pod
metadata:
 name: nginx
 annotations:
 v1.multus-cni.io/default-network: default/sriov-net1
spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
 resources:
 requests:
 yunsilicon.com/xsc_sriov: '1'
 limits:
 yunsilicon.com/xsc_sriov: '1'

```

## Offload

Pod ovs-ovn

```

ovs-appctl dptcl/dump-flows type=offloaded
flow-dump from pmd on cpu core: 9
ct_state(-new+est-rel+rpl+trk),ct_mark(0/0x3),recirc_id(0x2d277),in_port(15),packet_type(ns=0,id=0),eth(src=00:00:00:9d:fb:1a,dst=00:00:00:ce:cf:b9),eth_type(0x0800),ipv4(dst=10.16.0.14,frag=no),packets:6,bytes:588,used:7.276s,actions:ct(zone=4,nat),recirc(0x2d278)
ct_state(-new+est-rel+rpl+trk),ct_mark(0/0x3),recirc_id(0x2d275),in_port(8),packet_type(ns=0,id=0),eth(src=00:00:00:ce:cf:b9,dst=00:00:00:9d:fb:1a),eth_type(0x0800),ipv4(dst=10.16.0.18,frag=no),packets:5,bytes:490,used:7.434s,actions:ct(zone=6,nat),recirc(0x2d276)
ct_state(-new+est-rel-rpl+trk),ct_mark(0/0x1),recirc_id(0x2d276),in_port(8),packet_type(ns=0,id=0),eth(src=00:00:00:ce:cf:b9,dst=00:00:00:9d:fb:1a),eth_type(0x0800),eth_type(0x0800),ipv4(dst=10.16.0.18,frag=no),packets:5,bytes:490,used:7.434s,actions:15
recirc_id(0),in_port(15),packet_type(ns=0,id=0),eth(src=00:00:00:9d:fb:1a/01:00:00:00:00:00,dst=00:00:00:ce:cf:b9),eth_type(0x0800),ipv4(dst=10.16.0.14/255.192.0.0,frag=no),packets:6,bytes:588,used:7.277s,actions:ct(zone=6,nat),recirc(0x2d277)
recirc_id(0),in_port(8),packet_type(ns=0,id=0),eth(src=00:00:00:ce:cf:b9/01:00:00:00:00:00,dst=00:00:00:9d:fb:1a),eth_type(0x0800),ipv4(dst=10.16.0.18/255.192.0.0,frag=no),packets:6,bytes:588,used:7.434s,actions:ct(zone=4,nat),recirc(0x2d275)
ct_state(-new+est-rel+rpl+trk),ct_mark(0/0x1),recirc_id(0x2d278),in_port(15),packet_type(ns=0,id=0),eth(dst=00:00:00:ce:cf:b9/01:00:00:00:00:00),eth_type(0x0800),ipv4(frag=no),packets:6,bytes:588,used:7.277s,actions:8

```



PDF



Slack



Support

⌚ 2025 7 21

⌚ 2024 5 29



GitHub



7.17.3

## 7.18 Offload



### Note

2024

#### 7.18.1

- 2200E
- HADOS
- BIOS SR-IOV

#### 7.18.2

##### SR-IOV

1. 2200E vendor ID 1f47 ID 00:0a.0 00:0b.0 2200E

```
lspci | grep 1f47
00:0a.0 Ethernet controller: Device 1f47:1001 (rev 10)
00:0b.0 Ethernet controller: Device 1f47:1001 (rev 10)
```

1. VF

```
cat /sys/bus/pci/devices/0000\:00\:0a.0/sriov_totalvfs
256
```

1. VF VF

```
echo 7 > /sys/bus/pci/devices/0000\:00\:0a.0/sriov_numvfs
```

1. VF

```
lspci | grep 1f47
00:0a.0 Ethernet controller: Device 1f47:1001 (rev 10)
00:0a.1 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.2 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.3 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.4 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.5 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.6 Ethernet controller: Device 1f47:110f (rev 10)
00:0a.7 Ethernet controller: Device 1f47:110f (rev 10)
00:0b.0 Ethernet controller: Device 1f47:1001 (rev 10)
```

##### SR-IOV Device Plugin

1. SR-IOV Configmap SR-IOV Device Plugin VF Pod

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: sriovdp-config
 namespace: kube-system
data:
 config.json: |
 {
 "resourceList": [
 {
 "resourceName": "sriov_dpu",
 "resourcePrefix": "yusur.tech",
 "selectors": {
 "vendors": ["1f47"],
 "devices": ["110f"]
 }
 }
]
 }
```

```
]
}
```

## 1. SR-IOV Device Plugin

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/sriov-network-device-plugin/v3.6.2/deployments/sriovdp-daemonset.yaml
```

### 1. SR-IOV      kubernetes Node

```
kubectl describe node node1 | grep yusur
yusur.tech/sriov_dpu: 7
yusur.tech/sriov_dpu: 7
yusur.tech/sriov_dpu 0 0
```

## 7.18.3 Multus-CNI

Multus-CNI    Kube-OCN    SRIOV    Device ID

```
kubectl apply -f https://raw.githubusercontent.com/k8snetworkplumbingwg/multus-cni/v4.0.2/deployments/multus-daemonset-thick.yaml
```

### NetworkAttachmentDefinition

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: test
 namespace: kube-system
 annotations:
 k8s.v1.cni.cncf.io/resourceName: yusur.tech/sriov_dpu
spec:
 config: '{
 "cniVersion": "0.3.1",
 "name": "kube-ovn",
 "plugins": [
 {
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "test.kube-system.ovn"
 },
 {
 "type": "portmap",
 "capabilities": {
 "portMappings": true
 }
 }
]
 }
```

- provider : NetworkAttachmentDefinition {name}. {namespace}. ovn

## 7.18.4 Kube-OVN

### 1.

```
wget https://github.com/kubeovn/kube-ovn/blob/release-1.12/dist/images/install.sh
```

### 1. IFACE      IP

```
ENABLE_MIRROR=${ENABLE_MIRROR:-false}
HW_OFFLOAD=${HW_OFFLOAD:-true}
ENABLE_LB=${ENABLE_LB:-false}
IFACE="p0"
```

### 1. kube-ovn

```
bash install.sh
```

### VF    pod

yaml    VF    Pod

```

apiVersion: v1
kind: Pod
metadata:
 name: nginx
 namespace: default
 annotations:
 v1.multus-cni.io/default-network: kube-system/test
spec:
 containers:
 - name: nginx
 image: docker.io/library/nginx:alpine
 resources:
 requests:
 yusur.tech/sriov_dpu: '1'
 limits:
 yusur.tech/sriov_dpu: '1'

```

- v1.multus-cni.io/default-network : NetworkAttachmentDefinition {namespace}/{name}

## Offload

Pod ovs-ovn

```

ovs-appctl dpctl/dump-flows -m type=offloaded
ufid:67c2e10f-92d4-4574-be70-d072815ff166, skb_priority(0/0),skb_mark(0/0),ct_state(0/0x23),ct_zone(0/0),ct_mark(0/0),ct_label(0/0),recirc_id(0),dp_hash(0/0),in_port(d85b161b6840_h),packet_type(ns=0/0,id=0/0),eth(src=0:a:c9:1c:70:01:09,dst=8a:18:a4:22:b7:7d),eth_type(0x0800),ipv4(src=10.0.1.10,dst=10.0.1.6,proto=6,tos=0/0x3,ttl=0/0,frag=no),tcp(src=60774,dst=9001),packets:75021,bytes:109521630,offload_packets:75019,offload_bytes:109521498,used:3.990s,offloaded:yes,dp:tc,actions:set(tunnel(tun_id=0x5,dst=192.168.201.12,ttl=64,tp_dst=6081,geneve({class=0x102,type=0x80,len=4,0xa0006}),flags(csum(key))),genev_sys_6081
ufid:7940666e-a0bd-42a5-8116-1e84e81bb338, skb_priority(0/0),tunnel(tun_id=0x5,src=192.168.201.12,dst=192.168.201.11,ttl=0/0,tp_dst=6081,geneve({class=0x102,type=0x80,len=4,0x6000a}),flags(+key)),skb_mark(0/0),ct_state(0/0),ct_zone(0/0),ct_mark(0/0),ct_label(0/0),recirc_id(0),dp_hash(0/0),in_port(genev_sys_6081),packet_type(ns=0/0,id=0/0),eth(src=8a:18:a4:22:b7:7d,dst=0:a:c9:1c:70:01:09),eth_type(0x0800),ipv4(src=10.0.1.6,dst=10.0.1.10,proto=6,tos=0/0,ttl=0/0,frag=no),tcp(src=9001,dst=60774),packets:6946,bytes:459664,offload_packets:6944,offload_bytes:459532,used:4.170s,dp:tc,offloaded:yes,actions:d85b161b6840_h

```



PDF



Slack



Support

⌚2025 7 21

⌚2024 8 13

GitHub

7.18.5

7.19 DPDK

Kube-OVN      OVS-DPDK      KubeVirt      DPDK

KubeVirt OVS-DPDK patchVhostuser implementation KubeVirt KVM Device Plugin OVS-DPDK

## 7.19.1

- DPDK
  - Hugepages

## 7.19.2 DPDK

`driverctl` DPDK

```
driverctl set-override 0000:00:0b.0 uio_pci_generic
```

7.19.3

OVS-DPDK Kube-OVN

```
kubectl label nodes <node> ovn.kubernetes.io/ovs_dp_type="userspace"
```

OVS-DPDK /opt/ovs-config ovs-dpdk-config

ENCAP\_IP=192.168.122.193/24  
DPDK\_DEV=0000:00:0b.0

## 7.19.4 Kube-OVN

```
wget https://raw.githubusercontent.com/kubeovn/kube-ovn/release-1.14/dist/images/install.sh
```

PPDK

```
bash install.sh --with-hybrid-dpdk
```

7.19.5

## vhostuser QoS-DPDK

## KVM Device Plugin

```
kubectl apply -f https://raw.githubusercontent.com/kubevirt/kubernetes-device-plugins/master/manifests/kvm-ds.yaml
```

## NetworkAttachmentDefinition

```
apiVersion: k8s.cni.cncf.io/v1
kind: NetworkAttachmentDefinition
metadata:
 name: ovn-dpdk
 namespace: default
spec:
 config: >-
 {
 "cniVersion": "0.3.0",
 "tvsne": "kube-ovn"
```

```

 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "ovn-dpdk.default.ovn",
 "vhost_user_socket_volume_name": "vhostuser-sockets",
 "vhost_user_socket_name": "sock"
 }
}
```

## Dockerfile VM

```

FROM quay.io/kubevirt/virt-launcher:v0.46.1

wget http://cloud.centos.org/centos/7/images/CentOS-7-x86_64-GenericCloud.qcow2
COPY CentOS-7-x86_64-GenericCloud.qcow2 /var/lib/libvirt/images/CentOS-7-x86_64-GenericCloud.qcow2

```

```

apiVersion: v1
kind: ConfigMap
metadata:
 name: vm-config
data:
 start.sh: |
 chmod u+w /etc/libvirt/qemu.conf
 echo "hugepages_mount = \"/dev/hugepages\" >> /etc/libvirt/qemu.conf
 virtlogd &
 libvirtd &

 mkdir /var/lock

 sleep 5

 virsh define /root/vm/vm.xml
 virsh start vm

 tail -f /dev/null
 vm.xml: |
 <domain type='kvm'>
 <name>vm</name>
 <uuid>4a9b3f53-fa2a-47f3-a757-dd87720d9d1d</uuid>
 <memory unit='KiB'>2097152</memory>
 <currentMemory unit='KiB'>2097152</currentMemory>
 <memoryBacking>
 <hugepages>
 <page size='2' unit='M' nodeset='0' />
 </hugepages>
 </memoryBacking>
 <vcpu placement='static'>2</vcpu>
 <cpurtune>
 <shares>4096</shares>
 <vcpu pin vcpu='0' cpuset='4' />
 <vcpu pin vcpu='1' cpuset='5' />
 <emulatorpin cpuset='1,3' />
 </cpurtune>
 <os>
 <type arch='x86_64' machine='pc'>hvm</type>
 <boot dev='hd' />
 </os>
 <features>
 <acpi/>
 <apic/>
 </features>
 <cpu mode='host-model'>
 <model fallback='allow' />
 <topology sockets='1' cores='2' threads='1' />
 <numa>
 <cell id='0' cpus='0-1' memory='2097152' unit='KiB' memAccess='shared' />
 </numa>
 </cpu>
 <on_reboot>restart</on_reboot>
 <devices>
 <emulator>/usr/libexec/qemu-kvm</emulator>
 <disk type='file' device='disk'>
 <driver name='qemu' type='qcow2' cache='none' />
 <source file='/var/lib/libvirt/images/CentOS-7-x86_64-GenericCloud.qcow2' />
 <target dev='vda' bus='virtio' />
 </disk>
 <interface type='vhostuser'>
 <mac address='00:00:00:0A:30:89' />
 <source type='unix' path='/var/run/vm.sock' mode='server' />
 <model type='virtio' />
 <driver queues='2'>
 <host mrg_rxbuf='off' />
 </driver>
 </interface>
 <serial type='pty'>
 <target type='isa-serial' port='0'>
 <model name='isa-serial' />
 </target>
 </serial>
 <console type='pty'>
 <target type='serial' port='0' />
 </console>
 </devices>
 </domain>

```

```

</console>
<channel type='unix'>
 <source mode='bind' path='/var/lib/libvirt/qemu/channel/target/domain-1-vm/org.qemu.guest_agent.0' />
 <target type='virtio' name='org.qemu.guest_agent.0' state='connected' />
 <alias name='channel0' />
</channel>

</devices>
</domain>

apiVersion: apps/v1
kind: Deployment
metadata:
 name: vm-deployment
 labels:
 app: vm
spec:
 replicas: 1
 selector:
 matchLabels:
 app: vm
 template:
 metadata:
 labels:
 app: vm
 annotations:
 k8s.v1.cni.cncf.io/networks: default/ovn-dpdk
 ovn-dpdk.default.ovn.kubernetes.io/ip_address: 10.16.0.96
 ovn-dpdk.default.ovn.kubernetes.io/mac_address: 00:00:00:0A:30:89
 spec:
 nodeSelector:
 ovn.kubernetes.io/ovs_dp_type: userspace
 securityContext:
 runAsUser: 0
 volumes:
 - name: vhostuser-sockets
 emptyDir: {}
 - name: xml
 configMap:
 name: vm-config
 - name: hugepage
 emptyDir:
 medium: HugePages-2Mi
 - name: libvirt-runtime
 emptyDir: {}
 containers:
 - name: vm
 image: vm-vhostuser:latest
 command: ["bash", "/root/vm/start.sh"]
 securityContext:
 capabilities:
 add:
 - NET_BIND_SERVICE
 - SYS_NICE
 - NET_RAW
 - NET_ADMIN
 privileged: false
 runAsUser: 0
 resources:
 limits:
 cpu: '2'
 devices.kubevirt.io/kvm: '1'
 memory: '8784969729'
 hugepages-2Mi: 2Gi
 requests:
 cpu: 666m
 devices.kubevirt.io/kvm: '1'
 ephemeral-storage: 50M
 memory: '4490002433'
 volumeMounts:
 - name: vhostuser-sockets
 mountPath: /var/run/vm
 - name: xml
 mountPath: /root/vm/
 - mountPath: /dev/hugepages
 name: hugepage
 - name: libvirt-runtime
 mountPath: /var/run/libvirt

```

## Pod

```

virsh set-user-password vm root 12345
Password set successfully for root in vm

virsh console vm
Connected to domain 'vm'
Escape character is ^] (Ctrl +])

CentOS Linux 7 (Core)
Kernel 3.10.0-1127.el7.x86_64 on an x86_64

```

```
localhost login: root
Password:
Last login: Fri Feb 25 09:52:54 on ttys0
```

```
ip link set eth0 mtu 1400
ip addr add 10.16.0.96/16 dev eth0
ip ro add default via 10.16.0.1
ping 114.114.114.114
```

[PDF](#)[Slack](#)[Support](#)

⌚2022 7 16

⌚2022 5 24

⌚GitHub 

---

7.19.6

## 7.20 OpenStack



### 7.20.1

OVN-IC      Kubernetes      OpenStack    Kubernetes

1.      OpenStack    Kubernetes      CIDR
- 2.
3.      IP
4.      Kubernetes      OpenStack      VPC

#### OVN-IC

OVN-IC

```
docker run --name=ovn-ic-db -d --network=host -v /etc/ovn/:/etc/ovn -v /var/run/ovn:/var/run/ovn -v /var/log/ovn:/var/log/ovn kubeovn/kube-ovn:v1.14.9 bash start-ic-db.sh
```

#### Kubernetes

kube-system Namespace      ovn-ic-config ConfigMap

```
apiVersion: v1
kind: ConfigMap
metadata:
 name: ovn-ic-config
 namespace: kube-system
data:
 enable-ic: "true"
 az-name: "az1"
 ic-db-host: "192.168.65.3"
 ic-nb-port: "6645"
 ic-sb-port: "6646"
 gw-nodes: "az1-gw"
 auto-route: "true"
```

- enable-ic:
- az-name:
- ic-db-host:      OVN-IC
- ic-nb-port:      6645
- ic-sb-port:      OVN-IC      6646
- gw-nodes:
- auto-route:

#### OpenStack

Kubernetes

```
openstack router create router0
openstack router list
+-----+-----+-----+-----+
| ID | Name | Status | State | Project |
+-----+-----+-----+-----+
| d5b38655-249a-4192-8046-71aa4d2b4af1 | router0 | ACTIVE | UP | 98a29ab7388347e7b5ff8bdd181ba4f9 |
+-----+-----+-----+-----+
```

## OpenStack OVN

```
ovn-nbctl set NB_Global . name=op-az
```

OVN-IC OVN-IC

```
/usr/share/ovn/scripts/ovn-ctl --ovn-ic-nb-db=tcp:192.168.65.3:6645 \
--ovn-ic-sb-db=tcp:192.168.65.3:6646 \
--ovn-northd-nb-db=unix:/run/ovnnb_db.sock \
--ovn-northd-sb-db=unix:/run/ovn/ovnsb_db.sock \
start_ic
```

- ovn-ic-nb-db ovn-ic-sb-db : OVN-IC
- ovn-northd-nb-db ovn-northd-sb-db : OVN

```
ovs-vsctl set open_vswitch . external_ids:ovn-is-interconn=true
```

## OpenStack OVN

ts router0

```
ovn-nbctl lrp-add router0 lrp-router0-ts 00:02:ef:11:39:4f 169.254.100.73/24
ovn-nbctl lsp-add ts lsp-ts-router0 -- lsp-set-addresses lsp-ts-router0 router \
-- lsp-set-type lsp-ts-router0 router \
-- lsp-set-options lsp-ts-router0 router-port=lrp-router0-ts
ovn-nbctl lrp-set-gateway-chassis lrp-router0-ts {gateway chassis} 1000
ovn-nbctl set NB_Global . options:ic-route-adv=true options:ic-route-learn=true
```

## Kubernetes

```
ovn-nbctl lr-route-list router0
IPv4 Routes
 10.0.0.22 169.254.100.34 dst-ip (learned)
 10.16.0.0/16 169.254.100.34 dst-ip (learned)
```

router0 Kubernetes Pod

## 7.20.2 OVN

OpenStack Kubernetes OVN VPC Subnet

Kube-OVN OVN OpenStack Neutron OVN OpenStack networking-ovn Neutron

### Neutron

Neutron /etc/neutron/plugins/ml2/ml2\_conf.ini

```
[ovn]
...
ovn_nb_connection = tcp:[192.168.137.176]:6641,tcp:[192.168.137.177]:6641,tcp:[192.168.137.178]:6641
ovn_sb_connection = tcp:[192.168.137.176]:6642,tcp:[192.168.137.177]:6642,tcp:[192.168.137.178]:6642
ovn_l3_scheduler = OVN_L3_SCHEDULER
```

- ovn\_nb\_connection ovn\_sb\_connection : Kube-OVN ovn-central

### OVS

```
ovs-vsctl set open . external_ids:ovn-remote=tcp:[192.168.137.176]:6642,tcp:[192.168.137.177]:6642,tcp:[192.168.137.178]:6642
ovs-vsctl set open . external_ids:ovn-encap-type=geneve
ovs-vsctl set open . external_ids:ovn-encap-ip=192.168.137.200
```

- external\_ids:ovn-remote : Kube-OVN ovn-central
- ovn-encap-ip : IP

## Kubernetes OpenStack

Kubernetes    OpenStack    OpenStack    Pod



kube-ovn-controller args --enable-external-vpc=true

## OpenStack

```
openstack router list
+-----+-----+-----+
| ID | Name | Status | State | Project |
+-----+-----+-----+
| 22040ed5-0598-4f77-bffd-e7fd4db47e93 | router0 | ACTIVE | UP | 62381a21d569404aa236a5d8712449c |
+-----+-----+-----+
openstack network list
+-----+-----+
| ID | Name | Subnets |
+-----+-----+
| cd59e36a-37db-4c27-b709-d35379a7920f | provider | 01d73d9f-fdaa-426c-9b60-aa34abbfaeae |
+-----+-----+
openstack subnet list
+-----+-----+-----+
| ID | Name | Network | Subnet |
+-----+-----+-----+
| 01d73d9f-fdaa-426c-9b60-aa34abbfaeae | provider-v4 | cd59e36a-37db-4c27-b709-d35379a7920f | 192.168.1.0/24 |
+-----+-----+-----+
openstack server list
+-----+-----+-----+-----+-----+
| ID | Name | Status | Networks | Image | Flavor |
+-----+-----+-----+-----+-----+
| 8433d622-a8d6-41a7-8b31-49abfd64f639 | provider-instance | ACTIVE | provider=192.168.1.61 | ubuntu | m1 |
+-----+-----+-----+-----+-----+
```

## Kubernetes VPC

```
kubectl get vpc
NAME STANDBY SUBNETS
neutron-22040ed5-0598-4f77-bffd-e7fd4db47e93 true ["neutron-cd59e36a-37db-4c27-b709-d35379a7920f"]
ovn-cluster true ["join", "ovn-default"]
```

neutron-22040ed5-0598-4f77-bffd-e7fd4db47e93    OpenStack    VPC

Kube-OVN    VPC    Subnet    Pod

VPC, Subnet    Namespace net2    Pod:

```
apiVersion: v1
kind: Namespace
metadata:
 name: net2

apiVersion: kubeovn.io/v1
kind: Vpc
metadata:
 creationTimestamp: "2021-06-20T13:34:11Z"
 generation: 2
 labels:
 ovn.kubernetes.io/vpc_external: "true"
 name: neutron-22040ed5-0598-4f77-bffd-e7fd4db47e93
 resourceVersion: "583728"
 uid: 18d4c654-f511-4def-a3a0-a6434d237c1e
spec:
 namespaces:
 - net2

kind: Subnet
apiVersion: kubeovn.io/v1
metadata:
 name: net2
spec:
 vpc: neutron-22040ed5-0598-4f77-bffd-e7fd4db47e93
 namespaces:
 - net2
 cidrBlock: 12.0.1.0/24
 natOutgoing: false

apiVersion: v1
kind: Pod
metadata:
 name: ubuntu
```

```
namespace: net2
spec:
 containers:
 - image: docker.io/kubeovn/kube-ovn:v1.8.0
 command:
 - "sleep"
 - "604800"
 imagePullPolicy: IfNotPresent
 name: ubuntu
 restartPolicy: Always
```

[!\[\]\(72b4102a7f0c50652b1c0eac6181e615\_img.jpg\) PDF](#)[!\[\]\(e891cd23bae084630236c8c5e259a26b\_img.jpg\) Slack](#)[!\[\]\(2f34b0916eddf922c68ab45979b70ef5\_img.jpg\) Support](#)

 2025 8 4

 2022 5 24



7.20.3

---

## 7.21 IPsec

v1.13.0      UDP 500 4500

### 7.21.1

```
kube-ovn-cni certificatesigningrequest kube-ovn-controller kube-ovn-controller approve kube-ovn-cni ipsec
ipsec
```

### 7.21.2 IPsec

```
kube-ovn-controller kube-ovn-cni args --enable-ovn-ipsec=false --enable-ovn-ipsec=true
```

[!\[\]\(4ea2379173bd1bb922c00d50df3f1d87\_img.jpg\) PDF](#)[!\[\]\(129b5c615a454e46a1c5a569fea4d422\_img.jpg\) Slack](#)[!\[\]\(0241d1e18791ae3f7b7105f7f26b42cb\_img.jpg\) Support](#)

 2025 9 10

 2023 4 18



 GitHub



### 7.21.3

## 7.22 OVN

Pod	GRE/ERSPAN
Kube-OVN	v1.12

### 7.22.1 Multus-CNI

Multus-CNI      Multus

### 7.22.2

```
apiVersion: "k8s.cni.cncf.io/v1"
kind: NetworkAttachmentDefinition
metadata:
 name: attachnet
 namespace: default
spec:
 config: |
 {
 "cniVersion": "0.3.1",
 "type": "kube-ovn",
 "server_socket": "/run/openvswitch/kube-ovn-daemon.sock",
 "provider": "attachnet.default.ovn"
 }

```

provider <NAME>.<NAMESPACE>.ovn

### 7.22.3 Underlay

MTU      LSP/Pod      Underlay

Underlay

```
apiVersion: kubeovn.io/v1
kind: ProviderNetwork
metadata:
 name: net1
spec:
 defaultInterface: eth1

apiVersion: kubeovn.io/v1
kind: Vlan
metadata:
 name: vlan1
spec:
 id: 0
 provider: net1

apiVersion: kubeovn.io/v1
kind: Subnet
metadata:
 name: subnet1
spec:
 protocol: IPv4
 cidrBlock: 172.19.0.0/16
 excludeIps:
 - 172.19.0.2..172.19.0.20
 gateway: 172.19.0.1
 vlan: vlan1
 provider: attachnet.default.ovn

```

provider      provider

### 7.22.4 Pod

Pod

```
apiVersion: v1
kind: Pod

```

```

metadata:
 name: pod1
 annotations:
 k8s.v1.cni.cncf.io/networks: default/attachnet
spec:
 containers:
 - name: bash
 image: docker.io/kubeovn/kube-ovn:v1.14.9
 args:
 - bash
 - -c
 - sleep infinity
 securityContext:
 privileged: true

```

### Pod IP

```
$ kubectl get ips | grep pod1
pod1.default 10.16.0.12 00:00:00:FF:34:24 kube-ovn-worker ovn-default
pod1.default.attachnet.default.ovn 172.19.0.21 00:00:00:A0:30:68 kube-ovn-worker subnet1
```

IP 172.19.0.21

## 7.22.5 OVN

### OVN

```
kubectl ko nbctl mirror-add mirror1 gre 99 from-lport 172.19.0.21
kubectl ko nbctl lsp-attach-mirror coredns-787d4945fb-gpnkb.kube-system mirror1
```

coredns-787d4945fb-gpnkb.kube-system OVN LSP <POD\_NAME>.<POD\_NAMESPACE>

### OVN

```
ovn-nbctl mirror-add <NAME> <TYPE> <INDEX> <FILTER> <IP>

NAME - add a mirror with given name
TYPE - specify TYPE 'gre' or 'erspan'
INDEX - specify the tunnel INDEX value
 (indicates key if GRE, erSPAN_idx if ERSPAN)
FILTER - specify FILTER for mirroring selection
 ('to-lport' / 'from-lport')
IP - specify Sink / Destination i.e. Remote IP

ovn-nbctl mirror-del [NAME] remove mirrors
ovn-nbctl mirror-list print mirrors

ovn-nbctl lsp-attach-mirror PORT MIRROR attach source PORT to MIRROR
ovn-nbctl lsp-detach-mirror PORT MIRROR detach source PORT from MIRROR
```

## 7.22.6 Pod

### Pod

```
root@pod1:/kube-ovn# ip link add mirror1 type gretap local 172.19.0.21 key 99 dev net1
root@pod1:/kube-ovn# ip link set mirror1 up
```

### Pod

```
root@pod1:/kube-ovn# tcpdump -i mirror1 -nnve
tcpdump: listening on mirror1, link-type EN10MB (Ethernet), snapshot length 262144 bytes
05:13:30.328800 00:00:00:a3:f5:e2 > 00:00:00:97:0f:6e, ethertype ARP (0x0806), length 42: Ethernet (len 6), IPv4 (len 4), Request who-has 10.16.0.7 tell 10.16.0.4, length 28
05:13:30.559167 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 212: (tos 0x0, ttl 64, id 57364, offset 0, flags [DF], proto UDP (17), length 198)
 10.16.0.4.53 > 10.16.0.6.50472: 34511 NXDomain*- 0/1/1 (170)
05:13:30.560625 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 212: (tos 0x0, ttl 64, id 57365, offset 0, flags [DF], proto UDP (17), length 198)
 10.16.0.4.53 > 10.16.0.6.45177: 1659 NXDomain*- 0/1/1 (170)
05:13:30.562774 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 191: (tos 0x0, ttl 64, id 57368, offset 0, flags [DF], proto UDP (17), length 177)
 10.16.0.4.53 > 10.16.0.6.37755: 48737 NXDomain*- 0/1/1 (149)
05:13:30.563523 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 187: (tos 0x0, ttl 64, id 57369, offset 0, flags [DF], proto UDP (17), length 173)
 10.16.0.4.53 > 10.16.0.6.53887: 45519 NXDomain*- 0/1/1 (145)
05:13:30.564940 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 201: (tos 0x0, ttl 64, id 57370, offset 0, flags [DF], proto UDP (17),
```

```

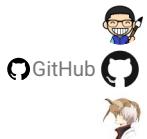
length 187)
 10.16.0.4.53 > 10.16.0.6.40846: 25745 NXDomain*- 0/1/1 (159)
05:13:30.565140 00:00:00:a3:f5:e2 > 00:00:00:89:d5:cc, ethertype IPv4 (0x0800), length 201: (tos 0x0, ttl 64, id 57371, offset 0, flags [DF], proto UDP (17),
length 187)
 10.16.0.4.53 > 10.16.0.6.45214: 61875 NXDomain*- 0/1/1 (159)
05:13:30.566023 00:00:00:a3:f5:e2 > 00:00:00:55:e4:4e, ethertype IPv4 (0x0800), length 80: (tos 0x0, ttl 64, id 45937, offset 0, flags [DF], proto UDP (17),
length 66)
 10.16.0.4.44116 > 172.18.0.1.53: 16025+ [1au] AAAA? kube-ovn.io. (38)

```

## 7.22.7

1.	ERSPAN	OVN	Linux	4.14	ERSPAN	IPv6	Linux	4.16
2.		OVN						

[PDF](#)
[Slack](#)
[Support](#)
 2025 7 2

 2023 4 20


## 7.22.8

## 7.23 DNS Kube-OVN

NodeLocal DNSCache      DaemonSet      DNS      DNS      Kube-OVN

### 7.23.1 DNS

Kubernetes      DNS

Kubernetes      Nodelocaldnscache

```
#!/bin/bash

localdns=169.254.20.10
domain=cluster.local
kubedns=10.96.0.10

wget https://raw.githubusercontent.com/kubernetes/kubernetes/master/cluster/addons/dns/nodelocaldns/nodelocaldns.yaml
sed -i "s/_PILLAR__LOCAL__DNS__/$localdns/g; s/_PILLAR__DNS__DOMAIN__/$domain/g; s/_PILLAR__DNS__SERVER__//g; s/_PILLAR__CLUSTER__DNS__/$kubedns/g"
nodelocaldns.yaml

kubectl apply -f nodelocaldns.yaml
```

kubelet      /var/lib/kubelet/config.yaml      clusterDNS      DNS IP 169.254.20.10      kubelet

### Kube-OVN DNS

Kubernetes      Nodelocal DNSCache      Kube-OVN

UNDERLAY SUBNET      U20

Underlay Subnet      DNS      U20      kubectl edit subnet {your subnet}      spec.u2oInterconnection = true ,      Overlay Subnet

KUBE-OVN-CONTROLLER      DNS IP

```
kubectl edit deployment kube-ovn-controller -n kube-system

spec.template.spec.containers.args --node-local-dns-ip=169.254.20.10
```

POD

Pod      /etc/resolv.conf      nameserver      DNS IP      Pod      nameserver      DNS ClusterIP      u2o      Pod      Pod

## 7.23.2 DNS

Pod      Pod      DNS      169.254.20.10

```
kubectl exec -it pod1 -- nslookup github.com
Server: 169.254.20.10
Address: 169.254.20.10:53
```

```
Name: github.com
Address: 20.205.243.166
```

DNS      ovn0      DNS      DNS

```
tcpdump -i any port 53

06:20:00.441889 659246098c56_h P ifindex 17 00:00:00:73:f1:06 ethertype IPv4 (0x0800), length 75: 10.16.0.2.40230 > 169.254.20.10.53: 1291+ A? baidu.com. (27)
06:20:00.441889 ovn0 In ifindex 7 00:00:00:50:32:cd ethertype IPv4 (0x0800), length 75: 10.16.0.2.40230 > 169.254.20.10.53: 1291+ A? baidu.com. (27)
06:20:00.441950 659246098c56_h P ifindex 17 00:00:00:73:f1:06 ethertype IPv4 (0x0800), length 75: 10.16.0.2.40230 > 169.254.20.10.53: 1611+ AAAA?
baidu.com. (27)
06:20:00.441950 ovn0 In ifindex 7 00:00:00:50:32:cd ethertype IPv4 (0x0800), length 75: 10.16.0.2.40230 > 169.254.20.10.53: 1611+ AAAA? baidu.com. (27)
06:20:00.442203 ovn0 Out ifindex 7 00:00:00:52:99:d8 ethertype IPv4 (0x0800), length 145: 169.254.20.10.53 > 10.16.0.2.40230: 1611* 0/1/0 (97)
```

```

06:20:00.442219 659246098c56_h Out ifindex 17 00:00:00:ea:b3:5e ethertype IPv4 (0x0800), length 145: 169.254.20.10.53 > 10.16.0.2.40230: 1611* 0/1/0 (97)
06:20:00.442273 ovn0 Out ifindex 7 00:00:00:52:99:d8 ethertype IPv4 (0x0800), length 125: 169.254.20.10.53 > 10.16.0.2.40230: 1291* 2/0/0 A 39.156.66.10, A
110.242.68.66 (77)
06:20:00.442278 659246098c56_h Out ifindex 17 00:00:00:ea:b3:5e ethertype IPv4 (0x0800), length 125: 169.254.20.10.53 > 10.16.0.2.40230: 1291* 2/0/0 A 39.
156.66.10, A 110.242.68.66 (77)

```

### 7.23.3



NetworkPolicy	NetworkPolicy	DNS IP	169.254.20.10	CIDR	NetworkPolicy	DNS	Pod
---------------	---------------	--------	---------------	------	---------------	-----	-----

#### NetworkPolicy

Pod	DNS	NetworkPolicy
-----	-----	---------------

```

apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
 name: allow-local-dns-and-node-cidr
 namespace: default #
spec:
 podSelector: {} # Pod
 policyTypes:
 - Ingress
 - Egress
 egress:
 # DNS
 - to:
 - ipBlock:
 cidr: 169.254.20.10/32
 # CIDR
 - to:
 - ipBlock:
 cidr: 10.0.0.0/8 # CIDR
 ingress:
 # DNS
 - from:
 - ipBlock:
 cidr: 169.254.20.10/32
 # CIDR
 - from:
 - ipBlock:
 cidr: 10.0.0.0/8 # CIDR

```

- 169.254.20.10/32 DNS IP
- 10.0.0.0/8 CIDR

[PDF](#)
[Slack](#)
[Support](#)

⌚2025 7 3

⌚2023 5 5



[GitHub](#)



### 7.23.4

## 7.24 VPC NAT

### 7.24.1

VPC	Overlay	natOutgoing	Subnet	Pod	SNAT	IP	Pod	SNAT
NAT		CIDR	IP		SNAT			

### 7.24.2

subnet.Spec    natOutgoing    natOutgoingPolicyRules

```
spec:
 natOutgoing: true
 natOutgoingPolicyRules:
 - action: forward
 match:
 srcIPs: 10.0.11.0/30,10.0.11.254
 - action: nat
 match:
 srcIPs: 10.0.11.128/26
 dstIPs: 114.114.114.114,8.8.8.8
```

NAT

1. IP 10.0.11.0/30 10.0.11.254 SNAT
2. IP 10.0.11.128/26 IP 114.114.114.114 8.8.8.8 SNAT

action	match	action, action	forward	nat	forward	SNAT, nat	SNAT	natOutgoingPolicyRules			
SNAT											
match		srcIPs	dstIPs			IP	IP	match.srcIPs	match.dstIPs	CIDR	IP
	match		natOutgoingPolicyRules								

[PDF](#)

[Slack](#)

[Support](#)

⌚2023 10 25

⌚2023 6 5



### 7.24.3

## 8.

## 8.1

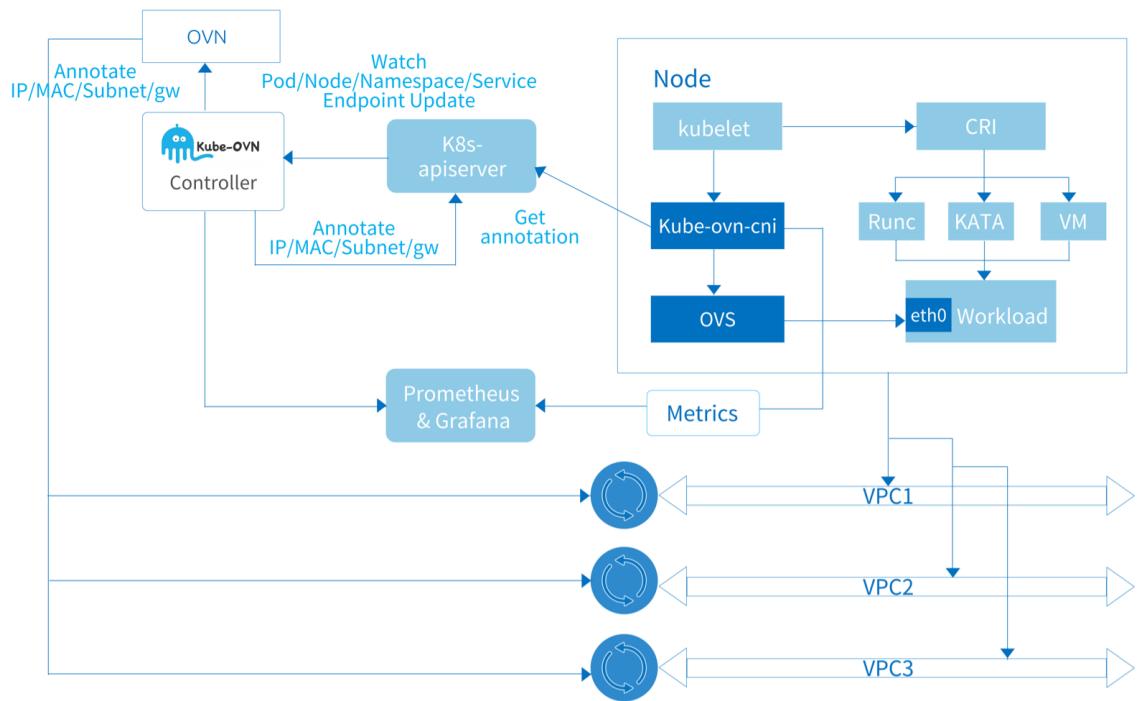
## Kube-OVN

Kube-OVN	Kubernetes	OVN	SDN	Kube-OVN	OVN	Kubernetes	CNI Service	Networkpolicy
SDN	VPC	QoS ACL						
Kube-OVN			Cilium Submariner Prometheus KubeVirt					

## 8.1.1

## Kube-OVN

- OVN/OVS
- Agent
- 



## OVN/OVS

OVN/OVS	Kube-OVN	OVN/OVS	SDN	Kube-OVN	ovn-architecture(7)
OVN	Kube-OVN	OVN	Kubernetes		
OVN/OVS		Kubernetes			

**OVN-CENTRAL**

```

ovn-central Deployment OVN ovn-nb , ovn-sb , ovn-northd
 • ovn-nb API kube-ovn-controller ovn-nb
 • ovn-sb ovn-nb
 • ovn-northd ovn-nb ovn-sb
ovn-central Raft

```

**OVS-OVN**

```

ovs-ovn DaemonSet Pod openvswitch, ovsdb, ovn-controller ovn-central Agent

```

**Agent**

```
Kube-OVN OVN Kubernetes
```

**KUBE-OVN-CONTROLLER**

Deployment	Kubernetes	OVN	Kube-OVN	kube-ovn-controller	OVN
Pod	Service	Endpoint	Node	NetworkPolicy	VPC
Pod	annotation	IPAM	ovn-central	Subnet	Vlan
CIDR		annotation	kube-ovn-cni		ProviderNetwork

**KUBE-OVN-CNI**

DaemonSet	CNI	OVS
DaemonSet	kube-ovn	kubelet  kube-ovn-cni
		CNI      kube-ovn-cni
		/opt/cni/bin

`kube-ovn-cni`

1. `ovn-controller vswitchd`
2. CNI add/del
  - a. veth OVS
  - b. OVS
  - c. iptables/ipset/route
3. QoS.
4. `ovn0`
5. Vlan/Underlay/EIP
- 6.

**Kube-OVN****KUBE-OVN-SPEAKER**

DaemonSet	Pod IP
	<b>BGP</b>

**KUBE-OVN-PINGER**

DaemonSet	OVS	Kube-OVN
-----------	-----	----------

**KUBE-OVN-MONITOR**

Deployment	OVN	Kube-OVN
------------	-----	----------

KUBECTL-KO

kubectl

kubectl

 PDF

 Slack

 Support

 2025 9 10

 2022 5 24



 GitHub 



8.1.2

---

## 8.2 What's Next

---

This document lists the features merged into the master branch for the next minor release.

### 8.2.1 Post-v1.14.0

- IPPool
- Multiple IPPools now can bind to the same Namespace. [#5731](#)
- Pods in a bound namespace will only get IPs from the bound pool(s), not other ranges in the subnet. [#5731](#)
- `AdminNetworkPolicy` now supports specify egress peers using FQDNs. [#5703](#)
- Using ARP for IPv4 network ready check: now you don't need ACL allow rules for gateway to make Pod running. [#5716](#)
- Non-primary CNI mode: you can run Kube-OVN as the secondary only network, without annoying unused annotations and logical switch port allocations. [#5618](#)
- VPC NAT Gateway:
  - No default EIP mode: the secondary interface can initialize without a default EIP to avoid the waste. [#5605](#)
  - Custom routes: you can control the route rules within the `vpc-nat-gateway` Pods to control traffic paths. [#5608](#)
  - Gratuitous ARP: VPC NAT Gateway automatically sends gratuitous ARP packets during initialization to accelerate network convergence. [#5607](#)
  - Healthchecks for static endpoints in `SwitchLBRules`: SLR with both selector or endpoints key can support healthchecks. [#5435](#)
- Underlay
  - Node Selectors for `ProviderNetwork`: instead of adding/removing nodes to the `ProviderNetwork` one by one, you can use node selectors to simplify the workflow. [#5518](#)
  - Different `NetworkProvider`s can now share the same VLAN. [#5471](#)
  - Adding `pod_name` and `pod_namespace` labels to interface metrics. [#5463](#)
- IPSec
  - Support `cert-manager` to issue certificates. [#5365](#)
  - Request new certificate if current certificate is not trusted. [#5710](#)
  - `kubectl-ko`
    - Collect IPSec and xfrm information. [#5472](#)
    - Replace `Endpoint` with `EndpointSlice`. [#5425](#)
  - NetworkAttachment caching: reduce APIServer load in large-scale deployments with Multus. [#5386](#)
  - Upgrade OVS to 3.5 and OVN to 25.03. [#5537](#)



2025 9 19

2025 9 15



### 8.2.2

---

## 8.3

---

### Kube-OVN

#### 8.3.1

##### Kube-OVN

```
1. default Pod IP CIDR 10.16.0.0/16 10.16.0.1
2. join Node Pod , CIDR 100.64.0.0/16 100.64.0.1
```

```
POD_CIDR="10.16.0.0/16"
POD_GATEWAY="10.16.0.1"
JOIN_CIDR="100.64.0.0/16"
EXCLUDE_IPS=""
```

```
EXCLUDE_IP POD_CIDR 192.168.10.20..192.168.10.30
```

Overlay              Service CIDR

[Join](#)

#### 8.3.2 Service

```
kube-proxy iptables Kube-OVN Kube-OVN Service CIDR
```

```
SVC_CIDR="10.96.0.0/12"
```

[kube-ovn-controller Deployment](#)

```
args:
- --service-cluster-ip-range=10.96.0.0/12
```

#### 8.3.3 Overlay

Kube-OVN        Kubernetes Node IP

```
IFACE=eth1
```

```
ens[a-z0-9]*,eth[a-z0-9]*
```

[kube-ovn-cni DaemonSet](#)

```
args:
- --iface=eth1
```

```
annotation ovn.kubernetes.io/tunnel_interface annotation iface annotation
```

```
kubectl annotate node no1 ovn.kubernetes.io/tunnel_interface=ethx
```

### 8.3.4 MTU

Overlay MTU	Kube-OVN	MTU	Overlay MTU	Pod MTU	MTU - 100 Underlay Pod
Overlay MTU	kube-ovn-cni DaemonSet				
			args: - --mtu=1333		

### 8.3.5

Kube-OVN	mirror0	tcpdump
----------	---------	---------

ENABLE_MIRROR=true
--------------------

kube-ovn-cni DaemonSet	:
------------------------	---

args: - --enable-mirror=true
---------------------------------

### 8.3.6 LB

Kube-OVN	OVN	L2 LB	Service	Overlay	kube-proxy	Service	,	Kube-OVN	LB
----------	-----	-------	---------	---------	------------	---------	---	----------	----

ENABLE_LB=false
-----------------

kube-ovn-controller Deployment
--------------------------------

args: - --enable-lb=false
------------------------------

LB

Kube-OVN v1.12.0	Deployment	subnet crd	spec enableLb	Kube-OVN	LB	LB	kube-ovn-controller
		enable-lb	load-balancer	enableLb	load-balancer	v1.12.0	enableLb

### 8.3.7 NetworkPolicy

Kube-OVN	OVN	ACL	NetworkPolicy	NetworkPolicy	Cilium	Chain	eBPF	NetworkPolicy	Kube-OVN
NetworkPolicy									

ENABLE_NP=false
-----------------

kube-ovn-controller Deployment
--------------------------------

args: - --enable-np=false
------------------------------

NetworkPolicy

### 8.3.8 EIP SNAT

EIP SNAT	kube-ovn-controller
----------	---------------------

```
ENABLE_EIP_SNAT=false
```

#### kube-ovn-controller Deployment

```
args:
- --enable-eip-snat=false
```

EIP SNAT

EIP SNAT

### 8.3.9 Load Balancer Service

VPC	Load Balancer Service
-----	-----------------------

LoadBalancer	Service
--------------	---------

```
ENABLE_LB_SVC=true
```

#### kube-ovn-controller Deployment

```
args:
- --enable-lb-svc=true
```

### 8.3.10 ECMP

ECMP	ECMP	kube-ovn-controller Deployment	:
------	------	--------------------------------	---

```
args:
- --enable-ecmp=true
```

Kube-OVN v1.12.0	subnet crd	spec enableEcmp	ECMP	ECMP	kube-ovn-controller
Deployment	enable-ecmp	v1.12.0			

### 8.3.11 Kubevirt VM

Kubevirt	VM	kube-ovn-controller	StatefulSet Pod	IP	VM
----------	----	---------------------	-----------------	----	----

1.10.6		kube-ovn-controller Deployment			
--------	--	--------------------------------	--	--	--

```
args:
- --keep-vm-ip=false
```

### 8.3.12 CNI

Kube-OVN	/opt/cni/bin	CNI	/etc/cni/net.d	CNI	01-kube-ovn.conflist	CNI
----------	--------------	-----	----------------	-----	----------------------	-----

```
CNI_CONF_DIR="/etc/cni/net.d"
CNI_BIN_DIR="/opt/cni/bin"
CNI_CONFIG_PRIORITY="01"
```

#### kube-ovn-cni DaemonSet Volume

```
volumes:
- name: cni-conf
 hostPath:
 path: "/etc/cni/net.d"
- name: cni-bin
 hostPath:
 path: "/opt/cni/bin"
...
args:
- --cni-conf-name=01-kube-ovn.conflist
```

### 8.3.13

Kube-OVN   Overlay      Geneve      Vxlan   STT

```
TUNNEL_TYPE="vxlan"
```

ovs-ovn DaemonSet

```
env:
- name: TUNNEL_TYPE
 value: "vxlan"
```

STT      ovs

### 8.3.14 SSL

OVN DB   API      SSL :

```
ENABLE_SSL=true
```

SSL

### 8.3.15 ip

kube-ovn-controller/kube-ovn-cni/kube-ovn-monitor      ip      0.0.0.0      ip

```
ENABLE_BIND_LOCAL_IP=true
```

kube-ovn-monitor      pod ip

```
netstat -tunlp |grep kube-ovn
tcp 0 0 172.18.0.5:10661 0.0.0.0:*
 LISTEN 2612/.kube-ovn-mon
```

deployment    daemonSet

```
env:
- name: ENABLE_BIND_LOCAL_IP
 value: "false"
```

 PDF

 Slack

 Support

⌚2025 9 17

⌚2022 5 24



### 8.3.16

## 8.4

---

Kube-OVN	Minor	Patch	Minor	OVN/OVS	API	Patch	Bug	API
----------	-------	-------	-------	---------	-----	-------	-----	-----

### 8.4.1

Kube-OVN		master, release-1.12	release-1.11		release-1.12		Bug
backport							
		release-1.11	backport	Bug			

### 8.4.2

Minor		Patch	Bug	Bug
-------	--	-------	-----	-----

### 8.4.3 Patch

Patch	<a href="#">hack/release.sh</a>							
1.	Build							
2.	tag	Docker Hub						
3.	tag	Github						
4.								
5.								
6.	Release Note PR							
7.	Release Note ( )							
8.	Merge github action	Release Note PR						
9.	Github Release							
10.	Github Release	Release	v1.12.12		Release Note	Release		

### 8.4.4 Minor

Minor									
1.	Github	release-1.13	( )						
2.	VERSION, dist/images/install.sh, charts/kube-ovn/values.yaml	charts/kube-ovn/Chart.yaml		Minor	v1.14.0	( )			
3.	tag	Docker Hub ( )							
4.	tag	Github ( )							
5.	v1.13	mkdocs.yml	version branch ( )						
6.	Release Note PR								
7.	Release Note ( )								
8.	Merge github action	Release Note PR							
9.	Github Release								
10.	Github Release	Release	v1.13.0		Release Note	Release			
11.	VERSION	Patch	v1.13.1						



PDF



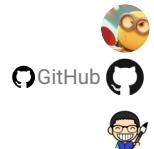
Slack



Support

⌚2024 5 16

⌚2024 5 8



8.4.5

---

## 8.5

---

Kube-OVN

**Alpha Beta GA**

### 8.5.1

---

#### **Alpha**

- 
- 
- API
- 
- 

#### **Beta**

- 
- 
- API
- 
- API

#### **GA**

- 
- 
- API
-

## 8.5.2

v1.8

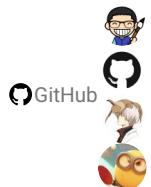
			Since	Until
Namespaced Subnet	true	GA	1.8	
	true	GA	1.8	
	true	GA	1.8	
ECMP	false	Beta	1.8	
ACL	true	Alpha	1.9	
( ACL )	true	Beta	1.8	
Underlay	true	GA	1.8	
	true	Beta	1.8	
DHCP	false	Alpha	1.10	
	false	Alpha	1.8	
OVN-IC	false	Beta	1.8	
Submariner	false	Alpha	1.9	
VIP	true	Alpha	1.10	
VPC	true	Beta	1.8	
VPC IP/SNAT/DNAT	true	Alpha	1.10	
VPC	true	Alpha	1.10	
VPC	true	Alpha	1.10	
VPC	true	Alpha	1.10	
QoS	true	GA	1.8	
linux-netem QoS	true	Alpha	1.9	
Grafana	false	GA	1.8	
	false	GA	1.8	
VPC EIP/SNAT	false	Beta	1.8	
	false	GA	1.8	
NetworkPolicy	true	Beta	1.8	
Webhook	false	Alpha	1.10	
	false	Beta	1.8	
Overlay	false	Alpha	1.8	
Overlay BGP	false	Alpha	1.9	
Cilium	false	Alpha	1.10	
VPC	false	Alpha	1.10	
Mellanox Offload	false	Alpha	1.8	
Offload	false	Alpha	1.10	
DPDK	false	Alpha	1.10	

			<b>Since</b>	<b>Until</b>
OpenStack		false	Alpha	1.9
Pod IP/Mac		true	GA	1.8
Workload IP		true	GA	1.8
StatefulSet IP		true	GA	1.8
VM IP		false	Beta	1.9
VPC Load Balancer Service		false	Alpha	1.11
VPC DNS		false	Alpha	1.11
Underlay Overlay		false	Beta	1.12
VPC Egress Gateway		true	Alpha	1.14

[!\[\]\(1aeedb8f6b9d7238c4caba7fbf5b5fdc\_img.jpg\) PDF](#)[!\[\]\(582ba375fc0a70dda840efab1d53e2a1\_img.jpg\) Slack](#)[!\[\]\(61ef16e7e0826837dbeb6e5059dc980a\_img.jpg\) Support](#)

⌚2025 9 17

⌚2022 8 31

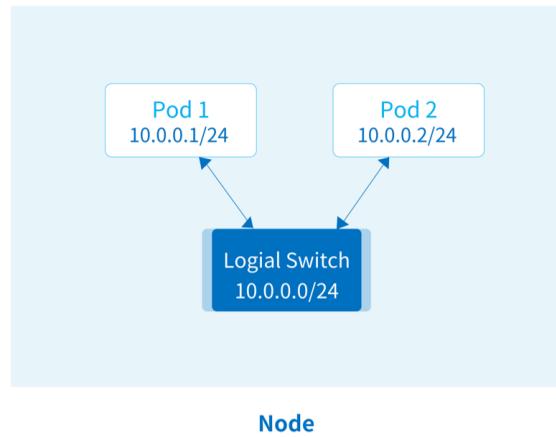


8.5.3

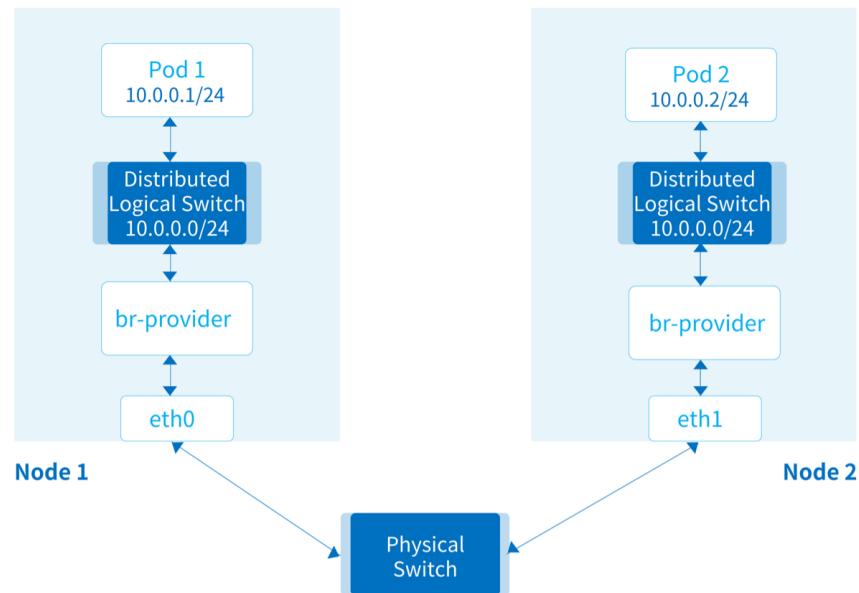
## 8.6 Underlay

### Underlay

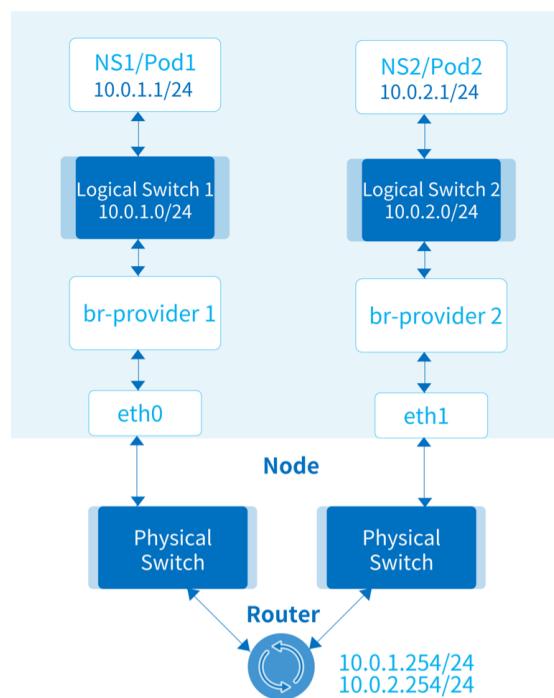
#### 8.6.1



#### 8.6.2



## 8.6.3

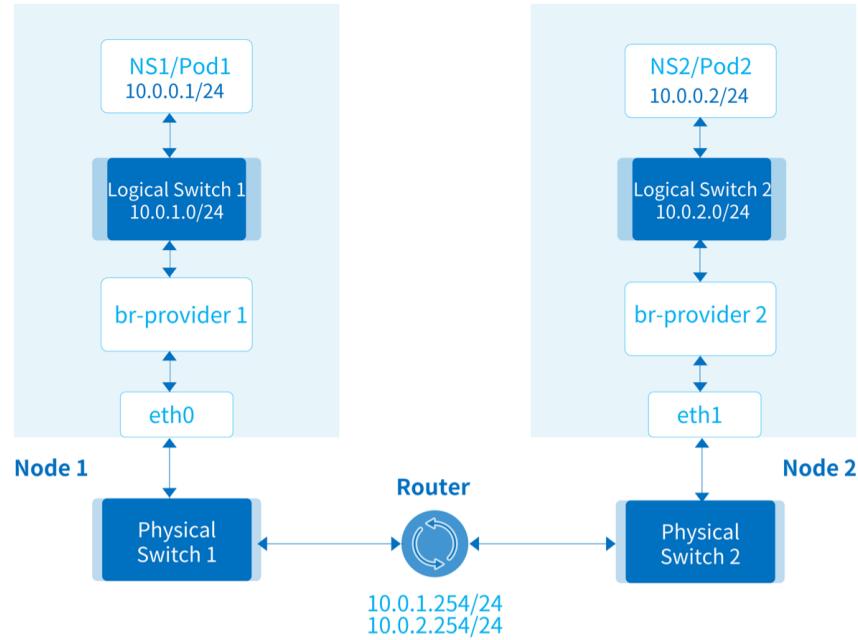


br-provider-1 br-provider-2

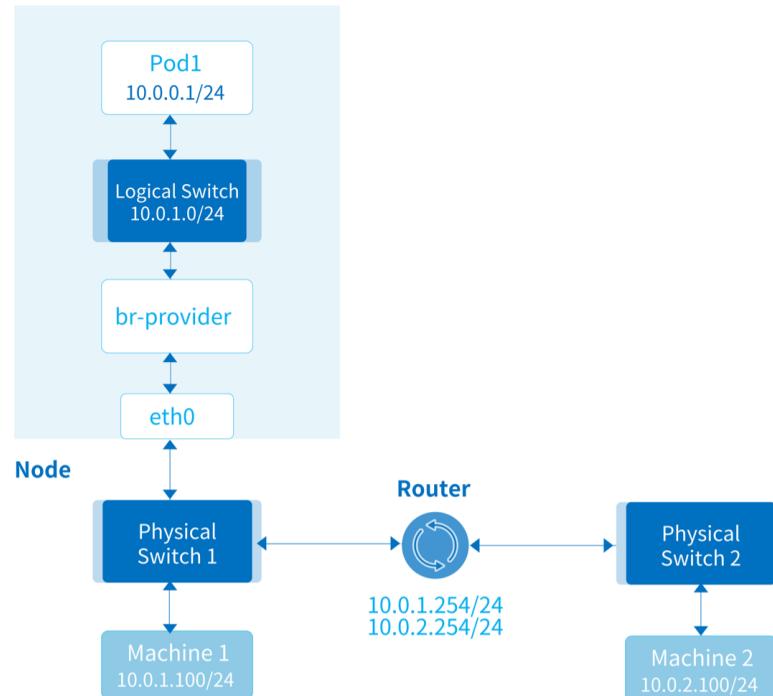
OVS

Provider Network

## 8.6.4

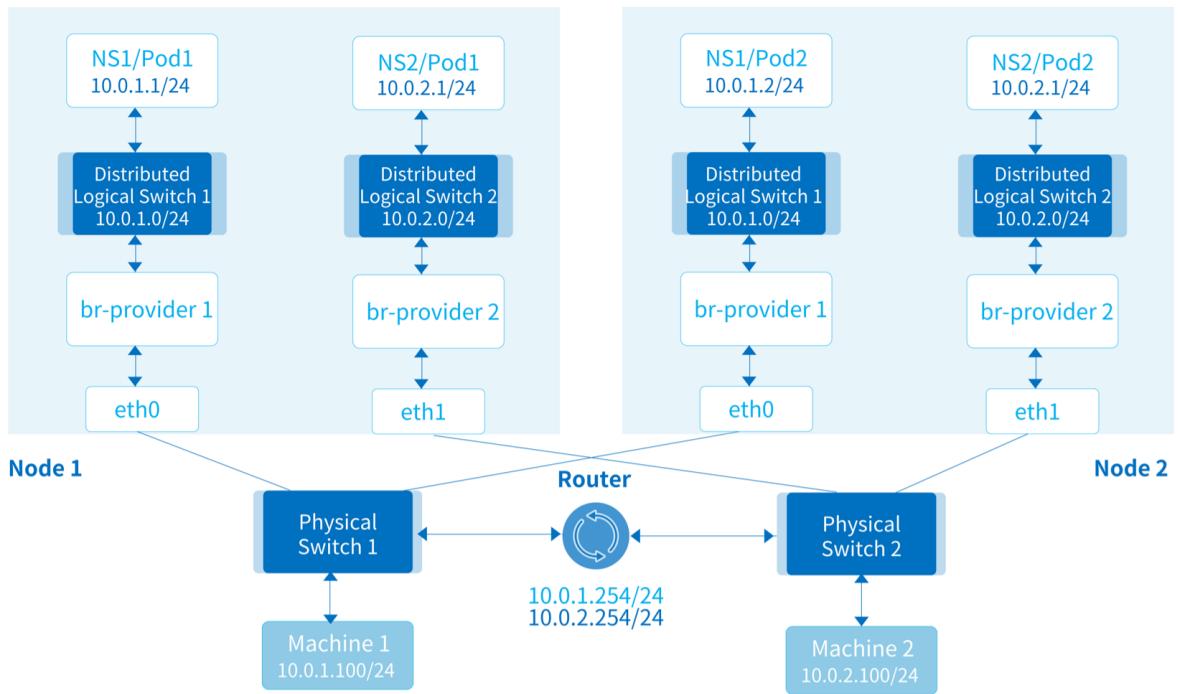


## 8.6.5

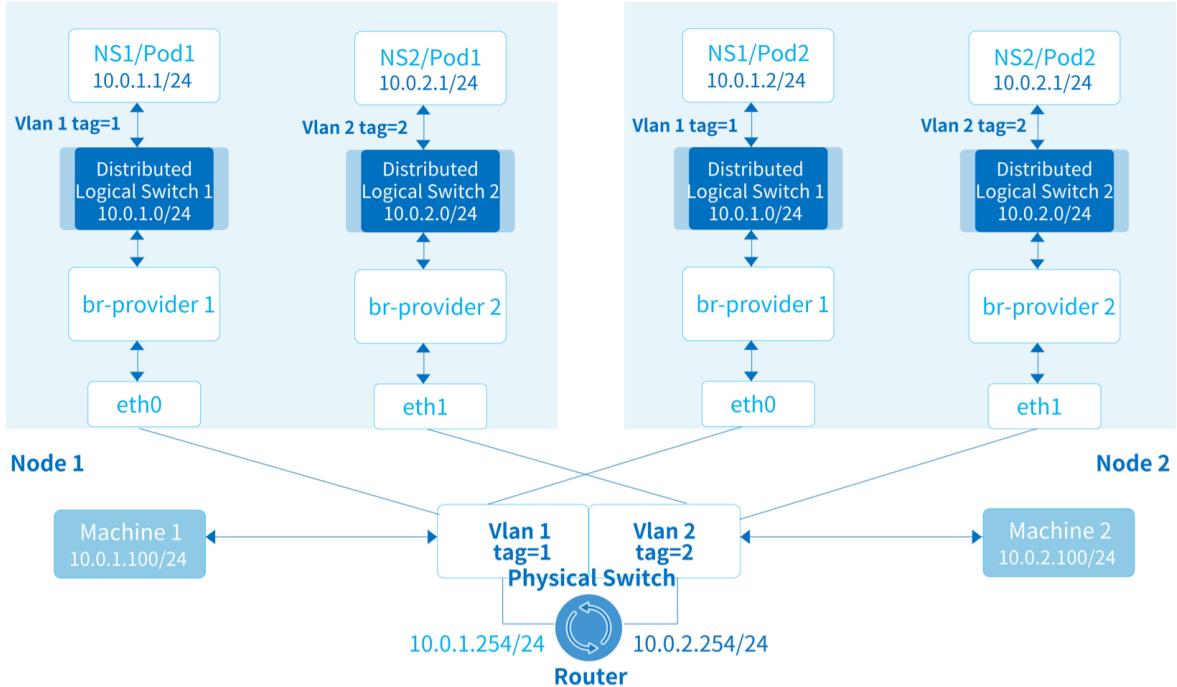


Pod

## 8.6.6 Vlan Tag

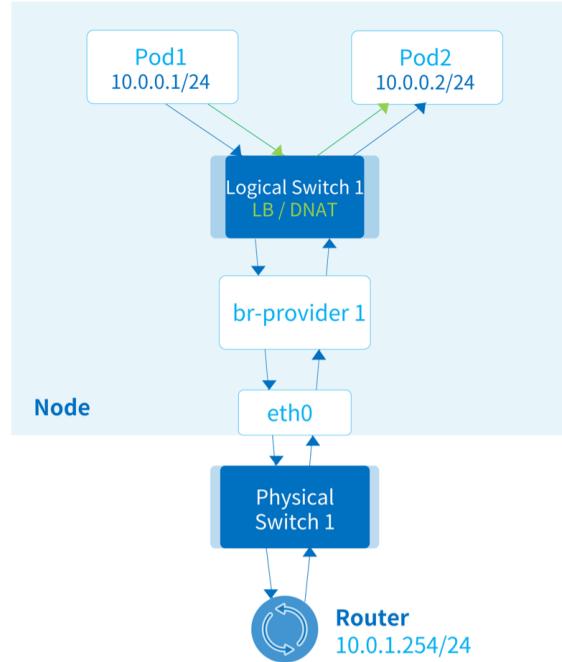
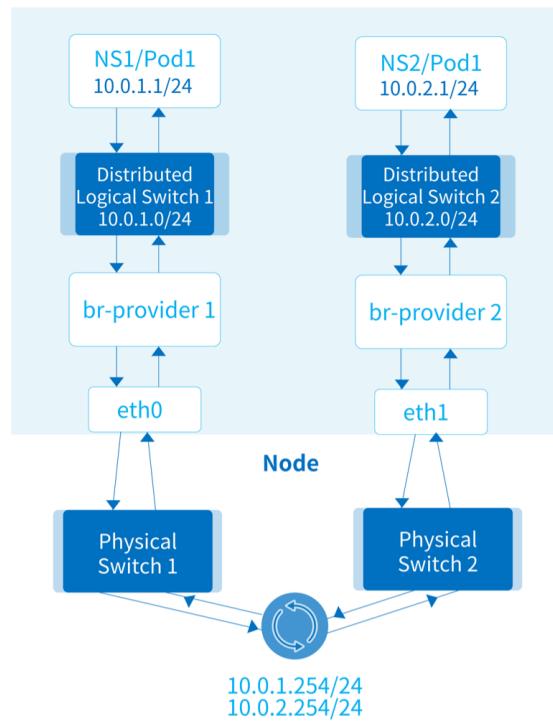


### 8.6.7 VLAN



### 8.6.8 Pod Service IP

Kube-OVN	Kubernetes Service	DNAT	IP	Service	Pod Endpoint	Service IP IP	Pod MAC	Service IP MAC	MAC
----------	--------------------	------	----	---------	--------------	---------------	---------	----------------	-----

**Service****Pod****Service****Pod**[PDF](#)[Slack](#)[Support](#)

⌚2022 9 23

⌚2022 5 20

⌚GitHub 🎉

8.6.9

---

## 8.7 Iptables

---

Kube-OVN    ipset    iptables    VPC    Overlay    NAT

ipset

IPv4/IPv6			
ovn40services/ovn60services	hash:net	Service	
ovn40subnets/ovn60subnets	hash:net	Overlay	NodeLocal DNS IP
ovn40subnets-nat/ovn60subnets-nat	hash:net	NatOutgoing	Overlay
ovn40subnets-distributed-gw/ovn60subnets-distributed-gw	hash:net	Overlay	
ovn40other-node/ovn60other-node	hash:net	IP	
ovn40local-pod-ip-nat/ovn60local-pod-ip-nat	hash:ip		
ovn40subnets-nat-policy	hash:net	natOutgoingPolicyRules	
ovn40natpr-418e79269dc5-dst	hash:net	natOutgoingPolicyRules	rule dstIPs
ovn40natpr-418e79269dc5-src	hash:net	natOutgoingPolicyRules	rule srcIPs

iptables IPv4

filter	INPUT	-m set --match-set ovn40services src -j ACCEPT	k8s Service Pod	--	
filter	INPUT	-m set --match-set ovn40services dst -j ACCEPT		--	
filter	INPUT	-m set --match-set ovn40subnets src -j ACCEPT		--	
filter	INPUT	-m set --match-set ovn40subnets dst -j ACCEPT		--	
filter	FORWARD	-m set --match-set ovn40services src -j ACCEPT		--	
filter	FORWARD	-m set --match-set ovn40services dst -j ACCEPT		--	
filter	FORWARD	-m set --match-set ovn40subnets src -j ACCEPT		--	
filter	FORWARD	-m set --match-set ovn40subnets dst -j ACCEPT		--	
filter	FORWARD	-s 10.16.0.0/16 -m comment --comment "ovn-subnet-gateway,ovn-default"	subnet	10.16.0.0/16 subnet cidr comment ovn-subnet-gateway subnet	iptables ovn-default subnet
filter	FORWARD	-d 10.16.0.0/16 -m comment --comment "ovn-subnet-gateway,ovn-default"	subnet		
filter	OUTPUT	-p udp -m udp --dport 6081 -j MARK --set-xmark 0x0	SNAT	UDP: bad checksum on VXLAN interface	
nat	PREROUTING	-m comment --comment "kube-ovn prerouting rules" -j OVN-PREROUTING	OVN-PREROUTING	--	
nat	POSTROUTING	-m comment --comment "kube-ovn postrouting rules" -j OVN-POSTROUTING	OVN-POSTROUTING	--	
nat	OVN-PREROUTING	-i ovn0 -m set --match-set ovn40subnets src -m set --match-set ovn40services dst -j MARK --set-xmark 0x4000/0x4000	Pod Service masquerade		LB
nat	OVN-PREROUTING	-p tcp -m addrtype --dst-type LOCAL -m set --match-		kube-proxy ipvs	

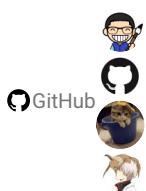
		set KUBE-NODE-PORT- LOCAL-TCP dst -j MARK -- set-xmark 0x80000/0x80000	ExternalTrafficPolicy Local Service TCP		
nat	OVN- PREROUTING	-p udp -m addrtype --dst- type LOCAL -m set --match- set KUBE-NODE-PORT- LOCAL-UDP dst -j MARK -- set-xmark 0x80000/0x80000	ExternalTrafficPolicy Local Service UDP		
nat	OVN- POSTROUTING	-m set --match-set ovn40services src -m set -- match-set ovn40subnets dst -m mark --mark 0x4000/0x4000 -j SNAT -- to-source	Service IP Overlay Pod IP		kube-proxy ipvs
nat	OVN- POSTROUTING	-m mark --mark 0x4000/0x4000 -j MASQUERADE	SNAT	--	
nat	OVN- POSTROUTING	-m set --match-set ovn40subnets src -m set -- match-set ovn40subnets dst -j MASQUERADE	Pod Service SNAT	--	
nat	OVN- POSTROUTING	-m mark --mark 0x80000/0x80000 -m set -- match-set ovn40subnets- distributed-gw dst -j RETURN	ExternalTrafficPolicy Local Service Endpoint SNAT	--	
nat	OVN- POSTROUTING	-m mark --mark 0x80000/0x80000 -j MASQUERADE	ExternalTrafficPolicy Local Service Endpoint SNAT	--	
nat	OVN- POSTROUTING	-p tcp -m tcp --tcp-flags SYN NONE -m conntrack -- ctstate NEW -j RETURN	Pod IP	SNAT	--
nat	OVN- POSTROUTING	-s 10.16.0.0/16 -m set ! -- match-set ovn40subnets dst -j SNAT --to-source 192.168.0.101	Pod NatOutgoing IP SNAT	10.16.0.0/16 IP	192.168.0.101
nat	OVN- POSTROUTING	-m set --match-set ovn40subnets-nat src -m set ! --match-set ovn40subnets dst -j MASQUERADE	Pod NatOutgoing SNAT	--	
nat	OVN- POSTROUTING	-m set --match-set ovn40subnets-nat-policy src -m set ! --match-set ovn40subnets dst -j OVN- NAT-POLICY	Pod natOutgoingPolicyRules SNAT	natOutgoingPolicyRules NAT-POLICY	OVN
nat					

	OVN- POSTROUTING	-m mark --mark 0x90001/0x90001 -j MASQUERADE --random- fully	OVN-NAT-POLICY 0x90001/0x90001	tag SNAT
nat	OVN- POSTROUTING	-m mark --mark 0x90002/0x90002 -j RETURN	OVN-NAT-POLICY 0x90002/0x90002	, tag SNAT
nat	OVN-NAT-POLICY	-s 10.0.11.0/24 -m comment --comment natPolicySubnet-net1 -j OVN-NAT-PSUBNET- aa98851157c5	10.0.11.0/24 net1 CIDR OVN-NAT- PSUBNET-aa98851157c5 natOutgoingPolicyRules	
nat	OVN-NAT- PSUBNET- xxxxxxxxxxxx	-m set --match-set ovn40natpr-418e79269dc5- src src -m set --match-set ovn40natpr-418e79269dc5- dst dst -j MARK --set-xmark 0x90002/0x90002	418e79269dc5 natOutgoingPolicyRules ID status.natOutgoingPolicyRules[index].Rule srcIPs ovn40natpr-418e79269dc5- src dstIPs ovn40natpr-418e79269dc5- dst tag 0x90002	
mangle	OVN-OUTPUT	-d 10.241.39.2/32 -p tcp -m tcp -dport 80 -j MARK --set- xmark 0x90003/0x90003	kubelet tproxy	
mangle	OVN- PREROUTING	-d 10.241.39.2/32 -p tcp -m tcp -dport 80 -j TPROXY -- on-port 8102 --on-ip 172.18.0.3 --tproxy-mark 0x90004/0x90004	kubelet tproxy	

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2022 9 6



8.7.1

## 8.8

---

### 8.8.1

Kube-OVN      Github      /      Github      Issue      PR      Maintainer      Review      Github Action

### 8.8.2

Kube-OVN      Go      Go Modules      G0111MODULE="on"

golangci-lint      local-installation

Kube-OVN      Docker buildx      Docker      buildx:

```
docker buildx create --use
```

### 8.8.3

Kube-OVN

```
git clone https://github.com/kubeovn/kube-ovn.git
cd kube-ovn
make release
```

ARM

```
make release-arm
```

### 8.8.4 base

OVS/OVN      base

base      Dockerfile      dist/images/Dockerfile.base

```
build x86 base image
make base-amd64

build arm base image
make base-arm64
```

### 8.8.5 E2E

Kube-OVN :

- KIND      Kubernetes      go install sigs.k8s.io/kind@latest
- jinjanator      : pip install jinjanator
- Ginkgo      go install github.com/onsi/ginkgo/v2/ginkgo; go get github.com/onsi/gomega/...

E2E

```
make kind-init
make kind-install
make e2e
```

Underlay E2E

```
make kind-init
make kind-install-underlay
make e2e-underlay-single-nic
```

ovn vpc nat gw eip, fip, snat, dnat

```
make kind-init
make kind-install
make ovn-vpc-nat-gw-conformance-e2e
```

iptables vpc nat gw eip, fip, snat, dnat

```
make kind-init
make kind-install-vpc-nat-gw
make iptables-vpc-nat-gw-conformance-e2e
```

loadbalancer service

```
make kind-init
make kind-install-lb-svc
make kube-ovn-lb-svc-conformance-e2e
```

```
make kind-clean
```

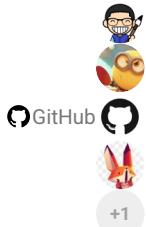
[!\[\]\(fea10d9bbf671e524c105b242c5b9338\_img.jpg\) PDF](#)

[!\[\]\(aec19455ee7b04cc6f70067cfbcdc99f\_img.jpg\) Slack](#)

[!\[\]\(903efed6be0d7ece1ae92593388eaed1\_img.jpg\) Support](#)

 2025 6 29

 2022 5 20



8.8.6

---

## 8.9 OVS/OVN

---

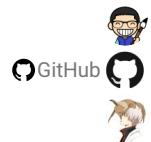
OVN/OVS	SDN	Kubernetes	Kube-OVN
OVN/OVS	Kube-OVN		

- [4228eab1d7](#) vswitchd ofport\_usage
- [54056ea65d](#) timer
- [6b4dccb311f](#) fdb
- [f627b7721e](#) hairpin fdb
- [3f3e3a436f](#) ovsdb-tool join-cluster Server ID
- [a6cb8215a8](#) QoS
- [d4d76ddb2e](#) ovsdb-tool fix-cluster
- [ffd2328d4a](#) netdev CPU
- [d088c5d8c2](#) ovs-router kube-ipvs0
- [1b31f07dc6](#)
- [54b7678229](#) ovs-sandbox docker run
- [9ee66bd91b](#)
- [e889d46924](#) Underlay resubmit
- [f9e97031b5](#) ovn-controller Kube-OVN localnet GARP
- [78cade0187](#) conntrack
- [85aa6263ad](#) northd DNS IP conntrack
- [34dc3e3fcf](#) lflow lport conntrack
- [a297b840c2](#) DNAT lsp
- [03e35ed9c5](#) ovn-controller
- [e7d3ba53cd](#) ACL DNS IP conntrack
- [9286e1fd57](#)
- [e5916eb53a](#) lr-lb DNAT
- [e4e6ea9c5f](#) BFD LRP
- [e76880e792](#) northd nb version\_compatibility
- [477695a010](#) northd localnet lrp arp/nd
  
- [20626ea909](#) LB ACL
- [a2d9ff3cccd](#) Deb



⌚2025 9 10

⌚2022 5 24



8.9.1

---

## 8.10

Kube-OVN OVN/OVS Geneve Vxlan STT OVN  
[OVN Architecture Design Decision](#)

### 8.10.1 Geneve

Geneve	Kube-OVN	OVN	Offload	Geneve	24bit
datapath		datapath 32768			
Mellanox	OVS	Geneve 5.4	backport		
UDP	TCP over UDP	TCP	CPU		

### 8.10.2 Vxlan

Vxlan	OVN	Offload	OVN	datapath	4096 datapath
datapath	4096	inport ACL			
Mellanox	OVS	Vxlan			
UDP	TCP over UDP	TCP	CPU		

### 8.10.3 STT

#### ⚠ Warning

OpenVswitch 3.6 STT Tunnel

STT	OVN	TCP	TCP	TCP	OVN	datapath
		OVS				
		OVS				

### 8.10.4

- [VXLAN vs GENEVE: Understand The Difference](#)
- [OVN FAQ](#)
- [What is Geneve](#)

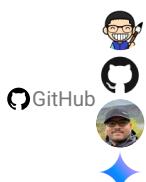
[!\[\]\(885cf0eed8cc7a871196b2e7848765d7\_img.jpg\) PDF](#)

[!\[\]\(e76f43154b16a42c81ccc316864065d4\_img.jpg\) Slack](#)

[!\[\]\(02ceceb12eaf9bd8ccf7c2a63b375a4c\_img.jpg\) Support](#)

 2025 9 10

 2022 6 17



## 8.10.5

---

## 8.11 Kube-OVN

---

### Kube-OVN

#### 8.11.1 ovn-monitor

##### OVN

Gauge	kube_ovn_ovn_status	OVN	(2) follower	(1) leader	(0)
Gauge	kube_ovn_failed_req_count	OVN			
Gauge	kube_ovn_log_file_size_bytes	OVN			
Gauge	kube_ovn_db_file_size_bytes	OVN			
Gauge	kube_ovn_chassis_info	OVN chassis	(1)	(0)	
Gauge	kube_ovn_db_status	OVN	, (1)	(0)	
Gauge	kube_ovn_logical_switch_info	OVN logical switch	(1)	logical switch	
Gauge	kube_ovn_logical_switch_external_id	OVN logical switch external_id	(1)	external-id	
Gauge	kube_ovn_logical_switch_port_binding	OVN logical switch logical switch port	(1)		
Gauge	kube_ovn_logical_switch_tunnel_key	OVN logical switch tunnel key			
Gauge	kube_ovn_logical_switch_ports_num	OVN logical switch logical port			
Gauge	kube_ovn_logical_switch_port_info	OVN logical switch port	(1)		
Gauge	kube_ovn_logical_switch_port_tunnel_key	OVN logical switch port tunnel key			
Gauge	kube_ovn_cluster_enabled	(1) OVN	(0) OVN		
Gauge	kube_ovn_cluster_role		(1)		
Gauge	kube_ovn_cluster_status		(1)		
Gauge	kube_ovn_cluster_term	RAFT term			
Gauge	kube_ovn_cluster_leader_self		leader (1)	(0)	
Gauge	kube_ovn_cluster_vote_self		leader (1)	(0)	
Gauge	kube_ovn_cluster_election_timer	election timer			
Gauge	kube_ovn_cluster_log_not_committed	commit RAFT			
Gauge	kube_ovn_cluster_log_not_applied	apply RAFT			
Gauge	kube_ovn_cluster_log_index_start	RAFT			
Gauge	kube_ovn_cluster_log_index_next	RAFT			
Gauge	kube_ovn_cluster_inbound_connections_total				
Gauge	kube_ovn_cluster_outbound_connections_total				
Gauge	kube_ovn_cluster_inbound_connections_error_total				
Gauge	kube_ovn_cluster_outbound_connections_error_total				

## 8.11.2 ovs-monitor

---

ovsdb vswitchd

Gauge	ovs_status	OVS	(1)	(0)
Gauge	ovs_info	OVS	(1)	
Gauge	failed_req_count	OVS		
Gauge	log_file_size	OVS		
Gauge	db_file_size	OVS		
Gauge	datapath	Datapath	(1)	
Gauge	dp_total	OVS datapath		
Gauge	dp_if	Datapath	(1)	
Gauge	dp_if_total	datapath port		
Gauge	dp_flows_total	Datapath flow		
Gauge	dp_flows_lookup_hit	Datapath flow		
Gauge	dp_flows_lookup_missed	Datapath flow		
Gauge	dp_flows_lookup_lost	Datapath userspace		
Gauge	dp_masks_hit	Datapath mask		
Gauge	dp_masks_total	Datapath mask		
Gauge	dp_masks_hit_ratio	Datapath mask		
Gauge	interface	OVS (1)		
Gauge	interface_admin_state	(0) down, (1) up, (2)		
Gauge	interface_link_state	(0) down, (1) up, (2)		
Gauge	interface_mac_in_use	OVS Interface MAC		
Gauge	interface_mtu	OVS Interface MTU		
Gauge	interface_of_port	OVS Interface OpenFlow Port ID		
Gauge	interface_if_index	OVS Interface Index		
Gauge	interface_tx_packets	OVS Interface		
Gauge	interface_tx_bytes	OVS Interface		
Gauge	interface_rx_packets	OVS Interface		
Gauge	interface_rx_bytes	OVS Interface		
Gauge	interface_rx_crc_err	OVS Interface		
Gauge	interface_rx_dropped	OVS Interface		
Gauge	interface_rx_errors	OVS Interface		
Gauge	interface_rx_frame_err	OVS Interface		
Gauge	interface_rx_missed_err	OVS Interface miss		
Gauge	interface_rx_over_err	OVS Interface overrun		
Gauge	interface_tx_dropped	OVS Interface		

Gauge	interface_tx_errors	OVS Interface
Gauge	interface_collisions	OVS interface

### 8.11.3 kube-ovn-pinger

Gauge	pinger_ovs_up	OVS
Gauge	pinger_ovs_down	OVS
Gauge	pinger_ovn_controller_up	ovn-controller
Gauge	pinger_ovn_controller_down	ovn-controller
Gauge	pinger_inconsistent_port_binding	OVN-SB portbinding OVS interface
Gauge	pinger_apiserver_healthy	kube-ovn-pinger apiserver
Gauge	pinger_apiserver_unhealthy	kube-ovn-pinger apiserver
Histogram	pinger_apiserver_latency_ms	kube-ovn-pinger apiserver
Gauge	pinger_internal_dns_healthy	kube-ovn-pinger
Gauge	pinger_internal_dns_unhealthy	kube-ovn-pinger
Histogram	pinger_internal_dns_latency_ms	kube-ovn-pinger
Gauge	pinger_external_dns_health	kube-ovn-pinger
Gauge	pinger_external_dns_unhealthy	kube-ovn-pinger
Histogram	pinger_external_dns_latency_ms	kube-ovn-pinger
Histogram	pinger_pod_ping_latency_ms	kube-ovn-pinger ping Pod
Gauge	pinger_pod_ping_lost_total	kube-ovn-pinger ping Pod
Gauge	pinger_pod_ping_count_total	kube-ovn-pinger ping Pod
Histogram	pinger_node_ping_latency_ms	kube-ovn-pinger ping Node
Gauge	pinger_node_ping_lost_total	kube-ovn-pinger ping Node
Gauge	pinger_node_ping_count_total	kube-ovn-pinger ping Node
Histogram	pinger_external_ping_latency_ms	kube-ovn-pinger ping
Gauge	pinger_external_lost_total	kube-ovn-pinger ping

## 8.11.4 kube-ovn-controller

`kube-ovn-controller`

Histogram	rest_client_request_latency_seconds	apiserver
Counter	rest_client_requests_total	apiserver
Counter	lists_total	API list
Summary	list_duration_seconds	API list
Summary	items_per_list	API list
Counter	watches_total	API watch
Counter	short_watches_total	API watch
Summary	watch_duration_seconds	API watch
Summary	items_per_watch	API watch
Gauge	last_resource_version	resource version
Histogram	ovs_client_request_latency_milliseconds	OVN
Gauge	subnet_available_ip_count	IP
Gauge	subnet_used_ip_count	IP

## 8.11.5 kube-ovn-cni

`kube-ovn-cni`

Histogram	cni_op_latency_seconds	CNI
Counter	cni_wait_address_seconds_total	CNI
Counter	cni_wait_connectivity_seconds_total	CNI
Counter	cni_wait_route_seconds_total	CNI
Histogram	rest_client_request_latency_seconds	apiserver
Counter	rest_client_requests_total	apiserver
Counter	lists_total	API list
Summary	list_duration_seconds	API list
Summary	items_per_list	API list
Counter	watches_total	API watch
Counter	short_watches_total	API watch
Summary	watch_duration_seconds	API watch
Summary	items_per_watch	API watch
Gauge	last_resource_version	resource version
Histogram	ovs_client_request_latency_milliseconds	OVN



⌚2022 10 18

⌚2022 6 21

GitHub 

8.11.6

---

## 8.12 Kube-OVN

---

Kube-OVN      Kube-OVN      CRD      CRD

### 8.12.1 Condition

type	String			
status	String	True	False	Unknown
reason	String			
message	String			
observedGeneration	Int64			
lastUpdateTime	Time			
lastTransitionTime	Time			

CRD      Status      Condition

### 8.12.2

#### Subnet

SUBNET

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	Subnet
metadata	ObjectMeta	Kubernetes	
spec	SubnetSpec	Subnet	
status	SubnetStatus	Subnet	

**SubnetSpec**

default	Bool				
vpc	String	VPC	ovn-cluster		
protocol	String	IP	IPv4	IPv6	Dual
namespaces	[]String		namespace		
cidrBlock	String		10.16.0.0/16		
gateway	String		CIDRBlock		
excludelps	[]String				
provider	String	OVN Subnet	NetworkAttachmentDefinition	. Kube-OVN	
gatewayType	String	Overlay	distributed	centralized	
gatewayNode	String		centralized		
natOutgoing	Bool	NAT	externalEgressGateway		
externalEgressGateway	String		natOutgoing		
policyRoutingPriority	Uint32				
policyRoutingTableID	Uint32		TableID		
mtu	Uint32		MTU		
private	Bool				
allowSubnets	[]String				
vlan	String		Vlan		
vips	[]String	virtual	Isp	virtual-ip	
logicalGateway	Bool				
disableGatewayCheck	Bool		Pod		
disableInterConnection	Bool				
enableDHCP	Bool		Isp	dhcp	
dhcpV4Options	String		Isp	dhcpv4_options	DHCP_Options
dhcpV6Options	String		Isp	dhcpv6_options	DHCP_Options
enableIPv6RA	Bool		Irp	ipv6_ra_configs	
ipv6RAConfigs	String		Irp	ipv6_ra_configs	
acls	[]Acl		logical-switch	acls	
allowEWTraffic	Bool				
natOutgoingPolicyRules	[]NatOutgoingPolicyRule	NAT			
u2oInterconnectionIP	String	Underlay/Overlay	IP		
u2oInterconnection	Bool	Overlay/Underlay			
enableLb	*Bool	logical-switch	load-balancer		
enableEcmp	Bool	ECMP			

enableMulticastSnoop	Bool
enableExternalLBAddress	Bool
routeTable	String
namespaceSelectors	[]LabelSelector

**Acl**

direction	String	Acl	from-lport	to-lport
priority	Int	Acl	0	32767
match	String	Acl		
action	String	Acl	allow-related	allow-stateless
			allow	drop
			reject	

**NatOutgoingPolicyRule**

match	NatOutGoingPolicyMatch
action	String

**NatOutGoingPolicyMatch**

srcIPs	String	IP
dstIPs	String	IP

## SubnetStatus

conditions	>[]SubnetCondition	Condition	
v4availableIPs	Float64	IPv4 IP	
v4availableIPrange	String	IPv4	
v4usingIPs	Float64	IPv4 IP	
v4usingIPrange	String	IPv4	
v6availableIPs	Float64	IPv6 IP	
v6availableIPrange	String	IPv6	
v6usingIPs	Float64	IPv6 IP	
v6usingIPrange	String	IPv6	
activateGateway	String		
dhcpV4OptionsUUID	String	lsp dhcpv4_options	DHCP_Options
dhcpV6OptionsUUID	String	lsp dhcpv6_options	DHCP_Options
u2oInterconnectionIP	String	Overlay/Underlay	IP
u2oInterconnectionMAC	String	Overlay/Underlay	MAC
u2oInterconnectionVPC	String	Overlay/Underlay	VPC
natOutgoingPolicyRules	[]NatOutgoingPolicyRuleStatus	NAT	
mcastQuerierIP	String	IP	
mcastQuerierMAC	String	MAC	

## IP

## IP

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IP
metadata	ObjectMeta	Kubernetes	
spec	IPSpec	IP	

**IPSpec**

IPSpec			
podName	String	Pod	
namespace	String	Pod Namespace	
subnet	String	IP Subnet	
attachSubnets	[]String	IP	
nodeName	String	Pod	
ipAddress	String	IP	v4IP v6IP
v4IpAddress	String	IPv4 IP	
v6IpAddress	String	IPv6 IP	
attachIps	[]String	IP	IP
macAddress	String	Pod	MAC
attachMacs	[]String	IP	MAC
containerID	String	Pod	Container ID
podType	String	Pod StatefulSet VirtualMachine	

**Vpc**

## VPC

VPC			
apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	Vpc
metadata	ObjectMeta	Kubernetes	
spec	VpcSpec	Vpc	
status	VpcStatus	Vpc	

## VpcSpec

VpcSpec			
defaultSubnet	String		
namespaces	[]String	Vpc	
staticRoutes	[]StaticRoute		
policyRoutes	[]PolicyRoute		
vpcPeerings	[]VpcPeering	VPC	
enableExternal	Bool		
extraExternalSubnets	[]String		
enableBfd	Bool	BFD ( )	
bfdPort	BFDPort	BFD	

## StaticRoute

policy	String	
cidr	String	
nextHopIP	String	IP
ecmpMode	String	ECMP
bfdId	String	BFD ID
routeTable	String	

## PolicyRoute

priority	Int	
match	String	
action	String	allow drop reroute
nextHopIP	String	IP action reroute

## VpcPeering

remoteVpc	String	VPC
localConnectIP	String	IP

## BFDPort

enabled	Bool	BFD
ip	String	BFD IP
nodeSelector	LabelSelector	BFD LRP

**VpcStatus**

conditions	IVpcCondition	Vpc	Condition
standby	Bool	VPC	
default	Bool	VPC	
defaultLogicalSwitch	String		
router	String		
tcpLoadBalancer	String	TCP	
udpLoadBalancer	String	UDP	
sctpLoadBalancer	String	SCTP	
tcpSessionLoadBalancer	String	TCP	
udpSessionLoadBalancer	String	UDP	
sctpSessionLoadBalancer	String	SCTP	
subnets	IString	VPC	
vpcPeerings	IString	VPC	
enableExternal	Bool		
extraExternalSubnets	IString		
enableBfd	Bool	BFD	

**8.12.3 Underlay****Vlan**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	Vlan
metadata	ObjectMeta	Kubernetes	
spec	VlanSpec	Vlan	
status	VlanStatus	Vlan	

**VLANSPEC**

id	Int	Vlan tag	0~4096
provider	String	Vlan	ProviderNetwork

**VLANSTATUS**

subnets	>[]String	Vlan	
conflict	Bool		
conditions	>[]VlanCondition	Vlan	Condition

**ProviderNetwork**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	ProviderNetwork
metadata	ObjectMeta	Kubernetes	
spec	ProviderNetworkSpec	ProviderNetwork	
status	ProviderNetworkStatus	ProviderNetwork	

**PROVIDERNETWORKSPEC**

defaultInterface	String	
customInterfaces	>[]CustomInterface	
excludeNodes	>[]String	
exchangeLinkName	Bool	OVS

**CustomInterface**

interface	String	Underlay
nodes	>[]String	

**PROVIDERNETWORKSTATUS**

ready	Bool		
readyNodes	>[]String		
notReadyNodes	>[]String		
vlans	>[]String	Vlan	
conditions	>[]ProviderNetworkCondition	ProviderNetwork	Condition

## 8.12.4

### SecurityGroup

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	SecurityGroup
metadata	ObjectMeta	Kubernetes	
spec	SecurityGroupSpec	SecurityGroup	
status	SecurityGroupStatus	SecurityGroup	

### SECURITYGROUPSPEC

ingressRules	[]SecurityGroupRule
egressRules	[]SecurityGroupRule
allowSameGroupTraffic	Bool

### SecurityGroupRule

ipVersion	String	IP	ipv4	ipv6		
protocol	SgProtocol		all	icmp	tcp	udp
priority	Int		1-200			
remoteType	SgRemoteType		address	securityGroup		
remoteAddress	String					
remoteSecurityGroup	String					
portRangeMin	Int		1			
portRangeMax	Int		65535			
policy	SgPolicy		allow	drop		

### SECURITYGROUPSTATUS

portGroup	String	
allowSameGroupTraffic	Bool	
ingressMd5	String	MD5
egressMd5	String	MD5
ingressLastSyncSuccess	Bool	
egressLastSyncSuccess	Bool	

## 8.12.5 IP

### Vip

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	Vip
metadata	ObjectMeta	Kubernetes	
spec	VipSpec	Vip	
status	VipStatus	Vip	

### VIPSPEC

namespace	String	VIP
subnet	String	VIP
type	String	VIP
v4ip	String	IPv4
v6ip	String	IPv6
macAddress	String	MAC
selector	[]String	
attachSubnets	[]String	

### VIPSTATUS

conditions	[]VipCondition	VIP	Condition
type	String	VIP	
v4ip	String	IPv4	
v6ip	String	IPv6	
mac	String	MAC	

### SwitchLBRule

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	SwitchLBRule
metadata	ObjectMeta	Kubernetes	
spec	SwitchLBRuleSpec	SwitchLBRule	
status	SwitchLBRuleStatus	SwitchLBRule	

**SWITCHLBRULESPEC**

vip	String	IP
namespace	String	
selector	[]String	
endpoints	[]String	
sessionAffinity	String	
ports	[]SwitchLBRulePort	

**SwitchLBRulePort**

name	String
port	Int32
targetPort	Int32
protocol	String

**SWITCHLBRULESTATUS**

conditions	[]SwitchLBRuleCondition	SwitchLBRule	Condition
ports	String	SwitchLBRule	
service	String	SwitchLBRule	service

---

**8.12.6 QoS IP****QoS Policy**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	QoS Policy
metadata	ObjectMeta	Kubernetes	
spec	QoS Policy Spec	QoS Policy	

**QOSPOLICYSPEC**

bandwidthLimitRules	QoS Policy Bandwidth Limit Rules
shared	Bool
bindingType	QoS Policy Binding Type

**IPPool**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IPPool
metadata	ObjectMeta	Kubernetes	
spec	IPPoolSpec	IPPool	
status	IPPoolStatus	IPPool	

**IPPOOLSPEC**

subnet	String		
namespaces	[]String		
ips	[]String	IP	

**IPPOOLSTATUS**

v4AvailableIPs	BigInt	IPv4	IP
v4AvailableIPRange	String	IPv4	IP
v4UsingIPs	BigInt	IPv4	IP
v4UsingIPRange	String	IPv4	IP
v6AvailableIPs	BigInt	IPv6	IP
v6AvailableIPRange	String	IPv6	IP
v6UsingIPs	BigInt	IPv6	IP
v6UsingIPRange	String	IPv6	IP
conditions	[]IPPoolCondition	IP	Condition

## 8.12.7 NAT IP

**IptablesEIP**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IptablesEIP
metadata	ObjectMeta	Kubernetes	
spec	IptablesEIPSpec	IptablesEIP	
status	IptablesEIPStatus	IptablesEIP	

## IPTABLESEIPSPEC

v4ip	String	IPv4
v6ip	String	IPv6
macAddress	String	MAC
natGwDp	String	NAT
qosPolicy	String	QoS
externalSubnet	String	

## IPTABLESEIPSTATUS

ready	Bool	IptablesEIP			
ip	String	IptablesEIP	IP	IPv4	
redo	String	IptablesEIP	CRD		
nat	String	IptablesEIP	fip	snat	dnat
qosPolicy	String	QoS			
conditions	[]IptablesEIPCondition	IptablesEIP		Condition	

## OvnEip

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	OvnEip
metadata	ObjectMeta	Kubernetes	
spec	OvnEipSpec	OvnEip	
status	OvnEipStatus	OvnEip	

## OVNEIPSPEC

externalSubnet	String	
v4Ip	String	IPv4
v6Ip	String	IPv6
macAddress	String	MAC
type	String	fip lsp nat

**IptablesFIPRule**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IptablesFIPRule
metadata	ObjectMeta	Kubernetes	
spec	IptablesFIPRuleSpec	IptablesFIPRule	

## IPTABLESFIPRULESPEC

eip	String	IP
internalIP	String	IP

**OvnFip**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	OvnFip
metadata	ObjectMeta	Kubernetes	
spec	OvnFipSpec	OvnFip	
status	OvnFipStatus	OvnFip	

## OVNFIPSPEC

ovnEip	String	OVN EIP	
ipType	String	IP	vip ip
ipName	String	IP	
vpc	String	VPC	
v4Ip	String	IPv4	
v6Ip	String	IPv6	
type	String	distributed	centralized

**IptablesDnatRule**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IptablesDnatRule
metadata	ObjectMeta	Kubernetes	
spec	IptablesDnatRuleSpec	IptablesDnatRule	

## IPTABLESDNATRULESPEC

eip	String	IP
externalPort	String	
protocol	String	
internalIP	String	IP
internalPort	String	

## OvnDnatRule

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	OvnDnatRule
metadata	ObjectMeta	Kubernetes	
spec	OvnDnatRuleSpec	OvnDnatRule	
status	OvnDnatRuleStatus	OvnDnatRule	

## OVNDNATRULESPEC

ovnEip	String	OVN EIP
ipType	String	IP vip ip
ipName	String	IP
internalPort	String	
externalPort	String	
protocol	String	
vpc	String	VPC
v4Ip	String	IPv4
v6Ip	String	IPv6

**OVNDNATRULESTATUS**

vpc	String	VPC	
v4Eip	String	IPv4 EIP	
v6Eip	String	IPv6 EIP	
externalPort	String		
v4Ip	String	IPv4	
v6Ip	String	IPv6	
internalPort	String		
protocol	String		
ipName	String	IP	
ready	Bool	DNAT	
conditions	[]OvnDnatRuleCondition	OVN DNAT	Condition

**IptablesSnatRule**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	IptablesSnatRule
metadata	ObjectMeta	Kubernetes	
spec	IptablesSnatRuleSpec	IptablesSnatRule	

**IPTABLESSNATRULESPEC**

eip	String	IP
internalCIDR	String	CIDR

**OvnSnatRule**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	OvnSnatRule
metadata	ObjectMeta	Kubernetes	
spec	OvnSnatRuleSpec	OvnSnatRule	
status	OvnSnatRuleStatus	OvnSnatRule	

## OVNSNATRULESPEC

ovnEip	String	OVN EIP
vpcSubnet	String	VPC
ipName	String	IP
vpc	String	VPC
v4IpCidr	String	IPv4 CIDR
v6IpCidr	String	IPv6 CIDR

## 8.12.8 VPC

## VpcNatGateway

apiVersion	String	Kubernetes	kubeovn.io/v1	
kind	String	Kubernetes	VpcNatGateway	
metadata	ObjectMeta	Kubernetes		
spec	VpcNatGatewaySpec	VpcNatGateway		
status	VpcNatGatewayStatus	VpcNatGateway		

## VPCNATGATEWAYSPEC

vpc	String	VPC	Pod	VPC
subnet	String	VPC	Pod	
externalSubnets	[]String			
lanIp	String	VPC	Pod	IP
selector	[]String	Kubernetes Selector		
tolerations	[]Toleration	Kubernetes		
affinity	Affinity	Kubernetes		
qosPolicy	String	QoS		
bgpSpeaker	VpcBgpSpeaker	BGP speaker		

**VpcBgpSpeaker**

enabled	Bool	BGP speaker
asn	Uint32	
remoteAsn	Uint32	
neighbors	[]String	BGP
holdTime	Duration	BGP
routerId	String	BGP ID
password	String	BGP
enableGracefulRestart	Bool	
extraArgs	[]String	

**Route**

cidr	String	CIDR
nextHopIP	String	IP

**VPCNATGATEWAYSTATUS**

qosPolicy	String	QoS
externalSubnets	[]String	
selector	[]String	Kubernetes Selector
tolerations	[]Toleration	Kubernetes
affinity	Affinity	Kubernetes

**VpcEgressGateway**

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	VpcEgressGateway
metadata	ObjectMeta	Kubernetes	
spec	VpcEgressGatewaySpec	VpcEgressGateway	
status	VpcEgressGatewayStatus	VpcEgressGateway	

## VPCEGRESSGATEWAYSPEC

vpc	String	VPC
replicas	Int32	
prefix	String	
image	String	
internalSubnet	String	
externalSubnet	String	
internalIPs	[]String	IP
externalIPs	[]String	IP
trafficPolicy	String	

## VpcDns

apiVersion	String	Kubernetes	kubeovn.io/v1
kind	String	Kubernetes	VpcDns
metadata	ObjectMeta	Kubernetes	
spec	VpcDNSSpec	VpcDns	
status	VpcDNSStatus	VpcDns	

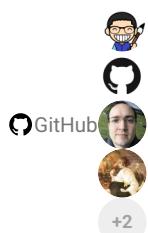
## VPCDNSSPEC

replicas	Int32	
vpc	String	VPC
subnet	String	



⌚2025 9 2

⌚2023 2 16



## 8.12.9

---

## 8.13 Annotation

---

Kube-OVN	Pod	Node Annotation	Annotation
			Annotation

### 8.13.1 Pod Annotation

Key	Value	Description									
ovn.kubernetes.io/allocated	true or false	Pod									
ovn.kubernetes.io/routed	true or false	Pod	OVN								
ovn.kubernetes.io/routes	String	Pod									
ovn.kubernetes.io/mac_address	String	Pod	Mac	Pod	Annotation	Mac					
ovn.kubernetes.io/ip_address	String	Pod	IP	Pod	Annotation	IP					
ovn.kubernetes.io/cidr	String	Pod	CIDR								
ovn.kubernetes.io/gateway	String	Pod	Gateway								
ovn.kubernetes.io/ip_pool	IP	Pod	Workload		IP						
ovn.kubernetes.io/bgp	true, cluster, local	BGP	Pod								
ovn.kubernetes.io/snat	String	Pod	SNAT								
ovn.kubernetes.io/eip	String	Pod	EIP								
ovn.kubernetes.io/vip	String	Pod	VIP	Annotation	VIP						
ovn.kubernetes.io/aaps	String	Pod	AAPs (Additional Allowed Addresses Pairs)								
ovn.kubernetes.io/virtualmachine	String	Pod	VirtualMachineInstance								
ovn.kubernetes.io/activation_strategy	String	Pod									
ovn.kubernetes.io/logical_router	String	Pod	VPC								
ovn.kubernetes.io/layer2_forward	true or false	Pod	OVN LSP	unknown							
ovn.kubernetes.io/port_security	true or false	Pod	Port Security								
ovn.kubernetes.io/logical_switch	String	Pod									
ovn.kubernetes.io/vlan_id	Int	Pod	Vlan ID								
ovn.kubernetes.io/ingress_rate	Int	Pod	Mbits/s								
ovn.kubernetes.io/egress_rate	Int	Pod	Mbits/s								
ovn.kubernetes.io/security_groups	String	Pod	Security Group								
ovn.kubernetes.io/default_route	true or false										
ovn.kubernetes.io/provider_network	String	Pod	ProviderNetwork								
ovn.kubernetes.io/mirror	true or false	Pod									
ovn.kubernetes.io/north_gateway	String	Pod									
ovn.kubernetes.io/latency	Int	Pod	ms								
ovn.kubernetes.io/limit	Int	Pod	qdisc								
ovn.kubernetes.io/loss	Float	Pod									
ovn.kubernetes.io/jitter	Int	Pod	ms								
ovn.kubernetes.io/generate-hash	true or false	Pod									
ovn.kubernetes.io/attachmentprovider	String	Pod									

### 8.13.2 Node Annotation

Key	Value	Description		
ovn.kubernetes.io/allocated	true or false	ovn0	join	
ovn.kubernetes.io/mac_address	String	Node ovn0	Mac	
ovn.kubernetes.io/ip_address	String	Node ovn0	IP	
ovn.kubernetes.io/cidr	String	Node ovn0	join	CIDR
ovn.kubernetes.io/gateway	String	Node ovn0	join	Gateway
ovn.kubernetes.io/chassis	String	Node OVN-SouthBoundDB	Chassis ID	
ovn.kubernetes.io/port_name	String	Node ovn0	OVN-NorthboundDB	LSP
ovn.kubernetes.io/logical_switch	String	Node ovn0		
ovn.kubernetes.io/tunnel_interface	String			

### 8.13.3 Namespace Annotation

Key	Value	Description	
ovn.kubernetes.io/cidr	CIDR	Namespace	CIDR
ovn.kubernetes.io/exclude_ips	excludelPs	Namespace	excludelPs

### 8.13.4 Subnet Annotation

Key	Value	Description
ovn.kubernetes.io/bgp	true, cluster, local	BGP

### 8.13.5 Service Annotation

Key	Value	Description		
ovn.kubernetes.io/bgp	true or false	BGP	Service	
ovn.kubernetes.io/switch_lb_vip	String	Service	Kube-OVN	VIP
ovn.kubernetes.io/vpc	String	Service	VPC	
ovn.kubernetes.io/service_external_ip_from_subnet	true or false	Service	IP	
ovn.kubernetes.io/service_health_check	true or false	Service		
ovn.kubernetes.io/lb_svc_img	String			

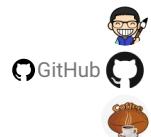
### 8.13.6 Networkpolicy Annotation

Key	Value	Description
ovn.kubernetes.io/enable_log	true or false	NetworkPolicy
ovn.kubernetes.io/log_acl_actions	"allow,drop,pass"	Action ACL

[PDF](#)[Slack](#)[Support](#)

⌚2025 9 10

⌚2024 6 12



8.13.7

---

## 8.14

---

### 8.14.1

Bad	Good

Bad	Good
Kube-OVN	1.10 Kube-OVN

  ■ ■

Bad	Good
wget 127.0.0.1	wget 127.0.0.1

### 8.14.2

yaml        yaml

Bad	Good
.... apiVersion: kubeovn.io/v1 kind: Subnet metadata: name: attach-subnet ....	....yaml apiVersion: kubeovn.io/v1 kind: Subnet metadata: name: attach-subnet ....

bash

Bad	Good
.... wget 127.0.0.1 ....	....bash wget 127.0.0.1 ....

#

**Bad**

```
oilbeater@macdeMac-3 ~ ping 114.114.114.114 -c 3
PING 114.114.114.114 (114.114.114.114): 56 data bytes
64 bytes from 114.114.114.114: icmp_seq=0 ttl=83 time=10.429 ms
64 bytes from 114.114.114.114: icmp_seq=1 ttl=79 time=11.360 ms
64 bytes from 114.114.114.114: icmp_seq=2 ttl=76 time=10.794 ms

--- 114.114.114.114 ping statistics ---
3 packets transmitted, 3 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 10.429/10.861/11.360/0.383 ms
```

**Good**

```
ping 114.114.114.114 -c 3
PING 114.114.114.114 (114.114.114.114): 56 data bytes
64 bytes from 114.114.114.114: icmp_seq=0 ttl=83 time=10.429 ms
64 bytes from 114.114.114.114: icmp_seq=1 ttl=79 time=11.360 ms
64 bytes from 114.114.114.114: icmp_seq=2 ttl=76 time=10.794 ms

--- 114.114.114.114 ping statistics ---
3 packets transmitted, 3 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 10.429/10.861/11.360/0.383 ms
```

#

**Bad**

```
mv /etc/origin/ovn/ovnnb_db.db /tmp
mv /etc/origin/ovn/ovnsb_db.db /tmp
```

**Good**

```
mv /etc/origin/ovn/ovnnb_db.db /tmp
mv /etc/origin/ovn/ovnsb_db.db /tmp
```

## 8.14.3

md

**Bad**

```
[](http://kubeovn.github.io/prepare)
```

**Good**

```
[](./prepare.md)
```

**Bad**

```
[Kubernetes](http://kubernetes.io)
```

**Good**

```
[Kubernetes](http://kubernetes.io){: target="_blank" }
```

## 8.14.4

**Bad**

```
```bash
wget 127.0.0.1
```

```

**Good**

```
```bash
wget 127.0.0.1
```

```

**Bad**

```
```bash
wget 127.0.0.1
```

```

**Good**

```
```bash
wget 127.0.0.1
```

```



PDF



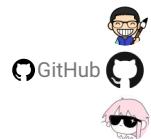
Slack



Support

⌚2025 9 10

⌚2022 7 19



8.14.5

---

9.

---

 [PDF](#) [Slack](#) [Support](#)

⌚2025 9 10

⌚2022 6 20



9.1

---